

PROBLEM for CS 179 :

by Prof. W. Kahan

Exhibit a program that starts from any three given floating-point numbers x , y and z , and computes $p := x \cdot y \cdot z$ in some order that avoids undeserved over/underflow. Do likewise for $q := x \cdot y / z$.

SOLUTIONS: The proofs that these programs work correctly depend upon the properties of three *Environmental Constants* associated with the floating-point formats in which x , y , z , p and q are represented, regardless of whether those constants appear in the programs. The *Overflow threshold* Ω is the biggest finite number in that format; the *Underflow threshold* η is the smallest *normalized* positive number. The magnitudes of x , y and z are presumed to lie between Ω and $\epsilon\eta$ inclusive where $\epsilon\eta$ is the smallest *nonzero* magnitude and may be far tinier than η if underflow is *gradual*; on machines that underflow abruptly to zero $\epsilon\eta = \eta$ except for CDC Cyber 17x's. $\epsilon\eta = 2\eta$ for these Cybers to cope with "partially underflowed" numbers between η and $\epsilon\eta$ that behave normally in *add*, *subtract* and *compare* but behave like zero in *multiply* and *divide*. Little is presumed about the product $\eta\Omega$, which lies very far from 1 on some machines.

An obvious program to compute p and q would first obtain their magnitudes using logarithms; $|p| = \exp(\ln|x| + \ln|y| + \ln|z|)$ and $|q| = \exp(\ln|x| + \ln|y| - \ln|z|)$. But these formulas lose accuracy badly when the data are very big or very small; the loss is caused by rounding each logarithm to working precision, and can be observed by comparing the computed values of $\exp(\ln|x|)$ and $|x|$ when it lies near Ω or η . And computing logarithms and exponentials wastes time. Our programs waste neither accuracy nor time.

Both programs start by Sorting $|x|$, $|y|$ and $|z|$ and continue thus:

Program for p :

Assume now that sorted $|x| \leq |y| \leq |z|$. Compute $x \cdot z$ first and then $p := (x \cdot z) \cdot y$ except on a machine with gradual underflow; on such a machine if $(x \cdot z)$ underflows recompute $p := (z \cdot y) \cdot x$.

Proof that p is correct.

If $x \cdot z$ overflowed, then $1 < |x| \leq |y| \leq \Omega < |x \cdot z| < |(x \cdot z) \cdot y|$ so p deserves to overflow too (except perhaps on a CRAY, which can overflow in certain cases when a product lies between $\Omega/2$ and Ω ; but that is too perverse to consider here). Similarly if $x \cdot z$ underflowed on a machine that underflows abruptly to zero, then

$$1 > |z| \geq |y| \geq |x| \geq \eta > |x \cdot z| > |(x \cdot z) \cdot y|$$

so p must underflow too. On a machine that underflows gradually conformity with IEEE standards 754/854 requires also the ability to detect underflow, and this should be exploited if any of the data can be subnormal (i.e., between $\epsilon\eta$ and η in magnitude). Then $x \cdot z$ underflows only when $1/\epsilon \geq |z| \geq |y| \geq |x| \geq \epsilon\eta$ and $\eta > |x \cdot z|$; since $\Omega > 1/\epsilon^2$ on those machines, $\Omega > z \cdot y$ so $z \cdot y$

cannot overflow and if it underflows too then either $|z| > 1$ and then $|x \cdot y \cdot z| = |(x \cdot z)(z \cdot y)/z| < \eta^2/|z| < \eta$, or else $|z| \leq 1$ and then $|x \cdot y \cdot z| < |x|\eta \leq \eta$, and p deserves to underflow either way.

Programs for q :

If we could treat q as a product $x \cdot y \cdot (1/z)$, we could compute it safely using the program for p ; but the risk that $1/z$ may over/underflow precludes that option. A safe and simple program works on machines that allow programs to branch on over/underflow:

```
First swap x and y if necessary to establish  $|x| \leq |y|$  ;
next compute  $p := x \cdot y$  ; subsequently
  if ( p overflowed and  $|z| > 1$  ) then  $q := (y/z) \cdot x$ 
  else if ( p underflowed and  $|z| < 1$  ) then
     $q := (((x/\varepsilon)/z) \cdot y) \cdot \varepsilon$ 
  else  $q := p/z$  . ( For Cybers use  $\varepsilon = 1$  here, not 2 .)
```

The validity of this program is easy to establish provided we may presume that $\sqrt{\eta}/\varepsilon^2 < \eta\Omega < \sqrt{\Omega}$, as appears to be true for all machines I know. But the ability to test for over/underflow and continue is not so common; what if over/underflow is silent? In the absence of a (portable) way to branch on over/underflow, we must produce a spaghetti-like code with branches that preclude spurious over/underflows. Such a program follows.

Two constants are needed. One is λ , the smallest power of the machine's radix no smaller than $\max\{1, 1/(\varepsilon\eta\Omega)\}$. The other is μ , the biggest power of the radix not exceeding $\min\{1, 1/(\eta\Omega)\}$. Multiplication by λ or μ is exact, so it cannot cause underflow on a machine that conforms to IEEE 754/854.

First sort $|x|$, $|y|$ and $|z|$, keeping track of z . This reduces the situation to one of three cases, depending upon whether $|z|$ is minimal, maximal, or neither:

```
In case  $|z|$  is minimal, say  $|z| \leq |x| \leq |y|$ , test  $|y|$  ;
  if  $|y| > 1$  then  $q := (x/z) \cdot y$ 
  else  $q := (x/(\lambda z)) \cdot (\lambda y)$  .
In case  $|z|$  is maximal, say  $|z| \geq |y| \geq |x|$ , test  $|x|$  ;
  if  $|x| < 1$  then  $q := (y/z) \cdot x$ 
  else  $q := (y/(\mu z)) \cdot (\mu x)$  .
In case  $|z|$  is neither, say  $|x| \leq |z| \leq |y|$ , test both;
  if  $|x| > 1$  then  $q := (y/z) \cdot x$ 
  else if  $|y| < 1$  then  $q := (x/z) \cdot y$ 
  else  $q := (x \cdot y)/z$  .
```

The proof that this program is correct is a tedious exercise in elementary inequalities, and is left to the reader.