

Notes on Kahan's *A Distillation Program*

Gideon Yuval

ABSTRACT

In the 8th set of Kahan course notes, there was a 1-page entry "A Distillation Program" (entry 3.5). My copy was almost illegible, which did not help any in trying to understand it.

The following is a recapitulation of the ideas in that Kahan note. It is more legible, and may or may not be clearer.

Exact 2-operand addition: Let Z be $p+q$ (machine precision, rounded), $\text{abs}(p) > \text{abs}(q)$. Let z be $(p-Z)+q$ (operations done at machine precision). Then $Z+z$ is exactly $p+q$; and z is as small as possible, among all pairs $Z+z$ for which this holds.

Exact N-operand addition: We try to span the nonzero bits of the exact N -operand sum with (approximately) as few terms as possible. This is called "distillation". We use two data-structures for this: one ("the list") holds all current terms, in order of absolute value. The other ("the heap") other holds pairs of successive list-entries, sorted in order of $\text{floor}(\log_2(p)) - \text{floor}(\log_2(q))$. In fact, all we need in the 2nd data-structure is the heap-property, i.e. the ability to find its smallest member fast.

A balanced tree (2-3 or AVL) will do fine for both data-structures ($\log N$ cost for any operation).

We can assume all entries are initially nonzero; otherwise, we just delete the zeros.

Repeat
 find smallest number in heap. Locate the two terms p, q , in list that correspond to it.
 Convert $p+q$ into $Z+z$, as above.
 Delete p and q from list (and update the heap accordingly)
 Insert Z into list (updating the heap accordingly)
 if $z \neq 0$, insert in list (updating the heap accordingly)
Until "convert" step is a no-op.

Termination: At every step, the values in the list become "more unequal": their sum remains constant, but the new list is "more exact" in the following sense: there are numbers such that

$$\sum X_i \text{ over } |X_i| > Z$$

becomes more accurate after the transformation than it was before. But the reverse does not occur.

Since this "more accurate" relation is a strict ordering, the iteration must terminate.

Final state: In the final state of the list, it contains N floating-point numbers, whose bit-positions never overlap (for otherwise, $p+q > Z+z$ would not have been a no-op). This number is thus within one of the optimum.