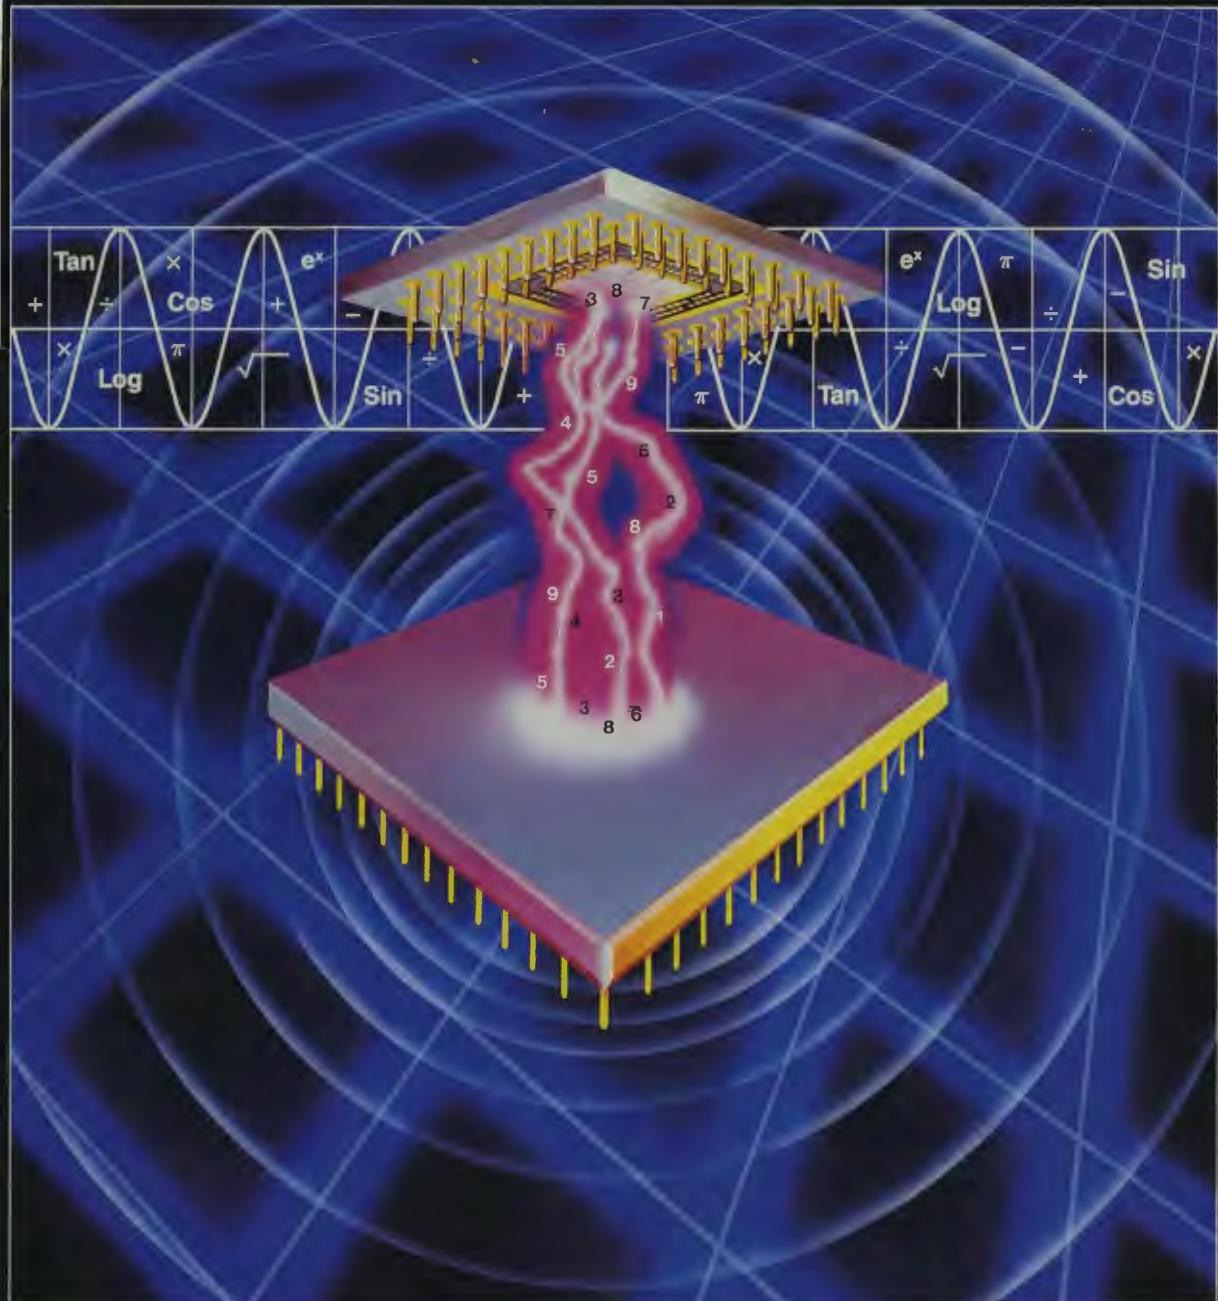


80387 Programmer's Reference Manual





LITERATURE SALES ORDER FORM

NAME: _____

COMPANY: _____

ADDRESS: _____

CITY: _____ STATE: _____ ZIP: _____

COUNTRY: _____

PHONE NO.: ()

ORDER NO.	TITLE	QTY.	PRICE	TOTAL
-----------	-------	------	-------	-------

ORDER NO.

TITLE

QT

PRICE

TOTAL

Subtotal _____

**Must Add Your
Local Sales Tax**

Must add appropriate postage to subtotal
(10% U.S. and Canada, 20% all other)

Postage

Postage

Total

Pay by Visa, MasterCard, American Express, Check, Money Order, or company purchase order payable to Intel Literature Sales. Allow 2-4 weeks for delivery.

Visa MasterCard American Express Expiration Date _____

Account No. _____

Mail To: Intel Literature Sales
P.O. Box 58130
Santa Clara, CA
95052-9130

International Customers outside the U.S. and Canada
should contact their local Intel Sales Office or Distributor
listed in the back of most Intel literature.

Call Toll Free: (800) 548-4725 for phone orders.

Prices good until 12/31/87

Source HR



**80387
PROGRAMMER'S REFERENCE
MANUAL**

1987

Intel Corporation makes no warranty for the use of its products and assumes no responsibility for any errors which may appear in this document nor does it make a commitment to update the information contained herein.

Intel retains the right to make changes to these specifications at any time, without notice.

Contact your local sales office to obtain the latest specifications before placing your order.

The following are trademarks of Intel Corporation and may only be used to identify Intel Products:

Above, BITBUS, COMMputer, CREDIT, Data Pipeline, FASTPATH, Genius, i, [△]i, ICE, iCEL, iCS, iDBP, iDIS, iICE, iLBX, im, iMDDX, iMMX, Inboard, Insite, Intel, intel, intelBOS, Intel Certified, Intelevision, intelligent Identifier, intelligent Programming, Intellec, Intellink, iOSP, iPDS, iPSC, iRMK, iRMX, iSBC, iSBX, iSDM, iSXM, KEPROM, Library Manager, MAPNET, MCS, Megachassis, MICROMAINFRAME, MULTIBUS, MULTICHANNEL, MULTIMODULE, MultiSERVER, ONCE, OpenNET, OTP, PC BUBBLE, Plug-A-Bubble, PROMPT, Promware, QUEST, QueX, Quick-Pulse Programming, Ripplemode, RMX/80, RUPI, Seamless, SLD, SugarCube, SupportNET, UPI, and VLSiCEL, and the combination of ICE, iCS, iRMX, iSBC, iSBX, iSXM, MCS, or UPI and a numerical suffix, 4-SITE.

MDS is an ordering code only and is not used as a product name or trademark. MDS® is a registered trademark of Mohawk Data Sciences Corporation.

*MULTIBUS is a patented Intel bus.

Unix is a trademark of AT&T Bell Labs.

MS-DOS, XENIX, and Multiplan are trademarks of Microsoft Corporation.

Lotus and 1-2-3 are registered trademarks of Lotus Development Corporation.

SuperCalc is a registered trademark of Computer Associates International.

Framework is a trademark of Ashton-Tate.

System 370 is a trademark of IBM Corporation.

AT is a registered trademark of IBM Corporation.

Additional copies of this manual or other Intel literature may be obtained from:

Intel Corporation
Literature Distribution
Mail Stop SC6-59
3065 Bowers Avenue
Santa Clara, CA 95051

PREFACE

This manual describes the 80387 Numeric Processor Extension (NPX) for the 80386 microprocessor. Understanding the 80387 requires an understanding of the 80386; therefore, a brief overview of 80386 concepts is presented first. A detailed discussion of the 80386 microprocessor can be found in the *80386 Programmer's Reference Manual*.

THE 80386 MICROSYSTEM

The 80386 is the basis of a new VLSI microprocessor system with exceptional capabilities for supporting large-system applications. This powerful microsystem is designed to support multiuser reprogrammable and real-time multitasking applications. Its dedicated system support circuits simplify system hardware; sophisticated hardware and software tools reduce both the time and the cost of product development. The 80386 microsystem offers a total-solution approach, enabling you to develop high-speed, interactive, multiuser, multitask— even multiprocessor—systems more rapidly and at higher performance than ever before.

- Reliability and system up-time are becoming increasingly important in all applications. Information must be protected from misuse or accidental loss. The 80386 includes a sophisticated and flexible four-level protection mechanism that can isolate layers of operating system programs from application programs to maintain a high degree of system integrity.
- The 80386 addresses up to 4 gigabytes of physical memory to support today's application requirements. This large physical memory enables the 80386 to keep many large programs and data structures simultaneously in memory for high-speed access.
- For applications with dynamically changing memory requirements, such as multiuser business systems, the 80386 CPU provides on-chip memory management and virtual memory support. On an 80386-based system, each user can have up to 64 terabytes of virtual-address space. This large address space virtually eliminates restrictions on the size of programs that may be part of the system. The memory management features are subject to control of systems software; therefore, systems software designers can choose among a variety of memory-organization models. Systems designers can choose to view memory in terms of fixed-length pages, in terms of variable length segments, or as a combination of pages and segments. The sizes of segments can range from one byte to 4 gigabytes. Virtual memory can be implemented either at the level of segments or at the level of pages.
- Large multiuser or real-time multitasking systems are easily supported by the 80386. High-performance features, such as a very high-speed task switch, fast interrupt-response time, intertask protection, page-oriented virtual memory, and a quick and direct operating system interface, make the 80386 highly suited to multiuser/multitasking applications.
- The 80386 has two primary operating modes: real-address mode and protected mode. In real-address mode, the 80386/80387 is fully upward compatible from the 8086, 8088, 80186, and 80188 microprocessors and from the 80286 real-address mode; all of the extensive libraries of 8086 and 8088 software execute 15 to 20 times faster on the 80386, without any modification.

- In protected-address mode, the advanced memory management and protection features of the 80386 become available, without any reduction in performance. Upgrading 8086 and 8088 application programs to use these new memory management and protection features usually requires only reassembly or recompilation (some programs may require minor modification). Entire 80286 protected-mode applications can run in this mode without modification.
- The virtual-8086 mode of the 80386 is available when the primary mode is protected mode. Virtual-8086 mode enables direct execution of multiple 8086/8088 programs within a protected-mode environment. Most 8086 and 8088 application programs can be executed in this environment without alteration (refer to the *80386 Programmer's Reference Manual* for differences from 8086). This high degree of compatibility between 80386 and earlier members of the 8086 processor family reduces both the time and the cost of software development.

THE ORGANIZATION OF THIS MANUAL

This manual describes the 80387 Numeric Processor Extension (NPX) for the 80386 microprocessor. The material in this manual is presented from the perspective of software designers, both at an applications and at a systems software level.

- Chapter 1, “Introduction to the 80387 Numerics Processor Extension,” gives an overview of the 80387 NPX and reviews the concepts of numeric computation using the 80387.
- Chapter 2, “80387 Numerics Processor Architecture,” presents the registers and data types of the 80387 to both applications and systems programmers.
- Chapter 3, “Special Computational Situations,” discusses the special values that can be represented in the 80387’s real formats—denormal numbers, zeros, infinities, NaNs (not a number)—as well as numerics exceptions. This chapter should be read thoroughly by systems programmers, but may be skimmed by applications programmers. Many of these special values and exceptions may never occur in applications programs.
- Chapter 4, “80387 Instruction Set,” provides functional information for software designers generating applications for systems containing an 80386 CPU with an 80387 NPX. The 80386/80387 instruction set mnemonics are explained in detail.
- Chapter 5, “Programming Numeric Applications,” provides a description of programming facilities for 80386/80387 systems. A comparative 80387 programming example is given.
- Chapter 6, “System-Level Numeric Programming,” provides information of interest to systems software writers, including details of the 80387 architecture and operational characteristics.
- Chapter 7, “Numeric Programming Examples,” provides several detailed programming examples for the 80387, including conditional branching, the conversion between floating-point values and their ASCII representations, and the use of trigonometric functions. These examples illustrate assembly-language programming on the 80387 NPX.
- Appendix A, “Machine Instruction Encoding and Decoding,” gives reference information on the encoding of NPX instructions. This information is useful to writers of debuggers, exception handlers, and compilers.

- Appendix B, "Exception Summary," provides a list of the exceptions that each instruction can cause. This list is valuable to both applications and systems programmers.
- Appendix C, "Compatibility between the 80387 and the 80287/8087," describes the differences from the 80387 that are common to the 80287 and the 8087.
- Appendix D, "Compatibility between the 80387 and the 8087," describes the additional differences between the 80387 and the 8087 that are of concern when porting 8086/8087 programs directly to the 80386/80387.
- Appendix E, "80387 80-Bit CMOS III Numeric Processor Extension," reproduces a data sheet of 80387 specifications that is separately available. The table of instruction timings in this appendix will be of interest to many readers of this manual. (The AC specifications have been deliberately left out.) The specifications in data sheets are subject to change; consult the most recent data sheet for design-in information.
- Appendix F, "PC/AT-Compatible 80387 Connection," documents a nonstandard method of connecting an 80387 to an 80386 to achieve compatibility with the IBM PC/AT.
- The Glossary defines 80387 and floating-point terminology. Refer to it as needed.

RELATED PUBLICATIONS

To best use the material in this manual, readers should be familiar with the operation and architecture of 80386 systems. The following manuals contain information related to the content of this manual and of interest to programmers of 80387 systems:

- *Introduction to the 80386*, order number 231252
- *80386 Data Sheet*, order number 231630
- *80386 Hardware Reference Manual*, order number 231732
- *80386 Programmer's Reference Manual*, order number 230985
- *80387 Data Sheet*, order number 231920

TABLE OF CONTENTS

	Page
CHAPTER 1	
INTRODUCTION TO THE 80387 NUMERICS PROCESSOR EXTENSION	
1.1 History	1-1
1.2 Performance	1-1
1.3 Ease of Use	1-3
1.4 Applications	1-4
1.5 Upgradability	1-5
1.6 Programming Interface	1-6
CHAPTER 2	
80387 NUMERICS PROCESSOR ARCHITECTURE	
2.1 80387 Registers	2-1
2.1.1 The NPX Register Stack	2-1
2.1.2 The NPX Status Word	2-3
2.1.3 Control Word	2-4
2.1.4 The NPX Tag Word	2-7
2.1.5 The NPX Instruction and Data Pointers	2-7
2.2 Computation Fundamentals	2-9
2.2.1 Number System	2-10
2.2.2 Data Types and Formats	2-11
2.2.2.1 Binary Integers	2-11
2.2.2.2 Decimal Integers	2-13
2.2.2.3 Real Numbers	2-13
2.2.3 Rounding Control	2-15
2.2.4 Precision Control	2-16
CHAPTER 3	
SPECIAL COMPUTATIONAL SITUATIONS	
3.1 Special Numeric Values	3-1
3.1.1 Denormal Real Numbers	3-1
3.1.1.1 Denormals and Gradual Underflow	3-4
3.1.2 Zeros	3-6
3.1.3 Infinity	3-9
3.1.4 NaN (Not-a-Number)	3-10
3.1.4.1 Signaling NaNs	3-11
3.1.4.2 Quiet NaNs	3-11
3.1.5 Indefinite	3-12
3.1.6 Encoding of Data Types	3-13
3.1.7 Unsupported Formats	3-13
3.2 Numeric Exceptions	3-18
3.2.1 Handling Numeric Exceptions	3-18
3.2.1.1 Automatic Exception Handling	3-18
3.2.1.2 Software Exception Handling	3-19

	Page
3.2.2 Invalid Operation	3-20
3.2.2.1 Stack Exception	3-20
3.2.2.2 Invalid Arithmetic Operation	3-21
3.2.3 Division by Zero	3-21
3.2.4 Denormal Operand	3-22
3.2.5 Numeric Overflow and Underflow	3-23
3.2.5.1 Overflow	3-23
3.2.5.2 Underflow	3-24
3.2.6 Inexact (Precision)	3-25
3.2.7 Exception Priority	3-26
3.2.8 Standard Underflow/Overflow Exception Handler	3-26

CHAPTER 4

THE 80387 INSTRUCTION SET

4.1 Compatibility with the 80287 and 8087	4-1
4.2 Numeric Operands	4-1
4.3 Data Transfer Instructions	4-2
4.3.1 FLD <i>source</i>	4-3
4.3.2 FST <i>destination</i>	4-3
4.3.3 FSTP <i>destination</i>	4-4
4.3.4 FXCH // <i>destination</i>	4-4
4.3.5 FILD <i>source</i>	4-4
4.3.6 FIST <i>destination</i>	4-4
4.3.7 FISTP <i>destination</i>	4-4
4.3.8 FBLD <i>source</i>	4-4
4.3.9 FBSTP <i>destination</i>	4-5
4.4 Nontranscendental Instructions	4-5
4.4.1 Addition	4-7
4.4.2 Normal Subtraction	4-8
4.4.3 Reversed Subtraction	4-8
4.4.4 Multiplication	4-8
4.4.5 Normal Division	4-8
4.4.6 Reversed Division	4-9
4.4.7 FSQRT	4-9
4.4.8 FSCALE	4-9
4.4.9 FPREM—Partial Remainder (80287/8087-Compatible)	4-9
4.4.10 FPREM1—Partial Remainder (IEEE Std. 754-Compatible)	4-10
4.4.11 FRNDINT	4-12
4.4.12 FXTRACT	4-12
4.4.13 FABS	4-13
4.4.14 FCHS	4-13
4.5 Comparison Instructions	4-13

	Page
4.5.1 FCOM //source	4-14
4.5.2 FCOMP //source	4-14
4.5.3 FCOMPP	4-14
4.5.4 FCOM source	4-14
4.5.5 FICOMP source	4-15
4.5.6 FTST	4-15
4.5.7 FUCOM //source	4-15
4.5.8 FUCOMP //source	4-15
4.5.9 FUCOMPP	4-15
4.5.10 FXAM	4-16
4.6 Transcendental Instructions	4-16
4.6.1 FCOS	4-17
4.6.2 FSIN	4-17
4.6.3 FSINCOS	4-17
4.6.4 FPTAN	4-17
4.6.5 FPATAN	4-18
4.6.6 F2XM1	4-18
4.6.7 FYL2X	4-19
4.6.8 FYL2XP1	4-19
4.7 Constant Instructions	4-19
4.7.1 FLDZ	4-20
4.7.2 FLD1	4-20
4.7.3 FLDPI	4-20
4.7.4 FLDL2T	4-20
4.7.5 FLDL2E	4-20
4.7.6 FLDLG2	4-20
4.7.7 FLDLN2	4-21
4.8 Processor Control Instructions	4-21
4.8.1 FINIT/FNINIT	4-22
4.8.2 FLDCW source	4-22
4.8.3 FSTCW/FNSTCW destination	4-22
4.8.4 FSTSW/FNSTSW destination	4-23
4.8.5 FSTSW AX/FNSTSW AX	4-23
4.8.6 FCLEX/FNCLEX	4-23
4.8.7 FSAVE/FNSAVE destination	4-23
4.8.8 FRSTOR source	4-25
4.8.9 FSTENV/FNSTENV destination	4-26
4.8.10 FLDENV source	4-26
4.8.11 FINCSTP	4-27
4.8.12 FDECSTP	4-28
4.8.13 FFREE destination	4-28
4.8.14 FNOP	4-28
4.8.15 FWAIT (CPU Instruction)	4-28

	Page
CHAPTER 5	
PROGRAMMING NUMERIC APPLICATIONS	
5.1 Programming Facilities	5-1
5.1.1 High-Level Languages	5-1
5.1.2 C Programs	5-1
5.1.3 PL/M-386	5-3
5.1.4 ASM386	5-4
5.1.4.1 Defining Data	5-4
5.1.4.2 Records and Structures	5-6
5.1.4.3 Addressing Methods	5-7
5.1.5 Comparative Programming Example	5-8
5.1.6 80387 Emulation	5-13
5.2 Concurrent Processing with the 80387	5-13
5.2.1 Managing Concurrency	5-14
5.2.1.1 Incorrect Exception Synchronization	5-16
5.2.1.2 Proper Exception Synchronization	5-16
CHAPTER 6	
SYSTEM-LEVEL NUMERIC PROGRAMMING	
6.1 80386/80387 Architecture	6-1
6.1.1 Instruction and Operand Transfer	6-1
6.1.2 Independent of CPU Addressing Modes	6-1
6.1.3 Dedicated I/O Locations	6-2
6.2 Processor Initialization and Control	6-2
6.2.1 System Initialization	6-2
6.2.2 Hardware Recognition of the NPX	6-2
6.2.3 Software Recognition of the NPX	6-3
6.2.4 Configuring the Numerics Environment	6-3
6.2.5 Initializing the 80387	6-5
6.2.6 80387 Emulation	6-6
6.2.7 Handling Numerics Exceptions	6-7
6.2.8 Simultaneous Exception Response	6-8
6.2.9 Exception Recovery Examples	6-8
CHAPTER 7	
NUMERIC PROGRAMMING EXAMPLES	
7.1 Conditional Branching Example	7-1
7.2 Exception Handling Examples	7-2
7.3 Floating-Point to ASCII Conversion Examples	7-6
7.3.1 Function Partitioning	7-18
7.3.2 Exception Considerations	7-18
7.3.3 Special Instructions	7-18
7.3.4 Description of Operation	7-19

	Page
7.3.5 Scaling the Value	7-19
7.3.5.1 Inaccuracy in Scaling	7-20
7.3.5.2 Avoiding Underflow and Overflow	7-20
7.3.5.3 Final Adjustments	7-20
7.3.6 Output Format	7-21
7.4 Trigonometric Calculation Examples (Not Tested)	7-21

**APPENDIX A
MACHINE INSTRUCTION ENCODING AND DECODING****APPENDIX B
EXCEPTION SUMMARY****APPENDIX C
COMPATIBILITY BETWEEN THE 80387 AND THE 80287/8087****APPENDIX D
COMPATIBILITY BETWEEN THE 80387 AND THE 8087****APPENDIX E
80387 80-BIT CMOS III NUMERIC PROCESSOR EXTENSION****APPENDIX F
PC/AT-COMPATIBLE 80387 CONNECTION****GLOSSARY OF 80387 AND FLOATING-POINT TERMINOLOGY****Figures**

Figure	Title	Page
1-1	Evolution and Performance of Numeric Processors	1-2
2-1	80387 Register Set	2-2
2-2	80387 Status Word	2-3
2-3	80387 Control Word Format	2-6
2-4	80387 Tag Word Format	2-7
2-5	Protected Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format	2-8
2-6	Real Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format	2-8

Figure	Title	Page
2-7	Protected Mode 80387 Instruction and Data Pointer Image in Memory, 16-Bit Format	2-9
2-8	Real Mode 80387 Instruction and Data Pointer Image in Memory, 16-Bit Format	2-9
2-9	80387 Double-Precision Number System	2-10
2-10	80387 Data Formats	2-12
3-1	Floating-Point System with Denormals	3-5
3-2	Floating-Point System without Denormals	3-5
3-3	Arithmetic Example Using Infinity	3-19
4-1	FSAVE/FRSTOR Memory Layout (32-Bit)	4-24
4-2	FSAVE/FRSTOR Memory Layout (16-Bit)	4-25
4-3	Protected Mode 80387 Environment, 32-Bit Format	4-26
4-4	Real Mode 80387 Environment, 32-Bit Format	4-27
4-5	Protected Mode 80387 Environment, 16-Bit Format	4-27
4-6	Real Mode 80387 Environment, 16-Bit Format	4-28
5-1	Sample C-386 Program	5-2
5-2	Sample 80387 Constants	5-5
5-3	Status Word Record Definition	5-6
5-4	Structure Definition	5-7
5-5	Sample PL/M-386 Program	5-9
5-6	Sample ASM386 Program	5-10
5-7	Instructions and Register Stack	5-12
5-8	Exception Synchronization Examples	5-15
6-1	Software Routine to Recognize the 80287	6-4
7-1	Conditional Branching for Compares	7-2
7-2	Conditional Branching for FXAM	7-3
7-3	Full-State Exception Handler	7-4
7-4	Reduced-Latency Exception Handler	7-5
7-5	Reentrant Exception Handler	7-6
7-6	Floating-Point to ASCII Conversion Routine	7-7
7-7	Relationships between Adjacent Joints	7-22
7-8	Robot Arm Kinematics Example	7-24

Tables

Table	Title	Page
1-1	Numeric Processing Speed Comparisons	1-2
1-2	Numeric Data Types	1-7
1-3	Principal NPX Instructions	1-8
2-1	Condition Code Interpretation	2-5
2-2	Correspondence between 80387 and 80386 Flag Bits	2-6

Table	Title	Page
2-3	Summary of Format Parameters	2-13
2-4	Real Number Notation	2-14
2-5	Rounding Modes	2-17
3-1	Arithmetic and Nonarithmetic Instructions	3-2
3-2	Denormalization Process	3-3
3-3	Zero Operands and Results	3-7
3-4	Infinity Operands and Results	3-9
3-5	Rules for Generating QNaNs	3-12
3-6	Binary Integer Encodings	3-14
3-7	Packed Decimal Encodings	3-15
3-8	Single and Double Real Encodings	3-16
3-9	Extended Real Encodings	3-17
3-10	Masked Responses to Invalid Operations	3-21
3-11	Masked Overflow Results	3-23
4-1	Data Transfer Instructions	4-3
4-2	Nontranscendental Instructions	4-6
4-3	Basic Nontranscendental Instructions and Operands	4-7
4-4	Condition Code Interpretation after FPREM and FPREM1 Instructions	4-11
4-5	Comparison Instructions	4-13
4-6	Condition Code Resulting from Comparisons	4-14
4-7	Condition Code Resulting from FTST	4-15
4-8	Condition Code Defining Operand Class	4-16
4-9	Transcendental Instructions	4-16
4-10	Results of FPATAN	4-18
4-11	Constant Instructions	4-20
4-12	Processor Control Instructions	4-21
5-1	PL/M-386 Built-In Procedures	5-3
5-2	ASM386 Storage Allocation Directives	5-4
5-3	Addressing Method Examples	5-7
6-1	NPX Processor State Following Initialization	6-6



CUSTOMER SUPPORT

CUSTOMER SUPPORT

Customer Support is Intel's complete support service that provides Intel customers with hardware support, software support, customer training, and consulting services. For more information contact your local sales offices.

After a customer purchases any system hardware or software product, service and support become major factors in determining whether that product will continue to meet a customer's expectations. Such support requires an international support organization and a breadth of programs to meet a variety of customer needs. As you might expect, Intel's customer support is quite extensive. It includes factory repair services and worldwide field service offices providing hardware repair services, software support services, customer training classes, and consulting services.

HARDWARE SUPPORT SERVICES

Intel is committed to providing an international service support package through a wide variety of service offerings available from Intel Hardware Support.

SOFTWARE SUPPORT SERVICES

Intel's software support consists of two levels of contracts. Standard support includes TIPS (Technical Information Phone Service), updates and subscription service (product-specific troubleshooting guides and COMMENTS Magazine). Basic support includes updates and the subscription service. Contracts are sold in environments which represent product groupings (i.e., iRMX environment).

CONSULTING SERVICES

Intel provides field systems engineering services for any phase of your development or support effort. You can use our systems engineers in a variety of ways ranging from assistance in using a new product, developing an application, personalizing training, and customizing or tailoring an Intel product to providing technical and management consulting. Systems Engineers are well versed in technical areas such as microcommunications, real-time applications, embedded microcontrollers, and network services. You know your application needs; we know our products. Working together we can help you get a successful product to market in the least possible time.

CUSTOMER TRAINING

Intel offers a wide range of instructional programs covering various aspects of system design and implementation. In just three to ten days a limited number of individuals learn more in a single workshop than in weeks of self-study. For optimum convenience, workshops are scheduled regularly at Training Centers worldwide or we can take our workshops to you for on-site instruction. Covering a wide variety of topics, Intel's major course categories include: architecture and assembly language, programming and operating systems, bitbus and LAN applications.

Introduction to the 80387 Numerics Processor Extension

1

CHAPTER 1

INTRODUCTION TO THE 80387 NUMERICS PROCESSOR EXTENSION

The 80387 NPX is a high-performance numerics processing element that extends the 80386 architecture by adding significant numeric capabilities and direct support for floating-point, extended-integer, and BCD data types. The 80386 CPU with 80387 NPX easily supports powerful and accurate numeric applications through its implementation of the IEEE Standard 754 for Binary Floating-Point Arithmetic. The 80387 provides floating-point performance comparable to that of large minicomputers while offering compatibility with object code for 8087 and 80287.

1.1 HISTORY

The 80387 Numeric Processor Extension (NPX) is compatible with its predecessors, the earlier Intel 8087 NPX and 80287 NPX. As the 80386 runs 8086 programs, so programs designed to use the 8087 and 80287 should run unchanged on the 80387.

The 8087 NPX was designed for use in 8086-family systems. The 8086 was the first microprocessor family to partition the processing unit to permit high-performance numeric capabilities. The 8087 NPX for this processor family implemented a complete numeric processing environment in compliance with an early proposal for the IEEE 754 Floating-Point Standard.

With the 80287 Numeric Processor Extension, high-speed numeric computations were extended to 80286 high-performance multitasking and multiuser systems. Multiple tasks using the numeric processor extension were afforded the full protection of the 80286 memory management and protection features.

The 80387 Numeric Processor Extension is Intel's third generation numerics processor. The 80387 implements the final IEEE standard, adds new trigonometric instructions, and uses a new design and CMOS-III process to allow higher clock rates and require fewer clocks per instruction. Together, the 80387 with additional instructions and the improved standard bring even more convenience and reliability to numerics programming and make this convenience and reliability available to applications that need the high-speed and large memory capacity of the 32-bit environment of the 80386 CPU.

Figure 1-1 illustrates the relative performance of 5-MHz 8086/8087, 8-MHz 80286/80287, and 20-MHz 80386/80387 systems in executing numerics-oriented applications.

1.2 PERFORMANCE

Table 1-1 compares the execution times of several 80387 instructions with the equivalent operations executed on an 8-MHz 80287. As indicated in the table, the 16-MHz 80387 NPX provides about 5 to 6 times the performance of an 8-MHz 80287 NPX. A 16-MHz

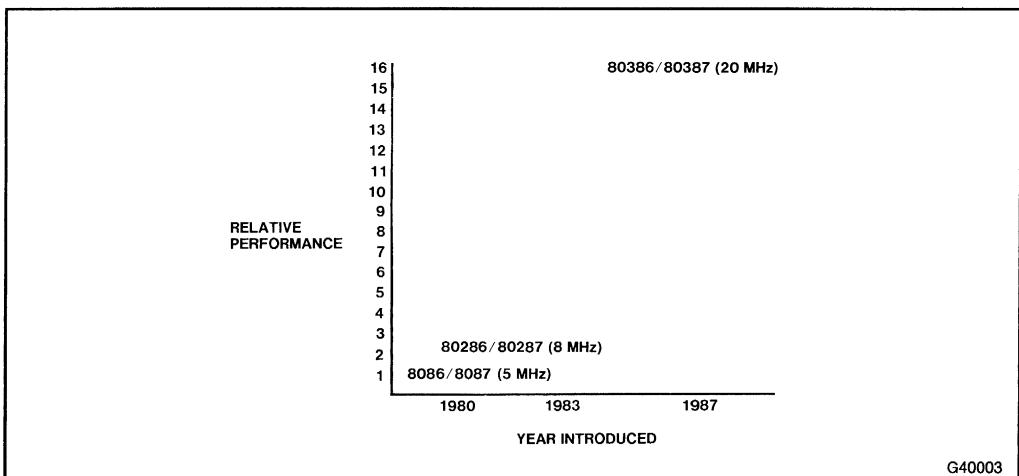


Figure 1-1. Evolution and Performance of Numeric Processors

Table 1-1. Numeric Processing Speed Comparisons

Floating-Point Instruction			Approximate Performance Ratios: 16 MHz 80386/80387 ÷ 8 MHz 80286/80287
FADD	ST, ST(i)	Addition	6.2
FDIV	dword_var	Division	4.7
FYL2X	stack (0), (1) assumed	Logarithm	6.0
FPATAX	stack (0) assumed	Arctangent	2.6*
F2XM1	stack (0) assumed	Exponentiation	2.7*

*The ratio is higher if the operand is not in range of the 80287 instruction.

80387 multiplies 32-bit and 64-bit floating-point numbers in about 1.9 and 2.8 microseconds, respectively. Of course, the actual performance of the NPX in a given system depends on the characteristics of the individual application.

Although the performance figures shown in Table 1-1 refer to operations on real (floating-point) numbers, the 80387 also manipulates fixed-point binary and decimal integers of up to 64 bits or 18 digits, respectively. The 80387 can improve the speed of multiple-precision software algorithms for integer operations by 10 to 100 times.

Because the 80387 NPX is an extension of the 80386 CPU, no software overhead is incurred in setting up the NPX for computation. The 80387 and 80386 processors coordinate their activities in a manner transparent to software. Moreover, built-in coordination facilities allow the 80386 CPU to proceed with other instructions while the 80387 NPX is simultaneously executing numeric instructions. Programs can exploit this concurrency of execution to further increase system performance and throughput.

1.3 EASE OF USE

The 80387 NPX offers more than raw execution speed for computation-intensive tasks. The 80387 brings the functionality and power of accurate numeric computation into the hands of the general user. These features are available in most high-level languages available for the 80386.

Like the 8087 and 80287 that preceded it, the 80387 is explicitly designed to deliver stable, accurate results when programmed using straightforward “pencil and paper” algorithms. The IEEE standard 754 specifically addresses this issue, recognizing the fundamental importance of making numeric computations both easy and safe to use.

For example, most computers can overflow when two single-precision floating-point numbers are multiplied together and then divided by a third, even if the final result is a perfectly valid 32-bit number. The 80387 delivers the correctly rounded result. Other typical examples of undesirable machine behavior in straightforward calculations occur when computing financial rate of return, which involves the expression $(1 + i)^n$ or when solving for roots of a quadratic equation:

$$\frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

If a does not equal 0, the formula is numerically unstable when the roots are nearly coincident or when their magnitudes are wildly different. The formula is also vulnerable to spurious over/underflows when the coefficients a , b , and c are all very big or all very tiny. When single-precision (4-byte) floating-point coefficients are given as data and the formula is evaluated in the 80387's normal way, keeping all intermediate results in its stack, the 80387 produces impeccable single-precision roots. This happens because, by default and with no effort on the programmer's part, the 80387 evaluates all those subexpressions with so much extra precision and range as to overwhelm any threat to numerical integrity.

If double-precision data and results were at issue, a better formula would have to be used, and once again the 80387's default evaluation of that formula would provide substantially enhanced numerical integrity over mere double-precision evaluation.

On most machines, straightforward algorithms will not deliver consistently correct results (and will not indicate when they are incorrect). To obtain correct results on traditional machines under all conditions usually requires sophisticated numerical techniques that are foreign to most programmers. General application programmers using straightforward algorithms will produce much more reliable programs using the 80387. This simple fact greatly reduces the software investment required to develop safe, accurate computation-based products.

Beyond traditional numerics support for scientific applications, the 80387 has built-in facilities for commercial computing. It can process decimal numbers of up to 18 digits without round-off errors, performing *exact arithmetic* on integers as large as 2^{64} or 10^{18} . Exact arithmetic is vital in accounting applications where rounding errors may introduce monetary losses that cannot be reconciled.

The NPX contains a number of optional facilities that can be invoked by sophisticated users. These advanced features include directed rounding, gradual underflow, and programmed exception-handling facilities.

These automatic exception-handling facilities permit a high degree of flexibility in numeric processing software, without burdening the programmer. While performing numeric calculations, the NPX automatically detects exception conditions that can potentially damage a calculation (for example, $X \div 0$ or \sqrt{X} when $X < 0$). By default, on-chip exception logic handles these exceptions so that a reasonable result is produced and execution may proceed without program interruption. Alternatively, the NPX can signal the CPU, invoking a software exception handler to provide special results whenever various types of exceptions are detected.

1.4 APPLICATIONS

The 80386's versatility and performance make it appropriate to a broad array of numeric applications. In general, applications that exhibit any of the following characteristics can benefit by implementing numeric processing on the 80387:

- Numeric data vary over a wide range of values, or include nonintegral values.
- Algorithms produce very large or very small intermediate results.
- Computations must be very precise; i.e., a large number of significant digits must be maintained.
- Performance requirements exceed the capacity of traditional microprocessors.
- Consistently safe, reliable results must be delivered using a programming staff that is not expert in numerical techniques.

Note also that the 80387 can reduce software development costs and improve the performance of systems that use not only real numbers, but operate on multiprecision binary or decimal integer values as well.

A few examples, which show how the 80387 might be used in specific numerics applications, are described below. In many cases, these types of systems have been implemented in the past with minicomputers or small mainframe computers. The advent of the 80387 brings the size and cost savings of microprocessor technology to these applications for the first time.

- Business data processing—The NPX's ability to accept decimal operands and produce *exact* decimal results of up to 18 digits greatly simplifies accounting programming. Financial calculations that use power functions can take advantage of the 80387's exponentiation and logarithmic instructions. Many business software packages can benefit from the speed and accuracy of the 80387; for example, Lotus* 1-2-3*, Multiplan*, SuperCalc*, and Framework*.

- Simulation—The large (32-bit) memory space of the 80386 coupled with the raw speed of the 80386 and 80387 processors make 80386/80387 microsystems suitable for attacking large simulation problems, which heretofore could only be executed on expensive mini and mainframe computers. For example, complex electronic circuit simulations using SPICE can now be performed on a microcomputer, the 80386/80387. Simulation of mechanical systems using finite element analysis can employ more elements, resulting in more detailed analysis or simulation of larger systems.
- Graphics transformations—The 80387 can be used in graphics terminals to locally perform many functions that normally demand the attention of a main computer; these include rotation, scaling, and interpolation. By also using an 82786 Graphics Display Controller to perform high-speed drawing and window management, very powerful and highly self-sufficient terminals can be built from a relatively small number of 80386 family parts.
- Process control—The 80387 solves dynamic range problems automatically, and its extended precision allows control functions to be fine-tuned for more accurate and efficient performance. Control algorithms implemented with the NPX also contribute to improved reliability and safety, while the 80387's speed can be exploited in real-time operations.
- Computer numerical control (CNC)—The 80387 can move and position machine tool heads with accuracy in real-time. Axis positioning also benefits from the hardware trigonometric support provided by the 80387.
- Robotics—Coupling small size and modest power requirements with powerful computational abilities, the 80387 is ideal for on-board six-axis positioning.
- Navigation—Very small, lightweight, and accurate inertial guidance systems can be implemented with the 80387. Its built-in trigonometric functions can speed and simplify the calculation of position from bearing data.
- Data acquisition—The 80387 can be used to scan, scale, and reduce large quantities of data as it is collected, thereby lowering storage requirements and time required to process the data for analysis.

The preceding examples are oriented toward *traditional* numerics applications. There are, in addition, many other types of systems that do not appear to the end user as *computational*, but can employ the 80387 to advantage. Indeed, the 80387 presents the imaginative system designer with an opportunity similar to that created by the introduction of the microprocessor itself. Many applications can be viewed as numerically-based if sufficient computational power is available to support this view (e.g., character generation for a laser printer). This is analogous to the thousands of successful products that have been built around “buried” microprocessors, even though the products themselves bear little resemblance to computers.

1.5 UPGRADABILITY

The architecture of the 80386 CPU is specifically adapted to allow easy upgradability to use an 80387, simply by plugging in the 80387 NPX. For this reason, designers of 80386 systems may wish to incorporate the 80387 NPX into their designs in order to offer two levels of price and performance at little additional cost.

Two features of the 80386 CPU make the design and support of upgradable 80386 systems particularly simple:

- The 80386 can be programmed to recognize the presence of an 80387 NPX; that is, software can recognize whether it is running on an 80386 with or without an 80387 NPX.
- After determining whether the 80387 NPX is available, the 80386 CPU can be instructed to let the NPX execute all numeric instructions. If an 80387 NPX is not available, the 80386 CPU can emulate all 80387 numeric instructions in software. This emulation is completely transparent to the application software—the same object code may be used by 80386 systems both with and without an 80387 NPX. No relinking or recompiling of application software is necessary; the same code will simply execute faster with the 80387 NPX than without.

To facilitate this design of upgradable 80386 systems, Intel provides a software emulator for the 80387 that provides the functional equivalent of the 80387 hardware, implemented in software on the 80386. Except for timing, the operation of this 80387 emulator (EMUL387) is the same as for the 80387 NPX hardware. When the emulator is combined as part of the systems software, the 80386 system with 80387 emulation and the 80386 with 80387 hardware are virtually indistinguishable to an application program. This capability makes it easy for software developers to maintain a single set of programs for both systems. System manufacturers can offer the NPX as a simple plug-in performance option without necessitating any changes in the user's software.

1.6 PROGRAMMING INTERFACE

The 80386/80387 pair is programmed as a single processor; all of the 80387 registers appear to a programmer as extensions of the basic 80386 register set. The 80386 has a class of instructions known as ESCAPE instructions, all having a common format. These ESC instructions are numeric instructions for the 80387 NPX. These numeric instructions for the 80387 are simply encoded into the instruction stream along with 80386 instructions.

All of the CPU memory-addressing modes may be used in programming the NPX, allowing convenient access to record structures, numeric arrays, and other memory-based data structures. All of the memory management and protection features of the CPU (both paging and segmentation) are extended to the NPX as well.

Numeric processing in the 80387 centers around the NPX register stack. Programmers can treat these eight 80-bit registers either as a fixed register set, with instructions operating on explicitly-designated registers, or as a classical stack, with instructions operating on the top one or two stack elements.

Internally, the 80387 holds all numbers in a uniform 80-bit extended format. Operands that may be represented in memory as 16-, 32-, or 64-bit integers, 32-, 64-, or 80-bit floating-point numbers, or 18-digit packed BCD numbers, are automatically converted into extended format as they are loaded into the NPX registers. Computation results are subsequently converted back into one of these destination data formats when they are stored into memory from the NPX registers.

Table 1-2 lists each of the seven data types supported by the 80387, showing the data format for each type. All operands are stored in memory with the least significant digits starting at the initial (lowest) memory address. Numeric instructions access and store memory operands using only this initial address. For maximum system performance, all operands should start at memory addresses divisible by four.

Table 1-3 lists the 80387 instructions by class. No special programming tools are necessary to use the 80387, because all of the NPX instructions and data types are directly supported by the ASM386 Assembler, by high-level languages from Intel, and by assemblers and compilers produced by many independent software vendors. Software routines for the 80387 may be written in ASM386 Assembler or any of the following higher-level languages from Intel:

PL/M-386
C-386

In addition, all of the development tools supporting the 8086/8087 and 80286/80287 can also be used to develop software for the 80386/80387.

All of these high-level languages provide programmers with access to the computational power and speed of the 80387 without requiring an understanding of the architecture of the 80386 and 80387 chips. Such architectural considerations as concurrency and synchronization are handled automatically by these high-level languages. For the ASM386 programmer, specific rules for handling these issues are discussed in a later section of this manual.

The following operating systems are known or expected to support the 80387: RMX-286/386, MS-DOS, Xenix-286/386, and Unix-286/386. Advanced in-circuit debugging support is provided by ICE-386.

Table 1-2. Numeric Data Types

Data Type	Bits	Significant Digits (Decimal)	Approximate Range (Decimal)
Word integer	16	4	$-32,768 \leq X \leq +32,767$
Short integer	32	9	$-2 \times 10^9 \leq X \leq +2 \times 10^9$
Long integer	64	18	$-9 \times 10^{18} \leq X \leq +9 \times 10^{18}$
Packed decimal	80	18	$-99\ldots99 \leq X \leq +99\ldots99$ (18 digits)
Single real	32	6-7	$1.18 \times 10^{-38} \leq X \leq 3.40 \times 10^{38}$
Double real	64	15-16	$2.23 \times 10^{-308} \leq X \leq 1.80 \times 10^{308}$
Extended real*	80	19	$3.30 \times 10^{-4932} \leq X \leq 1.21 \times 10^{4932}$

*Equivalent to *double extended* format of IEEE Std 754

Table 1-3. Principal NPX Instructions

Class	Instruction Types
Data Transfer	Load (all data types), Store (all data types), Exchange
Arithmetic	Add, Subtract, Multiply, Divide, Subtract Reversed, Divide Reversed, Square Root, Scale, Remainder, Integer Part, Change Sign, Absolute Value, Extract
Comparison	Compare, Examine, Test
Transcendental	Tangent, Arctangent, Sine, Cosine, Sine and Cosine, $2^x - 1$, $Y \cdot \log_2(X)$, $Y \cdot \log_2(X + 1)$
Constants	0, 1, π , $\log_{10}2$, \log_22 , $\log_{10}10$, \log_2e
Processor Control	Load Control Word, Store Control Word, Store Status Word, Load Environment, Store Environment, Save, Restore, Clear Exceptions, Initialize

*80387 Numerics
Processor Architecture*

2

CHAPTER 2

80387 NUMERICS PROCESSOR ARCHITECTURE

To the programmer, the 80387 NPX appears as a set of additional registers, data types, and instructions—all of which complement those of the 80386. Refer to Chapter 4 for detailed explanations of the 80387 instruction set. This chapter explains the new registers and data types that the 80387 brings to the architecture of the 80386.

2.1 80387 REGISTERS

The additional registers consist of

- Eight individually-addressable 80-bit numeric registers, organized as a register stack
- Three sixteen-bit registers containing:
 - the NPX status word
 - the NPX control word
 - the tag word
- Two 48-bit registers containing pointers to the current instruction and operand (these registers are actually located in the 80386)

All of the NPX numeric instructions focus on the contents of these NPX registers.

2.1.1 The NPX Register Stack

The 80387 register stack is shown in Figure 2-1. Each of the eight numeric registers in the 80387's register stack is 80 bits wide and is divided into fields corresponding to the NPX's extended real data type.

Numeric instructions address the data registers relative to the register on the top of the stack. At any point in time, this top-of-stack register is indicated by the TOP (stack TOP) field in the NPX status word. Load or push operations decrement TOP by one and load a value into the new top register. A store-and-pop operation stores the value from the current TOP register and then increments TOP by one. Like 80386 stacks in memory, the 80387 register stack grows *down* toward lower-addressed registers.

Many numeric instructions have several addressing modes that permit the programmer to implicitly operate on the top of the stack, or to explicitly operate on specific registers relative to the TOP. The ASM386 Assembler supports these register addressing modes, using the expression ST(0), or simply ST, to represent the current Stack Top and ST(*i*) to specify the

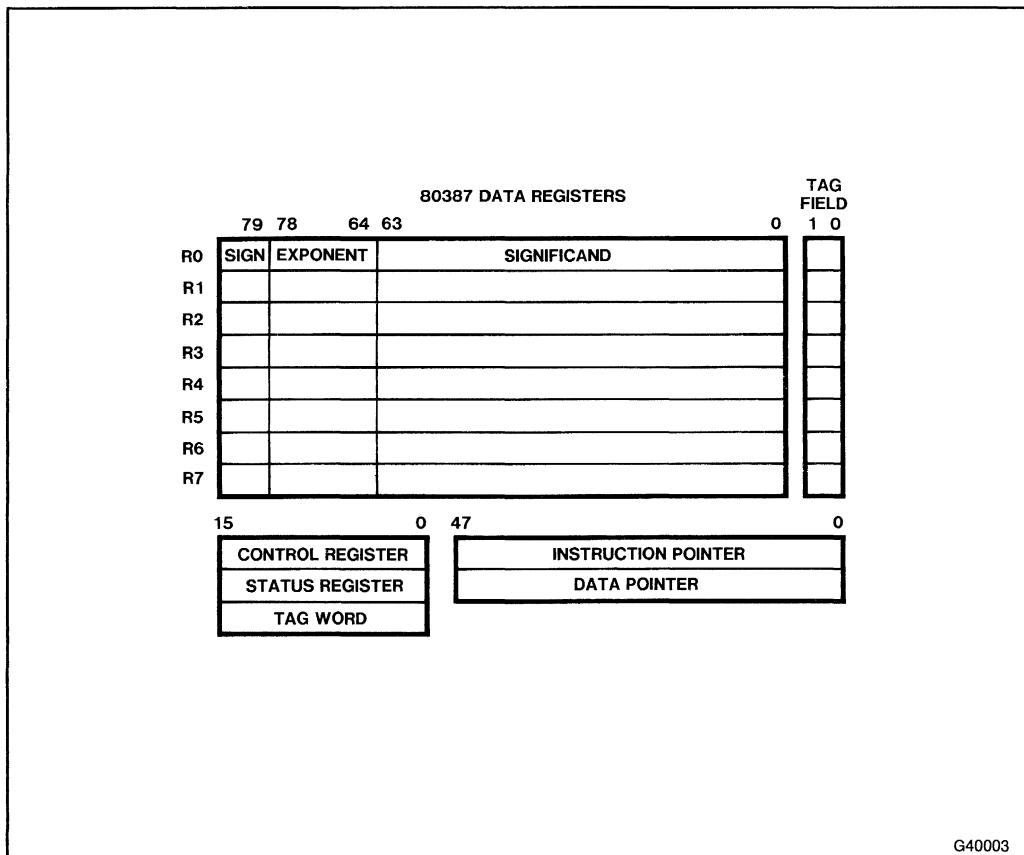


Figure 2-1. 80387 Register Set

G40003

ith register from TOP in the stack ($0 \leq i \leq 7$). For example, if TOP contains 011B (register 3 is the top of the stack), the following statement would add the contents of two registers in the stack (registers 3 and 5):

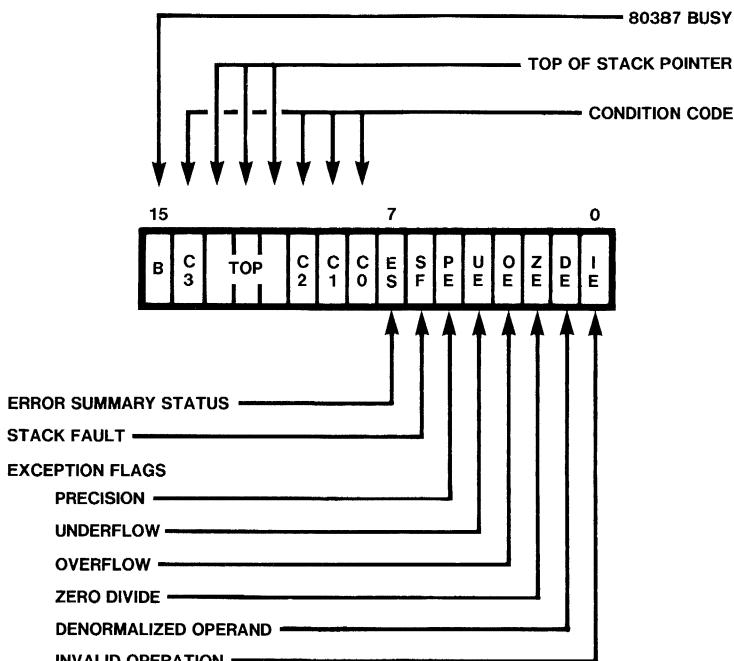
```
FADD ST, ST(2)
```

The stack organization and top-relative addressing of the numeric registers simplify subroutine programming by allowing routines to pass parameters on the register stack. By using the stack to pass parameters rather than using “dedicated” registers, calling routines gain more flexibility in how they use the stack. As long as the stack is not full, each routine simply loads the parameters onto the stack before calling a particular subroutine to perform a numeric calculation. The subroutine then addresses its parameters as ST, ST(1), etc., even though TOP may, for example, refer to physical register 3 in one invocation and physical register 5 in another.

2.1.2 The NPX Status Word

The 16-bit status word shown in Figure 2-2 reflects the overall state of the 80387. This status word may be stored into memory using the FSTSW/FNSTSW, FSTENV/FNSTENV, and FSAVE/FNSAVE instructions, and can be transferred into the 80386 AX register with the FSTSW AX/FNSTSW AX instructions, allowing the NPX status to be inspected by the CPU.

The B-bit (bit 15) is included for 8087 compatibility only. It reflects the contents of the ES bit (bit 7 of the status word), not the status of the BUSY# output of the 80387.



ES IS SET IF ANY UNMASKED EXCEPTION BIT IS SET; CLEARED OTHERWISE.
SEE TABLE 2-1 FOR INTERPRETATION OF CONDITION CODE.

TOP VALUES:

000 = REGISTER 0 IS TOP OF STACK
001 = REGISTER 1 IS TOP OF STACK

111 = REGISTER 7 IS TOP OF STACK

FOR DEFINITIONS OF EXCEPTIONS, REFER TO CHAPTER 3.

G40003

Figure 2-2. 80387 Status Word

The four NPX condition code bits (C_3-C_0) are similar to the flags in a CPU: the 80387 updates these bits to reflect the outcome of arithmetic operations. The effect of these instructions on the condition code bits is summarized in Table 2-1. These condition code bits are used principally for conditional branching. The FSTSW AX instruction stores the NPX status word directly into the CPU AX register, allowing these condition codes to be inspected efficiently by 80386 code. The 80386 SAHF instruction can copy C_3-C_0 directly to 80386 flag bits to simplify conditional branching. Table 2-2 shows the mapping of these bits to the 80386 flag bits.

Bits 12-14 of the status word point to the 80387 register that is the current Top of Stack (TOP). The significance of the stack top has been described in the prior section on the register stack.

Figure 2-2 shows the six exception flags in bits 0-5 of the status word. Bit 7 is the exception summary status (ES) bit. ES is set if any unmasked exception bits are set, and is cleared otherwise. If this bit is set, the ERROR# signal is asserted. Bits 0-5 indicate whether the NPX has detected one of six possible exception conditions since these status bits were last cleared or reset. They are “sticky” bits, and can only be cleared by the instructions FINIT, FCLEX, FLDENV, FSAVE, and FRSTOR.

Bit 6 is the stack fault (SF) bit. This bit distinguishes invalid operations due to stack overflow or underflow from other kinds of invalid operations. When SF is set, bit 9 (C_1) distinguishes between stack overflow ($C_1 = 1$) and underflow ($C_1 = 0$).

2.1.3 Control Word

The NPX provides the programmer with several processing options, which are selected by loading a word from memory into the control word. Figure 2-3 shows the format and encoding of the fields in the control word.

The low-order byte of this control word configures the 80387 exception masking. Bits 0-5 of the control word contain individual masks for each of the six exception conditions recognized by the 80387. The high-order byte of the control word configures the 80387 processing options, including

- Precision control
- Rounding control

The precision-control bits (bits 8-9) can be used to set the 80387 internal operating precision at less than the default precision (64-bit significand). These control bits can be used to provide compatibility with the earlier-generation arithmetic processors having less precision than the 80387. The precision-control bits affect the results of only the following five arithmetic instructions: ADD, SUB(R), MUL, DIV(R), and SQRT. No other operations are affected by PC.

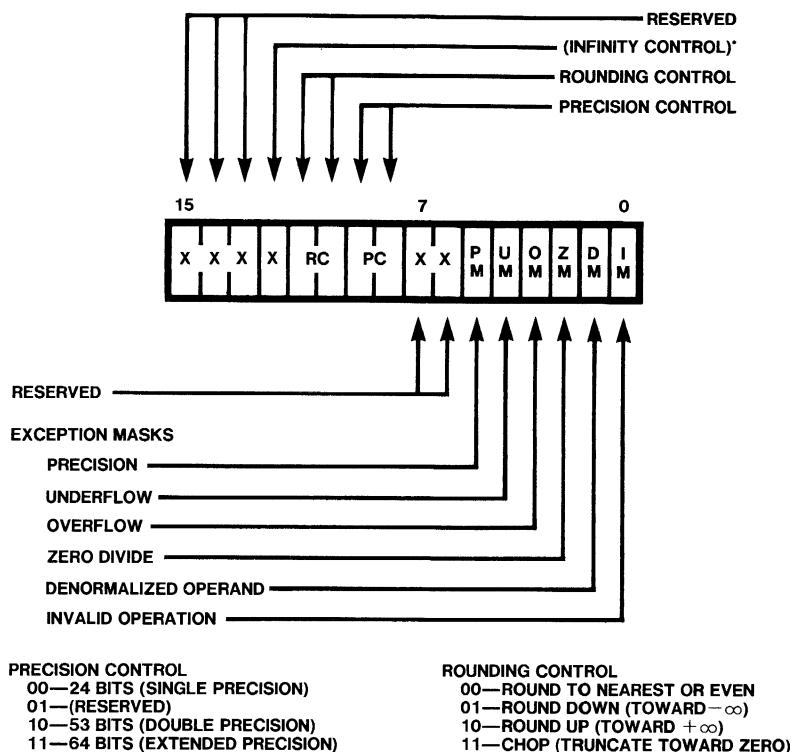
Table 2-1. Condition Code Interpretation

Instruction	C0 (S)	C3 (Z)	C1 (A)	C2 (C)
FPREM,FPREM1	Three least significant bits of quotient			Reduction 0=complete 1=incomplete
	Q2	Q0	Q1 or O/U#	
FCOM, FCOMP, FCOMPP, FTST, FUCOM, FUCOMP, FUCOMPP, FICOM, FICOMP	Result of comparison		Zero or O/U#	Operand is not comparable
FXAM	Operand class		Sign or O/U#	Operand class
FCHS, FABS, FXCH, FINCTOP, FDECSTOP, Constant loads, FXTRACT, FLD, FILD, FBLD, FSTP (ext real)	UNDEFINED		Zero or O/U#	UNDEFINED
FIST, FBSTP, FRNDINT, FST, FSTP, FADD, FMUL, FDIV, FDIVR, FSUB, FSUBR, FSCALE, FSQRT, FPATAN, F2XM1, FYL2X, FYL2XP1	UNDEFINED		Roundup or O/U#	UNDEFINED
FPTAN, FSIN, FCOS, FSINCOS	UNDEFINED		Roundup or O/U# undefined if C2=1	Reduction 0=complete 1=incomplete
FLDENV, FRSTOR	Each bit loaded from memory			
FLDCW, FSTENV, FSTCW, FSTSW, FCLEX, FINIT, FSAVE	UNDEFINED			

- O/U# When both IE and SF bits of status word are set, indicating a stack exception, this bit distinguishes between stack overflow (C1=1) and underflow (C1=0).
- Reduction If FPREM and FPREM1 produces a remainder that is less than the modulus, reduction is complete. When reduction is incomplete the value at the top of the stack is a partial remainder, which can be used as input to further reduction. For FPTAN, FSIN, FCOS, and FSINCOS, the reduction bit is set if the operand at the top of the stack is too large. In this case the original operand remains at the top of the stack.
- Roundup When the PE bit of the status word is set, this bit indicates whether the last rounding in the instruction was upward.
- UNDEFINED Do not rely on finding any specific value in these bits.

Table 2-2. Correspondence between 80387 and 80386 Flag Bits

80387 Flag	80386 Flag
C_0	CF
C_1	(none)
C_2	PF
C_3	ZF



*This “infinity control” bit is not meaningful to the 80387. To maintain compatibility with the 80287, this bit can be programmed; however, regardless of its value, the 80387 treats infinity in the affine sense ($-\infty < +\infty$).

G40003

Figure 2-3. 80387 Control Word Format

The rounding-control bits (bits 10–11) provide for the common round-to-nearest mode, as well as directed rounding and true chop. Rounding control affects only the arithmetic instructions (refer to Chapter 3 for lists of arithmetic and nonarithmetic instructions).

2.1.4 The NPX Tag Word

The tag word indicates the contents of each register in the register stack, as shown in Figure 2-4. The tag word is used by the NPX itself to distinguish between empty and nonempty register locations. Programmers of exception handlers may use this tag information to check the contents of a numeric register without performing complex decoding of the actual data in the register. The tag values from the tag word correspond to physical registers 0–7. Programmers must use the current top-of-stack (TOP) pointer stored in the NPX status word to associate these tag values with the relative stack registers ST(0) through ST(7).

The exact values of the tags are generated during execution of the FSTENV and FSAVE instructions according to the actual contents of the nonempty stack locations. During execution of other instructions, the 80387 updates the TW only to indicate whether a stack location is empty or nonempty.

2.1.5 The NPX Instruction and Data Pointers

The instruction and data pointers provide support for programmed exception-handlers. These registers are actually located in the 80386, but appear to be located in the 80387 because they are accessed by the ESC instructions FLDENV, FSTENV, FSAVE, and FRSTOR. Whenever the 80386 decodes an ESC instruction, it saves the instruction address, the operand address (if present), and the instruction opcode.

When stored in memory, the instruction and data pointers appear in one of four formats, depending on the operating mode of the 80386 (protected mode or real-address mode) and depending on the operand-size attribute in effect (32-bit operand or 16-bit operand). When the 80386 is in virtual-8086 mode, the real-address mode formats are used.

Figures 2-5 through 2-8 show these pointers as they are stored following an FSTENV instruction.

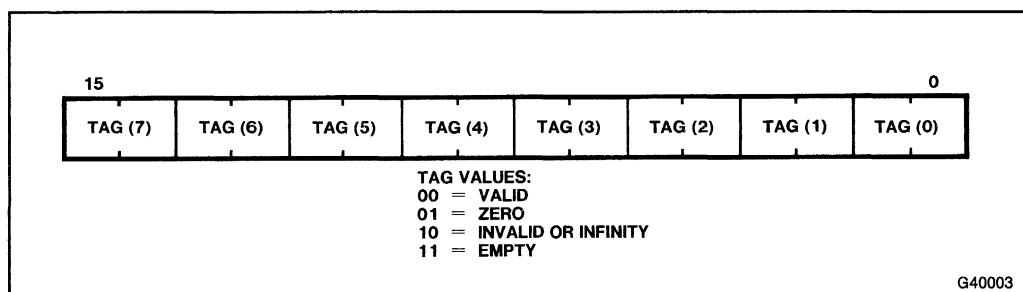


Figure 2-4. 80387 Tag Word Format

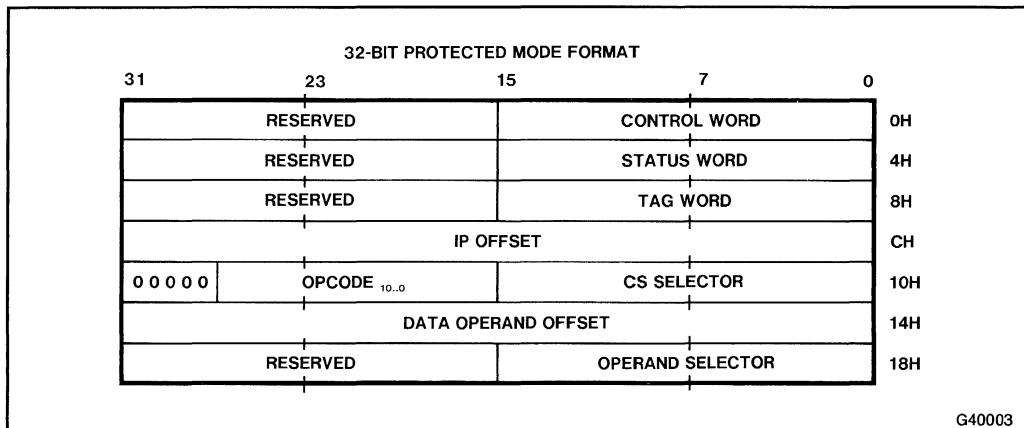


Figure 2-5. Protected Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format

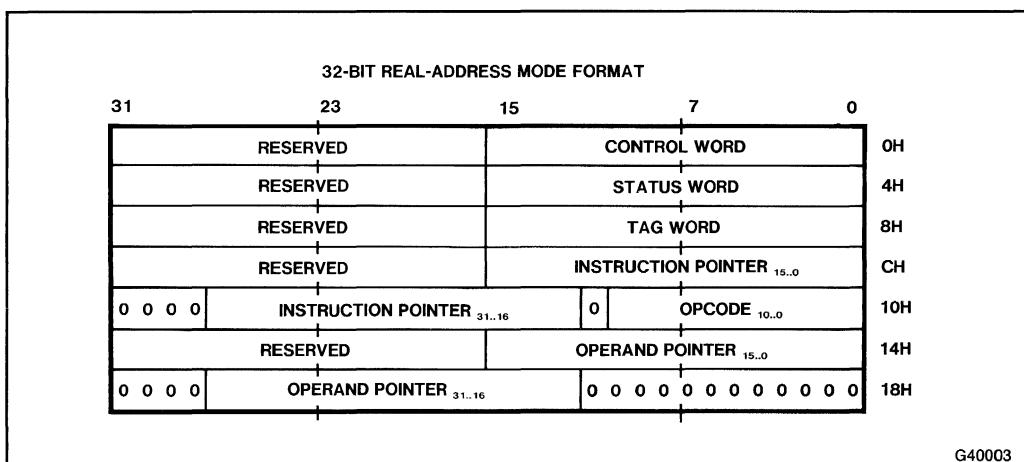
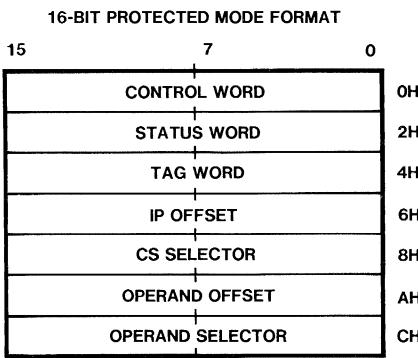


Figure 2-6. Real Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format

The FSTENV and FSAVE instructions store this data into memory, allowing exception handlers to determine the precise nature of any numeric exceptions that may be encountered.

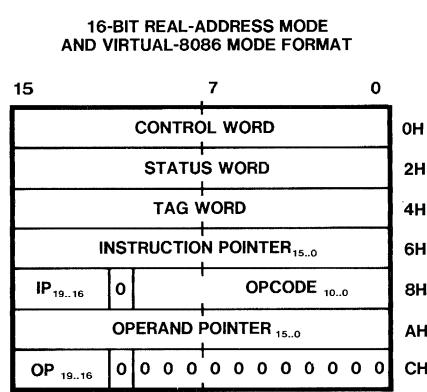
The instruction address saved in the 80386 (as in the 80287) points to any prefixes that preceded the instruction. This is different from the 8087, for which the instruction address points only to the ESC instruction opcode.

Note that the processor control instructions FINIT, FLDCW, FSTCW, FSTSW, FCLEX, FSTENV, FLDENV, FSAVE, FRSTOR, and FWAIT do not affect the data pointer. Note also that, except for the instructions just mentioned, the value of the data pointer is *undefined* if the prior ESC instruction did not have a memory operand.



G40003

**Figure 2-7. Protected Mode 80387 Instruction and Data Pointer Image in Memory,
16-Bit Format**



G40003

Figure 2-8. Real Mode 80387 Instruction and Data Pointer Image in Memory, 16-Bit Format

2.2 COMPUTATION FUNDAMENTALS

This section covers 80387 programming concepts that are common to all applications. It describes the 80387's internal number system and the various types of numbers that can be employed in NPX programs. The most commonly used options for rounding and precision (selected by fields in the control word) are described, with exhaustive coverage of less frequently used facilities deferred to later sections. Exception conditions that may arise during execution of NPX instructions are also described along with the options that are available for responding to these exceptions.

2.2.1 Number System

The system of real numbers that people use for pencil and paper calculations is conceptually infinite and continuous. There is no upper or lower limit to the magnitude of the numbers one can employ in a calculation, or to the precision (number of significant digits) that the numbers can represent. When considering any real number, there are always arbitrarily many numbers both larger and smaller. There are also arbitrarily many numbers between (i.e., with more significant digits than) any two real numbers. For example, between 2.5 and 2.6 are 2.51, 2.5897, 2.500001, etc.

While ideally it would be desirable for a computer to be able to operate on the entire real number system, in practice this is not possible. Computers, no matter how large, ultimately have fixed-size registers and memories that limit the system of numbers that can be accommodated. These limitations determine both the range and the precision of numbers. The result is a set of numbers that is finite and discrete, rather than infinite and continuous. This sequence is a subset of the real numbers that is designed to form a useful *approximation* of the real number system.

Figure 2-9 superimposes the basic 80387 real number system on a real number line (decimal numbers are shown for clarity, although the 80387 actually represents numbers in binary). The dots indicate the subset of real numbers the 80387 can represent as data and final results of calculations. The 80387's range of double-precision, normalized numbers is approximately $\pm 2.23 \times 10^{-308}$ to $\pm 1.80 \times 10^{308}$. Applications that are required to deal with data and final results outside this range are rare. For reference, the range of the IBM System 370* is about $\pm 0.54 \times 10^{-78}$ to $\pm 0.72 \times 10^{76}$.

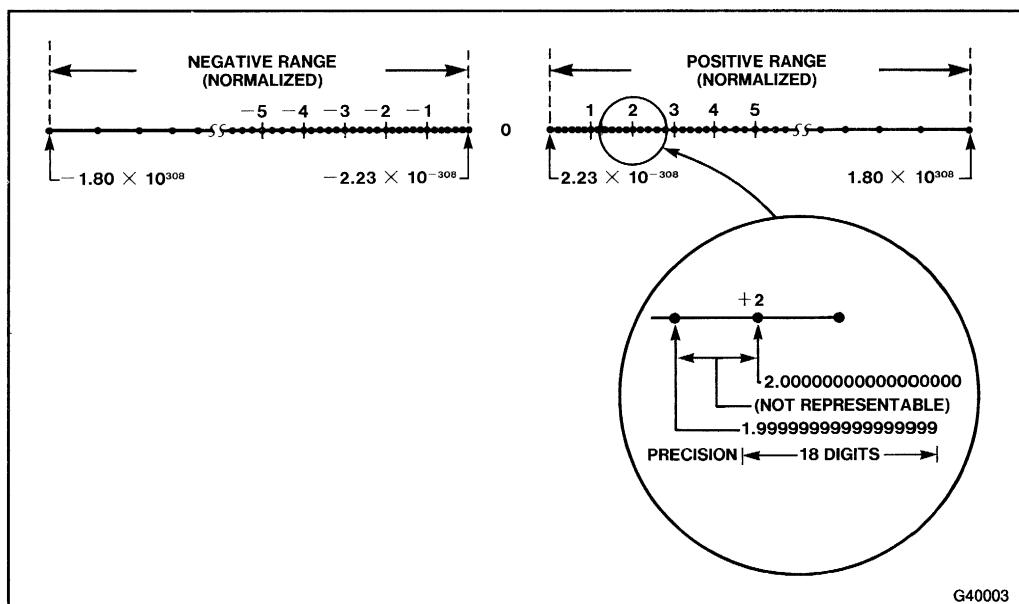


Figure 2-9. 80387 Double-Precision Number System

The finite spacing in Figure 2-9 illustrates that the NPX can represent a great many, but not all, of the real numbers in its range. There is always a gap between two adjacent 80387 numbers, and it is possible for the result of a calculation to fall in this space. When this occurs, the NPX rounds the true result to a number that it can represent. Thus, a real number that requires more digits than the 80387 can accommodate (e.g., a 20-digit number) is represented with some loss of accuracy. Notice also that the 80387's representable numbers are not distributed evenly along the real number line. In fact, an equal number of representable numbers exists between successive powers of 2 (i.e., as many representable numbers exist between 2 and 4 as between 65,536 and 131,072). Therefore, the gaps between representable numbers are larger as the numbers increase in magnitude. All integers in the range $\pm 2^{64}$ (approximately $\pm 10^{18}$), however, are exactly representable.

In its internal operations, the 80387 actually employs a number system that is a substantial superset of that shown in Figure 2-9. The internal format (called extended real) extends the 80387's range to about $\pm 3.30 \times 10^{-4932}$ to $\pm 1.21 \times 10^{4932}$, and its precision to about 19 (equivalent decimal) digits. This format is designed to provide extra range and precision for constants and intermediate results, and is not normally intended for data or final results.

From a practical standpoint, the 80387's set of real numbers is sufficiently large and dense so as not to limit the vast majority of microprocessor applications. Compared to most computers, including mainframes, the NPX provides a very good approximation of the real number system. It is important to remember, however, that it is not an exact representation, and that arithmetic on real numbers is inherently approximate.

Conversely, and equally important, the 80387 *does* perform exact arithmetic on integer operands. That is, if an operation on two integers is valid and produces a result that is in range, the result is exact. For example, $4 \div 2$ yields an exact integer, $1 \div 3$ does not, and $2^{40} \times 2^{30} + 1$ does not, because the result requires greater than 64 bits of precision.

2.2.2 Data Types and Formats

The 80387 recognizes seven numeric data types for memory-based values, divided into three classes: binary integers, packed decimal integers, and binary reals. A later section describes how these formats are stored in memory (the sign is always located in the highest-addressed byte).

Figure 2-10 summarizes the format of each data type. In the figure, the most significant digits of all numbers (and fields within numbers) are the leftmost digits.

2.2.2.1 BINARY INTEGERS

The three binary integer formats are identical except for length, which governs the range that can be accommodated in each format. The leftmost bit is interpreted as the number's sign: 0=positive and 1=negative. Negative numbers are represented in standard two's complement notation (the binary integers are the only 80387 format to use two's complement). The quantity zero is represented with a positive sign (all bits are 0). The 80387 word integer format is identical to the 16-bit signed integer data type of the 80386; the 80387 short integer format is identical to the 32-bit signed integer data type of the 80386.

DATA FORMATS	RANGE	PRECISION	MOST SIGNIFICANT BYTE								HIGHEST ADDRESSED BYTE											
			7	0	7	0	7	0	7	0	7	0	7	0	7	0	7	0				
WORD INTEGER	10^4	16 BITS													(TWO'S COMPLEMENT)							
			15	0																		
SHORT INTEGER	10^2	32 BITS													(TWO'S COMPLEMENT)							
			31	0																		
LONG INTEGER	10^{19}	64 BITS													(TWO'S COMPLEMENT)							
			63	0																		
PACKED BCD	10^{18}	18 DIGITS	S	X	d ₁₇	d ₁₆	d ₁₅	d ₁₄	d ₁₃	d ₁₂	d ₁₁	d ₁₀	d ₉	d ₈	d ₇	d ₆	d ₅	d ₄	d ₃	d ₂	d ₁	d ₀
			79	72											MAGNITUDE							C
SINGLE PRECISION	10^{-38}	24BITS	S	BIASED EXPONENT	SIGNIFICAND								31	23	0							
DOUBLE PRECISION	10^{-308}	53 BITS	S	BIASED EXPONENT	SIGNIFICAND								63	52	0							
EXTENDED PRECISION	10^{-4932}	64 BITS	S	BIASED EXPONENT	1	SIGNIFICAND								79	64 63Δ	0						

- (1) S = SIGN BIT (0 = positive, 1 = negative)
 - (2) d_n = DECIMAL DIGIT (TWO PER TYPE)
 - (3) X = BITS HAVE NO SIGNIFICANCE; 80387 IGNORES WHEN LOADING, ZEROS WHEN STORING
 - (4) Δ = POSITION OF IMPLICIT BINARY POINT
 - (5) I = INTEGER BIT OF SIGNIFICAND; STORED IN TEMPORARY REAL, IMPLICIT IN SINGLE AND DOUBLE PRECISION
 - (6) EXPONENT BIAS (NORMALIZED VALUES):
 - SINGLE: 127 (7FH)
 - DOUBLE: 1023 (3FFH)
 - EXTENDED REAL: 16383 (3FFFH)
 - (7) PACKED BCD: $(-1)^s (D_1 \dots D_6)$
 - (8) REAL: $(-1)^s (2^{E-BIAS}) (F_0 F_1 \dots)$

G40003

Figure 2-10. 80387 Data Formats

The binary integer formats exist in memory only. When used by the 80387, they are automatically converted to the 80-bit extended real format. All binary integers are exactly representable in the extended real format.

2.2.2.2 DECIMAL INTEGERS

Decimal integers are stored in packed decimal notation, with two decimal digits “packed” into each byte, except the leftmost byte, which carries the sign bit (0=positive, 1=negative). Negative numbers are not stored in two’s complement form and are distinguished from positive numbers only by the sign bit. The most significant digit of the number is the leftmost digit. All digits must be in the range 0–9.

The decimal integer format exists in memory only. When used by the 80387, it is automatically converted to the 80-bit extended real format. All decimal integers are exactly representable in the extended real format.

2.2.2.3 REAL NUMBERS

The 80387 represents real numbers of the form:

$$(-1)^s 2^E (b_0 b_1 b_2 b_3 \dots b_{p-1})$$

...where...

$s = 0$ or 1

E = any integer between E_{min} and E_{max} , inclusive

$b_i = 0$ or 1

p = number of bits of precision

Table 2-3 summarizes the parameters for each of the three real-number formats.

Table 2-3. Summary of Format Parameters

Parameter	Format		
	Single	Double	Extended
Format width in bits	32	64	80
p (bits of precision)	24	53	64
Exponent width in bits	8	11	15
E_{max}	+127	+1023	+16383
E_{min}	-126	-1022	-16382
Exponent bias	+127	+1023	+16383

The 80387 stores real numbers in a three-field binary format that resembles scientific, or exponential, notation. The format consists of the following fields:

- The number's significant digits are held in the *significand* field, $b_{0\Delta}b_1b_2b_3..b_{p-1}$. (The term "significand" is analogous to the term "mantissa" used to describe floating point numbers on some computers.)
- The *exponent* field, $e = E + \text{bias}$, locates the binary point within the significant digits (and therefore determines the number's magnitude). (The term "exponent" is analogous to the term "characteristic" used to describe floating point numbers on some computers.)
- The 1-bit *sign* field indicates whether the number is positive or negative. Negative numbers differ from positive numbers only in the sign bits of their significands.

Table 2-4 shows how the real number 178.125 (decimal) is stored in the 80387 single real format. The table lists a progression of equivalent notations that express the same value to show how a number can be converted from one form to another. (The ASM386 and PL/M-386 language translators perform a similar process when they encounter programmer-defined real number constants.) Note that not every decimal fraction has an exact binary equivalent. The decimal number 1/10, for example, cannot be expressed exactly in binary (just as the number 1/3 cannot be expressed exactly in decimal). When a translator encounters such a value, it produces a rounded binary approximation of the decimal value.

The NPX usually carries the digits of the significand in normalized form. This means that, except for the value zero, the significand contains an *integer bit* and *fraction bits* as follows:

$1_{\Delta}fff...ff$

where Δ indicates an assumed binary point. The number of fraction bits varies according to the real format: 23 for single, 52 for double, and 63 for extended real. By normalizing real numbers so that their integer bit is always a 1, the 80387 eliminates leading zeros in small

Table 2-4. Real Number Notation

Notation	Value		
Ordinary Decimal	178.125		
Scientific Decimal	$1_{\Delta}78125E2$		
Scientific Binary	$1_{\Delta}0110010001E111$		
Scientific Binary (Biased Exponent)	$1_{\Delta}0110010001E10000110$		
80387 Single Format (Normalized)	Sign	Biased Exponent	Significand
	0	10000110	0110010001000000000000 $1_{\Delta}(\text{implicit})$

values ($|X| < 1$). This technique maximizes the number of significant digits that can be accommodated in a significand of a given width. Note that, in the single and double formats, the integer bit is *implicit* and is not actually stored; the integer bit is physically present in the extended format only.

If one were to examine only the significand with its assumed binary point, all normalized real numbers would have values greater than or equal to 1 and less than 2. The exponent field locates the *actual* binary point in the significant digits. Just as in decimal scientific notation, a positive exponent has the effect of moving the binary point to the right, and a negative exponent effectively moves the binary point to the left, inserting leading zeros as necessary. An unbiased exponent of zero indicates that the position of the assumed binary point is also the position of the actual binary point. The exponent field, then, determines a real number's magnitude.

In order to simplify comparing real numbers (e.g., for sorting), the 80387 stores exponents in a biased form. This means that a constant is added to the *true exponent* described above. As Table 2-3 shows, the value of this *bias* is different for each real format. It has been chosen so as to force the *biased exponent* to be a positive value. This allows two real numbers (of the same format and sign) to be compared as if they are unsigned binary integers. That is, when comparing them bitwise from left to right (beginning with the leftmost exponent bit), the first bit position that differs orders the numbers; there is no need to proceed further with the comparison. A number's true exponent can be determined simply by subtracting the bias value of its format.

The single and double real formats exist in memory only. If a number in one of these formats is loaded into an 80387 register, it is automatically converted to extended format, the format used for all internal operations. Likewise, data in registers can be converted to single or double real for storage in memory. The extended real format may be used in memory also, typically to store intermediate results that cannot be held in registers.

Most applications should use the double format to store real-number data and results; it provides sufficient range and precision to return correct results with a minimum of programmer attention. The single real format is appropriate for applications that are constrained by memory, but it should be recognized that this format provides a smaller margin of safety. It is also useful for the debugging of algorithms, because roundoff problems will manifest themselves more quickly in this format. The extended real format should normally be reserved for holding intermediate results, loop accumulations, and constants. Its extra length is designed to shield final results from the effects of rounding and overflow/underflow in intermediate calculations. However, the range and precision of the double format are adequate for most microcomputer applications.

2.2.3 Rounding Control

Internally, the 80387 employs three extra bits (guard, round, and sticky bits) that enable it to round numbers in accord with the infinitely precise true result of a computation; these bits are not accessible to programmers. Whenever the destination can represent the infinitely precise true result, the 80387 delivers it. Rounding occurs in arithmetic and store operations when the format of the destination cannot exactly represent the infinitely precise true result.

For example, a real number may be rounded if it is stored in a shorter real format, or in an integer format. Or, the infinitely precise true result may be rounded when it is returned to a register.

The NPX has four rounding modes, selectable by the RC field in the control word (see Figure 2-3). Given a true result b that cannot be represented by the target data type, the 80387 determines the two representable numbers a and c that most closely bracket b in value ($a < b < c$). The processor then rounds (changes) b to a or to c according to the mode selected by the RC field as shown in Table 2-5. Rounding introduces an error in a result that is less than one unit in the last place to which the result is rounded.

- “Round to nearest” is the default mode and is suitable for most applications; it provides the most accurate and statistically unbiased estimate of the true result.
- The “chop” or “round toward zero” mode is provided for integer arithmetic applications.
- “Round up” and “round down” are termed *directed rounding* and can be used to implement interval arithmetic. Interval arithmetic generates a certifiable result independent of the occurrence of rounding and other errors. The upper and lower bounds of an interval may be computed by executing an algorithm twice, rounding up in one pass and down in the other.

Rounding control affects only the arithmetic instructions (refer to Chapter 3 for lists of arithmetic and nonarithmetic instructions).

2.2.4 Precision Control

The 80387 allows results to be calculated with either 64, 53, or 24 bits of precision in the significand as selected by the precision control (PC) field of the control word. The default setting, and the one that is best suited for most applications, is the full 64 bits of significance provided by the extended real format. The other settings are required by the IEEE standard and are provided to obtain compatibility with the specifications of certain existing programming languages. Specifying less precision nullifies the advantages of the extended format's extended fraction length. When reduced precision is specified, the rounding of the fractional value clears the unused bits on the right to zeros.

Table 2-5. Rounding Modes

RC Field	Rounding Mode	Rounding Action
00	Round to nearest	Closer to b of a or c ; if equally close, select even number (the one whose least significant bit is zero).
01	Round down (toward $-\infty$)	a
10	Round up (toward $+\infty$)	c
11	Chop (toward 0)	Smaller in magnitude of a or c .

NOTE: $a < b < c$; a and c are successive representable numbers; b is not representable.

Special Computational Situations **3**

CHAPTER 3

SPECIAL COMPUTATIONAL SITUATIONS

Besides being able to represent positive and negative numbers, the 80387 data formats may be used to describe other entities. These special values provide extra flexibility, but most users will not need to understand them in order to use the 80387 successfully. This section describes the special values that may occur in certain cases and the significance of each. The 80387 exceptions are also described, for writers of exception handlers and for those interested in probing the limits of computation using the 80387.

The material presented in this section is mainly of interest to programmers concerned with writing exception handlers. Many readers will only need to skim this section.

When discussing these special computational situations, it is useful to distinguish between *arithmetic instructions* and *nonarithmetic instructions*. Nonarithmetic instructions are those that have no operands or transfer their operands without substantial change; arithmetic instructions are those that make significant changes to their operands. Table 3-1 defines these two classes of instructions.

3.1 SPECIAL NUMERIC VALUES

The 80387 data formats encompass encodings for a variety of special values in addition to the typical real or integer data values that result from normal calculations. These special values have significance and can express relevant information about the computations or operations that produced them. The various types of special values are

- Denormal real numbers
- Zeros
- Positive and negative infinity
- NaN (Not-a-Number)
- Indefinite
- Unsupported formats

The following sections explain the origins and significance of each of these special values. Tables 3-6 through 3-9 at the end of this section show how each of these special values is encoded for each of the numeric data types.

3.1.1 Denormal Real Numbers

The 80387 generally stores nonzero real numbers in normalized floating-point form; that is, the integer (leading) bit of the significand is always a one. (Refer to Chapter 2 for a review of operand formats.) This bit is explicitly stored in the extended format, and is implicitly

Table 3-1. Arithmetic and Nonarithmetic Instructions

Nonarithmetic Instructions	Arithmetic Instructions
FABS	F2XM1
FCHS	FADD(P)
FCLEX	FBLD
FDECSTP	FBSTP
FFREE	FCOMP(P)(P)
FINCSTP	FCOS
FINIT	FDIV(R)(P)
FLD (register-to-register)	FIADD
FLD (extended format from memory)	FICOM(P)
FLD constant	FIDIV(R)
FLDCW	FILD
FLDENV	FIMUL
FNOP	FIST(P)
FRSTOR	FISUB(R)
FSAVE	FLD (conversion)
FST(P) (register-to-register)	FMUL(P)
FSTP (extended format to memory)	FPATAN
FSTCW	FPREM
FSTENV	FPREM1
FSTSW	FPTAN
FWAIT	FRNDINT
FXAM	FSCALE
FXCH	FSIN
	FSINCOS
	FSQRT
	FST(P) (conversion)
	FSUB(R)(P)
	FTST
	FUCOM(P)(P)
	FXTRACT
	FYL2X
	FYL2XP1

assumed to be a one (1_{Δ}) in the single and double formats. Since leading zeros are eliminated, normalized storage allows the maximum number of significant digits to be held in a significand of a given width.

When a numeric value becomes very close to zero, normalized floating-point storage cannot be used to express the value accurately. The term *tiny* is used here to precisely define what values require special handling by the 80387. A number R is said to be *tiny* when $-2^{E_{\min}} < R < 0$ or $0 < R < +2^{E_{\min}}$. (As defined in Chapter 2, E_{\min} is -126 for single format, -1022 for double format, and -16382 for extended format.) In other words, a nonzero number is *tiny* if its exponent would be too negative to store in the destination format.

To accommodate these instances, the 80387 can store and operate on reals that are not normalized, i.e., whose significands contain one or more leading zeros. Denormals typically arise when the result of a calculation yields a value that is *tiny*.

Denormal values have the following properties:

- The biased floating-point exponent is stored at its smallest value (zero)
- The integer bit of the significand (whether explicit or implicit) is zero

The leading zeros of denormals permit smaller numbers to be represented, at the possible cost of some lost precision (the number of significant bits is reduced by the leading zeros). In typical algorithms, extremely small values are most likely to be generated as intermediate, rather than final, results. By using the NPX's extended real format for holding intermediate values, quantities as small as $\pm 3.4 \times 10^{-4932}$ can be represented; this makes the occurrence of denormal numbers a rare phenomenon in 80387 applications. Nevertheless, the NPX can load, store, and operate on denormalized real numbers when they do occur.

Denormals receive special treatment by the 80387 in three respects:

- The 80387 avoids creating denormals whenever possible. In other words, it always normalizes real numbers except in the case of tiny numbers.
- The 80387 provides the unmasked underflow exception to permit programmers to detect cases when denormals would be created.
- The 80387 provides the denormal exception to permit programmers to detect cases when denormals enter into further calculations.

Denormalizing means incrementing the true result's exponent and inserting a corresponding leading zero in the significand, shifting the rest of the significand one place to the right. Denormal values may occur in any of the single, double, or extended formats. Table 3-2 illustrates how a result might be denormalized to fit a single format destination.

Denormalization produces either a denormal or a zero. Denormals are readily identified by their exponents, which are always the minimum for their formats; in biased form, this is always the bit string: 00..00. This same exponent value is also assigned to the zeros, but a denormal has a nonzero significand. A denormal in a register is tagged *special*. Tables 3-8 and 3-9 later in this chapter show how denormal values are encoded in each of the real data formats.

The denormalization process causes loss of significance if low-order one-bit bits are shifted off the right of the significand. In a severe case, *all* the significand bits of the true result are shifted out and replaced by the leading zeros. In this case, the result of denormalization is a true zero, and, if the value is in a register, it is tagged as a zero.

Table 3-2. Denormalization Process

Operation	Sign	Exponent	Significand
True Result	0	-129	1 _Δ 01011100..00
Denormalize	0	-128	0 _Δ 101011100..00
Denormalize	0	-127	0 _Δ 0101011100..00
Denormalize	0	-126	0 _Δ 00101011100..00
Denormal Result	0	-126	0 _Δ 00101011100..00

Denormals are rarely encountered in most applications. Typical debugged algorithms generate extremely small results during the evaluation of intermediate subexpressions; the final result is usually of an appropriate magnitude for its single or double format real destination. If intermediate results are held in temporary real, as is recommended, the great range of this format makes underflow very unlikely. Denormals are likely to arise only when an application generates a great many intermediates, so many that they cannot be held on the register stack or in extended format memory variables. If storage limitations force the use of single or double format reals for intermediates, and small values are produced, underflow may occur, and, if masked, may generate denormals.

When a denormal number is single or double format is used as a source operand and the denormal exception is masked, the 80387 automatically *normalizes* the number when it is converted to extended format.

3.1.1.1 DENORMALS AND GRADUAL UNDERFLOW

Floating-point arithmetic cannot carry out all operations exactly for all operands; approximation is unavoidable when the exact result is not representable as a floating-point variable. To keep the approximation mathematically tractable, the hardware is made to conform to accuracy standards that can be modeled by certain inequalities instead of equations. Let the assignment

$$X \leftarrow Y @ Z \quad (\text{where } @ \text{ is some operation})$$

represent a typical operation. In the default rounding mode (round to nearest), each operation is carried out with an absolute error no larger than half the separation between the two floating-point numbers closest to the exact results. Let x be the value stored for the variable whose name in the program is X , and similarly y for Y , and z for Z . Normally y and z will differ by accumulated errors from what is desired and from what would have been obtained in the absence of error. For the calculation of x we assume that y and z are the best approximations available, and we seek to compute x as well as we can. If $y@z$ is representable exactly, then we expect $x = y@z$, and that is what we get for every algebraic operation on the 80387 (i.e., when $y@z$ is one of $y+z$, $y-z$, $y\times z$, $y\div z$, \sqrt{z}). But if $y@z$ must be approximated, as is usually the case, then x must differ from $y@z$ by no more than half the difference between the two representable numbers that straddle $y@z$. That difference depends on two factors:

1. The precision to which the calculation is carried out, as determined either by the precision control bits or by the format used in memory. On the 80387, the precisions are single (24 significant bits), double (53 significant bits), and extended (64 significant bits).
2. How close $y@z$ is to zero. In this respect the presence of denormal numbers on the 80387 provides a distinct advantage over systems that do not admit denormal numbers.

In any floating-point number system, the density of representable numbers is greater near zero than near the largest representable magnitudes. However, machines that do not use denormal numbers suffer from an enormous gap between zero and its closest neighbors. Figures 3-1 and 3-2 show what happens near zero in two kinds of floating-point number systems.

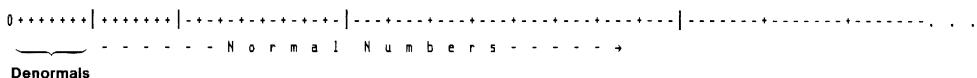


Figure 3-1. Floating-Point System with Denormals

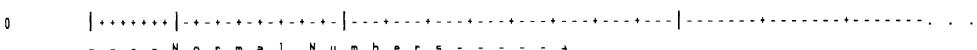


Figure 3-2. Floating-Point System without Denormals

Figure 3-1 shows a floating-point number system that (like the 80387) admits denormal numbers. For simplicity, only the non-negative numbers appear and the figure illustrates a number system that carries just four significant bits instead of the 24, 53, or 64 significant bits that the 80387 offers.

Each vertical mark stands for a number representable in four significant bits, and the bolder marks stand for the normal powers of 2. The denormal numbers lie between 0 and the nearest normal power of 2. They are no less dense than the remaining normal nonzero numbers.

Figure 3-2 shows a floating-point number system that (unlike the 80387) does not admit denormal numbers. There are two yawning gaps, one on the positive side of zero (as illustrated) and one on the negative side of zero (not illustrated). The gap between zero and the nearest neighbor of zero differs from the gap between that neighbor and the next bigger number by a factor of about 8.4×10^6 for single, 4.5×10^{15} for double, and 9.2×10^{18} for extended format. Those gaps would horribly complicate error analysis.

The advantage of denormal numbers is apparent when one considers what happens in either case when the underflow exception is masked and $y@z$ falls into the space between zero and the smallest normal magnitude. The 80387 returns the nearest denormal number. This action might be called “gradual underflow.” The effect is no different than the rounding that can occur when $y@z$ falls in the normal range.

On the other hand, the system that does not have denormal numbers returns zero as the result, an action that can be much more inaccurate than rounding. This action could be called “abrupt underflow.”

3.1.2 Zeros

The value *zero* in the real and decimal integer formats may be signed either positive or negative, although the sign of a binary integer zero is always positive. For computational purposes, the value of zero always behaves identically, regardless of sign, and typically the fact that a zero may be signed is transparent to the programmer. If necessary, the FXAM instruction may be used to determine a zero's sign.

If a zero is loaded or generated in a register, the register is tagged zero. Table 3-3 lists the results of instructions executed with zero operands and also shows how a zero may be created from nonzero operands.

Table 3-3. Zero Operands and Results

Operation	Operands	Result
FLD,FBLD	+0 -0	+0 -0
FILD	+0	+0
FST,FSTP	+0 -0 +X -X	+0 -0 +0 ¹ -0 ¹
FBSTP	+0 -0	+0 -0
FIST,FISTP	+0 -0 +X -X	+0 -0 +0 ³ -0 ³
Addition	+0 plus +0 -0 plus -0 +0 plus -0, -0 plus +0 -X plus +X, +X plus -X ±0 plus ±X, ±X plus ±0	+0 -0 ±0 ² ±0 ² #X
Subtraction	+0 minus -0 -0 minus +0 +0 minus +0, -0 minus -0 +X minus +X, -X minus -X ±0 minus ±X ±X minus ±0	+0 -0 ±0 ² ±0 ² -#X #X
Multiplication	+0 × +0, -0 × -0 +0 × -0, -0 × +0 +0 × +X, +X × +0 +0 × -X, -X × +0 -0 × +X, -X × +0 -0 × -X, -X × -0	+0 -0 +0 -0 -0 +0
Multiplication	+X × +Y, -X × -Y +X × -Y, -X × +Y	+0 ¹ -0 ¹
Division	±0 ÷ ±0 ±X ÷ ±0 +0 ÷ +X, -0 ÷ -X +0 ÷ -X, -0 ÷ +X -X ÷ -Y, +X ÷ +Y -X ÷ +Y, +X ÷ -Y	Invalid Operation ⊕∞ (Zero Divide) +0 -0 +0 ¹ -0 ¹
FPREM, FPREM1	±0 rem ±0 ±X rem ±0 +0 rem ±X -0 rem ±X	Invalid Operation Invalid Operation +0 -0
FPREM	+X rem ±Y -X rem ±Y	+0 Y exactly divides X -0 Y exactly divides X
FPREM1	+X rem ±Y -X rem ±Y	+0 Y exactly divides X -0 Y exactly divides X

X and Y denote nonzero positive operands.

1 When extreme underflow denormalizes the result to zero.

2 Sign determined by rounding mode: + for nearest, up, or chop, - for down.

3 When $0 < X < 1$ and rounding mode is not up.

* Sign of original zero operand.

Sign of original X operand.

-# Complement of sign of original X operand.

⊕ Exclusive OR of the signs of the operands.

Table 3-3. Zero Operands and Results (Cont'd.)

Operation	Operands	Result
FSQRT	+0 -0	+0 -0
Compare	$\pm 0 : +X$ $\pm 0 : \pm 0$ $\pm 0 : -X$	$\pm 0 < +X$ $\pm 0 = \pm 0$ $\pm 0 > -X$ $\pm 0 = 0$
FTST	+0 +0 -0	$C_3=1; C_2=C_1=C_0=0$ $C_3=C_1=1; C_2=C_0=0$
FCHS	+0 -0	-0 +0
FABS	+0	+0
F2XM1	+0 -0	+0 -0
FRNDINT	+0 -0	+0 -0
FSCALE	± 0 scaled by $-\infty$ ± 0 scaled by $+\infty$ ± 0 scaled by X	*0 Invalid Operation *0
FXTRACT	+0 -0	ST = +0, ST(1) = $-\infty$, Zero divide ST = -0, ST(1) = $-\infty$, Zero divide
FPTAN	+0	*0
FSIN (or SIN result of FSINCOS)	± 0	*0
FCOS (or COS result of FSINCOS)	± 0	+1
FPATAN	$\pm 0 \div +X$ $\pm 0 \div -X$ $\pm X \div \pm 0$ $\pm 0 \div +0$ $\pm 0 \div -0$ $+\infty \div \pm 0$ $-\infty \div \pm 0$ $\pm 0 \div +\infty$ $\pm 0 \div -\infty$	*0 * π $\#\pi/2$ *0 * π $+\pi/2$ $-\pi/2$ *0 * π
FYL2X	$\pm Y \times \log(\pm 0)$ $\pm 0 \times \log(\pm 0)$	Zero Divide Invalid Operation
FYL2XP1	$+Y \times \log(\pm 0 + 1)$ $-Y \times \log(\pm 0 + 1)$	*0 - *0

X and Y denote nonzero positive operands.

* Sign of original zero operand.

Sign of original X operand.

-# Complement of sign of original X operand.

3.1.3 Infinity

The real formats support signed representations of infinities. These values are encoded with a biased exponent of all ones and a significand of $1_Δ 00..00$; if the infinity is in a register, it is tagged special.

A programmer may code an infinity, or it may be created by the NPX as its masked response to an overflow or a zero divide exception. Note that depending on rounding mode, the masked response may create the largest valid value representable in the destination rather than infinity.

The signs of the infinities are observed, and comparisons are possible. Infinities are always interpreted in the affine sense; that is, $-\infty < (\text{any finite number}) < +\infty$. Arithmetic on infinities is always exact and, therefore, signals no exceptions, except for the invalid operations specified in Table 3-4.

Table 3-4. Infinity Operands and Results

Operation	Operands	Result
Addition	$+\infty$ plus $+\infty$	$+\infty$
	$-\infty$ plus $-\infty$	$-\infty$
	$+\infty$ plus $-\infty$	Invalid Operation
	$-\infty$ plus $+\infty$	Invalid Operation
	$\pm\infty$ plus $\pm X$	$*\infty$
	$\pm X$ plus $\pm\infty$	$*\infty$
Subtraction	$+\infty$ minus $-\infty$	$+\infty$
	$-\infty$ minus $+\infty$	$-\infty$
	$+\infty$ minus $+\infty$	Invalid Operation
	$-\infty$ minus $-\infty$	Invalid Operation
	$\pm\infty$ minus $\pm X$	$*\infty$
	$\pm X$ minus $\pm\infty$	$-\infty$
Multiplication	$\pm\infty$ \times $\pm\infty$	$\oplus\infty$
	$\pm\infty$ \times $\pm Y$, $\pm Y$ \times $\pm\infty$	$\oplus\infty$
	± 0 \times $\pm\infty$, $\pm\infty$ \times ± 0	Invalid Operation
	$\pm\infty$ \div $\pm\infty$	Invalid Operation
Division	$\pm\infty$ \div $\pm X$	$\oplus\infty$
	$\pm X$ \div $\pm\infty$	$\oplus 0$
	$\pm\infty$ \div ± 0	$\oplus\infty$
	$-\infty$	Invalid Operation
FSQRT	$+\infty$	$+\infty$
	$\pm\infty$	Invalid Operation
	± 0	$\oplus 0$
FPREM, FPREM1	$\pm\infty$ rem $\pm\infty$	Invalid Operation
	$\pm\infty$ rem $\pm X$	Invalid Operation
	$\pm X$ rem $\pm\infty$	$$X, Q = 0$
FRNDINT	$\pm\infty$	$*\infty$

X Zero or nonzero positive operand.

Y Nonzero positive operand.

* Sign of original infinity operand.

-* Complement of sign of original infinity operand.

\$ Sign of original operand.

⊕ Exclusive OR of signs of operands.

Table 3-4. Infinity Operands and Results (Cont'd.)

Operation	Operands	Result
FSCALE	$\pm\infty$ scaled by $-\infty$ $\pm\infty$ scaled by $+\infty$ $\pm\infty$ scaled by $\pm X$ ± 0 scaled by $-\infty$ ± 0 scaled by $+\infty$ $\pm Y$ scaled by $+\infty$ $\pm Y$ scaled by $-\infty$	Invalid Operation $*\infty$ $*\infty$ $\pm 0^1$ Invalid Operation $\#\infty$ $\#0$
FXTRACT	$\pm\infty$	$ST = *\infty, ST(1) = +\infty$
Compare	$+\infty : +\infty$ $-\infty : -\infty$ $+\infty : -\infty$ $-\infty : +\infty$ $+\infty : \pm X$ $-\infty : \pm X$ $\pm X : +\infty$ $\pm X : -\infty$	$+\infty = +\infty$ $-\infty = -\infty$ $+\infty > -\infty$ $-\infty < +\infty$ $+\infty > X$ $-\infty < X$ $X < +\infty$ $X > +\infty$
FTST	$+\infty$ $-\infty$	$+\infty > 0$ $-\infty < 0$
FPATAN	$\pm\infty \div \pm X$ $\pm Y \div +\infty$ $\pm Y \div -\infty$ $\pm\infty \div +\infty$ $\pm\infty \div -\infty$ $\pm\infty \div \pm 0$ $+0 \div +\infty$ $+0 \div -\infty$ $-0 \div +\infty$ $-0 \div -\infty$	$*\pi/2$ $\#0$ $\#\pi$ $*\pi/4$ $*3\pi/4$ $*\pi/2$ $+0$ $+\pi$ -0 $-\pi$ $+\infty$ -1
F2XM1	$+\infty$ $-\infty$	
FYL2X, FYL2XP1	$\pm\infty \times \log(1)$ $\pm\infty \times \log(Y>1)$ $\pm\infty \times \log(0<Y<1)$ $\pm Y \times \log(+\infty)$ $\pm 0 \times \log(+\infty)$ $\pm Y \times \log(-\infty)$	Invalid Operation $*\infty$ $-*\infty$ $\#\infty$ Invalid Operation Invalid Operation

X Zero or nonzero positive operand.

Y Nonzero positive operand.

* Sign of original infinity operand.

-* Complement of sign of original infinity operand.

Sign of the original Y operand.

1 Sign of original zero operand.

3.1.4 NaN (Not-a-Number)

A NaN (Not a Number) is a member of a class of special values that exists in the real formats only. A NaN has an exponent of 11..11B, may have either sign, and may have any significand except 1_A00..00B, which is assigned to the infinities. A NaN in a register is tagged special.

There are two classes of NaNs: signaling (SNaN) and quiet (QNaN). Among the QNaNs, the value *real indefinite* is of special interest.

3.1.4.1 SIGNALING NaNs

A signaling NaN is a NaN that has a zero as the most significant bit of its significand. The rest of the significand may be set to any value. The 80387 never generates a signaling NaN as a result; however, it recognizes signaling NaNs when they appear as operands. Arithmetic operations (as defined at the beginning of this chapter) on a signaling NaN cause an invalid-operation exception (except for load operations, FXCH, FCHS, and FABS).

By unmasking the invalid operation exception, the programmer can use signaling NaNs to trap to the exception handler. The generality of this approach and the large number of NaN values that are available provide the sophisticated programmer with a tool that can be applied to a variety of special situations.

For example, a compiler could use signaling NaNs as references to uninitialized (real) array elements. The compiler could preinitialize each array element with a signaling NaN whose significand contained the index (relative position) of the element. If an application program attempted to access an element that it had not initialized, it would use the NaN placed there by the compiler. If the invalid operation exception were unmasked, an interrupt would occur, and the exception handler would be invoked. The exception handler could determine which element had been accessed, since the operand address field of the exception pointers would point to the NaN, and the NaN would contain the index number of the array element.

3.1.4.2 QUIET NaNs

A quiet NaN is a NaN that has a one as the most significant bit of its significand. The 80387 creates the quiet NaN *real indefinite* (defined below) as its default response to certain exceptional conditions. The 80387 may derive other QNaNs by converting an SNaN. The 80387 converts a SNaN by setting the most significant bit of its significand to one, thereby generating an QNaN. The remaining bits of the significand are not changed; therefore, diagnostic information that may be stored in these bits of the SNaN is propagated into the QNaN.

The 80387 will generate the special QNaN, *real indefinite*, as its masked response to an invalid operation exception. This NaN is signed negative; its significand is encoded $1_{\Delta}100..00$. All other NaNs represent values created by programmers or derived from values created by programmers.

Both quiet and signaling NaNs are supported in all operations. A QNaN is generated as the masked response for invalid-operation exceptions and as the result of an operation in which at least one of the operands is a QNaN. The 80387 applies the rules shown in Table 3-5 when generating a QNaN:

Note that handling of a QNaN operand has greater priority than all exceptions except certain invalid-operation exceptions (refer to the section "Exception Priority" in this chapter).

Table 3-5. Rules for Generating QNaNs

Operation	Action
Real operation on an SNaN and a QNaN	Deliver the QNaN operand.
Real operation on two SNaNs	Deliver the QNaN that results from converting the SNaN that has the larger significand.
Real operation on two QNaNs	Deliver the QNaN that has the larger significand.
Real operation on an SNaN and another number	Deliver the QNaN that results from converting the SNaN.
Real operation on a QNaN and another number	Deliver the QNaN.
Invalid operation that does not involve NaNs	Deliver the default QNaN <i>real indefinite</i> .

Quiet NaNs could be used, for example, to speed up debugging. In its early testing phase, a program often contains multiple errors. An exception handler could be written to save diagnostic information in memory whenever it was invoked. After storing the diagnostic data, it could supply a quiet NaN as the result of the erroneous instruction, and that NaN could point to its associated diagnostic area in memory. The program would then continue, creating a different NaN for each error. When the program ended, the NaN results could be used to access the diagnostic data saved at the time the errors occurred. Many errors could thus be diagnosed and corrected in one test run.

3.1.5 Indefinite

For every 80387 numeric data type, one unique encoding is reserved for representing the special value *indefinite*. The 80387 produces this encoding as its response to a masked invalid-operation exception.

In the case of reals, the *indefinite* value is a QNaN as discussed in the prior section.

Packed decimal *indefinite* may be stored by the NPX in a FBSTP instruction; attempting to use this encoding in a FBLD instruction, however, will have an undefined result; thus *indefinite* cannot be loaded from a packed decimal integer.

In the binary integers, the same encoding may represent either *indefinite* or the largest negative number supported by the format (-2^{15} , -2^{31} , or -2^{63}). The 80387 will store this encoding as its masked response to an invalid operation, or when the value in a source register represents or rounds to the largest negative integer representable by the destination. In situations where its origin may be ambiguous, the invalid-operation exception flag can be examined to see if the value was produced by an exception response. When this encoding is loaded or used by an integer arithmetic or compare operation, it is always interpreted as a negative number; thus *indefinite* cannot be loaded from a binary integer.

3.1.6 Encoding of Data Types

Tables 3-6 through 3-9 show how each of the special values just described is encoded for each of the numeric data types. In these tables, the least-significant bits are shown to the right and are stored in the lowest memory addresses. The sign bit is always the left-most bit of the highest-addressed byte.

3.1.7 Unsupported Formats

The extended format permits many bit patterns that do not fall into any of the previously mentioned categories. Some of these encodings were supported by the 80287 NPX; however, most of them are not supported by the 80387 NPX. These changes are required due to changes made in the final version of the IEEE 754 standard that eliminated these data types.

The categories of encodings formerly known as pseudozeros, pseudo-NaNs, pseudoinfinities, and unnormal numbers are not supported by the 80387. The 80387 raises the invalid-operation exception when they are encountered as operands.

The encodings formerly known as pseudodenormal numbers are not generated by the 80387; however, they are correctly utilized when encountered in operands to 80387 instructions. The exponent is treated as if it were 00..01 and the mantissa is unchanged. The denormal exception is raised.

Table 3-6. Binary Integer Encodings

Class		Sign	Magnitude
Positives	(Largest)	0	11..11
		•	•
		•	•
		•	•
		•	•
	(Smallest)	0	00..01
	Zero	0	00..00
Negatives	(Smallest)	1	11..11
		•	•
		•	•
		•	•
	(Largest/Indefinite*)	1	00..00
			Word: _____ 15 bits Short: _____ 31 bits Long: _____ 63 bits

*If this encoding is used as a source operand (as in an integer load or integer arithmetic instruction), the 80387 interprets it as the largest negative number representable in the format... -2^{15} , -2^{31} , or -2^{63} . The 80387 delivers this encoding to an integer destination in two cases:

1. If the result is the largest negative number.
2. As the response to a masked invalid operation exception, in which case it represents the special value *integer indefinite*.

Table 3-7. Packed Decimal Encodings

Class		Sign		Magnitude							
				digit	digit	digit	digit	...	digit		
Positives	(Largest)	0	0000000	1 0 0 1	1 0 0 1	1 0 0 1	1 0 0 1	...	1 0 0 1		
					
					
					
	(Smallest)	0	0000000	0 0 0 0	0 0 0 0	0 0 0 0	0 0 0 0	...	0 0 0 1		
	Zero	0	0000000	0 0 0 0	0 0 0 0	0 0 0 0	0 0 0 0	...	0 0 0 0		
Negatives	Zero	1	0000000	0 0 0 0	0 0 0 0	0 0 0 0	0 0 0 0	...	0 0 0 0		
	(Smallest)	1	0000000	0 0 0 0	0 0 0 0	0 0 0 0	0 0 0 0	...	0 0 0 1		
					
					
(Largest)	1	0000000	1 0 0 1	1 0 0 1	1 0 0 1	1 0 0 1	1 0 0 1	...	1 0 0 1		
	Indefinite*	1	1111111	1 1 1 1	1 1 1 1	U U U U**	U U U U	...	U U U U		
				—1 byte—		—9 bytes—					

* The *packed decimal indefinite* is stored by FBSTP in response to a masked invalid operation exception. Attempting to load this value via FBLD produces an undefined result.

** UUUU means bit values are undefined and may contain any value.

Table 3-8. Single and Double Real Encodings

Class			Sign	Biased Exponent	Significand ff-ff*
Positives	NaNs	Quiet	0	11..11 • •	11..11 • •
			0	11..11	10..00
	Signalings		0	11..11 • •	01..11 • •
			0	11..11	00..01
	Infinity		0	11..11	00..00
	Reals	Normals	0	11..10 • •	11..11 • •
			0	00..01	00..00
		Denormals	0	00..00 • •	11..11 • •
			0	00..00	00..01
	Zero		0	00..00	00..00
Negatives	Reals	Zero	1	00..00	00..00
			1	00..00 • •	00..01 • •
		Denormals	1	00..00	11..11
			1	00..01 • •	00..00 • •
	NaNs	Normals	1	00..01 • •	00..00 • •
			1	11..10	11..11
		Infinity		11..11	00..00
		Signalings	1	11..11 • •	00..01 • •
			1	11..11	01..11
	Quiet	Indefinite	1	11..11 • •	10..00 • •
			1	11..11	11..11
			Single: Double:	8 bits 11 bits	23 bits 52 bits

*Integer bit is implied and not stored.

Table 3-9. Extended Real Encodings

Class			Sign	Biased Exponent	Significand i.ff-ff
Positives	NaNs	Quiet	0 • 0	11..11 • 11..11	1 11..11 • 1 10..00
		Signaling	0 • 0	11..11 • 11..11	1 01..11 • 1 00..01
Infinity			0	11..11	1 00..00
Positives	Reals	Normals	0 • 0	11..10 • 00..01	1 11..11 • 1 00..00
		Unsupported 8087 Unnormals	0 • 0	11..10 • 00..01	0 11..11 • 0 00..00
		Pseudodenormals	0 • 0	00..00 • 00..00	1 11..11 • 1 00..00
		Denormals	0 • 0	00..00 • 00..00	0 11..11 • 0 00..01
		Zero	0	00..00	000..00
		Zero	1	00..00	0 00..00
		Denormals	1 • 1	00..00 • 00..00	0 00..01 • 0 11..11
		Pseudodenormals	1 • 1	00..00 • 00..00	1 00..00 • 1 11..11
		Unsupported 8087 Unnormals	1 • 1	00..00 • 11..10	0 00..00 • 0 11..11
		Normals	1 • 1	00..01 • 11..10	1 00..00 • 1 11..11
Infinity			1	11..11	1 00..00
Negatives	NaNs	Signaling	1 • 1	11..11 • 11..11	1 00..01 • 1 01..11
		Indefinite	1 • 1	11..11 • 11..11	1 10..00 • 1 11..11
		Quiet	1 • 1	11..11 • 11..11	1 10..00 • 1 11..11
				15 bits	64 bits

3.2 NUMERIC EXCEPTIONS

The 80387 can recognize six classes of numeric exception conditions while executing numeric instructions:

1. I— Invalid operation
 - Stack fault
 - IEEE standard invalid operation
2. Z— Divide-by-zero
3. D— Denormalized operand
4. O— Numeric overflow
5. U— Numeric underflow
6. P— Inexact result (precision)

3.2.1 Handling Numeric Exceptions

When numeric exceptions occur, the NPX takes one of two possible courses of action:

- The NPX can itself handle the exception, producing the most reasonable result and allowing numeric program execution to continue undisturbed.
- A software exception handler can be invoked by the CPU to handle the exception.

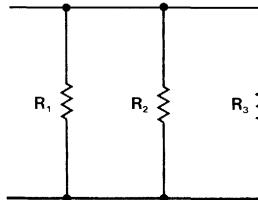
Each of the six exception conditions described above has a corresponding flag bit in the 80387 status word and a mask bit in the 80387 control word. If an exception is masked (the corresponding mask bit in the control word = 1), the 80387 takes an appropriate default action and continues with the computation. If the exception is unmasked (mask=0), the 80387 asserts the ERROR# output to the 80386 to signal the exception and invoke a software exception handler.

Note that when exceptions are masked, the NPX may detect multiple exceptions in a single instruction, because it continues executing the instruction after performing its masked response. For example, the 80387 could detect a denormalized operand, perform its masked response to this exception, and then detect an underflow.

3.2.1.1 AUTOMATIC EXCEPTION HANDLING

The 80387 NPX has a default fix-up activity for every possible exception condition it may encounter. These masked-exception responses are designed to be safe and are generally acceptable for most numeric applications.

As an example of how even severe exceptions can be handled safely and automatically using the NPX's default exception responses, consider a calculation of the parallel resistance of several values using only the standard formula (Figure 3-3). If R1 becomes zero, the circuit resistance becomes zero. With the divide-by-zero and precision exceptions masked, the 80387 NPX will produce the correct result.



$$\text{EQUIVALENT RESISTANCE} = \frac{1}{\frac{1}{R_1} + \frac{1}{R_2} + \frac{1}{R_3}}$$

122164-11

Figure 3-3. Arithmetic Example Using Infinity

By masking or unmasking specific numeric exceptions in the NPX control word, NPX programmers can delegate responsibility for most exceptions to the NPX, reserving the most severe exceptions for programmed exception handlers. Exception-handling software is often difficult to write, and the NPX's masked responses have been tailored to deliver the most reasonable result for each condition. For the majority of applications, masking all exceptions other than invalid-operation yields satisfactory results with the least programming effort. An invalid-operation exception normally indicates a program error that must be corrected; this exception should not normally be masked.

The exception flags in the NPX status word provide a cumulative record of exceptions that have occurred since these flags were last cleared. Once set, these flags can be cleared only by executing the FCLEX (clear exceptions) instruction, by reinitializing the NPX, or by overwriting the flags with an FRSTOR or FLDENV instruction. This allows a programmer to mask all exceptions (except invalid operation), run a calculation, and then inspect the status word to see if any exceptions were detected at any point in the calculation.

3.2.1.2 SOFTWARE EXCEPTION HANDLING

If the NPX encounters an unmasked exception condition, it signals the exception to the 80386 CPU using the ERROR# status line between the two processors.

The next time the 80386 CPU encounters a WAIT or ESC instruction in its instruction stream, the 80386 will detect the active condition of the ERROR# status line and automatically trap to an exception response routine using interrupt #16, the "processor extension error" exception.

This exception response routine is normally a part of the systems software. Typical exception responses may include:

- Incrementing an exception counter for later display or printing
- Printing or displaying diagnostic information (e.g., the 80387 environment and registers)
- Aborting further execution
- Using the exception pointers to build an instruction that will run without exception and executing it

For 80386 systems having systems software support for the 80387 NPX, applications programmers should consult the operating system's reference manuals for the appropriate system response to NPX exceptions. For systems programmers, specific details on writing software exception handlers are included in Chapter 6.

3.2.2 Invalid Operation

This exception may occur in response to two general classes of operations:

1. Stack operations
2. Arithmetic operations

The stack flag (SF) of the status word indicates which class of operation caused the exception. When SF is 1 a stack operation has resulted in stack overflow or underflow; when SF is 0, an arithmetic instruction has encountered an invalid operand.

3.2.2.1 STACK EXCEPTION

When SF is 1, indicating a stack operation, the O/U# bit of the condition code (bit C₁) distinguishes between stack overflow and underflow as follows:

O/U# = 1 Stack overflow— an instruction attempted to push down a nonempty stack location.

O/U# = 0 Stack underflow— an instruction attempted to read an operand from an empty stack location.

When the invalid-operation exception is masked, the 80387 returns the QNaN *indefinite*. This value overwrites the destination register, destroying its original contents.

When the invalid-operation exception is not masked, the 80386 exception “processor extension error” is triggered. TOP is not changed, and the source operands remain unaffected.

3.2.2.2 INVALID ARITHMETIC OPERATION

This class includes the invalid operations defined in IEEE Std 754. The 80387 reports an invalid operation in any of the cases shown in Table 3-10. Also shown in this table are the 80387's responses when the invalid exception is masked. When unmasked, the 80386 exception "processor extension error" is triggered, and the operands remain unaltered. An invalid operation generally indicates a program error.

3.2.3 Division by Zero

If an instruction attempts to divide a finite nonzero operand by zero, the 80387 will report a zero-divide exception. This is possible for F(I)DIV(R)(P) as well as the other instructions

Table 3-10. Masked Responses to Invalid Operations

Condition	Masked Response
Any arithmetic operation on an unsupported format.	Return the QNaN <i>indefinite</i> .
Any arithmetic operation on a signaling NaN.	Return a QNaN (refer to the section "Rules for Generating QNaNs").
Compare and test operations: one or both operands is a NaN.	Set condition codes "not comparable."
Addition of opposite-signed infinities or subtraction of like-signed infinities.	Return the QNaN <i>indefinite</i> .
Multiplication: $\infty \times 0$; or $0 \times \infty$.	Return the QNaN <i>indefinite</i> .
Division: $\infty \div \infty$; or $0 \div 0$.	Return the QNaN <i>indefinite</i> .
Remainder instructions FPREM, FPREM1 when modulus (divisor) is zero or dividend is ∞ .	Return the QNaN <i>indefinite</i> ; set C ₂ .
Trigonometric instructions FCOS, FPTAN, FSIN, FSINCOS when argument is ∞ .	Return the QNaN <i>indefinite</i> ; set C ₂ .
FSQRT of negative operand (except FSQRT (-0) = -0), FYL2X of negative operand (except FYL2X (-0) = $-\infty$), FYL2XP1 of operand more negative than -1 .	Return the QNaN <i>indefinite</i> .
FIST(P) instructions when source register is empty, a NaN, ∞ , or exceeds representable range of destination.	Store integer <i>indefinite</i> .
FBSTP instruction when source register is empty, a NaN, ∞ , or exceeds 18 decimal digits.	Store packed decimal <i>indefinite</i> .
FXCH instruction when one or both registers are tagged empty.	Change empty registers to the QNaN <i>indefinite</i> and then perform exchange.

that perform division internally: FYL2X and FXTRACT. The masked response for FDIV and FYL2X is to return an infinity signed with the exclusive OR of the signs of the operands. For FXTRACT, ST(1) is set to $-\infty$; ST is set to zero with the same sign as the original operand. If the divide-by-zero exception is unmasked, the 80386 exception "processor extension error" is triggered; the operands remain unaltered.

3.2.4 Denormal Operand

If an arithmetic instruction attempts to operate on a denormal operand, the NPX reports the denormal-operand exception. Denormal operands may have reduced significance due to lost low-order bits, therefore it may be advisable in certain applications to preclude operations on these operands. This can be accomplished by an exception handler that responds to unmasked denormal exceptions. Most users will mask this exception so that computation may proceed; any loss of accuracy will be analyzed by the user when the final result is delivered.

When this exception is masked, the 80387 sets the D-bit in the status word, then proceeds with the instruction. Gradual underflow and denormal numbers as handled on the 80387 will produce results at least as good as, and often better than what could be obtained from a machine that flushes underflows to zero. In fact, a denormal operand in single- or double-precision format will be normalized to the extended-real format when loaded into the 80387. Subsequent operations will benefit from the additional precision of the extended-real format used internally.

When this exception is not masked, the D-bit is set and the exception handler is invoked. The operands are not changed by the instruction and are available for inspection by the exception handler.

If an 8087/80287 program uses the denormal exception to automatically normalize denormal operands, then that program can run on an 80387 by masking the denormal exception. The 8087/80287 denormal exception handler would not be used by the 80387 in this case. A numerics program runs faster when the 80387 performs normalization of denormal operands. A program can detect at run-time whether it is running on an 80387 or 8087/80287 and disable the denormal exception when an 80387 is used. The following code sequence is recommended to distinguish between an 80387 and an 8087/80287.

```
FINIT          ; Use default infinity mode:  
              ; projective for 8087/80287,  
              ; affine for 80387  
FLD1           ; Generate infinity  
FLDZ  
FDIV  
FLD    ST       ; Form negative infinity  
FCHS  
FCOMPP         ; Compare +infinity with -infinity  
FSTSW temp     ; 8087/80287 will say they are equal  
MOV    AX, temp  
SAHF  
JNZ   Using_80387
```

The denormal-operand exception of the 80387 permits emulation of arithmetic on unnormal operands as provided by the 8087/80287. The standard does not require the denormal exception nor does it recognize the unnormal data type.

3.2.5 Numeric Overflow and Underflow

If the exponent of a numeric result is too large for the destination real format, the 80387 signals a numeric overflow. Conversely, if the exponent of a result is too small to be represented in the destination format, a numeric underflow is signaled. If either of these exceptions occur, the result of the operation is outside the range of the destination real format.

Typical algorithms are most likely to produce extremely large and small numbers in the calculation of intermediate, rather than final, results. Because of the great range of the extended-precision format (recommended as the destination format for intermediates), overflow and underflow are relatively rare events in most 80387 applications.

3.2.5.1 OVERFLOW

The overflow exception can occur whenever the rounded true result would exceed in magnitude the largest finite number in the destination format. The exception can occur in the execution of most of the arithmetic instructions and in some of the conversion instructions; namely, FST(P), F(I)ADD(P), F(I)SUB(R)(P), F(I)MUL(P), FDIV(R)(P), FSCALE, FYL2X, and FYL2XP1.

The response to an overflow condition depends on whether the overflow exception is masked:

- Overflow exception masked. The value returned depends on the rounding mode as Table 3-11 illustrates.

Table 3-11. Masked Overflow Results

Rounding Mode	Sign of True Result	Result
To nearest	+ -	+∞ -∞
Toward -∞	+ -	Largest finite positive number -∞
Toward +∞	+ -	+∞ Largest finite negative number
Toward zero	+ -	Largest finite positive number Largest finite negative number

- Overflow exception not masked. The unmasked response depends on whether the instruction is supposed to store the result on the stack or in memory:
 - Destination is the stack. The true result is divided by $2^{24,576}$ and rounded. (The bias 24,576 is equal to 3×2^{13} .) The significand is rounded to the appropriate precision (according to the precision control (PC) bit of the control word, for those instructions controlled by PC, otherwise to extended precision). The roundup bit (C_1) of the status word is set if the significand was rounded upward.

The biasing of the exponent by 24,576 normally translates the number as nearly as possible to the middle of the exponent range so that, if desired, it can be used in subsequent scaled operations with less risk of causing further exceptions. With the instruction FSCALE, however, it can happen that the result is too large and overflows even after biasing. In this case, the unmasked response is exactly the same as the masked round-to-nearest response, namely \pm infinity. The intention of this feature is to ensure the trap handler will discover that a translation of the exponent by -24574 would not work correctly without obliging the programmer of Decimal-to-Binary or Exponential functions to determine which trap handler, if any, should be invoked.

- Destination is memory (this can occur only with the store instructions). No result is stored in memory. Instead, the operand is left intact in the stack. Because the data in the stack is in extended-precision format, the exception handler has the option either of reexecuting the store instruction after proper adjustment of the operand or of rounding the significand on the stack to the destination's precision as the standard requires. The exception handler should ultimately store a value into the destination location in memory if the program is to continue.

3.2.5.2 UNDERFLOW

Underflow can occur in the execution of the instructions FST(P), FADD(P), FSUB(RP), FMUL(P), F(I)DIV(RP), FSCALE, FPREM(1), FPTAN, FSIN, FCOS, FSINCOS, FPATAN, F2XM1, FYL2X, and FYL2XP1.

Two related events contribute to underflow:

1. Creation of a tiny result which, because it is so small, may cause some other exception later (such as overflow upon division).
2. Creation of an inexact result; i.e. the delivered result differs from what would have been computed were both the exponent range and precision unbounded.

Which of these events triggers the underflow exception depends on whether the underflow exception is masked:

1. Underflow exception masked. The underflow exception is signaled when the result is both tiny and inexact.
2. Underflow exception not masked. The underflow exception is signaled when the result is tiny, regardless of inexactness.

The response to an underflow exception also depends on whether the exception is masked:

1. Masked response. The result is denormal or zero. The precision exception is also triggered.
2. Unmasked response. The unmasked response depends on whether the instruction is supposed to store the result on the stack or in memory:
 - Destination is the stack. The true result is multiplied by $2^{24,576}$ and rounded. (The bias 24,576 is equal to 3×2^{13} .) The significand is rounded to the appropriate precision (according to the precision control (PC) bit of the control word, for those instructions controlled by PC, otherwise to extended precision). The roundup bit (C_1) of the status word is set if the significand was rounded upward.

The biasing of the exponent by 24,576 normally translates the number as nearly as possible to the middle of the exponent range so that, if desired, it can be used in subsequent scaled operations with less risk of causing further exceptions. With the instruction FSCALE, however, it can happen that the result is too tiny and underflows even after biasing. In this case, the unmasked response is exactly the same as the masked round-to-nearest response, namely ± 0 . The intention of this feature is to ensure the trap handler will discover that a translation by +24576 would not work correctly without obliging the programmer of Decimal-to-Binary or Exponential functions to determine which trap handler, if any, should be invoked.

- Destination is memory (this can occur only with the store instructions). No result is stored in memory. Instead, the operand is left intact in the stack. Because the data in the stack is in extended-precision format, the exception handler has the option either of reexecuting the store instruction after proper adjustment of the operand or of rounding the significand on the stack to the destination's precision as the standard requires. The exception handler should ultimately store a value into the destination location in memory if the program is to continue.

3.2.6 Inexact (Precision)

This exception condition occurs if the result of an operation is not exactly representable in the destination format. For example, the fraction $1/3$ cannot be precisely represented in binary form. This exception occurs frequently and indicates that some (generally acceptable) accuracy has been lost.

All the transcendental instructions are inexact by definition; they always cause the inexact exception.

The C_1 (roundup) bit of the status word indicates whether the inexact result was rounded up ($C_1 = 1$) or chopped ($C_1 = 0$).

The inexact exception accompanies the underflow exception when there is also a loss of accuracy. When underflow is masked, the underflow exception is signaled only when there is a loss of accuracy; therefore the precision flag is always set as well. When underflow is unmasked, there may or may not have been a loss of accuracy; the precision bit indicates which is the case.

This exception is provided for applications that need to perform exact arithmetic only. Most applications will mask this exception. The 80387 delivers the rounded or over/underflowed result to the destination, regardless of whether a trap occurs.

3.2.7 Exception Priority

The 80387 deals with exceptions according to a predetermined precedence. Precedence in exception handling means that higher-priority exceptions are flagged and results are delivered according to the requirements of that exception. Lower-priority exceptions may not be flagged even if they occur. For example, dividing an SNaN by zero causes an invalid-operand exception (due to the SNaN) and not a zero-divide exception; the masked result is the QNaN *real indefinite*, not ∞ . A denormal or inexact (precision) exception, however, can accompany a numeric underflow or overflow exception.

The exception precedence is as follows:

1. Invalid operation exception, subdivided as follows:
 - a. Stack underflow.
 - b. Stack overflow.
 - c. Operand of unsupported format.
 - d. SNaN operand.
2. QNaN operand. Though this is not an exception, if one operand is a QNaN, dealing with it has precedence over lower-priority exceptions. For example, a QNaN divided by zero results in a QNaN, not a zero-divide exception.
3. Any other invalid-operation exception not mentioned above or zero divide.
4. Denormal operand. If masked, then instruction execution continues, and a lower-priority exception can occur as well.
5. Numeric overflow and underflow. Inexact result (precision) can be flagged as well.
6. Inexact result (precision).

3.2.8 Standard Underflow/Overflow Exception Handler

As long as the underflow and overflow exceptions are masked, no additional software is required to cause the output of the 80387 to conform to the requirements of IEEE Std 754. When unmasked, these exceptions give the exception handler an additional option in the case of store instructions. No result is stored in memory; instead, the operand is left intact on the stack. The handler may round the significand of the operand on the stack to the destination's precision as the standard requires, or it may adjust the operand and reexecute the faulting instruction.

The 80387 Instruction Set

4

CHAPTER 4

THE 80387 INSTRUCTION SET

This chapter describes the operation of all 80387 instructions. Within this section, the instructions are divided into six functional classes:

- Data Transfer instructions
- Nontranscendental instructions
- Comparison instructions
- Transcendental instructions
- Constant instructions
- Processor Control instructions

Throughout this chapter, the instruction set is described as it appears to the ASM386 programmer who is coding a program. Not included in this chapter are details of instruction format, encoding, and execution times. This detailed information may be found in Appendix A and Appendix E. Refer also to Appendix B for a summary of the exceptions caused by each instruction.

4.1 COMPATIBILITY WITH THE 80287 AND 8087

The instruction set for the 80387 NPX is largely the same as that for the 80287 NPX (used with 80286 systems) and that for the 8087 NPX (used with 8086 and 8088 systems). Most object programs generated for the 80287 or 8087 will execute without change on the 80387. Several instructions are new to the 80387, and several 80287 and 8087 instructions perform no useful function on the 80387. Appendix C and Appendix D give details of these instruction set differences.

4.2 NUMERIC OPERANDS

The typical NPX instruction accepts one or two operands as inputs, operates on these, and produces a result as an output. An operand is most often the contents of a register or of a memory location. The operands of some instructions are predefined; for example, FSQRT always takes the square root of the number in the top NPX stack element. Others allow, or require, the programmer to explicitly code the operand(s) along with the instruction mnemonic. Still others accept one explicit operand and one implicit operand, which is usually the top NPX stack element. All 80387 instructions that have a data operand use ST as one operand or as the only operand.

Whether supplied by the programmer or utilized automatically, the two basic types of operands are *sources* and *destinations*. A source operand simply supplies one of the inputs to an instruction; it is not altered by the instruction. Even when an instruction converts the source operand from one format to another (e.g., real to integer), the conversion is actually performed in an internal work area to avoid altering the source operand. A destination

operand may also provide an input to an instruction. It is distinguished from a source operand, however, because its content may be altered when it receives the result produced by the operation; that is, the destination is replaced by the result.

Many instructions allow their operands to be coded in more than one way. For example, FADD (add real) may be written without operands, with only a source or with a destination and a source. The instruction descriptions in this section employ the simple convention of separating alternative operand forms with slashes; the slashes, however, are not coded. Consecutive slashes indicate an option of no explicit operands. The operands for FADD are thus described as

//*source/destination, source*

This means that FADD may be written in any of three ways:

Written Form	Action
FADD	Add ST to ST(1), put result in ST(1), then pop ST
FADD <i>source</i>	Add source to ST(0)
FADD <i>destination, source</i>	Add source to destination

The assembler can allow the same instruction to be specified in different ways; for example:

FADD = FADDP ST(1), ST
FADD ST(1) = FADD ST, ST(1)

When reading this section, it is important to bear in mind that memory operands may be coded with any of the CPU's memory addressing methods provided by the ModR/M byte. To review these methods (BASE + (INDEX × SCALE) + DISPLACEMENT) refer to the *80386 Programmer's Reference Manual*. Chapter 5 also provides several addressing mode examples.

4.3 DATA TRANSFER INSTRUCTIONS

These instructions (summarized in Table 4-1) move operands among elements of the register stack, and between the stack top and memory. Any of the seven data types can be converted to extended real and loaded (pushed) onto the stack in a single operation; they can be stored to memory in the same manner. The data transfer instructions automatically update the 80387 tag word to reflect whether the register is empty or full following the instruction.

Table 4-1. Data Transfer Instructions

Real Transfers	
FLD FST FSTP FXCH	Load Real Store real Store real and pop Exchange registers
Integer Transfers	
FILD FIST FISTP	Integer load Integer store Integer store and pop
Packed Decimal Transfers	
FBLD FBSTP	Packed decimal (BCD) load Packed decimal (BCD) store and pop

4.3.1 FLD source

FLD (load real) loads (pushes) the source operand onto the top of the register stack. This is done by decrementing the stack pointer by one and then copying the content of the source to the new stack top. ST(7) must be empty to avoid causing an invalid-operation exception. The new stack top is tagged nonempty. The source may be a register on the stack (ST(i)) or any of the real data types in memory. If the source is a register, the register number used is that before TOP is decremented by the instruction. Coding FLD ST(0) duplicates the stack top. Single and double real source operands are converted to extended real automatically. Loading an extended real operand does not require conversion; therefore, the I and D exceptions do not occur in this case.

4.3.2 FST destination

FST (store real) copies the NPX stack top to the destination, which may be another register on the stack or a single or double (but not extended-precision) memory operand. If the destination is single or double real, the copy of the significand is rounded to the width of the destination according to the RC field of the control word, and the copy of the exponent is converted to the width and bias of the destination format. The over/underflow condition is checked for as well.

If, however, the stack top contains zero, $\pm\infty$, or a NaN, then the stack top's significand is not rounded but is chopped (on the right) to fit the destination. Neither is the exponent converted, rather it also is chopped on the right and transferred "as is". This preserves the value's identification as ∞ or a NaN (exponent all ones) so that it can be properly loaded and used later in the program if desired.

Note that the 80387 does not signal the invalid-operation exception when the destination is a nonempty stack element.

4.3.3 FSTP destination

FSTP (store real and pop) operates identically to FST except that the NPX stack is popped following the transfer. This is done by tagging the top stack element empty and then incrementing TOP. FSTP also permits storing to an extended-precision real memory variable, whereas FST does not. If the source operand is a register, the register number used is that before TOP is incremented by the instruction. Coding FSTP ST(0) is equivalent to popping the stack with no data transfer.

4.3.4 FXCH //destination

FXCH (exchange registers) swaps the contents of the destination and the stack top registers. If the destination is not coded explicitly, ST(1) is used. Many 80387 instructions operate only on the stack top; FXCH provides a simple means of effectively using these instructions on lower stack elements. For example, the following sequence takes the square root of the third register from the top (assuming that ST is nonempty):

```
FXCH ST(3)  
FSQRT  
FXCH ST(3)
```

4.3.5 FILD source

FILD (integer load) converts the source memory operand from its binary integer format (word, short, or long) to extended real and pushes the result onto the NPX stack. ST(7) must be empty to avoid causing an exception. The (new) stack top is tagged nonempty. FILD is an exact operation; the source is loaded with no rounding error.

4.3.6 FIST destination

FIST (integer store) stores the content of the stack top to an integer according to the RC field (rounding control) of the control word and transfers the result to the destination, leaving the stack top unchanged. The destination may define a word or short integer variable. Negative zero is stored in the same encoding as positive zero: 0000...00.

4.3.7 FISTP destination

FISTP (integer and pop) operates like FIST except that it also pops the NPX stack following the transfer. The destination may be any of the binary integer data types.

4.3.8 FBLD source

FBLD (packed decimal (BCD) load) converts the content of the source operand from packed decimal to extended real and pushes the result onto the NPX stack. ST(7) must be empty to avoid causing an exception. The sign of the source is preserved, including the case where

the value is negative zero. FBLD is an exact operation; the source is loaded with no rounding error.

The packed decimal digits of the source are assumed to be in the range 0–9. The instruction does not check for invalid digits (A–FH), and the result of attempting to load an invalid encoding is undefined.

4.3.9 FBSTP *destination*

FBSTP (packed decimal (BCD) store and pop) converts the content of the stack top to a packed decimal integer, stores the result at the destination in memory, and pops the stack. FBSTP rounds a nonintegral value according to the RC (rounding control) field of the control word.

4.4 NONTRANSCENDENTAL INSTRUCTIONS

The 80387's nontranscendental instruction set (Table 4-2) provides a wealth of variations on the basic add, subtract, multiply, and divide operations, and a number of other useful functions. These range from a simple absolute value to a square root instruction that executes faster than ordinary division; 80387 programmers no longer need to spend valuable time eliminating square roots from algorithms because they run too slowly. Other nontranscendental instructions perform exact modulo division, round real numbers to integers, and scale values by powers of two.

The 80387's basic nontranscendental instructions (addition, subtraction, multiplication, and division) are designed to encourage the development of very efficient algorithms. In particular, they allow the programmer to reference memory as easily as the NPX register stack.

Table 4-3 summarizes the available operation/operand forms that are provided for basic arithmetic. In addition to the four normal operations, two “reversed” instructions make subtraction and division “symmetrical” like addition and multiplication. The variety of instruction and operand forms give the programmer unusual flexibility:

- Operands may be located in registers or memory.
- Results may be deposited in a choice of registers.
- Operands may be a variety of NPX data types: extended real, double real, single real, short integer or word integer, with automatic conversion to extended real performed by the 80387.

Five basic instruction forms may be used across all six operations, as shown in Table 4-3. The classical stack form may be used to make the 80387 operate like a classical stack machine. No operands are coded in this form, only the instruction mnemonic. The NPX picks the source operand from the stack top and the destination from the next stack element. It then pops the stack, performs the operation, and returns the result to the new stack top, effectively replacing the operands by the result.

Table 4-2. Nontranscendental Instructions

Addition	
FADD FADDP FIADD	Add real Add real and pop Integer add
Subtraction	
FSUB FSUBP FISUB FSUBR FSUBRP FISUBR	Subtract real Subtract real and pop Integer subtract Subtract real reversed Subtract real reversed and pop Integer subtract reversed
Multiplication	
FMUL FMULP FIMUL	Multiply real Multiply real and pop Integer multiply
Division	
FDIV FDIVP FIDIV FDIVR FDIVRP FIDIVR	Divide real Divide real and pop Integer divide Divide real reversed Divide real reversed and pop Integer divide reversed
Other Operations	
FSQRT FSCALE FPREM FPREM1 FRNDINT FXTRACT FABS FCHS	Square root Scale Partial remainder IEEE standard partial remainder Round to integer Extract exponent and significand Absolute value Change sign

The register form is a generalization of the classical stack form; the programmer specifies the stack top as one operand and any register on the stack as the other operand. Coding the stack top as the destination provides a convenient way to access a constant, held elsewhere in the stack, from the stack top. The destination need not always be ST, however. All two operand instructions allow use of another register as the destination. This coding (ST is the source operand) allows, for example, adding the stack top into a register used as an accumulator.

Often the operand in the stack top is needed for one operation but then is of no further use in the computation. The register pop form can be used to pick up the stack top as the source

Table 4-3. Basic Nontranscendental Instructions and Operands

Instruction Form	Mnemonic Form	Operand Forms destination, source	ASM386 Example
Classical stack	Fop	{ ST(1), ST }	FADD
Classical stack, extra pop	FopP	{ ST(1), ST }	FADDP
Register	Fop	ST(i), ST or ST, ST(i)	FSUB
Register pop	FopP	ST(i), ST	FMULP
Real memory	Fop	{ ST, } single/double	FDIV
Integer memory	Flop	{ ST, } word-integer/short-integer	FIDIV
			PULSES

NOTES:

Braces ({ }) surround *implicit* operands; these are not coded, and are shown here for information only.

op =	ADD	destination \leftarrow destination + source
	SUB	destination \leftarrow destination - source
	SUBR	destination \leftarrow source - destination
	MUL	destination \leftarrow destination • source
	DIV	destination \leftarrow destination \div source
	DIVR	destination \leftarrow source \div destination

operand, and then discard it by popping the stack. Coding operands of ST(1), ST with a register pop mnemonic is equivalent to a classical stack operation: the top is popped and the result is left at the new top.

The two memory forms increase the flexibility of the 80387's nontranscendental instructions. They permit a real number or a binary integer in memory to be used directly as a source operand. This is useful in situations where operands are not used frequently enough to justify holding them in registers. Note that any memory addressing method may be used to define these operands, so they may be elements in arrays, structures, or other data organizations, as well as simple scalars.

The six basic operations are discussed further in the next paragraphs, and descriptions of the remaining seven operations follow.

4.4.1 Addition

FADD	//source/destination,source
FADDP	//destination,source
FIADD	source

The addition instructions (add real, add real and pop, integer add) add the source and destination operands and return the sum to the destination. The operand at the stack top may be doubled by coding:

FADD ST, ST(0)

If the source operand is in memory, conversion of an integer, a single real, or a double real operand to extended real is performed automatically.

4.4.2 Normal Subtraction

FSUB //source/destination,source
FSUBP //destination,source
FISUB source

The normal subtraction instructions (subtract real, subtract real and pop, integer subtract) subtract the source operand from the destination and return the difference to the destination.

4.4.3 Reversed Subtraction

FSUBR //source/destination,source
FSUBRP //destination,source
FISUBR source

The reversed subtraction instructions (subtract real reversed, subtract real reversed and pop, integer subtract reversed) subtract the destination from the source and return the difference to the destination. For example, FSUBR ST, ST(1) means subtract ST from ST(1) and leave the result in ST.

4.4.4 Multiplication

FMUL //source/destination,source
FMULP //destination,source
FIMUL source

The multiplication instructions (multiply real, multiply real and pop, integer multiply) multiply the source and destination operands and return the product to the destination. Coding FMUL ST, ST(0) squares the content of the stack top.

4.4.5 Normal Division

FDIV //source/destination,source
FDIVP //destination,source
FIDIV source

The normal division instructions (divide real, divide real and pop, integer divide) divide the destination by the source and return the quotient to the destination.

4.4.6 Reversed Division

FDIVR //source/destination,source
FDIVRP //destination,source
FIDIVR source

The reversed division instructions (divide real reversed, divide real reversed and pop, integer divide reversed) divide the source operand by the destination and return the quotient to the destination.

4.4.7 FSQRT

FSQRT (square root) replaces the content of the top stack element with its square root. (Note: The square root of -0 is defined to be -0 .)

4.4.8 FSCALE

FSCALE (scale) interprets the value contained in ST(1) as an integer and adds this value to the exponent of the number in ST. This is equivalent to

$$\text{ST} \leftarrow \text{ST} \cdot 2^{\text{ST}(1)}$$

Thus, FSCALE provides rapid multiplication or division by integral powers of 2. It is particularly useful for scaling the elements of a vector.

There is no limit on the range of the scale factor in ST(1). If the value is not integral, FSCALE uses the nearest integer smaller in magnitude; i.e., it chops the value toward 0. If the resulting integer is zero, the value in ST is not changed.

4.4.9 FPREM — Partial Remainder (80287/8087-Compatible)

FPREM computes the remainder of division of ST by ST(1) and leaves the result in ST. FPREM finds a remainder REM and a quotient Q such that

$$\text{REM} = \text{ST} - \text{ST}(1)*\text{Q}$$

The quotient Q is chosen to be the integer obtained by chopping the exact value of ST(ST(1)) toward zero. The sign of the remainder is the same as the sign of the original dividend from ST.

By ignoring precision control, the 80387 produces an exact result with FPREM. The precision (inexact) exception does not occur and the rounding control has no effect.

The FPREM instruction is not the remainder operation specified in the IEEE standard. To get that remainder, the FPREM1 instruction should be used.

The FPREM instruction is designed to be executed iteratively in a software-controlled loop. It operates by performing successive scaled subtractions; therefore, obtaining the exact remainder when the operands differ greatly in magnitude can consume large amounts of execution time. Because the 80387 can only be preempted between instructions, the remainder function could seriously increase interrupt latency in these cases. For this reason, the maximum number of iterations is limited. The instruction may terminate before it has completely terminated the calculation. The C2 bit of the status word indicates whether the calculation is complete or whether the instruction must be executed again.

FPREM can reduce the exponent of ST by up to (but not including) 64 in one execution. If FPREM produces a remainder that is less than the modulus (i.e., the divisor), the function is complete and bit C2 of the status word condition code is cleared. If the function is incomplete, C2 is set to 1; the result in ST is then called the *partial* remainder. Software can inspect C2 by storing the status word following execution of FPREM, reexecuting the instruction (using the partial remainder in ST as the dividend) until C2 is cleared. A higher priority interrupting routine that needs the 80387 can force a context switch between the instructions in the remainder loop.

An important use for FPREM is to reduce arguments (operands) of transcendental functions to the range permitted by these instructions. For example, the FPTAN (tangent) instruction requires its argument ST to be less than 2^{63} . For $\pi/4 < |ST| < 2^{63}$, FPTAN (as well as the other trigonometric instructions) performs an internal reduction of ST to a value less than $\pi/4$ using an internally stored $\pi/4$ divisor that has 67 significant bits. Because of its greater accuracy, this method of reduction is recommended when the argument is within the required range.

However, when $|ST| \geq 2^{63}$, FPREM can be employed to reduce ST. With $\pi/4$ as a modulus, FPREM can reduce an argument so that it is within range of FPTAN and so that no further reduction is required by FPTAN.

Because FPREM produces an exact result, the argument reduction does *not* introduce roundoff error into the calculation, even if several iterations are required to bring the argument into range. However, π is never accurate. The rounding of π , when it is used by FPREM to reduce an argument for a periodic trigonometric function, does not create the effect of a rounded argument, but of a rounded period.

When reduction is complete, FPREM provides the least-significant three bits of the quotient generated by FPREM (in C_3 , C_1 , C_0). This is also important for transcendental argument reduction, because it locates the original angle in the correct one of eight $\pi/4$ segments of the unit circle (see Table 4-4).

4.4.10 FPREM1—Partial Remainder (IEEE Std. 754-Compatible)

FPREM1 computes the remainder of division of ST by ST(1) and leaves the result in ST. FPREM1 finds a remainder REM1 and a quotient Q1 such that

$$\text{REM1} = \text{ST} - \text{ST}(1)*\text{Q1}$$

Table 4-4. Condition Code Interpretation after FPREM and FPREM1 Instructions

Condition Code				Interpretation after FPREM and FPREM1	
C2(PF)	C3	C1	C0		
1	X	X	X	Incomplete Reduction: further iteration required or complete reduction	
0	Q1	Q0	Q2	Q MOD 8	Complete Reduction: C0, C3, C1 contain three least significant bits of quotient
	0	0	0	0	
	0	1	0	1	
	1	0	0	2	
	1	1	0	3	
	0	0	1	4	
	0	1	1	5	
	1	0	1	6	
	1	1	1	7	

The quotient Q1 is chosen to be the integer nearest to the exact value of ST/ST(1). When ST/ST(1) is exactly $N + 1/2$ (for some integer N), there are two integers equally close to ST/ST(1). In this case the value chosen for Q1 is the even integer.

The result produced by FPREM1 is always *exact*; no rounding is necessary, and therefore the precision exception does not occur and the rounding control has no effect.

The FPREM1 instruction is designed to be executed iteratively in a software-controlled loop. FPREM1 operates by performing successive scaled subtractions; therefore, obtaining the exact remainder when the operands differ greatly in magnitude can consume large amounts of execution time. Because the 80387 can only be preempted between instructions, the remainder function could seriously increase interrupt latency in these cases. For this reason, the maximum number of iterations is limited. The instruction may terminate before it has completely terminated the calculation. The C2 bit of the status word indicates whether the calculation is complete or whether the instruction must be executed again.

FPREM1 can reduce the exponent of ST by up to (but not including) 64 in one execution. If FPREM1 produces a remainder that is less than the modulus (i.e., the divisor), the function is complete and bit C2 of the status word condition code is cleared. If the function is incomplete, C2 is set to 1; the result in ST is then called the partial remainder. Software can inspect C2 by storing the status word following execution of FPREM1, reexecuting the instruction (using the partial remainder in ST as the dividend) until C2 is cleared. When C2 is cleared, FPREM1 also provides the least-significant three bits of the quotient generated by FPREM1 (in C₃, C₁, C₀).

The uses for FPREM1 are the same as those for FPREM.

FPREM1 differs from FPREM in these respects:

- FPREM and FPREM1 choose the value of the quotient differently; the low-order three bits of the quotient as reported in bits C3,C1,C0 of the status word may differ by one in some cases.
- FPREM and FPREM1 may produce different remainders. FPREM produces a remainder R such that $0 \leq R < |ST(1)|$ or $-|ST(1)| < R \leq 0$, depending on the sign of the dividend. FPREM1 produces a remainder R1 such that $-|ST(1)|/2 < R1 < +|ST(1)|/2$.

4.4.11 FRNDINT

FRNDINT (round to integer) rounds the top stack element to an integer according to the RC bits of the control word. For example, assume that ST contains the 80387 real number encoding of the decimal value 155.625. FRNDINT will change the value to 155 if the RC field of the control word is set to down or chop, or to 156 if it is set to up or nearest.

4.4.12 FXTRACT

FXTRACT (extract exponent and significand) performs a superset of the IEEE-recommended **logb**(*x*) function by “decomposing” the number in the stack top into two numbers that represent the actual value of the operand’s exponent and significand fields. The “exponent” replaces the original operand on the stack and the “significand” is pushed onto the stack. (ST(7) must be empty to avoid causing the invalid-operation exception.) Following execution of FXTRACT, ST (the new stack top) contains the value of the original significand expressed as a real number: its sign is the same as the operand’s, its exponent is 0 true (16,383 or 3FFFH biased), and its significand is identical to the original operand’s. ST(1) contains the value of the original operand’s true (unbiased) exponent expressed as a real number.

If the original operand is zero, FXTRACT leaves $-\infty$ in ST(1) (the exponent) while ST is assigned the value zero with a sign equal to that of the original operand. The zero-divide exception is raised in this case, as well.

To illustrate the operation of FXTRACT, assume that ST contains a number whose true exponent is +4 (i.e., its exponent field contains 4003H). After executing FXTRACT, ST(1) will contain the real number +4.0; its sign will be positive, its exponent field will contain 4001H (+2 true) and its significand field will contain 1_A00...00B. In other words, the value in ST(1) will be $1.0 \times 2^2 = 4$. If ST contains an operand whose true exponent is -7 (i.e., its exponent field contains 3FF8H), then FXTRACT will return an “exponent” of -7.0; after the instruction executes, ST(1)’s sign and exponent fields will contain C001H (negative sign, true exponent of 2), and its significand will be 1_A1100...00B. In other words, the value in ST(1) will be $-1.75 \times 2^2 = -7.0$. In both cases, following FXTRACT, ST’s sign and significand fields will be the same as the original operand’s, and its exponent field will contain 3FFFH (0 true).

FXTRACT is useful for power and range scaling operations. Both FXTRACT and the base 2 exponential instruction F2XM1 are needed to perform a general power operation. Converting numbers in 80387 extended real format to decimal representations (e.g., for printing or displaying) requires not only FBSTP but also FXTRACT to allow scaling that does not overflow the range of the extended format. FXTRACT can also be useful for debugging, because it allows the exponent and significand parts of a real number to be examined separately.

4.4.13 FABS

FABS (absolute value) changes the top stack element to its absolute value by making its sign positive. Note that the invalid-operation exception is not signaled even if the operand is a signaling NaN or has a format that is not supported.

4.4.14 FCHS

FCHS (change sign) complements (reverses) the sign of the top stack element. Note that the invalid-operation exception is not signaled even if the operand is a signaling NaN or has a format that is not supported.

4.5 COMPARISON INSTRUCTIONS

The instructions of this class allow comparison of numbers of all supported real and integer data types. Each of these instructions (Table 4-5) analyzes the top stack element, often in relationship to another operand, and reports the result as a condition code in the status word.

The basic operations are compare, test (compare with zero), and examine (report type, sign, and normalization). Special forms of the compare operation are provided to optimize algorithms by allowing direct comparisons with binary integers and real numbers in memory, as well as popping the stack after a comparison.

The FSTSW (store status word) instruction may be used following a comparison to transfer the condition code to memory or to the 80386 AX register for inspection. The 80386 SAHF

Table 4-5. Comparison Instructions

FCOM	Compare real
FCOMP	Compare real and pop
FCOMPP	Compare real and pop twice
FICOM	Integer compare
FICOMP	Integer compare and pop
FTST	Test
FUCOM	Unordered compare real
FUCOMP	Unordered compare real and pop
FUCOMPP	Unordered compare real and pop twice
FXAM	Examine

instruction is recommended for copying the 80387 flags from AX to the 80386 flags for easy conditional branching.

Note that instructions other than those in the comparison group may update the condition code. To ensure that the status word is not altered inadvertently, store it immediately following a comparison operation.

4.5.1 FCOM //source

FCOM (compare real) compares the stack top to the source operand. The source operand may be a register on the stack, or a single or double real memory operand. If an operand is not coded, ST is compared to ST(1). The sign of zero is ignored, so that $+0 = -0$. Following the instruction, the condition codes reflect the order of the operands as shown in Table 4-6.

If either operand is a NaN (either quiet or signaling) or an undefined format, or if a stack fault occurs, the invalid-operation exception is raised and the condition bits are set to “unordered.”

4.5.2 FCOMP //source

FCOMP (compare real and pop) operates like FCOM, and in addition pops the stack.

4.5.3 FCOMPP

FCOMPP (compare real and pop twice) operates like FCOM and additionally pops the stack twice, discarding both operands. FCOMPP always compares ST to ST(1); no operands may be explicitly specified.

4.5.4 FICOM source

FICOM (integer compare) converts the source operand, which may reference a word or short binary integer variable, to extended real and compares the stack top to it. The condition code bits in the status word are set as for FCOM.

Table 4-6. Condition Code Resulting from Comparisons

Order	C3 (ZF)	C2 (PF)	C0 (CF)	80386 Conditional Branch
ST > Operand	0	0	0	JA
ST < Operand	0	0	1	JB
ST = Operand	1	0	0	JE
Unordered	1	1	1	JP

4.5.5 FICOMP source

FICOMP (integer compare and pop) operates identically to FICOM and additionally discards the value in ST by popping the NPX stack.

4.5.6 FTST

FTST (test) tests the top stack element by comparing it to zero. The result is posted to the condition codes as shown in Table 4-7.

4.5.7 FUCOM //source

FUCOM (unordered compare real) operates like FCOM, with two differences:

1. It does not cause an invalid-operation exception when one of the operands is a NaN. If either operand is a NaN, the condition bits of the status word are set to *unordered* as shown in Table 4-6.
2. Only operands on the NPX stack can be compared.

4.5.8 FUCOMP //source

FUCOMP (unordered compare real and pop) operates like FUCOM and in addition pops the NPX stack.

4.5.9 FUCOMPP

FUCOMPP (unordered compare real and pop) operates like FUCOM and in addition pops the NPX stack twice, discarding both operands. FUCOMPP always compares ST to ST(1); no operands can be explicitly specified.

Table 4-7. Condition Code Resulting from FTST

Order	C3 (ZF)	C2 (ZF)	C0 (ZF)	83086 Conditional Branch
ST > 0.0	0	0	0	JA
ST < 0.0	0	0	1	JB
ST = 0.0	1	0	0	JE
Unordered	1	1	1	JP

4.5.10 FXAM

FXAM (examine) reports the content of the top stack element as positive/negative and NaN, denormal, normal, zero, infinity, unsupported, or empty. Table 4-8 lists and interprets all the condition code values that FXAM generates.

4.6 TRANSCENDENTAL INSTRUCTIONS

The instructions in this group (Table 4-9) perform the time-consuming *core calculations* for all common trigonometric, inverse trigonometric, hyperbolic, inverse hyperbolic, logarithmic, and exponential functions. The transcendentals operate on the top one or two stack elements, and they return their results to the stack. The trigonometric operations assume their arguments are expressed in radians. The logarithmic and exponential operations work in base 2.

The results of transcendental instructions are highly accurate. The absolute value of the relative error of the transcendental instructions is guaranteed to be less than 2^{-62} . (Relative error is the ratio between the absolute error and the exact value.)

Table 4-8. Condition Code Defining Operand Class

C3	C2	C1	C0	Value at TOP
0	0	0	0	+Unsupported
0	0	0	1	+NaN
0	0	1	0	-Unsupported
0	0	1	1	-NaN
0	1	0	0	+Normal
0	1	0	1	+Infinity
0	1	1	0	-Normal
0	1	1	1	-Infinity
1	0	0	0	+0
1	0	0	1	+Empty
1	0	1	0	-0
1	0	1	1	-Empty
1	1	0	0	+Denormal
1	1	1	0	-Denormal

Table 4-9. Transcendental Instructions

FSIN FCOS FSINCOS FPTAN FPATAN F2XM1 FYL2X FYL2XP1	Sine Cosine Sine and cosine Tangent of ST Arctangent of ST(1)/ST $2^x - 1$ $Y \cdot \log_2 X$; Y is ST(1), X is ST $Y \cdot \log_2(X + 1)$; Y is ST(1), X is ST
---	--

The trigonometric functions accept a practically unrestricted range of operands, whereas the other transcendental instructions require that arguments be more restricted in range. FPREM or FPREM1 may be used to bring the otherwise valid operand of a periodic function into range. Prologue and epilogue software may be used to reduce arguments for other instructions to the expected range and to adjust the result to correspond to the original arguments if necessary. The instruction descriptions in this section document the allowed operand range for each instruction.

4.6.1 FCOS

When complete, this function replaces the contents of ST with COS(ST). ST, expressed in radians, must lie in the range $|θ| < 2^{63}$ (for most practical purposes unrestricted). If ST is in range, C2 of the status word is cleared and the result of the operation is produced.

If the operand is outside of the range, C2 is set to one (function incomplete) and ST remains intact (i.e., no reduction of the operand is performed). It is the programmers responsibility to reduce the operand to an absolute value smaller than 2^{63} . The instructions FPREM1 and FPREM are available for this purpose.

4.6.2 FSIN

When complete, this function replaces the contents of ST with SIN(ST). FSIN is equivalent to FCOS in the way it reduces the operand. ST is expressed in radians.

4.6.3 FSINCOS

When complete, this instruction replaces the contents of ST with SIN(ST), then pushes COS(ST) onto the stack. (ST(7) must be empty to avoid an invalid exception.) FSINCOS is equivalent to FCOS in the way it reduces the operand. ST is expressed in radians.

4.6.4 FPTAN

When complete, FPTAN (partial tangent) computes the function $Y = \tan(ST)$. ST is expressed in radians. Y replaces ST, then the value 1 is pushed, becoming the new stack top. (ST(7) must be empty to avoid an invalid exception.) When the function is complete $ST(1) = \tan(\text{arg})$ and $ST = 1$. FPTAN is equivalent to FCOS in the way it reduces the operand.

The fact that FPTAN places two results on the stack maintains compatibility with the 8087/80287 and aids the calculation of other trigonometric functions that can be derived from **tan** via standard trigonometric identities. For example, the **cot** function is given by this identity:

$$\cot x = 1 / \tan x .$$

Therefore, simply executing the reverse divide instruction FDIVR after FPTAN yields the **cot** function.

4.6.5 FPATAN

FPATAN (arctangent) computes the function $\Theta = \text{ARCTAN}(Y/X)$. X is taken from ST(0) and Y from ST(1). The instruction pops the NPX stack and returns Θ to the (new) stack top, overwriting the Y operand. The result is expressed in radians. The range of operands is not restricted; however, the range of the result depends on the relationship between the operands according to Table 4-10.

The fact that the argument of FPATAN is a ratio aids calculation of other trigonometric functions, including Arcsin and Arccos. These can be derived from Arctan via standard trigonometric identities. For example, the Arcsin function can be easily calculated using this identity:

$$\text{Arcsin } x = \text{Arctan} (x / \sqrt{1 - x^2}).$$

Thus, to find Arcsin (Y), push Y onto the NPX stack, then calculate $X = \sqrt{1 - Y^2}$, pushing the result X onto the stack. Executing FPATAN then leaves Arcsin (Y) at the top of the stack.

4.6.6 F2XM1

F2XM1 (2 to the X minus 1) calculates the function $Y = 2^X - 1$. X is taken from the stack top and must be in the range $-1 \leq X \leq 1$. The result Y replaces the argument X at the stack top. If the argument is out of range, the results are undefined.

This instruction is designed to produce a very accurate result even when X is close to 0. For values of the argument very close in magnitude to 1, a larger error will be incurred. To obtain $Y = 2^X$, add 1 to the result delivered by F2XM1.

Table 4-10. Results of FPATAN

Sign(Y)	Sign(X)	$ Y < X ?$	Final Result
+	+	Yes	$0 < \text{atan}(Y/X) < \pi/4$
+	+	No	$\pi/4 < \text{atan}(Y/X) < \pi/2$
+	-	No	$\pi/2 < \text{atan}(Y/X) < 3 \cdot \pi/4$
+	-	Yes	$3 \cdot \pi/4 < \text{atan}(Y/X) < \pi$
-	+	Yes	$-\pi/4 < \text{atan}(Y/X) < 0$
-	+	No	$-\pi/2 < \text{atan}(Y/X) < -\pi/4$
-	-	No	$-3 \cdot \pi/4 < \text{atan}(Y/X) < -\pi/2$
-	-	Yes	$-\pi < \text{atan}(Y/X) < -3 \cdot \pi/4$

The following formulas show how values other than 2 may be raised to a power of X:

$$10^X = 2^{X \cdot \text{LOG}_2(10)}$$

$$e^X = 2^{X \cdot \text{LOG}_2(e)}$$

$$y^X = 2^{X \cdot \text{LOG}_2(Y)}$$

As shown in the next section, the 80387 has built-in instructions for loading the constants $\text{LOG}_2 10$ and $\text{LOG}_2 e$, and the FYL2X instruction may be used to calculate $X \cdot \text{LOG}_2 Y$.

4.6.7 FYL2X

FYL2X (Y log base 2 of X) calculates the function $Z = Y \cdot \text{LOG}_2 X$. X is taken from the stack top and Y from ST(1). The operands must be in the following ranges:

$$\begin{aligned} 0 \leq X &< +\infty \\ -\infty < Y &< +\infty \end{aligned}$$

The instruction pops the NPX stack and returns Z at the (new) stack top, replacing the Y operand. If the operand is out of range (i.e., in negative) the invalid-operation exception occurs.

This function optimizes the calculations of log to any base other than two, because a multiplication is always required:

$$\text{LOG}_N X = (\text{LOG}_2 N)^{-1} \cdot \text{LOG}_2 X$$

4.6.8 FYL2XP1

FYL2XP1 (Y log base 2 of $(X + 1)$) calculates the function $Z = Y \cdot \text{LOG}_2(X + 1)$. X is taken from the stack top and must be in the range $-(1 - \text{SQRT}(2)/2) < X < 1 - \text{SQRT}(2)/2$. Y is taken from ST(1) and is unlimited in range $(-\infty < Y < +\infty)$. FYL2XP1 pops the stack and returns Z at the (new) stack top, replacing Y. If the argument is out of range, the results are undefined.

This instruction provides improved accuracy over FYL2X when computing the logarithm of a number very close to 1, for example $1 + \epsilon$ where $\epsilon << 1$. Providing ϵ rather than $1 + \epsilon$ as the input to the function allows more significant digits to be retained.

4.7 CONSTANT INSTRUCTIONS

Each of these instructions (Table 4-11) loads (pushes) a commonly used constant onto the stack. (ST(7) must be empty to avoid an invalid exception.) The values have full extended real precision (64 bits) and are accurate to approximately 19 decimal digits. Because an external real constant occupies 10 memory bytes, the constant instructions, which are only

Table 4-11. Constant Instructions

FLDZ	Load + 0.0
FLD1	Load + 1.0
FLDPi	Load π
FLDL2T	Load $\log_{10}10$
FLDL2E	Load \log_2e
FLDLG2	Load $\log_{10}2$
FLDLN2	Load \log_22

two bytes long, save storage and improve execution speed, in addition to simplifying programming.

The constants used by these instructions are stored internally in a format more precise even than extended real. When loading the constant, the 80387 rounds the more precise internal constant according the RC (rounding control) bit of the control word. However, in spite of this rounding, the precision exception is not raised (to maintain compatibility). When the rounding control is set to round to nearest on the 80387, the 80387 produces the same constant that is produced by the 80287.

4.7.1 FLDZ

FLDZ (load zero) loads (pushes) +0.0 onto the NPX stack.

4.7.2 FLD1

FLD1 (load one) loads (pushes) +1.0 onto the NPX stack.

4.7.3 FLDPi

FLDPi (load π) loads (pushes) π onto the NPX stack.

4.7.4 FLDL2T

FLDL2T (load log base 2 of 10) loads (pushes) the value \log_210 onto the NPX stack.

4.7.5 FLDL2E

FLDL2E (load log base 2 of e) loads (pushes) the value \log_2e onto the NPX stack.

4.7.6 FLDLG2

FLDLG2 (load log base 10 of 2) loads (pushes) the value $\log_{10}2$ onto the NPX stack.

4.7.7 FLDLN2

FLDLN2 (load log base e of 2) loads (pushes) the value $\text{LOG}_e 2$ onto the NPX stack.

4.8 PROCESSOR CONTROL INSTRUCTIONS

The processor control instructions are shown in Table 4-12. The instruction FSTSW is commonly used for conditional branching. The remaining instructions are not typically used in calculations; they provide control over the 80387 NPX for system-level activities. These activities include initialization, exception handling, and task switching.

As shown in Table 4-12, many of the NPX processor control instructions have two forms of assembler mnemonic:

1. A *wait* form, where the mnemonic is prefixed only with an F, such as FSTSW. This form checks for unmasked numeric exceptions.
2. A *no-wait* form, where the mnemonic is prefixed with an FN, such as FNSTSW. This form ignores unmasked numeric exceptions.

When the control instruction is coded using the *no-wait* form of the mnemonic, the ASM386 assembler does not precede the ESC instruction with a *wait* instruction, and the CPU does not test the ERROR# status line from the NPX before executing the processor control instruction.

Only the *processor control* class of instructions have this alternate no-wait form. All numeric instructions are automatically synchronized by the 80386; the CPU transfers all operands before initiating the next instruction. Because of this automatic synchronization by the 80386, *numeric* instructions for the 80387 need not be preceded by a CPU wait instruction in order to execute correctly.

Table 4-12. Processor Control Instructions

FINIT/FNINIT FLDCW FSTCW/FNSTCW FSTSW/FNSTSW FSTSW AX/FNSTSW AX FCLEX/FNCLEX FSTENV/FNSTENV FLDENV FSAVE/FNSAVE FRSTOR FINCSTP FDECSTP FFREE FNOP FWAIT	Initialize processor Load control word Store control word Store status word Store status word to AX Clear exceptions Store environment Load environment Save state Restore state Increment stack pointer Decrement stack pointer Free register No operation CPU Wait
---	--

It should also be noted that the 8087 instructions FENI and FDISI and the 80287 instruction FSETPF perform no function in the 80387. If these opcodes are detected in an 80386/80387 instruction stream, the 80387 performs no specific operation and no internal states are affected. For programmers interested in porting numeric software from 80287 or 8087 environments to the 80386, however, it should be noted that program sections containing these exception-handling instructions are not likely to be completely portable to the 80387. Appendix C and Appendix D contains a more complete description of the differences between the 80387 and the 80287/8087.

4.8.1 FINIT/FNINIT

FINIT/FNINIT (initialize processor) sets the 80387 NPX into a known state, unaffected by any previous activity. It sets the control word to its default value 037FH (round to nearest, all exceptions masked, 64 bits of precision), clears the status word, and empties all floating-point stack registers. The no-wait form of this instruction causes the 80387 to abort any previous numeric operations currently executing in the NEU.

This instruction performs the functional equivalent of a hardware RESET, with one exception: RESET causes the IM bit of the control word to be reset and the ES and IE bits of the status word to be set as a means of signaling the presence of an 80387; FINIT puts the opposite values in these bits.

FINIT checks for unmasked numeric exceptions, FNINIT does not. Note that if FNINIT is executed while a previous 80387 memory-referencing instruction is running, 80387 bus cycles in progress are aborted. This instruction may be necessary to clear the 80387 if a processor-extension segment-overrun exception (interrupt 9) is detected by the CPU.

4.8.2 FLDCW source

FLDCW (load control word) replaces the current processor control word with the word defined by the source operand. This instruction is typically used to establish or change the 80387's mode of operation. Note that if an exception bit in the status word is set, loading a new control word that unmasks that exception will activate the ERROR# output of the 80387. When changing modes, the recommended procedure is to first clear any exceptions and then load the new control word.

4.8.3 FSTCW/FNSTCW destination

FSTCW/FNSTCW (store control word) writes the processor control word to the memory location defined by the destination. FSTCW checks for unmasked numeric exceptions; FNSTCW does not.

4.8.4 FSTSW/FNSTSW *destination*

FSTSW/FNSTSW (store status word) writes the current value of the 80387 status word to the destination operand in memory. The instruction is used to

- Implement conditional branching following a comparison, FPREM, or FPREM1 instruction (FSTSW).
- Invoke exception handlers (by polling the exception bits) in environments that do not use interrupts (FSTSW).

FSTSW checks for unmasked numeric exceptions, FNSTSW does not.

4.8.5 FSTSW AX/FNSTSW AX

FSTSW AX/FNSTSW AX (store status word to AX) is a special 80387 instruction that writes the current value of the 80387 status word directly into the 80386 AX register. This instruction optimizes conditional branching in numeric programs, where the 80386 CPU must test the condition of various NPX status bits. The waited form FSTSW AX checks for unmasked numeric exceptions, the non-waited form FNSTSW AX does not.

When this instruction is executed, the 80386 AX register is updated with the NPX status word before the CPU executes any further instructions. The status stored is that from the completion of the prior ESC instruction.

4.8.6 FCLEX/FNCLEX

FCLEX/FNCLEX (clear exceptions) clears all exception flags, the exception status flag and the busy flag in the status word. As a consequence, the 80387's ERROR# line goes inactive. FCLEX checks for unmasked numeric exceptions, FNCLEX does not.

4.8.7 FSAVE/FNSAVE *destination*

FSAVE/FNSAVE (save state) writes the full 80387 state—environment plus register stack—to the memory location defined by the destination operand. Figure 4-1 and Figure 4-2 show the layout of the save area; the size and layout of the save area depends on the operating mode of the 80386 (real-address mode or protected mode) and on the operand-size attribute in effect for the instruction (32-bit operand or 16-bit operand). When the 80386 is in virtual-8086 mode, the real-address mode formats are used. Typically the instruction is coded to save this image on the CPU stack.

The values in the tag word in memory are determined during the execution of FSAVE/FNSAVE. If the tag in the status register indicates that the corresponding register is nonempty, the 80387 examines the data in the register and stores the appropriate tag in memory. Thus the tag that is stored always reflects the actual content of the register.

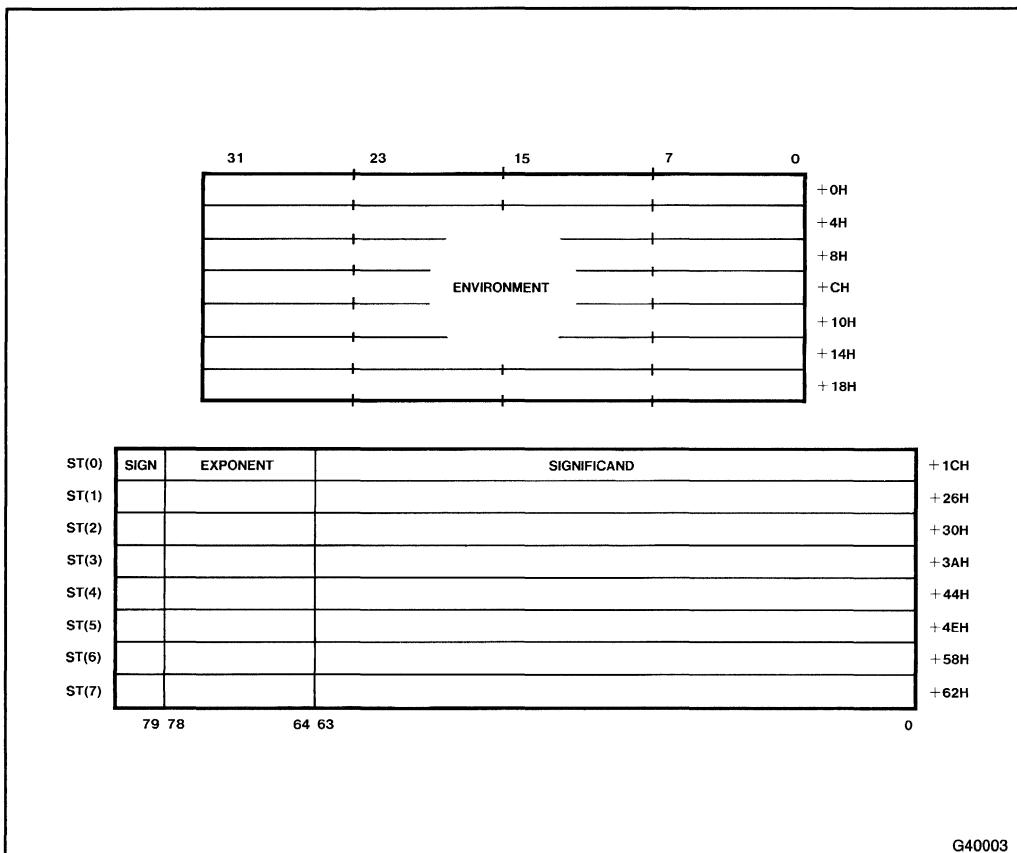


Figure 4-1. FSAVE/FRSTOR Memory Layout (32-Bit)

FNSAVE delays its execution until all NPX activity completes normally. Thus, the save image reflects the state of the NPX following the completion of any running instruction. After writing the state image to memory, FSAVE/FNSAVE initializes the 80387 as if FINIT/FNINIT had been executed.

FSAVE/FNSAVE is useful whenever a program wants to save the current state of the NPX and initialize it for a new routine. Three examples are

1. An operating system needs to perform a context switch (suspend the task that had been running and give control to a new task).
2. An exception handler needs to use the 80387.
3. An application task wants to pass a “clean” 80387 to a subroutine.

FSAVE checks for unmasked numeric exceptions before executing, FNSAVE does not.

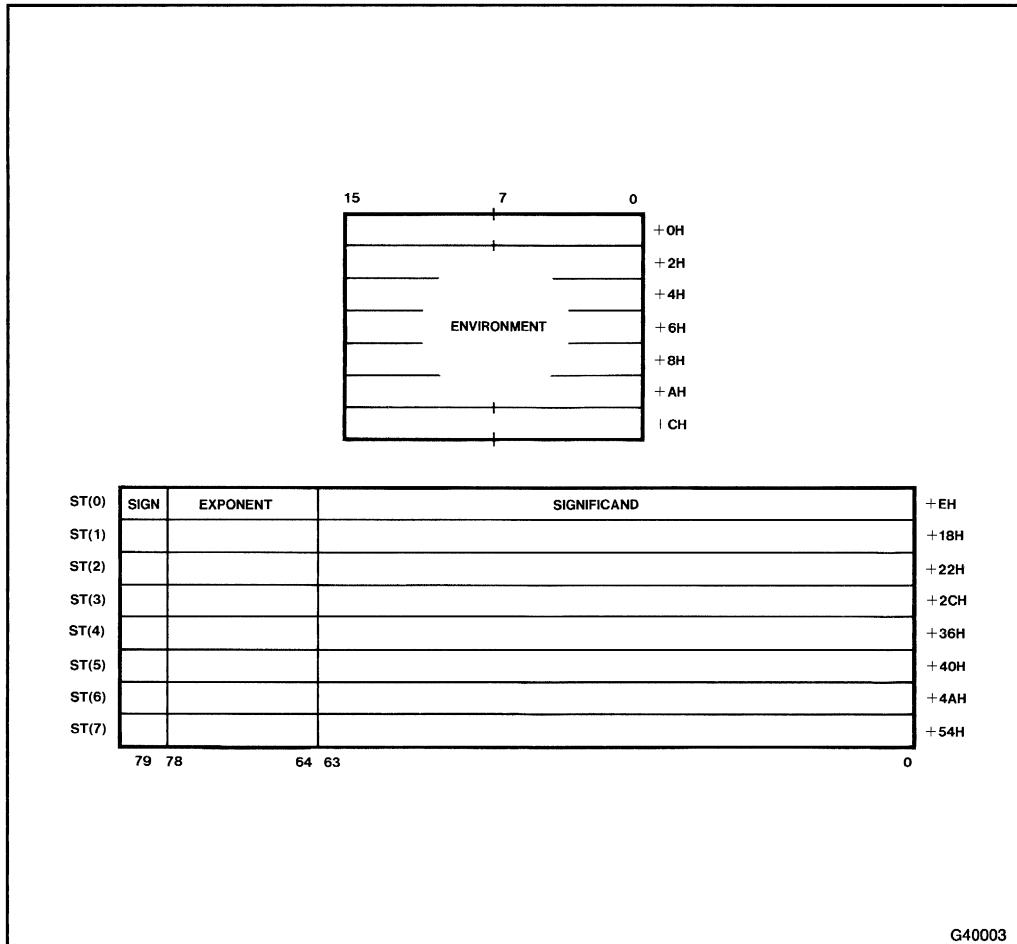


Figure 4-2. FSAVE/FRSTOR Memory Layout (16-Bit)

4.8.8 FRSTOR source

FRSTOR (restore state) reloads the 80387 state from the memory area defined by the source operand. This information should have been written by a previous FSAVE/FNSAVE instruction and not altered by any other instruction. FRSTOR automatically waits checking for interrupts until all data transfers are completed before continuing to the next instruction.

Note that the 80387 “reacts” to its new state at the conclusion of the FRSTOR. It generates an exception request, for example, if the exception and mask bits in the memory image so indicate when the next WAIT or exception-checking ESC instruction is executed.

4.8.9 FSTENV/FNSTENV destination

FSTENV/FNSTENV (store environment) writes the 80387's basic status—control, status, and tag words, and exception pointers—to the memory location defined by the destination operand. Typically, the environment is saved on the CPU stack. FSTENV/FNSTENV is often used by exception handlers because it provides access to the exception pointers that identify the offending instruction and operand. After saving the environment, FSTENV/FNSTENV sets all exception masks in the 80387 control word (i.e., masks all exceptions). FSTENV checks for pending exceptions before executing, FNSTENV does not.

Figures 4-3 through 4-6 shows the format of the environment data in memory; the size and layout of the save area depends on the operating mode of the 80386 (real-address mode or protected mode) and on the operand-size attribute in effect for the instruction (32-bit operand or 16-bit operand). When the 80386 is in virtual-8086 mode, the real-address mode formats are used. FNSTENV does not store the environment until all NPX activity has completed. Thus, the data saved by the instruction reflects the 80387 after any previously decoded instruction has been executed.

The values in the tag word in memory are determined during the execution of FNSTENV/FSTENV. If the tag in the status register indicates that the corresponding register is nonempty, the 80387 examines the data in the register and stores the appropriate tag in memory. Thus the tag that is stored always reflects the actual content of the register.

4.8.10 FLDENV source

FLDENV (load environment) reloads the environment from the memory area defined by the source operand. This data should have been written by a previous FSTENV/FNSTENV

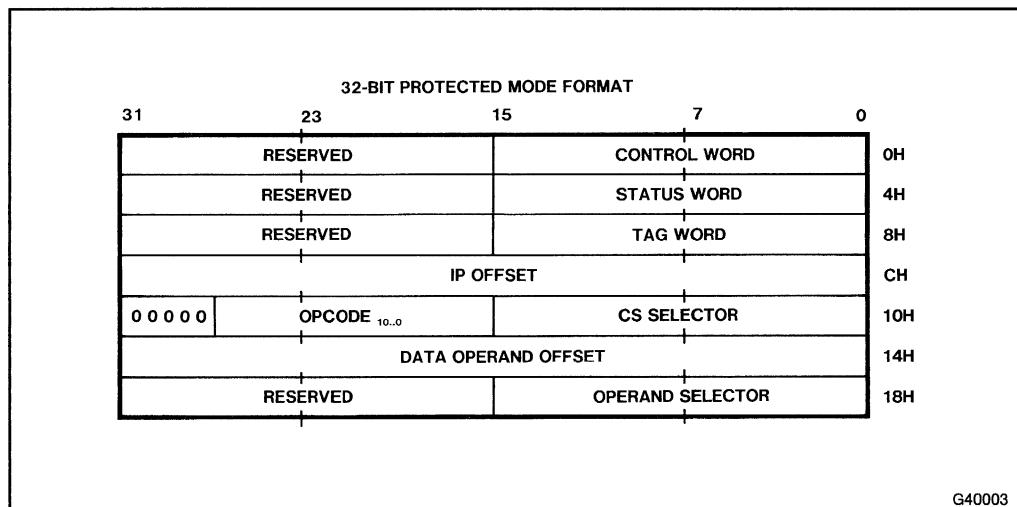
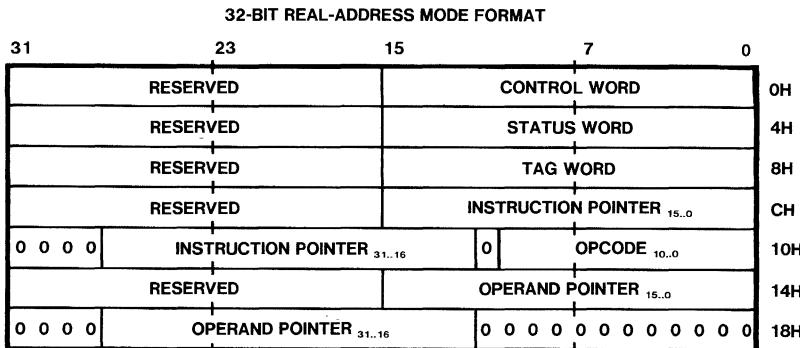
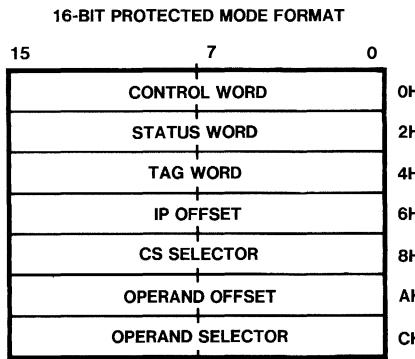


Figure 4-3. Protected Mode 80387 Environment, 32-Bit Format



G40003

Figure 4-4. Real Mode 80387 Environment, 32-Bit Format



G40003

Figure 4-5. Protected Mode 80387 Environment, 16-Bit Format

instruction. CPU instructions (that do not reference the environment image) may immediately follow FLDENV. FLDENV automatically waits for all data transfers to complete before executing the next instruction.

Note that loading an environment image that contains an unmasked exception causes a numeric exception when the next WAIT or exception-checking ESC instruction is executed.

4.8.11 FINCSTP

FINCSTP (increment NPX stack pointer) adds 1 to the stack top pointer (TOP) in the status word. It does not alter tags or register contents, nor does it transfer data. It is not

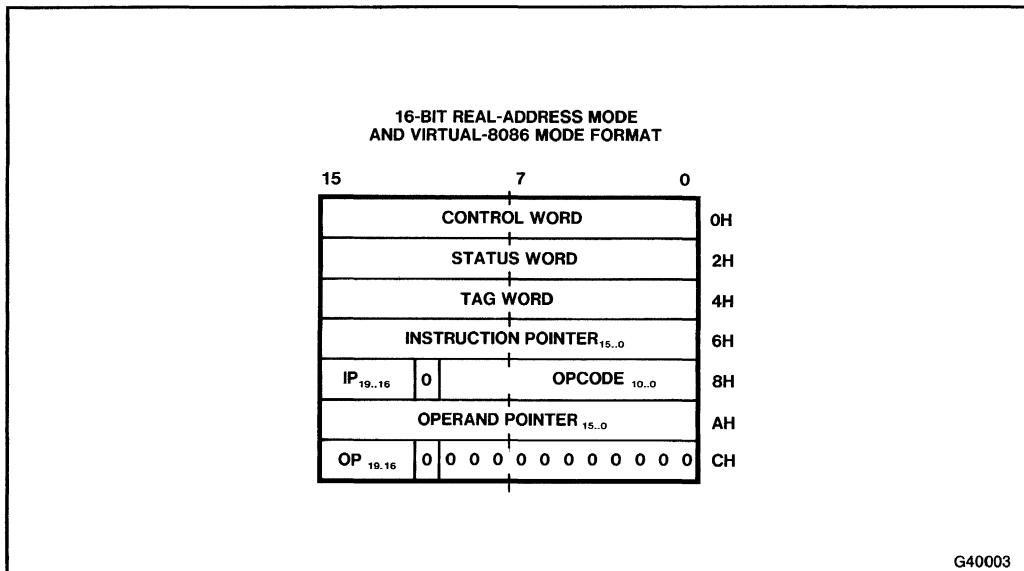


Figure 4-6. Real Mode 80387 Environment, 16-Bit Format

equivalent to popping the stack, because it does not set the tag of the previous stack top to empty. Incrementing the stack pointer when ST=7 produces ST=0.

4.8.12 FDECSTP

FDECSTP (decrement NPX stack pointer) subtracts 1 from ST, the stack top pointer in the status word. No tags or registers are altered, nor is any data transferred. Executing FDECSTP when ST=0 produces ST=7.

4.8.13 FFREE *destination*

FFREE (free register) changes the destination register's tag to *empty*; the content of the register is unaffected.

4.8.14 FNOP

FNOP (no operation) effectively performs no operation.

4.8.15 FWAIT (CPU Instruction)

FWAIT is not actually an 80387 instruction, but an alternate mnemonic for the 80386 WAIT instruction. The FWAIT or WAIT mnemonic should be coded whenever the programmer

wants to check for a pending error before modifying a variable used in the previous floating-point instruction. Coding an FWAIT instruction after an 80387 instruction ensures that unmasked numeric exceptions occur and exception handlers are invoked before the next instruction has a chance to examine the results of the 80387 instruction.

More information on when to code an FWAIT instruction is given in Chapter 5 in the section “Concurrent Processing with the 80387.”



Programming Numeric Applications

5



CHAPTER 5

PROGRAMMING NUMERIC APPLICATIONS

5.1 PROGRAMMING FACILITIES

As described previously, the 80387 NPX is programmed simply as an extension of the 80386 CPU. This section describes how programmers in ASM386 and in a variety of higher-level languages can work with the 80387.

The level of detail in this section is intended to give programmers a basic understanding of the software tools that can be used with the 80387, but this information does not document the full capabilities of these facilities. Complete documentation is available with each program development product.

5.1.1 High-Level Languages

For programmers using high-level languages, the programming and operation of the NPX is handled automatically by the compiler. A variety of Intel high-level languages are available that automatically make use of the 80387 NPX when appropriate. These languages include C-386 and PL/M-386. In addition many high-level language compilers are available from independent software vendors.

Each of these high-level languages has special numeric libraries allowing programs to take advantage of the capabilities of the 80387 NPX. No special programming conventions are necessary to make use of the 80387 NPX when programming numeric applications in any of these languages.

Programmers in PL/M-386 and ASM386 can also make use of many of these library routines by using routines contained in the 80387 Support Library. These libraries implement many of the functions provided by higher-level languages, including exception handlers, ASCII-to-floating-point conversions, and a more complete set of transcendental functions than that provided by the 80387 instruction set.

5.1.2 C Programs

C programmers automatically cause the C compiler to generate 80387 instructions when they use the **double** and **float** data types. The **float** type corresponds to the 80387's single real format; the **double** type corresponds to the 80387's double real format. The statement `#include <math.h>` causes mathematical functions such as **sin** and **sqrt** to return values of type **double**. Figure 5-1 illustrates the ease with which C programs interface with the 80387.

```
XENIX286 C386 COMPILER, V0.2 COMPILE OF MODULE SAMPLE
OBJECT MODULE PLACED IN sample.obj
COMPILER INVOKED BY: c386 sample.c

stmt level

1      ****
2      *
3      *          SAMPLE C PROGRAM
4      *
5      ****
6
7      /** Include /usr/include/stdio.h if necessary ***/
8      /** Include math declarations for transcendentals and others ***/
9
10     #include </usr/include/math.h>
11     #define PI 3.141592654
12
13     main()
14     {
15     double      sin_result, cos_result;
16     double      angle_deg = 0.0, angle_rad;
17     int         i, no_of_trial = 4;
18
19     for( i = 1; i <= no_of_trial; i++){
20         angle_rad = angle_deg * PI / 180.0;
21         sin_result = sin (angle_rad);
22         cos_result = cos (angle_rad);
23         printf("sine of %f degrees equals %f\n", angle_deg, sin_result);
24         printf("cosine of %f degrees equals %f\n\n", angle_deg, cos_result);
25         angle_deg = angle_deg + 30.0;
26     }
27     /** etc. */
28 }

C386 COMPILE COMPLETE. 0 WARNINGS, 0 ERRORS
```

Figure 5-1. Sample C-386 Program

5.1.3 PL/M-386

Programmers in PL/M-386 can access a very useful subset of the 80387's numeric capabilities. The PL/M-386 REAL data type corresponds to the NPX's single real (32-bit) format. This data type provides a range of about $8.43 \times 10^{-37} \leq |X| \leq 3.38 \times 10^{38}$, with about seven significant decimal digits. This representation is adequate for the data manipulated by many microcomputer applications.

The utility of the REAL data type is extended by the PL/M-386 compiler's practice of holding intermediate results in the 80387's extended real format. This means that the full range and precision of the processor are utilized for intermediate results. Underflow, overflow, and rounding exceptions are most likely to occur during intermediate computations rather than during calculation of an expression's final result. Holding intermediate results in extended-precision real format greatly reduces the likelihood of overflow and underflow and eliminates roundoff as a serious source of error until the final assignment of the result is performed.

The compiler generates 80387 code to evaluate expressions that contain REAL data types, whether variables or constants or both. This means that addition, subtraction, multiplication, division, comparison, and assignment of REALs will be performed by the NPX. INTEGER expressions, on the other hand, are evaluated on the CPU.

Five built-in procedures (Table 5-1) give the PL/M-386 programmer access to 80387 functions manipulated by the processor control instructions. Prior to any arithmetic operations, a typical PL/M-386 program will set up the NPX using the INIT\$REAL\$MATH\$UNIT procedure and then issue SET\$REAL\$MODE to configure the NPX. SET\$REAL\$MODE loads the 80387 control word, and its 16-bit parameter has the format shown for the control word in Chapter 1. The recommended value of this parameter is 033EH (round to nearest, 64-bit precision, all exceptions masked except invalid operation). Other settings may be used at the programmer's discretion.

If any exceptions are unmasked, an exception handler must be provided in the form of an interrupt procedure that is designated to be invoked via CPU interrupt vector number 16. The exception handler can use the GET\$REAL\$ERROR procedure to obtain the low-order

Table 5-1. PL/M-386 Built-In Procedures

Procedure	80387 Instruction	Description
INIT\$REAL\$MATH\$UNIT ⁽¹⁾	FINIT	Initialize processor.
SET\$REAL\$MODE	FLDCW	Set exception masks, rounding precision, and infinity controls.
GET\$REAL\$ERROR ⁽²⁾	FNSTSW & FNCLEX	Store, then clear, exception flags.
SAVE\$REAL\$STATUS	FNSAVE	Save processor state.
RESTORE\$REAL\$STATUS	FRSTOR	Restore processor state.

byte of the 80387 status word and to then clear the exception flags. The byte returned by GET\$REAL\$ERROR contains the exception flags; these can be examined to determine the source of the exception.

The SAVE\$REAL\$STATUS and RESTORE\$REAL\$STATUS procedures are provided for multitasking environments where a running task that uses the 80387 may be preempted by another task that also uses the 80387. It is the responsibility of the operating system to issue SAVE\$REAL\$STATUS before it executes any statements that affect the 80387; these include the INIT\$REAL\$MATH\$UNIT and SET\$REAL\$MODE procedures as well as arithmetic expressions. SAVE\$REAL\$STATUS saves the 80387 state (registers, status, and control words, etc.) on the CPU's stack. RESTORE\$REAL\$STATUS reloads the state information; the preempting task must invoke this procedure before terminating in order to restore the 80387 to its state at the time the running task was preempted. This enables the preempted task to resume execution from the point of its preemption.

5.1.4 ASM386

The ASM386 assembly language provides programmers with complete access to all of the facilities of the 80386 and 80387 processors.

The programmer's view of the 80386/80387 hardware is a single machine with these resources:

- 160 instructions
- 12 data types
- 8 general registers
- 6 segment registers
- 8 floating-point registers, organized as a stack

5.1.4.1 DEFINING DATA

The ASM386 directives shown in Table 5-2 allocate storage for 80387 variables and constants. As with other storage allocation directives, the assembler associates a type with any variable defined with these directives. The type value is equal to the length of the storage unit in bytes (10 for DT, 8 for DQ, etc.). The assembler checks the type of any variable coded in an instruction to be certain that it is compatible with the instruction. For example, the coding FIADD ALPHA will be flagged as an error if ALPHA's type is not 2 or 4,

Table 5-2. ASM386 Storage Allocation Directives

Directive	Interpretation	Data Types
DW	Define Word	Word integer
DD	Define Doubleword	Short integer, short real
DQ	Define Quadword	Long integer, long real
DT	Define Tenbyte	Packed decimal, temporary real

because integer addition is only available for word and short integer (doubleword) data types. The operand's type also tells the assembler which machine instruction to produce; although to the programmer there is only an FIADD instruction, a different machine instruction is required for each operand type.

On occasion it is desirable to use an instruction with an operand that has no declared type. For example, if register BX points to a short integer variable, a programmer may want to code FIADD [BX]. This can be done by informing the assembler of the operand's type in the instruction, coding FIADD DWORD PTR [BX]. The corresponding overrides for the other storage allocations are WORD PTR, QWORD PTR, and TBYTE PTR.

The assembler does not, however, check the types of operands used in processor control instructions. Coding FRSTOR [BP] implies that the programmer has set up register BP to point to the location (probably in the stack) where the processor's 94-byte state record has been previously saved.

The initial values for 80387 constants may be coded in several different ways. Binary integer constants may be specified as bit strings, decimal integers, octal integers, or hexadecimal strings. Packed decimal values are normally written as decimal integers, although the assembler will accept and convert other representations of integers. Real values may be written as ordinary decimal real numbers (decimal point required), as decimal numbers in scientific notation, or as hexadecimal strings. Using hexadecimal strings is primarily intended for defining special values such as infinities, NaNs, and denormalized numbers. Most programmers will find that ordinary decimal and scientific decimal provide the simplest way to initialize 80387 constants. Figure 5-2 compares several ways of setting the various 80387 data types to the same initial value.

```
; THE FOLLOWING ALL ALLOCATE THE CONSTANT: -126
; NOTE TWO'S COMPLETE STORAGE OF NEGATIVE BINARY INTEGERS.
;
; EVEN
WORD_INTEGER    DW   11111111000010B      ; FORCE WORD ALIGNMENT
SHORT_INTEGER    DD   0FFFFFFFFFF82H      ; BIT STRING
                 ; HEX STRING MUST START
                 ; WITH DIGIT
LONG_INTEGER     DQ   -126                ; ORDINARY DECIMAL
SINGLE_REAL      DD   -126.0              ; NOTE PRESENCE OF '.'
DOUBLE_REAL     DD   -1.26E2             ; "SCIENTIFIC"
PACKED_DECIMAL   DT   -126                ; ORDINARY DECIMAL INTEGER
;
; IN THE FOLLOWING, SIGN AND EXPONENT IS 'C005'
; SIGNIFICAND IS '7E00...00', 'R' INFORMS ASSEMBLER THAT
; THE STRING REPRESENTS A REAL DATA TYPE.
;
EXTENDED_REAL    DT   0C0057E000000000000000000R ; HEX STRING
```

Figure 5-2. Sample 80387 Constants

Note that preceding 80387 variables and constants with the ASM386 EVEN directive ensures that the operands will be word-aligned in memory. The best performance is obtained when data transfers are double-word aligned. All 80387 data types occupy integral numbers of words so that no storage is “wasted” if blocks of variables are defined together and preceded by a single EVEN declarative.

5.1.4.2 RECORDS AND STRUCTURES

The ASM386 RECORD and STRUC (structure) declaratives can be very useful in NPX programming. The record facility can be used to define the bit fields of the control, status, and tag words. Figure 5-3 shows one definition of the status word and how it might be used in a routine that polls the 80387 until it has completed an instruction.

Because structures allow different but related data types to be grouped together, they often provide a natural way to represent “real world” data organizations. The fact that the structure template may be “moved” about in memory adds to its flexibility. Figure 5-4 shows a simple structure that might be used to represent data consisting of a series of test score samples. A structure could also be used to define the organization of the information stored and loaded by the FSTENV and FLDENV instructions.

```
; RESERVE SPACE FOR STATUS WORD
STATUS_WORD
; LAY OUT STATUS WORD FIELDS
STATUS RECORD
& BUSY:           1,
& COND_CODE3:     1,
& STACK_TOP:      3,
& COND_CODE2:     1,
& COND_CODE1:     1,
& COND_CODE0:     1,
& INT_REQ:        1,
& S_FLAG:         1,
& P_FLAG:         1,
& U_FLAG:         1,
& O_FLAG:         1,
& Z_FLAG:         1,
& D_FLAG:         1,
& I_FLAG:         1
; REDUCE UNTIL COMPLETE
REDUCE: FPREM1
    FNSTS  STATUS_WORD
    TEST   STATUS_WORD, MASK_COND_CODE2
    JNZ    REDUCE
```

Figure 5-3. Status Word Record Definition

```

SAMPLE      STRUC
    N_OBS     DD   ?    ; SHORT INTEGER
    MEAN      DQ   ?    ; DOUBLE REAL
    MODE      DW   ?    ; WORD INTEGER
    STD_DEV   DQ   ?    ; DOUBLE REAL
    ; ARRAY OF OBSERVATIONS -- WORD INTEGER
    TEST_SCORES DW 1000 DUP (?)
SAMPLE      ENDS

```

Figure 5-4. Structure Definition

Table 5-3. Addressing Method Examples

Coding	Interpretation
FIADD ALPHA	ALPHA is a simple scalar (mode is direct).
FDIVR ALPHA.BETA	BETA is a field in a structure that is "overlaid" on ALPHA (mode is direct).
FMUL QWORD PTR [BX]	BX contains the address of a long real variable (mode is register indirect).
FSUB ALPHA [SI]	ALPHA is an array and SI contains the offset of an array element from the start of the array (mode is indexed).
FILD [BP].BETA	BP contains the address of a structure on the CPU stack and BETA is a field in the structure (mode is based).
FBLD TBYTE PTR [BX] [DI]	BX contains the address of a packed decimal array and DI contains the offset of an array element (mode is based indexed).

5.1.4.3 Addressing Methods

80387 memory data can be accessed with any of the memory addressing methods provided by the ModR/M byte and (optionally) the SIB byte. This means that 80387 data types can be incorporated in data aggregates ranging from simple to complex according to the needs of the application. The addressing methods and the ASM386 notation used to specify them in instructions make the accessing of structures, arrays, arrays of structures, and other organizations direct and straightforward. Table 5-3 gives several examples of 80387 instructions coded with operands that illustrate different addressing methods.

5.1.5 Comparative Programming Example

Figures 5-5 and 5-6 show the PL/M-386 and ASM386 code for a simple 80387 program, called ARRSUM. The program references an array (X\$ARRAY), which contains 0–100 single real values; the integer variable N\$OF\$X indicates the number of array elements the program is to consider. ARRSUM steps through X\$ARRAY accumulating three sums:

- SUM\$X, the sum of the array values
- SUM\$INDEXES, the sum of each array value times its index, where the index of the first element is 1, the second is 2, etc.
- SUM\$SQUARES, the sum of each array element squared

(A true program, of course, would go beyond these steps to store and use the results of these calculations.) The control word is set with the recommended values: round to nearest, 64-bit precision, interrupts enabled, and all exceptions masked except invalid operation. It is assumed that an exception handler has been written to field the invalid operation if it occurs, and that it is invoked by interrupt pointer 16. Either version of the program will run on an actual or an emulated 80387 without altering the code shown.

The PL/M-386 version of ARRSUM (Figure 5-5) is very straightforward and illustrates how easily the 80387 can be used in this language. After declaring variables, the program calls built-in procedures to initialize the processor (or its emulator) and to load to the control word. The program clears the sum variables and then steps through X\$ARRAY with a DO-loop. The loop control takes into account PL/M-386's practice of considering the index of the first element of an array to be 0. In the computation of SUM\$INDEXES, the built-in procedure FLOAT converts I+1 from integer to real because the language does not support "mixed mode" arithmetic. One of the strengths of the NPX, of course, is that it *does* support arithmetic on mixed data types (because all values are converted internally to the 80-bit extended-precision real format).

The ASM386 version (Figure 5-6) defines the external procedure INIT387, which makes the different initialization requirements of the processor and its emulator transparent to the source code. After defining the data and setting up the segment registers and stack pointer, the program calls INIT387 and loads the control word. The computation begins with the next three instructions, which clear three registers by loading (pushing) zeros onto the stack. As shown in Figure 5-7, these registers remain at the bottom of the stack throughout the computation while temporary values are pushed on and popped off the stack above them.

The program uses the CPU LOOP instruction to control its iteration through X_ARRAY; register ECX, which LOOP automatically decrements, is loaded with N_OF_X, the number of array elements to be summed. Register ESI is used to select (index) the array elements. The program steps through X_ARRAY from back to front, so ESI is initialized to point at the element just beyond the first element to be processed. The ASM386 TYPE operator is used to determine the number of bytes in each array element. This permits changing X_ARRAY to a double-precision real array by simply changing its definition (DD to DQ) and reassembling.

XENIX286 PL/M-386 DEBUG X291a COMPILATION OF MODULE ARRAYSUM
OBJECT MODULE PLACED IN arraysum.obj
COMPILER INVOKED BY: plm386 arraysum.plm

```
*****  
*          *  
*          ARRAYSUM  MODULE          *  
*          *  
*****  
  
1      array$sum:    do;  
  
2  1      declare (sum$x, sum$indexes, sum$squares) real;  
3  1      declare x$array(100) real;  
4  1      declare (n$of$x, i) integer;  
5  1      declare control$387 literally '033eh';  
  
6  1      /* Assume x$array and n$of$x are initialized */  
7  1      call init$real$math$unit;  
7  1      call set$real$mode(control$387);  
  
8  1      /* Clear sums */  
8  1      sum$x, sum$indexes, sum$squares = 0.0;  
  
9  1      /* Loop through array, accumulating sums */  
10 2      do i = 0 to n$of$x - 1;  
11 2      sum$x = sum$x + x$array(i);  
12 2      sum$indexes = sum$indexes + (x$array(i)*float(i+1));  
13 2      sum$squares = sum$squares + (x$array(i)*x$array(i));  
13 2      end;  
  
14 1      /* etc. */  
14 1      end array$sum;
```

MODULE INFORMATION:

CODE AREA SIZE	= 000000A0H	1600
CONSTANT AREA SIZE	= 00000004H	40
VARIABLE AREA SIZE	= 00001A4H	4200
MAXIMUM STACK SIZE	= 00000004H	40
32 LINES READ		
0 PROGRAM WARNINGS		
0 PROGRAM ERRORS		

DICTIONARY SUMMARY:

8KB MEMORY USED
0KB DISK SPACE USED

END OF PL/M-386 COMPILATION

Figure 5-5. Sample PL/M-386 Program

```

XENIX286 80386 MACRO ASSEMBLER V1.0, ASSEMBLY OF MODULE ARRAYSUM
OBJECT MODULE PLACED IN arraysum.obj
ASSEMBLER INVOKED BY: asm386 arraysum.asm

LOC     OBJ          LINE    SOURCE
      1   name      arraysum
      2
      3 ; Define initialization routine
      4
      5   extrn    init387:far
      6
      7 ; Allocate space for data
      8
      9   data      segment rw public
00000000 3E03 10   control_387    dw 033eh
00000002 ????????
00000006 (100 11   n_of_x      dd ?
??????????
)
00000196 ????????
0000019A ????????
0000019E ????????
      12   x_array    dd 100 dup (?)
      13
      14
      15
      16   data      ends
      17
      18 ; Allocate CPU stack space
      19
      20   stack     stackseg 400
      21
      22 ; Begin code
      23
      24   code      segment er public
      25
      26   assume  ds:data, ss:stack
      27
00000000 28 start:
00000000 66B8---- R   29   mov      ax, data
00000004 8ED8 30   mov      ds, ax
00000006 66B8---- R   31   mov      ax, stack
0000000A BB000000000 32   mov      eax, 0h
0000000F 8ED0 33   mov      ss, ax
00000011 BC00000000 34   mov      esp, stackstart stack
35
36 ; Assume x_array and n_of_x have
37 ; been initialized
38
39 ; Prepare the 80387 or its emulator
40
00000016 9A00000000---- E   41   call    init387
0000001D D92000000000 R   42   fldcw   control_387
43
44 ; Clear three registers to hold
45 ; running sums
46
00000023 D9EE 47   fldz
00000025 D9EE 48   fldz
00000027 D9EE 49   fldz

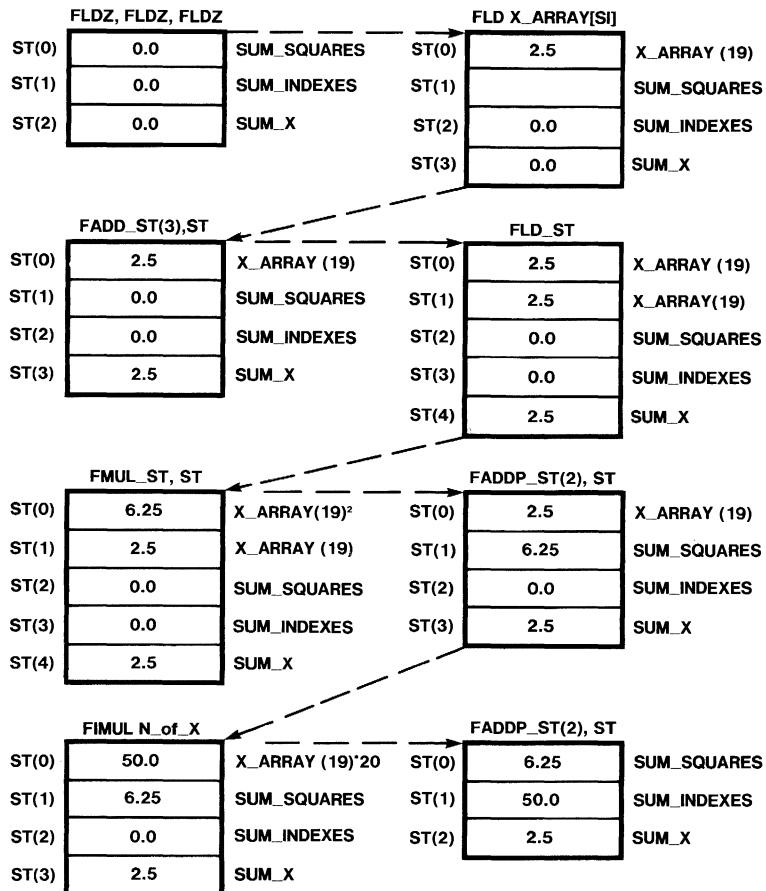
```

Figure 5-6. Sample ASM386 Program

LOC	OBJ	LINE	SOURCE
		50	
		51	; Setup ECX as loop counter and ESI
		52	; as index into x_array
		53	
00000029	8B0D02000000	R 54	mov ecx, n_of_x
0000002F	F7E9		imul ecx
00000031	8BF0		mov esi, eax
		57	
		58	; ESI now contains index of last
		59	; element + 1
		60	; Loop through x_array and
		61	; accumulate sum
		62	
00000033		63	sum_next:
		64	; backup one element and push on
		65	; the stack
		66	
00000033	83EE04	R 67	sub esi, type x_array
00000036	D98606000000	R 68	fld x_array[esi]
		69	
		70	; add to the sum and duplicate x
		71	; on the stack
		72	
0000003C	DCC3	73	fadd st(3), st
0000003E	D9C0	74	fld st
		75	
		76	; square it and add into the sum of
		77	; (index+1) and discard
		78	
00000040	DCC8	79	fmul st, st
00000042	DEC2	80	faddp st(2), st
		81	
		82	; reduce index for next iteration
		83	
00000044	FF0D02000000	R 84	dec n_of_x
0000004A	E2E7	85	loop sum_next
		86	
		87	; Pop sums into memory
		88	
0000004C		89	pop_results:
0000004C	D91D96010000	R 90	fstp sum_squares
00000052	D91D9A010000	R 91	fstp sum_indexes
00000058	D91D9E010000	R 92	fstp sum_x
0000005E	9B	93	fwait
		94	
		95	;
		96	; Etc.
		97	;
-----		98	code ends
		99	end start, ds:data, ss:stack

ASSEMBLY COMPLETE, NO WARNINGS, NO ERRORS.

Figure 5-6. Sample ASM386 Program (Cont'd.).



122164-14

Figure 5-7. Instructions and Register Stack

Figure 5-7 shows the effect of the instructions in the program loop on the NPX register stack. The figure assumes that the program is in its first iteration, that N_OF_X is 20, and that X_ARRAY(19) (the 20th element) contains the value 2.5. When the loop terminates, the three sums are left as the top stack elements so that the program ends by simply popping them into memory variables.

5.1.6 80387 Emulation

The programming of applications to execute on both 80386 with an 80387 and 80386 systems without an 80387 is made much easier by the existence of an 80387 emulator for 80386 systems. The Intel EMUL387 emulator offers a complete software counterpart to the 80387 hardware; NPX instructions can be simply emulated in software rather than being executed in hardware. With software emulation, the distinction between 80386 systems with or without an 80387 is reduced to a simple performance differential. Identical numeric programs will simply execute more slowly (using software emulation of NPX instructions) on 80386 systems without an 80387 than on an 80386/80387 system executing NPX instructions directly.

When incorporated into the systems software, the emulation of NPX instructions on the 80386 systems is completely transparent to the applications programmer. Applications software needs no special libraries, linking, or other activity to allow it to run on an 80386 with 80387 emulation.

To the applications programmer, the development of programs for 80386 systems is the same whether the 80387 NPX hardware is available or not. The full 80387 instruction set is available for use, with NPX instructions being either emulated or executed directly. Applications programmers need not be concerned with the hardware configuration of the computer systems on which their applications will eventually run.

For systems programmers, details relating to 80387 emulators are described in Chapter 6.

The EMUL387 software emulator for 80386 systems is available from Intel as a separate program product.

5.2 CONCURRENT PROCESSING WITH THE 80387

Because the 80386 CPU and the 80387 NPX have separate execution units, it is possible for the NPX to execute numeric instructions in parallel with instructions executed by the CPU. This simultaneous execution of different instructions is called concurrency.

No special programming techniques are required to gain the advantages of concurrent execution; numeric instructions for the NPX are simply placed in line with the instructions for the CPU. CPU and numeric instructions are initiated in the same order as they are encountered by the CPU in its instruction stream. However, because numeric operations performed by the NPX generally require more time than operations performed by the CPU, the CPU can often execute several of its instructions before the NPX completes a numeric instruction previously initiated.

This concurrency offers obvious advantages in terms of execution performance, but concurrency also imposes several rules that must be observed in order to assure proper synchronization of the 80386 CPU and 80387 NPX.

All Intel high-level languages automatically provide for and manage concurrency in the NPX. Assembly-language programmers, however, must understand and manage some areas of concurrency in exchange for the flexibility and performance of programming in assembly language. This section is for the assembly-language programmer or well-informed high-level-language programmer.

5.2.1 Managing Concurrency

Concurrent execution of the host and 80387 is easy to establish and maintain. The activities of numeric programs can be split into two major areas: program control and arithmetic. The program control part performs activities such as deciding what functions to perform, calculating addresses of numeric operands, and loop control. The arithmetic part simply adds, subtracts, multiplies, and performs other operations on the numeric operands. The NPX and host are designed to handle these two parts separately and efficiently.

Concurrency management is required to check for an exception before letting the 80386 change a value just used by the 80387. Almost any numeric instruction can, under the wrong circumstances, produce a numeric exception. For programmers in higher-level languages, all required synchronization is automatically provided by the appropriate compiler. For assembly-language programmers exception synchronization remains the responsibility of the assembly-language programmer.

A complication is that a programmer may not expect his numeric program to cause numeric exceptions, but in some systems, they may regularly happen. To better understand these points, consider what can happen when the NPX detects an exception.

Depending on options determined by the software system designer, the NPX can perform one of two things when a numeric exception occurs:

- The NPX can provide a default fix-up for selected numeric exceptions. Programs can mask individual exception types to indicate that the NPX should generate a safe, reasonable result whenever that exception occurs. The default exception fix-up activity is treated by the NPX as part of the instruction causing the exception; no external indication of the exception is given. When exceptions are detected, a flag is set in the numeric status register, but no information regarding where or when is available. If the NPX performs its default action for all exceptions, then the need for exception synchronization is not manifest. However, as will be shown later, this is not sufficient reason to ignore exception synchronization when designing programs that use the 80387.
- As an alternative to the NPX default fix-up of numeric exceptions, the 80386 CPU can be notified whenever an exception occurs. When a numeric exception is unmasked and the exception occurs, the NPX stops further execution of the numeric instruction and signals this event to the CPU. On the next occurrence of an ESC or WAIT instruction,

the CPU traps to a software exception handler. The exception handler can then implement any sort of recovery procedures desired for any numeric exception detectable by the NPX. Some ESC instructions do not check for exceptions. These are the nonwaiting forms FNINIT, FNSTENV, FNSAVE, FNSTSW, FNSTCW, and FNCLEX.

When the NPX signals an unmasked exception condition, it is requesting help. The fact that the exception was unmasked indicates that further numeric program execution under the arithmetic and programming rules of the NPX is unreasonable.

If concurrent execution is allowed, the state of the CPU when it recognizes the exception is undefined. The CPU may have changed many of its internal registers and be executing a totally different program by the time the exception occurs. To handle this situation, the NPX has special registers updated at the start of each numeric instruction to describe the state of the numeric program when the failed instruction was attempted.

Exception synchronization ensures that the NPX is in a well-defined state after an unmasked numeric exception occurs. Without a well-defined state, it would be impossible for exception recovery routines to determine why the numeric exception occurred, or to recover successfully from the exception.

The following two sections illustrate the need to always consider exception synchronization when writing 80387 code, even when the code is initially intended for execution with exceptions masked. If the code is later moved to an environment where exceptions are unmasked, the same code may not work correctly. An example of how some instructions written without exception synchronization will work initially, but fail when moved into a new environment is shown in Figure 5-8.

INCORRECT ERROR SYNCHRONIZATION

```
FILD  COUNT    ; NPX instruction
INC   COUNT    ; CPU instruction alters operand
FSQRT COUNT   ; subsequent NPX instruction -- error from
                ; previous NPX instruction detected here
```

PROPER ERROR SYNCHRONIZATION

```
FILD  COUNT    ; NPX instruction
FSQRT COUNT   ; subsequent NPX instruction -- error from
                ; previous NPX instruction detected here
INC   COUNT    ; CPU instruction alters operand
```

Figure 5-8. Exception Synchronization Examples

5.2.1.1 INCORRECT EXCEPTION SYNCHRONIZATION

In Figure 5-8, three instructions are shown to load an integer, calculate its square root, then increment the integer. The 80386-to-80387 interface and synchronous execution of the NPX emulator will allow this program to execute correctly when no exceptions occur on the FILD instruction.

This situation changes if the 80387 numeric register stack is extended to memory. To extend the NPX stack to memory, the invalid exception is unmasked. A push to a full register or pop from an empty register sets SF and causes an invalid exception.

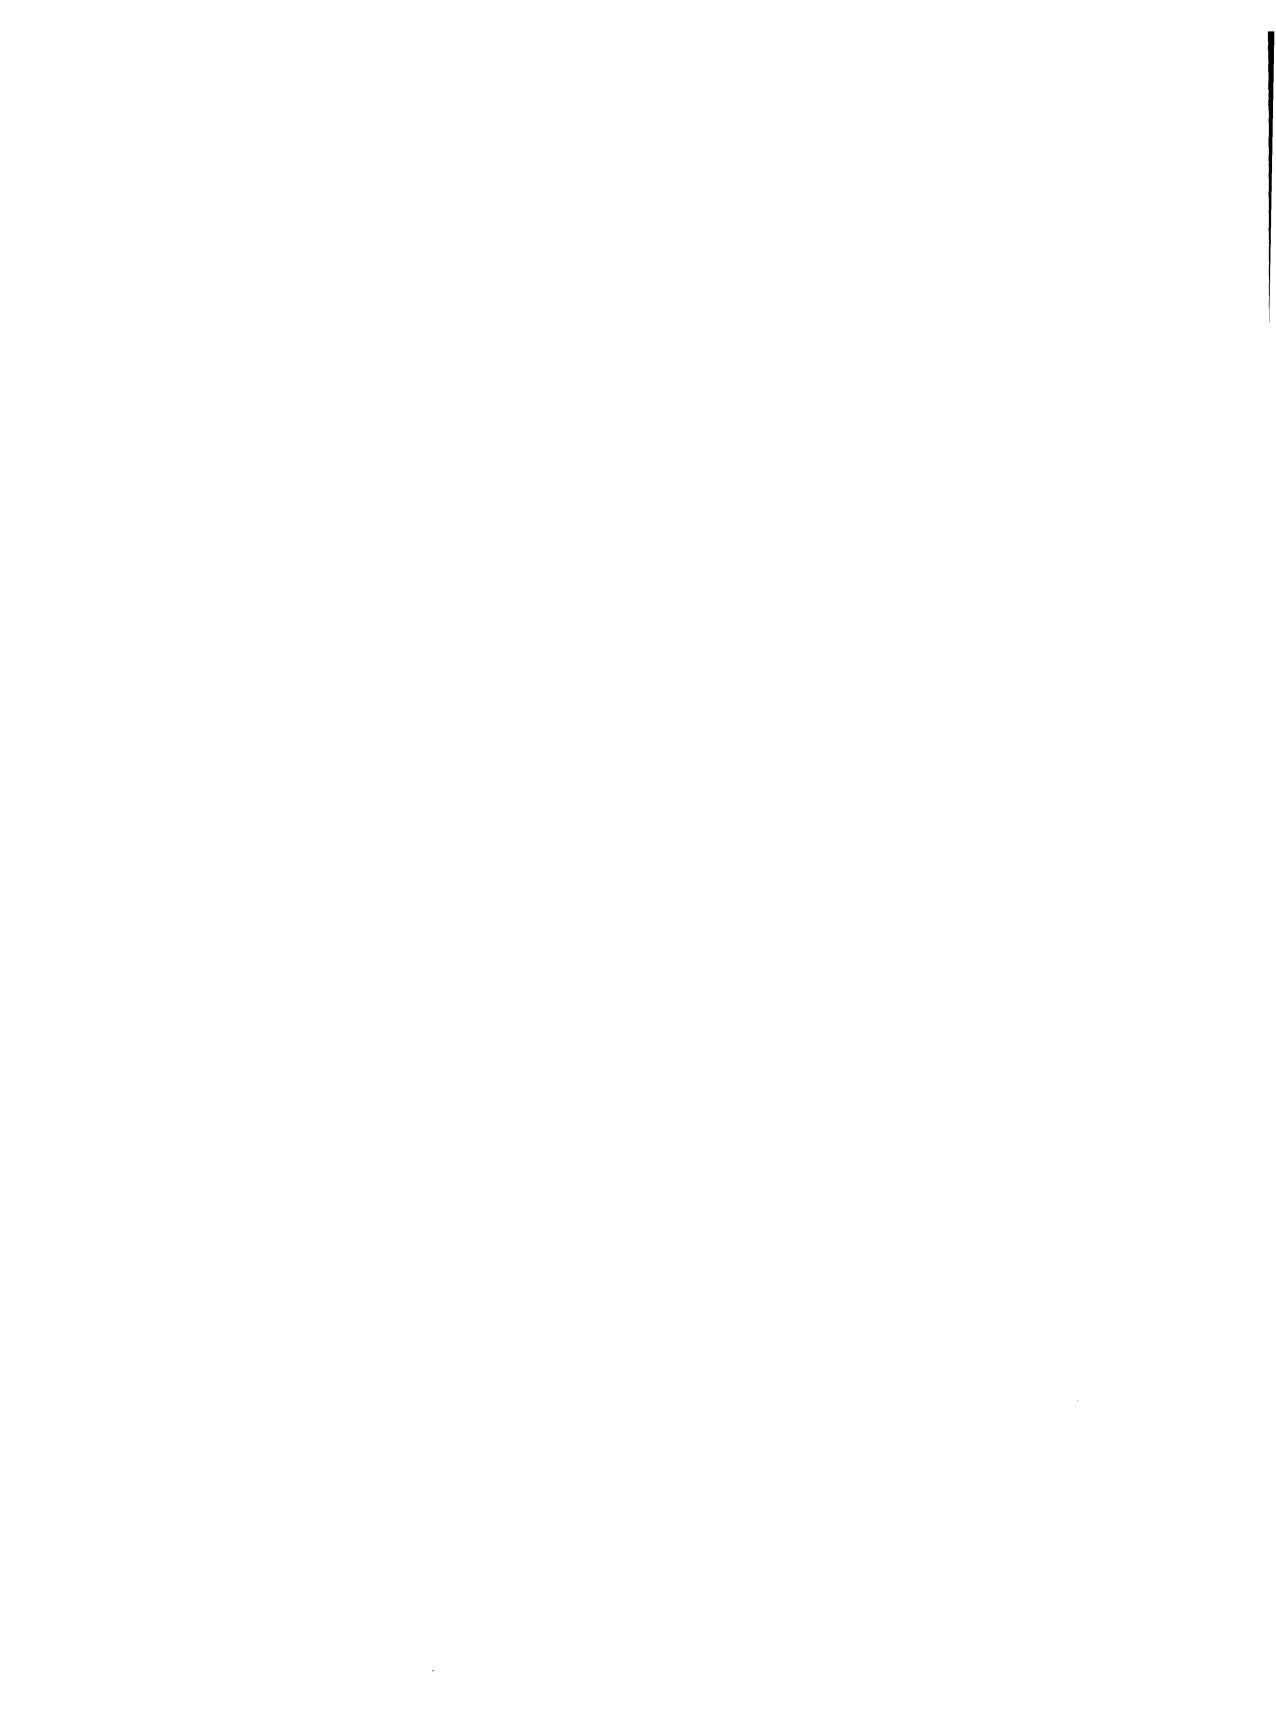
The recovery routine for the exception must recognize this situation, fix up the stack, then perform the original operation. The recovery routine will not work correctly in the first example shown in the figure. The problem is that the value of COUNT is incremented before the NPX can signal the exception to the CPU. Because COUNT is incremented before the exception handler is invoked, the recovery routine will load an incorrect value of COUNT, causing the program to fail or behave unreliably.

5.2.1.2 PROPER EXCEPTION SYNCHRONIZATION

Exception synchronization relies on the WAIT instruction and the BUSY# and ERROR# signals of the 80387. When an unmasked exception occurs in the 80387, it asserts the ERROR# signal, signaling to the CPU that a numeric exception has occurred. The next time the CPU encounters a WAIT instruction or an exception-checking ESC instruction, the CPU acknowledges the ERROR# signal by trapping automatically to Interrupt #16, the processor-extension exception vector. If the following ESC or WAIT instruction is properly placed, the CPU will not yet have disturbed any information vital to recovery from the exception.

System-Level Numeric Programming

6



CHAPTER 6

SYSTEM-LEVEL NUMERIC PROGRAMMING

System programming for 80387 systems requires a more detailed understanding of the 80387 NPX than does application programming. Such things as emulation, initialization, exception handling, and data and error synchronization are all the responsibility of the systems programmer. These topics are covered in detail in the sections that follow.

6.1 80386/80387 ARCHITECTURE

On a software level, the 80387 NPX appears as an extension of the 80386 CPU. On the hardware level, however, the mechanisms by which the 80386 and 80387 interact are more complex. This section describes how the 80387 NPX and 80386 CPU interact and points out features of this interaction that are of interest to systems programmers.

6.1.1 Instruction and Operand Transfer

All transfers of instructions and operands between the 80387 and system memory are performed by the 80386 using I/O bus cycles. The 80387 appears to the CPU as a special peripheral device. It is special in two respects: the CPU initiates I/O automatically when it encounters ESC instructions, and the CPU uses reserved I/O addresses to communicate with the 80387. These I/O operations are completely transparent to software.

Because the 80386 actually performs all transfers between the 80387 and memory, no additional bus drivers, controllers, or other components are necessary to interface the 80387 NPX to the local bus. The 80387 can utilize instructions and operands located in any memory accessible to the 80386 CPU.

6.1.2 Independent of CPU Addressing Modes

Unlike the 80287, the 80387 is not sensitive to the addressing and memory management of the CPU. The 80387 operates the same regardless of whether the 80386 CPU is operating in real-address mode, in protected mode, or in virtual 8086 mode.

The instruction FSETPM that was necessary in 80286/80287 systems to set the 80287 into protected mode is not needed for the 80387. The 80387 treats this instruction as a no-op.

Because the 80386 actually performs all transfers between the 80387 and memory, 80387 instructions can utilize any memory location accessible by the task currently executing on the 80386. When operating in protected mode, all references to memory operands are automatically verified by the 80386's memory management and protection mechanisms as for any other memory references by the currently-executing task. Protection violations associated with NPX instructions automatically cause the 80386 to trap to an appropriate exception handler.

To the numerics programmer, the operating modes of the 80386 affect only the manner in which the NPX instruction and data pointers are represented in memory following an FSAVE or FSTENV instruction. Each of these instructions produces one of four formats depending on both the operating mode and on the operand-size attribute in effect for the instruction. The differences are detailed in the discussion of the FSAVE and FSTENV instructions in Chapter 4.

6.1.3 Dedicated I/O Locations

The 80387 NPX does not require that any memory addresses be set aside for special purposes. The 80387 does make use of I/O port addresses, but these are 32-bit addresses with the high-order bit set (i.e. > 80000000H); therefore, these I/O operations are completely transparent to the 80386 software. Because these addresses are beyond the 64 Kbyte I/O addressing limit of I/O instructions, 80386 programs cannot reference these reserved I/O addresses directly.

6.2 PROCESSOR INITIALIZATION AND CONTROL

One of the principal responsibilities of systems software is the initialization, monitoring, and control of the hardware and software resources of the system, including the 80387 NPX. In this section, issues related to system initialization and control are described, including recognition of the NPX, emulation of the 80387 NPX in software if the hardware is not available, and the handling of exceptions that may occur during the execution of the 80387.

6.2.1 System Initialization

During initialization of an 80386 system, systems software must

- Recognize the presence or absence of the NPX.
- Set flags in the 80386 MSW to reflect the state of the numeric environment.

If an 80387 NPX is present in the system, the NPX must be initialized. All of these activities can be quickly and easily performed as part of the overall system initialization.

6.2.2 Hardware Recognition of the NPX

The 80386 identifies the type of its coprocessor (80287 or 80387) by sampling its ERROR# input some time after the falling edge of RESET and before executing the first ESC instruction. The 80287 keeps its ERROR# output in inactive state after hardware reset; the 80387 keeps its ERROR# output in active state after hardware reset. The 80386 records this difference in the ET bit of control register zero (CR0). The 80386 subsequently uses ET to control its interface with the coprocessor. If ET is set, it employs the 32-bit protocol of the 80387; if ET is not set, it employs the 16-bit protocol of the 80287.

Systems software can (if necessary) change the value of ET. There are three reasons that ET may not be set:

1. An 80287 is actually present.
2. No coprocessor is present.
3. An 80387 is present but it is connected in a nonstandard manner that does not trigger the setting of ET.

An example of case three is the PC/AT-compatible design described in Appendix F. In such cases, initialization software may need to change the value of ET.

6.2.3 Software Recognition of the NPX

Figure 6-1 shows an example of a recognition routine that determines whether an NPX is present, and distinguishes between the 80387 and the 8087/80287. This routine can be executed on any 80386, 80286, or 8086 hardware configuration that has an NPX socket.

The example guards against the possibility of accidentally reading an expected value from a floating data bus when no NPX is present. Data read from a floating bus is undefined. By expecting to read a specific bit pattern from the NPX, the routine protects itself from the indeterminate state of the bus. The example also avoids depending on any values in reserved bits, thereby maintaining compatibility with future numerics coprocessors.

6.2.4 Configuring the Numerics Environment

Once the 80386 CPU has determined the presence or absence of the 80387 or 80287 NPX, the 80386 must set either the MP or the EM bit in its own control register zero (CR0) accordingly. The initialization routine can either

- Set the MP bit in CR0 to allow numeric instructions to be executed directly by the NPX.
- Set the EM bit in the CR0 to permit software emulation of the numeric instructions.

The MP (monitor coprocessor) flag of CR0 indicates to the 80386 whether an NPX is physically available in the system. The MP flag controls the function of the WAIT instruction. When executing a WAIT instruction, the 80386 tests the task switched (TS) bit only if MP is set; if it finds TS set under these conditions, the CPU traps to exception #7.

The Emulation Mode (EM) bit of CR0 indicates to the 80386 whether NPX functions are to be emulated. If the CPU finds EM set when it executes an ESC instruction, program control is automatically trapped to exception #7, giving the exception handler the opportunity to emulate the functions of an 80387.

For correct 80386 operation, the EM bit must never be set concurrently with MP. The EM and MP bits of the 80386 are described in more detail in the *80386 Programmer's Reference Manual*. More information on software emulation for the 80387 NPX is described in the "80387 Emulation" section later in this chapter. In any case, if ESC instructions are to be executed, either the MP or EM bit must be set, but not both.

8086/87/88/186 MACRO ASSEMBLER Test for presence of a Numerics Chip, Revision 1.0

PAGE 1

DOS 3.20 (033-N) 8086/87/88/186 MACRO ASSEMBLER V2.0 ASSEMBLY OF MODULE TEST_NPX
OBJECT MODULE PLACED IN FINDNPX.OBJ

LOC	OBJ	LINE	SOURCE
		1	+1 \$title('Test for presence of a Numerics Chip, Revision 1.0')
		2	
		3	name Test_NPX
		4	
----		5	stack segment stack 'stack'
0000 (100		6	dw 100 dup (?)
???			
)			
00C8 ????		7	sst dw ?
----		8	stack ends
		9	
----		10	data segment public 'data'
0000 0000		11	temp dw 0h
----		12	data ends
		13	
		14	dgroup group data, stack
		15	cgroup group code
		16	
----		17	code segment public 'code'
		18	assume cs:cgroup, ds:dgroup
		19	
0000		20	start:
		21	;
		22	; Look for an 8087, 80287, or 80387 NPX.
		23	; Note that we cannot execute WAIT on 8086/88 if no 8087 is present.
		24	;
0000		25	test_npx:
0000 90DBE3	R	26	finit ; Must use non-wait form
0003 BE0000		27	mov si,offset dgroup:temp
0006 C7045A5A		28	mov word ptr [si],5A5AH ; Initialize temp to non-zero value
000A 90DD3C		29	fnstsw [si] ; Must use non-wait form of fstsw
		30	; It is not necessary to use a WAIT instruction
		31	; after fnstsw or fnstcw. Do not use one here.
000D 803C00		32	cmp byte ptr [si],0 ; See if correct status with zeroes was read
0010 752A		33	jne no_npx ; Jump if not a valid status word, meaning no NPX
		34	;
		35	; Now see if ones can be correctly written from the control word.
		36	;
0012 90D93C		37	fnstcw [si] ; Look at the control word; do not use WAIT form
		38	; Do not use a WAIT instruction here!
0015 8804		39	mov ax,[si] ; See if ones can be written by NPX
0017 253F10		40	and ax,103fh ; See if selected parts of control word look OK
001A 3D3F00		41	cmp ax,3fh ; Check that ones and zeroes were correctly read
001D 7510		42	jne no_npx ; Jump if no NPX is installed
		43	;
		44	; Some numerics chip is installed. NPX instructions and WAIT are now safe.
		45	; See if the NPX is an 8087, 80287, or 80387.
		46	; This code is necessary if a denormal exception handler is used or the
		47	; new 80387 instructions will be used.
		48	;

Figure 6-1. Software Routine to Recognize the 80287

8086/87/88/186 MACRO ASSEMBLER Test for presence of a Numerics Chip, Revision 1.0			PAGE	2
LOC	OBJ	LINE	SOURCE	
001F	9BD9E8	49	fld1	; Must use default control word from FNINIT
0022	9BD9EE	50	fldz	; Form infinity
0025	9BDEF9	51	fdiv	; 8087/287 says +inf = -inf
0028	9BD9C0	52	fld st	; Form negative infinity
002B	9BD9E0	53	fchs	; 80387 says +inf <> -inf
002E	9BDED9	54	fcompp	; See if they are the same and remove them
0031	9BDD3C	55	fstsw [si]	; Look at status from FCOMPP
0034	8804	56	mov ax,[si]	
0036	9E	57	sahf	; See if the infinities matched
0037	7406	58	je found_87_287	; Jump if 8087/287 is present
		59	;	
		60	;	An 80387 is present. If denormal exceptions are used for an 8087/287,
		61	;	they must be masked. The 80387 will automatically normalize denormal
		62	;	operands faster than an exception handler can.
		63	;	
0039	EB0790	64	jmp found_387	
003C		65	no_npx:	
		66	;	set up for no NPX
		67	;	...
		68	;	
003C	EB0490	69	jmp exit	
003F		70	found_87_287:	
		71	;	set up for 87/287
		72	;	...
		73	;	
003F	EB0190	74	jmp exit	
0042		75	found_387:	
		76	;	set up for 387
		77	;	...
		78	;	
0042		79	exit:	
----		80	code ends	
		81	end start,ds:dgroup,ss:dgroup:sst	

ASSEMBLY COMPLETE, NO ERRORS FOUND

Figure 6-1. Software Routine to Recognize the 80287 (Cont'd.).

6.2.5 Initializing the 80387

Initializing the 80387 NPX simply means placing the NPX in a known state unaffected by any activity performed earlier. A single FNINIT instruction performs this initialization. All the error masks are set, all registers are tagged empty, TOP is set to zero, and default rounding and precision controls are set. Table 6-1 shows the state of the 80387 NPX following FINIT or FNINIT. This state is compatible with that of the 80287 after FINIT or after hardware RESET.

The FNINIT instruction *does not* leave the 80387 in the same state as that which results from the hardware RESET signal. Following a hardware RESET signal, such as after initial power-up, the state of the 80387 differs in the following respects:

1. The mask bit for the invalid-operation exception is reset.
2. The invalid-operation exception flag is set.
3. The exception-summary bit is set (along with its mirror image, the B-bit).

Table 6-1. NPX Processor State Following Initialization

Field	Value	Interpretation
Control Word (Infinity Control)* Rounding Control Precision Control Exception Masks	0 00 11 111111	Affine Round to nearest 64 bits All exceptions masked
Status Word (Busy) Condition Code Stack Top Exception Summary Stack Flag Exception Flags	0 0000 000 0 0 000000	— — Register 0 is stack top No exceptions — No exceptions
Tag Word Tags	11	Empty
Registers	N.C.	Not changed
Exception Pointers Instruction Code Instruction Address Operand Address	N.C. N.C. N.C.	Not changed Not changed Not changed

*The 80387 does not have infinity control. This value is listed to emphasize that programs written for the 80287 may not behave the same on the 80387 if they depend on this bit.

These settings cause assertion of the ERROR# signal as described previously. The FNINIT instruction must be used to change the 80387 state to one compatible with the 80287.

6.2.6 80387 Emulation

If it is determined that no 80387 NPX is available in the system, systems software may decide to emulate ESC instructions in software. This emulation is easily supported by the 80386 hardware, because the 80386 can be configured to trap to a software emulation routine whenever it encounters an ESC instruction in its instruction stream.

Whenever the 80386 CPU encounters an ESC instruction, and its MP and EM status bits are set appropriately (MP=0, EM=1), the 80386 automatically traps to interrupt #7, the “processor extension not available” exception. The return link stored on the stack points to the first byte of the ESC instruction, including the prefix byte(s), if any. The exception handler can use this return link to examine the ESC instruction and proceed to emulate the numeric instruction in software.

The emulator must step the return pointer so that, upon return from the exception handler, execution can resume at the first instruction following the ESC instruction.

To an application program, execution on an 80386 system with 80387 emulation is almost indistinguishable from execution on a system with an 80387, except for the difference in execution speeds.

There are several important considerations when using emulation on an 80386 system:

- When operating in protected mode, numeric applications using the emulator must be executed in execute-readable code segments. Numeric software cannot be emulated if it is executed in execute-only code segments. This is because the emulator must be able to examine the particular numeric instruction that caused the emulation trap.
- Only privileged tasks can place the 80386 in emulation mode. The instructions necessary to place the 80386 in emulation mode are privileged instructions, and are not typically accessible to an application.

An emulator package (EMUL387) that runs on 80386 systems is available from Intel. This emulation package operates in both real and protected mode as well as in virtual 8086 mode, providing a complete functional equivalent for the 80387 emulated in software.

When using the EMUL387 emulator, writers of numeric exception handlers should be aware of one slight difference between the emulated 80387 and the 80387 hardware:

- On the 80387 hardware, exception handlers are invoked by the 80386 at the first WAIT or ESC instruction following the instruction causing the exception. The return link, stored on the 80386 stack, points to this second WAIT or ESC instruction where execution will resume following a return from the exception handler.
- Using the EMUL387 emulator, numeric exception handlers are invoked from within the emulator itself. The return link stored on the stack when the exception handler is invoked will therefore point back to the EMUL387 emulator, rather than to the program code actually being executed (emulated). An IRET return from the exception handler returns to the emulator, which then returns immediately to the emulated program. This added layer of indirection should not cause confusion, however, because the instruction causing the exception can always be identified from the 80387's instruction and data pointers.

6.2.7 Handling Numerics Exceptions

Once the 80387 has been initialized and normal execution of applications has been commenced, the 80387 NPX may occasionally require attention in order to recover from numeric processing exceptions. This section provides details for writing software exception handlers for numeric exceptions. Numeric processing exceptions have already been introduced in Chapter 3.

The 80387 NPX can take one of two actions when it recognizes a numeric exception:

- If the exception is masked, the NPX will automatically perform its own masked exception response, correcting the exception condition according to fixed rules, and then continuing with its instruction execution.
- If the exception is unmasked, the NPX signals the exception to the 80386 CPU using the ERROR# status line between the two processors. Each time the 80386 encounters an ESC or WAIT instruction in its instruction stream, the CPU checks the condition of this ERROR# status line. If ERROR# is active, the CPU automatically traps to Interrupt vector #16, the Processor Extension Error trap.

Interrupt vector #16 typically points to a software exception handler, which may or may not be a part of systems software. This exception handler takes the form of an 80386 interrupt procedure.

When handling numeric errors, the CPU has two responsibilities:

- The CPU must not disturb the numeric context when an error is detected.
- The CPU must clear the error and attempt recovery from the error.

Although the manner in which programmers may treat these responsibilities varies from one implementation to the next, most exception handlers will include these basic steps:

- Store the NPX environment (control, status, and tag words, operand and instruction pointers) as it existed at the time of the exception.
- Clear the exception bits in the status word.
- Enable interrupts on the CPU.
- Identify the exception by examining the status and control words in the saved environment.
- Take some system-dependent action to rectify the exception.
- Return to the interrupted program and resume normal execution.

6.2.8 Simultaneous Exception Response

In cases where multiple exceptions arise simultaneously, the 80387 signals one exception according to the precedence shown at the end of Chapter 3. This means, for example, that an SNaN divided by zero results in an invalid operation, not in a zero divide exception.

6.2.9 Exception Recovery Examples

Recovery routines for NPX exceptions can take a variety of forms. They can change the arithmetic and programming rules of the NPX. These changes may redefine the default fix-up for an error, change the appearance of the NPX to the programmer, or change how arithmetic is defined on the NPX.

A change to an exception response might be to automatically normalize all denormals loaded from memory. A change in appearance might be extending the register stack into memory to provide an “infinite” number of numeric registers. The arithmetic of the NPX can be changed to automatically extend the precision and range of variables when exceeded. All these functions can be implemented on the NPX via numeric exceptions and associated recovery routines in a manner transparent to the application programmer.

Some other possible application-dependent actions might include:

- Incrementing an exception counter for later display or printing
- Printing or displaying diagnostic information (e.g., the 80387 environment and registers)
- Aborting further execution
- Storing a diagnostic value (a NaN) in the result and continuing with the computation

Notice that an exception may or may not constitute an error, depending on the application. Once the exception handler corrects the condition causing the exception, the floating-point instruction that caused the exception can be restarted, if appropriate. This cannot be accomplished using the IRET instruction, however, because the trap occurs at the ESC or WAIT instruction following the offending ESC instruction. The exception handler must obtain (using FSAVE or FSTENV) the address of the offending instruction in the task that initiated it, make a copy of it, execute the copy in the context of the offending task, and then return via IRET to the current CPU instruction stream.

In order to correct the condition causing the numeric exception, exception handlers must recognize the precise state of the NPX at the time the exception handler was invoked, and be able to reconstruct the state of the NPX when the exception initially occurred. To reconstruct the state of the NPX, programmers must understand when, during the execution of an NPX instruction, exceptions are actually recognized.

Invalid operation, zero divide, and denormalized exceptions are detected before an operation begins, whereas overflow, underflow, and precision exceptions are not raised until a true result has been computed. When a *before* exception is detected, the NPX register stack and memory have not yet been updated, and appear as if the offending instructions has not been executed.

When an *after* exception is detected, the register stack and memory appear as if the instruction has run to completion; i.e., they may be updated. (However, in a store or store-and-pop operation, unmasked over/underflow is handled like a *before* exception; memory is not updated and the stack is not popped.) The programming examples contained in Chapter 7 include an outline of several exception handlers to process numeric exceptions for the 80387.



Numeric Programming Examples

7



CHAPTER 7

NUMERIC PROGRAMMING EXAMPLES

The following sections contain examples of numeric programs for the 80387 NPX written in ASM386. These examples are intended to illustrate some of the techniques for programming the 80386/80387 computing system for numeric applications.

7.1 CONDITIONAL BRANCHING EXAMPLE

As discussed in Chapter 2, several numeric instructions post their results to the condition code bits of the 80387 status word. Although there are many ways to implement conditional branching following a comparison, the basic approach is as follows:

- Execute the comparison.
- Store the status word. (80387 allows storing status directly into AX register.)
- Inspect the condition code bits.
- Jump on the result.

Figure 7-1 is a code fragment that illustrates how two memory-resident double-format real numbers might be compared (similar code could be used with the FTST instruction). The numbers are called A and B, and the comparison is A to B.

The comparison itself requires loading A onto the top of the 80387 register stack and then comparing it to B, while popping the stack with the same instruction. The status word is then written into the 80386 AX register.

A and B have four possible orderings, and bits C3, C2, and C0 of the condition code indicate which ordering holds. These bits are positioned in the upper byte of the NPX status word so as to correspond to the CPU's zero, parity, and carry flags (ZF, PF, and CF), when the byte is written into the flags. The code fragment sets ZF, PF, and CF of the CPU status word to the values of C3, C2, and C0 of the NPX status word, and then uses the CPU conditional jump instructions to test the flags. The resulting code is extremely compact, requiring only seven instructions.

The FXAM instruction updates all four condition code bits. Figure 7-2 shows how a jump table can be used to determine the characteristics of the value examined. The jump table (FXAM_TBL) is initialized to contain the 32-bit displacement of 16 labels, one for each possible condition code setting. Note that four of the table entries contain the same value, "EMPTY." The first two condition code settings correspond to "EMPTY." The two other table entries that contain "EMPTY" will never be used on the 80387, but may be used if the code is executed with an 80287.

The program fragment performs the FXAM and stores the status word. It then manipulates the condition code bits to finally produce a number in register BX that equals the condition

```
.  
. .  
A DQ ?  
B DQ ?  
. .  
fld A ; LOAD A ONTO TOP OF 387 STACK  
fcomp B ; COMPARE A:B, POP A  
fstsw ax ; STORE RESULT TO CPU AX REGISTER  
;  
; CPU AX REGISTER CONTAINS CONDITION CODES  
; (RESULTS OF COMPARE)  
; LOAD CONDITION CODES INTO CPU FLAGS  
;  
sahf  
;  
; USE CONDITIONAL JUMPS TO DETERMINE ORDERING OF A TO B  
;  
jp a_b_unordered ; TEST C2 (PF)  
jb a_less ; TEST C0 (CF)  
je a_equal ; TEST C3 (ZF)  
a_greater: ; C0 (CF) = 0, C3 (ZF) = 0  
. .  
a_equal: ; C0 (CF) = 0, C3 (ZF) = 1  
. .  
a_less: ; C0 (CF) = 1, C3 (ZF) = 0  
. .  
a_b_unordered: ; C2 (PF) = 1  
. .
```

Figure 7-1. Conditional Branching for Comparisons

code times 2. This involves zeroing the unused bits in the byte that contains the code, shifting C3 to the right so that it is adjacent to C2, and then shifting the code to multiply it by 2. The resulting value is used as an index that selects one of the displacements from FXAM_TBL (the multiplication of the condition code is required because of the 2-byte length of each value in FXAM_TBL). The unconditional JMP instruction effectively vectors through the jump table to the labeled routine that contains code (not shown in the example) to process each possible result of the FXAM instruction.

7.2 EXCEPTION HANDLING EXAMPLES

There are many approaches to writing exception handlers. One useful technique is to consider the exception handler procedure as consisting of “prologue,” “body,” and “epilogue” sections of code. This procedure is invoked via interrupt number 16.

```
; JUMP TABLE FOR EXAMINE ROUTINE
; FXAM_TBL DD POS_UNNORM, POS_NAN, NEG_UNNORM, NEG_NAN,
&           POS_NORM, POS_INFINITY, NEG_NORM,
&           NEG_INFINITY, POS_ZERO, EMPTY, NEG_ZERO,
&           EMPTY, POS_DENORM, EMPTY, NEG_DENORM, EMPTY

; EXAMINE ST AND STORE RESULT (CONDITION CODES)

FXAM
XOR EAX,EAX ; CLEAR EAX
FSTSW AX

; CALCULATE OFFSET INTO JUMP TABLE

AND AX,0100011100000000B ; CLEAR ALL BITS EXCEPT C3, C2-C0
SHR EAX,6      ; SHIFT C2-C0 INTO PLACE      (0000XXX0)
SAL AH,5       ; POSITION C3            (000X0000)
OR AL,AH       ; DROP C3 IN ADJACENT TO C2 (000XXXX0)
XOR AH,AH       ; CLEAR OUT THE OLD COPY OF C3

; JUMP TO THE ROUTINE 'ADDRESSED' BY CONDITION CODE

JMP FXAM_TBL[EAX]

; HERE ARE THE JUMP TARGETS, ONE TO HANDLE
; EACH POSSIBLE RESULT OF FXAM

POS_UNNORM:
.
POS_NAN:
.
NEG_UNNORM:
.
NEG_NAN:
.
POS_NORM:
.
POS_INFINITY:
.
NEG_NORM:
.
NEG_INFINITY:
.
POS_ZERO:
.
EMPTY:
.
NEG_ZERO:
.
POS_DENORM:
.
NEG_DENORM:
```

Figure 7-2. Conditional Branching for FXAM

At the beginning of the prologue, CPU interrupts have been disabled. The prologue performs all functions that must be protected from possible interruption by higher-priority sources. Typically, this involves saving CPU registers and transferring diagnostic information from the 80387 to memory. When the critical processing has been completed, the prologue may enable CPU interrupts to allow higher-priority interrupt handlers to preempt the exception handler.

The body of the exception handler examines the diagnostic information and makes a response that is necessarily application-dependent. This response may range from halting execution, to displaying a message, to attempting to repair the problem and proceed with normal execution.

The epilogue essentially reverses the actions of the prologue, restoring the CPU and the NPX so that normal execution can be resumed. The epilogue must *not* load an unmasked exception flag into the 80387 or another exception will be requested immediately.

Figures 7-3 through 7-5 show the ASM386 coding of three skeleton exception handlers. They show how prologues and epilogues can be written for various situations, but provide comments indicating only where the application dependent exception handling body should be placed.

```
SAVE_ALL      PROC
;
; SAVE CPU REGISTERS, ALLOCATE STACK SPACE
; FOR 80387 STATE IMAGE
    PUSH EBP
    MOV  EBP,ESP
    SUB  ESP,108
;
; SAVE FULL 80387 STATE, ENABLE CPU INTERRUPTS
    FNSAVE [EBP-108]
    STI
;
; APPLICATION-DEPENDENT EXCEPTION HANDLING
; CODE GOES HERE
;
; CLEAR EXCEPTION FLAGS IN STATUS WORD
; (WHICH IS IN MEMORY)
; RESTORE MODIFIED STATE IMAGE
    MOV  BYTE PTR [EBP-104], 0H
    FRSTOR [EBP-108]
;
; DEALLOCATE STACK SPACE, RESTORE CPU REGISTERS
    MOVE ESP,EBP
    .
    .
    POP EBP
;
; RETURN TO INTERRUPTED CALCULATION
    IRET
SAVE_ALL      ENDP
```

Figure 7-3. Full-State Exception Handler

```
SAVE_ENVIRONMENT PROC
;
; SAVE CPU REGISTERS, ALLOCATE STACK SPACE
; FOR 80387 ENVIRONMENT
; PUSH    EBP
;
MOV     EBP,ESP
SUB    ESP,28
; SAVE ENVIRONMENT, ENABLE CPU INTERRUPTS
FNSTENV [EBP-28]
STI
;
; APPLICATION EXCEPTION-HANDLING CODE GOES HERE
;
; CLEAR EXCEPTION FLAGS IN STATUS WORD
; (WHICH IS IN MEMORY)
; RESTORE MODIFIED ENVIRONMENT IMAGE
MOV    BYTE PTR [EBP-24], 0H
FLDENV [EBP-28]
; DE-ALLOCATE STACK SPACE, RESTORE CPU REGISTERS
MOV    ESP,EBP
POP    EBP
;
; RETURN TO INTERRUPTED CALCULATION
IRET
SAVE_ENVIRONMENT ENDP
```

Figure 7-4. Reduced-Latency Exception Handler

Figures 7-3 and 7-4 are very similar; their only substantial difference is their choice of instructions to save and restore the 80387. The tradeoff here is between the increased diagnostic information provided by FNSAVE and the faster execution of FNSTENV. For applications that are sensitive to interrupt latency or that do not need to examine register contents, FNSTENV reduces the duration of the “critical region,” during which the CPU does not recognize another interrupt request.

After the exception handler body, the epilogues prepare the CPU and the NPX to resume execution from the point of interruption (i.e., the instruction following the one that generated the unmasked exception). Notice that the exception flags in the memory image that is loaded into the 80387 are cleared to zero prior to reloading (in fact, in these examples, the entire status word image is cleared).

The examples in Figures 7-3 and 7-4 assume that the exception handler itself will not cause an unmasked exception. Where this is a possibility, the general approach shown in Figure 7-5 can be employed. The basic technique is to save the full 80387 state and then to load a new control word in the prologue. Note that considerable care should be taken when designing an exception handler of this type to prevent the handler from being reentered endlessly.

```
        .  
        LOCAL_CONTROL DW ? ; ASSUME INITIALIZED  
        .  
  
REENTRANT          PROC  
; SAVE CPU REGISTERS, ALLOCATE STACK SPACE FOR  
; 80387 STATE IMAGE  
    PUSH    EBP  
    .  
    .  
    MOV     EBP,ESP  
    SUB     ESP,108  
; SAVE STATE, LOAD NEW CONTROL WORD,  
; ENABLE CPU INTERRUPTS  
    FNSAVE  [EBP-108]  
    FLDCW   LOCAL_CONTROL  
    STI  
    .  
    .  
; APPLICATION EXCEPTION HANDLING CODE GOES HERE.  
; AN UNMASKED EXCEPTION GENERATED HERE WILL  
; CAUSE THE EXCEPTION HANDLER TO BE REENTERED.  
; IF LOCAL STORAGE IS NEEDED, IT MUST BE  
; ALLOCATED ON THE CPU STACK.  
    .  
    .  
; CLEAR EXCEPTION FLAGS IN STATUS WORD  
; (WHICH IS IN MEMORY)  
; RESTORE MODIFIED STATE IMAGE  
    MOV     BYTE PTR [EBP-104], 0H  
    FRSTOR  [EBP-108]  
; DE-ALLOCATE STACK SPACE, RESTORE CPU REGISTERS  
    MOV     ESP,EBP  
    .  
    .  
    POP    EBP  
; RETURN TO POINT OF INTERRUPTION  
    IRET  
REENTRANT          ENDP
```

Figure 7-5. Reentrant Exception Handler

7.3 FLOATING-POINT TO ASCII CONVERSION EXAMPLES

Numeric programs must typically format their results at some point for presentation and inspection by the program user. In many cases, numeric results are formatted as ASCII strings for printing or display. This example shows how floating-point values can be converted to decimal ASCII character strings. The function shown in Figure 7-6 can be invoked from PL/M-386, Pascal-386, FORTRAN-386, or ASM386 routines.

```
XENIX286 80386 MACRO ASSEMBLER V1.0, ASSEMBLY OF MODULE FLOATING_TO_ASCII
OBJECT MODULE PLACED IN fpasc.obj
ASSEMBLER INVOKED BY: asm386 fpasc.asm

LOC    OBJ          LINE      SOURCE
00000000          1  +1 $title('Convert a floating point number to ASCII')
                   2
                   3          name      floating_to_ascii
                   4
                   5          public    floating_to_ascii
                   6          extrn    get_power_10:near,tos_status:near
                   7          ;
                   8          ; This subroutine will convert the floating point
                   9          ; number in the top of the FPU stack to an ASCII
                  10         ; string and separate power of 10 scaling value
                  11         ; (in binary). The maximum width of the ASCII string
                  12         ; formed is controlled by a parameter which must be
                  13         ; > 1. Unnormal values, denormal values, and pseudo
                  14         ; zeroes will be correctly converted. However, unnormals
                  15         ; and pseudo zeros are no longer supported formats on the
                  16         ; 80387 in conformance with the IEEE floating point
                  17         ; standard) and hence not generated internally. A
                  18         ; returned value will indicate how many binary bits
                  19         ; of precision were lost in an unnormal or denormal
                  20         ; value. The magnitude (in terms of binary power)
                  21         ; of a pseudo zero will also be indicated. Integers
                  22         ; less than 10**18 in magnitude are accurately converted
                  23         ; if the destination ASCII string field is wide enough
                  24         ; to hold all the digits. Otherwise the value is converted
                  25         ; to scientific notation.
                  26         ;
                  27         ; The status of the conversion is identified by the
                  28         ; return value, it can be:
                  29         ;
                  30         ;     0      conversion complete, string_size is defined
                  31         ;     1      invalid arguments
                  32         ;     2      exact integer conversion, string_size is defined
                  33         ;     3      indefinite
                  34         ;     4      + NAN (Not A Number)
                  35         ;     5      - NAN
                  36         ;     6      + Infinity
                  37         ;     7      - Infinity
                  38         ;     8      pseudo zero found, string_size is defined
                  39         ;
                  40         ; The PLM/386 calling convention is:
                  41         ;
                  42         ; floating_to_ascii:
                  43         ;     procedure (number,denormal_ptr,string_ptr,size_ptr,
                  44         ;             field_size,power_ptr) word external;
                  45         ;     declare (denormal_ptr,string_ptr,power_ptr,size_ptr)
                  46         ; pointer;
                  47         ;     declare field_size word,
                  48         ;     string_size based size_ptr word;
                  49         ;     declare number real;
                  50         ;     declare denormal integer based denormal_ptr;
```

Figure 7-6. Floating-Point to ASCII Conversion Routine

LOC	OBJ	LINE	SOURCE
		51	; declare power integer based power_ptr;
		52	; end floating_to_ascii;
		53	;
		54	; The floating point value is expected to be
		55	on the top of the NPX stack. This subroutine
		56	expects 3 free entries on the NPX stack and
		57	will pop the passed value off when done. The
		58	generated ASCII string will have a leading
		59	character either '-' or '+' indicating the sign
		60	of the value. The ASCII decimal digits will
		61	immediately follow. The numeric value of the
		62	ASCII string is (ASCII STRING.)*10**POWER. If
		63	the given number was zero, the ASCII string will
		64	contain a sign and a single zero character. The
		65	value string_size indicates the total length of
		66	the ASCII string including the sign character.
		67	String(0) will always hold the sign. It is
		68	possible for string_size to be less than
		69	field_size. This occurs for zeroes or integer
		70	values. A pseudo zero will return a special
		71	return code. The denormal count will indicate
		72	the power of two originally associated with the
		73	value. The power of ten and ASCII string will
		74	be as if the value was an ordinary zero.
		75	;
		76	This subroutine is accurate up to a maximum of
		77	18 decimal digits for integers. Integer values
		78	will have a decimal power of zero associated
		79	with them. For non integers, the result will be
		80	accurate to within 2 decimal digits of the 16th
		81	decimal place(double precision). The exponentiate
		82	instruction is also used for scaling the value into
		83	the range acceptable for the BCD data type. The
		84	rounding mode in effect on entry to the
		85	subroutine is used for the conversion.
		86	;
		87	The following registers are not transparent:
		88	;
		89	eax ebx ecx edx esi edi eflags
		90	;
		91	;
		92	Define the stack layout.
		93	;
00000000()		94	ebp_save equ dword ptr [ebp]
00000004()		95	es_save equ ebp_save + size ebp_save
00000008()		96	return_ptr equ es_save + size es_save
0000000C()		97	power_ptr equ return_ptr + size return_ptr
00000010()		98	field_size equ power_ptr + size power_ptr
00000014()		99	size_ptr equ field_size + size field_size
00000018()		100	string_ptr equ size_ptr + size size_ptr
0000001C()		101	denormal_ptr equ string_ptr + size string_ptr
		102	
0014		103	parms_size equ size power_ptr + size field_size +
		104	& size size_ptr + size string_ptr +
		105	& size denormal_ptr

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.).

LOC	OBJ	LINE	SOURCE
		106	;
		107	; Define constants used
		108	;
0012		109	BCD_DIGITS equ 18 ; Number of digits in bcd_value
0004		110	WORD_SIZE equ 4
000A		111	BCD_SIZE equ 10
0001		112	MINUS equ 1 ; Define return values
0004		113	NAN equ 4 ; The exact values chosen
0006		114	INFINITY equ 6 ; here are important. They must
0003		115	INDEFINITE equ 3 ; correspond to the possible return
0008		116	PSEUDO_ZERO equ 8 ; values and be in the same numeric
-0002		117	INVALID equ -2 ; order as tested by the program.
-0004		118	ZERO equ -4
-0006		119	DENORMAL equ -6
-0008		120	UNNORMAL equ -8
0000		121	NORMAL equ 0
0002		122	EXACT equ 2
		123	;
		124	; Define layout of temporary storage area.
		125	;
FFFFFFFFFFC[]		126	power_two equ word ptr [ebp - WORD_SIZE]
FFFFFFFFFF2[]		127	bcd_value equ tbyte ptr power_two - BCD_SIZE
FFFFFFFFFF2[]		128	bcd_byte equ byte ptr bcd_value
FFFFFFFFFF2[]		129	fraction equ bcd_value
		130	
000C		131	local_size equ size power_two + size bcd_value
		132	;
		133	; Allocate stack space for the temporaries so
		134	; the stack will be big enough
		135	;
-----		136	stack stackseg (local_size+6) ; Allocate stack
		137	; space for locals
		138	+1 \$eject

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

LOC	OBJ	LINE	SOURCE
		139	code
		140	segment public er
		141	extrn power_table:qword
		142	;
		143	;
		144	;
00000000 0A00		145	even
		146	dw 10
		147	;
		148	;
		149	;
		150	;
00000002 F8		151	status_table db UNNORMAL, NAN, UNNORMAL + MINUS,
00000003 04		152	& NAN + MINUS, NORMAL, INFINITY,
00000004 F9		153	& NORMAL + MINUS, INFINITY + MINUS,
00000005 05		154	& ZERO, INVALID, ZERO + MINUS, INVALID,
00000006 00		155	& DENORMAL, INVALID, DENORMAL + MINUS, INVALID
00000007 06			
00000008 01			
00000009 07			
0000000A FC			
0000000B FE			
0000000C FD			
0000000D FE			
0000000E FA			
0000000F FE			
00000010 FB			
00000011 FE			
00000012		156	floating_to_ascii proc
00000012 E800000000	E	157	
		158	
		159	call tos_status ; Look at status of ST(0)
		160	
00000017 2E0FB68002000000	R	161	;
0000001F 3CFE		162	get descriptor from table
00000021 7527		163	movzx eax, status_table[eax]
		164	cmp al,INVALID ; Look for empty ST(0)
		165	;
		166	jne not_empty
		167	;
00000023 C21400		168	ST(0) is empty! Return the status value.
		169	ret parms_size
		170	;
		171	Remove infinity from stack and exit.
00000026		172	;
00000026 D0D8		173	found_infinity:
00000028 EB02		174	fstp st(0) ; OK to leave fstp running
		175	jmp short exit_proc
		176	;
		177	String space is too small!
		178	;
0000002A		179	Return invalid code.
0000002A B0FE		180	small_string:
0000002C		181	mov al,INVALID
0000002C C9		182	exit_proc:
			leave ; Restore stack setup

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

LOC	OBJ	LINE	SOURCE
		183	pop es
0000002D 07		184	ret parms_size
		185	;
		186	; ST(0) is NAN or indefinite. Store the
		187	; value in memory and look at the fraction
		188	; field to separate indefinite from an ordinary NAN.
		189	;
00000031		190	NAN_or_indefinite:
00000031 DB7DF2		191	fstp fraction ; Remove value from stack
		192	; for examination
00000034 A801		193	test al,MINUS ; Look at sign bit
00000036 9B		194	fwait ; Insure store is done
00000037 74F3		195	jz exit_proc ; Can't be indefinite if
		196	; positive
		197	
00000039 BB000000C0		198	mov ebx,0C000000H ; Match against upper 32
		199	;bits of fraction
		200	
0000003E 2B5DF6		201	; Compare bits 63-32
		202	sub ebx,dword ptr fraction + 4
		203	
00000041 0B5DF2		204	; Bits 31-0 must be zero
00000044 75E6		205	or ebx,dword ptr fraction
		206	jnz exit_proc
		207	
00000046 B003		208	; Set return value for indefinite value
00000048 EBE2		209	mov al,INDEFINITE
		210	jmp exit_proc
		211	
		212	; Allocate stack space for local variables
		213	; and establish parameter addressability.
		214	
0000004A		215	not_empty:
0000004A 06		216	push es ; Save working register
0000004B C80C0000		217	enter local_size,0 ; Setup stack addressing
		218	
		219	
0000004F 884D10		220	; Check for enough string space
00000052 83F902		221	mov ecx,field_size
00000055 7CD3		222	cmp ecx,2
		223	jl small_string
		224	
00000057 49		225	dec ecx ; Adjust for sign character
		226	
00000058 83F912		227	; See if string is too large for BCD
00000058 7605		228	cmp ecx,BCD_DIGITS
		229	jbe size_ok
		230	
00000050 B912000000		231	; Else set maximum string size
00000062		232	mov ecx,BCD_DIGITS
00000062 3C06		233	size_ok:
		234	cmp al,INFINITY ; Look for infinity
		235	
00000064 7DC0		236	; Return status value for + or - inf
		237	jge found_infinity

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.).

LOC	OBJ	LINE	SOURCE
		238	
00000066 3C04		239	cmp al,NAN ; Look for NAN or INDEFINITE
00000068 7DC7		240	jge NAN_or_indefinite
		241	;
		242	; Set default return values and check that
		243	the number is normalized.
		244	;
0000006A D9E1		245	fabs ; Use positive value only
		246	; sign bit in al has true sign of value
0000006C 31D2		247	xor edx,edx ; Form 0 constant
0000006E 887D1C		248	mov edi,denormal_ptr; Zero denormal count
00000071 668917		249	mov [edi],dx
00000074 885D0C		250	mov ebx,power_ptr ; Zero power of ten value
00000077 668913		251	mov [ebx],dx
0000007A 88C2		252	mov dl,al
0000007C 80E201		253	and dl,1
0000007F 80C202		254	add dl,EXACT
00000082 3FCF		255	cmp al,ZERO ; Test for zero
00000084 0F83BC000000		256	jae convert_integer ; Skip power code if value
		257	; is zero
0000008A DB7DF2		258	fstp fraction
0000008B 98		259	fwait
0000008E 8A45F9		260	mov al,bcd_byte+7
00000091 804DF980		261	or byte ptr bcd_byte+7,80h
00000095 DB6DF2		262	fld fraction
00000098 D9F4		263	fxtract
0000009A A880		264	test al,80h
0000009C 7524		265	jnz normal_value
		266	
0000009E D9E8		267	fldl
000000A0 DEE9		268	fsub
000000A2 D9E4		269	ftst
000000A4 98DFE0		270	fstsw ax
000000A7 9E		271	sahf
000000A8 7510		272	jnz set_unnormal_count
		273	;
		274	; Found a pseudo zero
		275	;
000000AA D9EC		276	fldlg2 ; Develop power of ten estimate
000000AC 80C206		277	add dl,PSEUDO_ZERO-EXACT
000000AF DECA		278	fmulp st(2),st
000000B1 D9C9		279	fxch
000000B3 DF1B		280	fistp word ptr [ebx]; Set power of ten
000000B5 E98C000000		281	jmp convert_integer
		282	
000000BA		283	set_unnormal_count:
000000BA D9F4		284	fxtract ; Get original fraction,
		285	; now normalized
000000BC D9C9		286	fxch ; Get unnormal count
000000BE D9E0		287	fchs
000000C0 DF1F		288	fistp word ptr [edi]; Set unnormal count
		289	
		290	;
		291	; Calculate the decimal magnitude associated
		292	; with this number to within one order. This

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

LOC	OBJ	LINE	SOURCE
		293	; error will always be inevitable due to
		294	; rounding and lost precision. As a result,
		295	; we will deliberately fail to consider the
		296	; LOG10 of the fraction value in calculating
		297	; the order. Since the fraction will always
		298	; be $1 \leq F < 2$, its LOG10 will not change
		299	; the basic accuracy of the function. To
		300	; get the decimal order of magnitude, simply
		301	; multiply the power of two by LOG10(2) and
		302	; truncate the result to an integer.
		303	;
		304	normal_value:
		305	fstp fraction ; Save the fraction field
		306	; for later use
		307	fist power_two ; Save power of two
		308	fldlg2 ; Get LOG10(2)
		309	; Power_two is now safe to use
		310	fmul ; Form LOG10(of exponent of number)
		311	fstp word ptr [ebx] ; Any rounding mode
		312	; will work here
		313	;
		314	Check if the magnitude of the number rules
		315	out treating it as an integer.
		316	;
		317	CX has the maximum number of decimal digits
		318	allowed.
		319	;
		320	fwait ; Wait for power_ten to be valid
		321	;
		322	; Get power of ten of value
		323	movsx si, word ptr [ebx]
		324	sub esi,ecx ; Form scaling factor
		325	; necessary in ax
		326	ja adjust_result ; Jump if number will not fit
		327	;
		328	The number is between 1 and $10^{**(\text{field_size})}$.
		329	Test if it is an integer.
		330	;
		331	fld power_two ; Restore original number
		332	sub dl,NORMAL-EXACT ; Convert to exact return
		333	; value
		334	fld fraction ; Form full value, this
		335	fscale ; is safe here
		336	;
		337	fst st(1) ; Copy value for compare
		338	frndint ; Test if its an integer
		339	fcomp ; Compare values
		340	fstsw ax ; Save status
		341	sahf ; C3=1 implies it was
		342	; an integer
		343	jnz convert_integer
		344	;
		345	fstp st(0) ; Remove non integer value
		346	add dl,NORMAL-EXACT ; Restore original return value
		347	;

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

LOC	OBJ	LINE	SOURCE
		348	; Scale the number to within the range allowed
		349	; by the BCD format. The scaling operation should
		350	; produce a number within one decimal order of
		351	; magnitude of the largest decimal number
		352	; representable within the given string width.
		353	;
		354	;
		355	The scaling power of ten value is in si.
		356	;
000000F2		adjust_result:	
000000F2 8BC6		357	mov eax,esi ; Setup for pow10
000000F4 668903		358	mov word ptr [ebx],ax ; Set initial power
000000F7 F7D8		359	;
		360	neg eax ; of ten return value
		361	; Subtract one for each order of
000000F9 E800000000	E	362	; magnitude the value is scaled by
		363	call get_power_10 ; Scaling factor is
		364	; returned as
000000FE DB6DF2		365	; exponent and fraction
00000101 DEC9		366	fld fraction ; Get fraction
00000103 8BF1		367	fmul ; Combine fractions
00000105 C1E603		368	mov esi,ecx ; Form power of ten of
00000108 DF45FC		369	; the maximum
0000010B DEC2		370	shl esi,3 ; BCD value to fit in
0000010D D9FD		371	;
		372	fild power_two ; Combine powers of two
		373	faddp st(2),st ; Form full value,
		374	; exponent was safe
0000010F DDD9		375	fstp st(1) ; Remove exponent
		376	;
		377	; Test the adjusted value against a table
		378	; of exact powers of ten. The combined errors
		379	; of the magnitude estimate and power function
		380	; can result in a value one order of magnitude
		381	; too small or too large to fit correctly in
		382	; the BCD field. To handle this problem, pretest
		383	; the adjusted value, if it is too small or
		384	; large, then adjust it by ten and adjust the
		385	; power of ten value.
		386	;
00000111		387	test_power:
		388	;
		389	; Compare against exact power entry. Use the next
		390	; entry since cx has been decremented by one
00000111 2EDC9608000000	E	391	fcom power_table[esi]+type power_table
00000118 9BDFF0		392	fstsw ax ; No wait is necessary
0000011B 9E		393	sahf ; If C3 = C0 = 0 then
0000011C 720F		394	jb test_for_small ; too big
0000011E 2EDE3500000000	R	395	;
00000125 80E2FD		396	fidiv const10 ; Else adjust value
00000128 66FF03		397	and dl,not EXACT ; Remove exact flag
0000012B EB17		398	inc word ptr [ebx] ; Adjust power of ten value
		399	jmp short_in_range ; Convert the value to a BCD
		400	; integer
0000012D		401	test_for_small:
0000012D 2EDC9600000000	E	402	fcom power_table[esi] ; Test relative size

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.).

OC	OBJ	LINE	SOURCE
		403	fstsw ax ; No wait is necess
0000134 9BDFE0		404	ary sahf ; If CO = 0 then
0000137 9E		405	; st(0) >= lower bound
0000138 720A		406	jc in_range ; Convert the value
		407	to a ; BCD integer
		408	
000013A 2EDE0D00000000	R	409	fimul const10 ; Adjust value into range
0000141 66FF0B		410	dec word ptr [ebx] ; Adjust power of ten value
0000144		411	in_range:
0000144 D9FC		412	frndint ; Form integer value
		413	;
		414	; Assert: 0 <= TOS <= 999,999,999,999,999,999
		415	; The TOS number will be exactly representable
		416	; in 18 digit BCD format.
		417	;
0000146		418	convert_integer:
0000146 DF75F2		419	fbstp bcd_value ; Store as BCD format number
		420	;
		421	While the store BCD runs, setup registers
		422	for the conversion to ASCII.
		423	;
0000149 BE08000000		424	mov esi,BCD_SIZE-2 ; Initial BCD index value
000014E 66B9040F		425	mov cx,0F04h ; Set shift count and mask
0000152 B801000000		426	mov ebx,1 ; Set initial size of ASCII
		427	; field for sign
0000157 887D18		428	mov edi,string_ptr ; Get address of start of
		429	; ASCII string
000015A 8CD8		430	mov ax,ds ; Copy ds to es
000015C 8EC0		431	mov es,ax
000015E FC		432	cld ; Set autoincrement mode
000015F B02B		433	mov al,'+' ; Clear sign field
0000161 F6C201		434	test dl,MINUS ; Look for negative value
0000164 7402		435	jz positive_result
		436	
0000166 B02D		437	mov al,'-'
0000168 AA		438	positive_result:
		439	stosb ; Bump string pointer
		440	; past sign
0000169 80E2FE		441	and dl,not MINUS ; Turn off sign bit
000016C 9B		442	fwait ; Wait for fbstp to finish
		443	;
		444	Register usage:
		445	ah: BCD byte value in use
		446	al: ASCII character value
		447	dx: Return value
		448	ch: BCD mask = 0fh
		449	cl: BCD shift count = 4
		450	bx: ASCII string field width
		451	esi: BCD field index
		452	di: ASCII string field pointer
		453	ds,es: ASCII string segment base
		454	
		455	Remove leading zeroes from the number.

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.).

LOC	OBJ	LINE	SOURCE
		456	;
00000160	8A6435F2	457	skip_leading_zeroes:
00000170	88E0	458	mov ah,bcd_byte[esi] ; Get BCD byte
00000173	D2E8	459	mov al,ah ; Copy value
00000175	240F	460	shr al,cl ; Get high order digit
00000177	7517	461	and al,0fh ; Set zero flag
		462	jnz enter_odd ; Exit loop if leading
		463	; non zero found
		464	
00000179	88E0	465	mov al,ah ; Get BCD byte again
0000017B	240F	466	and al,0fh ; Get low order digit
0000017D	7519	467	jnz enter_even ; Exit loop if non zero
		468	; digit found
		469	
0000017F	4E	470	dec esi ; Decrement BCD index
00000180	79EB	471	jns skip_leading_zeroes
		472	;
		473	The significand was all zeroes.
		474	;
00000182	B030	475	mov al,'0' ; Set initial zero
00000184	AA	476	stosb
00000185	43	477	inc ebx ; Bump string length
00000186	EB17	478	jmp short exit_with_value
		479	;
		480	Now expand the BCD string into digit
		481	per byte values 0-9.
		482	;
00000188		483	digit_loop:
00000188	8A6435F2	484	mov ah,bcd_byte[esi] ; Get BCD byte
0000018C	88E0	485	mov al,ah
0000018E	D2E8	486	shr al,cl ; Get high order digit
00000190		487	enter_odd:
00000190	0430	488	add al,'0' ; Convert to ASCII
00000190	AA	489	stosb ; Put digit into ASCII
		490	; string area
00000193	88E0	491	mov al,ah ; Get low order digit
00000195	240F	492	and al,0fh
00000197	43	493	inc ebx ; Bump field size counter
00000198		494	enter_even:
00000198	0430	495	add al,'0' ; Convert to ASCII
00000198	AA	496	stosb ; Put digit into ASCII area
00000198	43	497	inc ebx ; Bump field size counter
0000019C	4E	498	dec esi ; Go to next BCD byte
0000019D	79E9	499	jns digit_loop
		500	;
		501	Conversion complete. Set the string
		502	size and remainder.
		503	;
0000019F		504	exit_with_value:
0000019F	8B7D14	505	mov edi,size_ptr
000001A2	66891F	506	mov word ptr [edi],bx
000001A5	8BC2	507	mov eax,edx ; Set return value
000001A7	E980FFFF	508	jmp exit_proc
		509	
000001AC		510	floating_to_ascii endp
		511	
-----		512	code ends
		513	end

ASSEMBLY COMPLETE, NO WARNINGS, NO ERRORS.

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

XENIX286 80386 MACRO ASSEMBLER V1.0, ASSEMBLY OF MODULE GET_POWER_10
 OBJECT MODULE PLACED IN power10.obj
 ASSEMBLER INVOKED BY: asm386 power10.asm

LOC	OBJ	LINE	SOURCE
		1	+1 \$title(Calculate the value of 10**ax)
		2	; ;
		3	; This subroutine will calculate the
		4	; value of 10**eax. For values of
		5	; 0 <= eax < 19, the result will exact.
		6	; All 80386 registers are transparent
		7	; and the value is returned on the TOS
		8	; as two numbers, exponent in ST(1) and
		9	; fraction in ST(0). The exponent value
		10	; can be larger than the largest
		11	; exponent of an extended real format
		12	; number. Three stack entries are used.
		13	;
00000000		14	name get_power_10
		15	public get_power_10,power_table
		16	
-----		17	stack stackseg 8
-----		18	
		19	code segment public er
		20	;
		21	; Use exact values from 1.0 to 1e18.
		22	;
00000000	000000000000F03F	23	power_table even ; Optimize 16 bit access
00000008	0000000000002440	24	dq 1.0,1e1,1e2,1e3
00000010	0000000000005940		
00000018	0000000000040BF40		
00000020	000000000088C340	25	dq 1e4,1e5,1e6,1e7
00000028	000000000006AF840		
00000030	00000000080842E41		
00000038	00000000D0126341		
00000040	0000000084D79741	26	dq 1e8,1e9,1e10,1e11
00000048	0000000065CDCD41		
00000056	0000000205FA00242		
00000058	00000008764A83742		
00000060	000000A2941A6D42	27	dq 1e12,1e13,1e14,1e15
00000068	000040E59C30A242		
00000070	0000901EC4BCD642		
00000078	00003426F56B0C43		
00000080	0080E03779C34143	28	dq 1e16,1e17,1e18
00000088	00AD88557347643		
00000090	00CB4E676DC1AB43		
00000098	3D12000000	29	
00000099	770B	30	get_power_10 proc
0000009A		31	
		32	cmp eax,18 ; Test for 0 <= ax < 19
		33	ja out_of_range
		34	
0000009F	2EDD04C500000000	R	fld power_table[eax*8]; Get exact value
000000A7	D9F4	35	fxtract ; Separate power
		36	

Figure 7-6. Floating-Point to ASCII Conversion Routine (Cont'd.)

Shortness, speed, and accuracy were chosen rather than providing the maximum number of significant digits possible. An attempt is made to keep integers in their own domain to avoid unnecessary conversion errors.

Using the extended precision real number format, this routine achieves a worst case accuracy of three units in the 16th decimal position for a noninteger value or integers greater than 10^{18} . This is double precision accuracy. With values having decimal exponents less than 100 in magnitude, the accuracy is one unit in the 17th decimal position.

Higher precision can be achieved with greater care in programming, larger program size, and lower performance.

7.3.1 Function Partitioning

Three separate modules implement the conversion. Most of the work of the conversion is done in the module FLOATING_TO_ASCII. The other modules are provided separately, because they have a more general use. One of them, GET_POWER_10, is also used by the ASCII to floating-point conversion routine. The other small module, TOS_STATUS, identifies what, if anything, is in the top of the numeric register stack.

7.3.2 Exception Considerations

Care is taken inside the function to avoid generating exceptions. Any possible numeric value is accepted. The only possible exception is insufficient space on the numeric register stack.

The value passed in the numeric stack is checked for existence, type (NaN or infinity), and status (denormal, zero, sign). The string size is tested for a minimum and maximum value. If the top of the register stack is empty, or the string size is too small, the function returns with an error code.

Overflow and underflow is avoided inside the function for very large or very small numbers.

7.3.3 Special Instructions

The functions demonstrate the operation of several numeric instructions, different data types, and precision control. Shown are instructions for automatic conversion to BCD, calculating the value of 10 raised to an integer value, establishing and maintaining concurrency, data synchronization, and use of directed rounding on the NPX.

Without the extended precision data type and built-in exponential function, the double precision accuracy of this function could not be attained with the size and speed of the shown example.

The function relies on the numeric BCD data type for conversion from binary floating-point to decimal. It is not difficult to unpack the BCD digits into separate ASCII decimal digits. The major work involves scaling the floating-point value to the comparatively limited range of BCD values. To print a 9-digit result requires accurately scaling the given value to an

integer between 10^8 and 10^9 . For example, the number +0.123456789 requires a scaling factor of 10^9 to produce the value +123456789.0, which can be stored in 9 BCD digits. The scale factor must be an exact power of 10 to avoid changing any of the printed digit values.

These routines should exactly convert all values exactly representable in decimal in the field size given. Integer values that fit in the given string size are not be scaled, but directly stored into the BCD form. Noninteger values exactly representable in decimal within the string size limits are also exactly converted. For example, 0.125 is exactly representable in binary or decimal. To convert this floating-point value to decimal, the scaling factor is 1000, resulting in 125. When scaling a value, the function must keep track of where the decimal point lies in the final decimal value.

7.3.4 Description of Operation

Converting a floating-point number to decimal ASCII takes three major steps: identifying the magnitude of the number, scaling it for the BCD data type, and converting the BCD data type to a decimal ASCII string.

Identifying the magnitude of the result requires finding the value X such that the number is represented by $I \times 10^X$, where $1.0 \leq I < 10.0$. Scaling the number requires multiplying it by a scaling factor 10^S , so that the result is an integer requiring no more decimal digits than provided for in the ASCII string.

Once scaled, the numeric rounding modes and BCD conversion put the number in a form easy to convert to decimal ASCII by host software.

Implementing each of these three steps requires attention to detail. To begin with, not all floating-point values have a numeric meaning. Values such as infinity, indefinite, or NaN may be encountered by the conversion routine. The conversion routine should recognize these values and identify them uniquely.

Special cases of numeric values also exist. Denormals have numeric values, but should be recognized because they indicate that precision was lost during some earlier calculations.

Once it has been determined that the number has a numeric value, and it is normalized (setting appropriate denormal flags, if necessary, to indicate this to the calling program), the value must be scaled to the BCD range.

7.3.5 Scaling the Value

To scale the number, its magnitude must be determined. It is sufficient to calculate the magnitude to an accuracy of 1 unit, or within a factor of 10 of the required value. After scaling the number, a check is made to see if the result falls in the range expected. If not, the result can be adjusted one decimal order of magnitude up or down. The adjustment test after the scaling is necessary due to inevitable inaccuracies in the scaling value.

Because the magnitude estimate for the scale factor need only be close, a fast technique is used. The magnitude is estimated by multiplying the power of 2, the unbiased floating-point exponent, associated with the number by $\log_{10}2$. Rounding the result to an integer produces an estimate of sufficient accuracy. Ignoring the fraction value can introduce a maximum error of 0.32 in the result.

Using the magnitude of the value and size of the number string, the scaling factor can be calculated. Calculating the scaling factor is the most inaccurate operation of the conversion process. The relation $10^x = 2^{(x \cdot \log_2 10)}$ is used for this function. The exponentiate instruction F2XM1 is used.

Due to restrictions on the range of values allowed by the F2XM1 instruction, the power of 2 value is split into integer and fraction components. The relation $2^{(I + F)} = 2^I \times 2^F$ allows using the FSCALE instruction to recombine the 2^F value, calculated through F2XM1, and the 2^I part.

7.3.5.1 INACCURACY IN SCALING

The inaccuracy in calculating the scale factor arises because of the trailing zeros placed into the fraction value of the power of two when stripping off the integer valued bits. For each integer valued bit in the power of 2 value separated from the fraction bits, one bit of precision is lost in the fraction field due to the zero fill occurring in the least significant bits.

Up to 14 bits may be lost in the fraction because the largest allowed floating point exponent value is $2^{14} - 1$. These bits directly reduce the accuracy of the calculated scale factor, thereby reducing the accuracy of the scaled value. For numbers in the range of $10^{\pm 30}$, a maximum of 8 bits of precision are lost in the scaling process.

7.3.5.2 AVOIDING UNDERFLOW AND OVERFLOW

The fraction and exponent fields of the number are separated to avoid underflow and overflow in calculating the scaling values. For example, to scale 10^{-4932} to 10^8 requires a scaling factor of 10^{4950} , which cannot be represented by the NPX.

By separating the exponent and fraction, the scaling operation involves adding the exponents separate from multiplying the fractions. The exponent arithmetic involves small integers, all easily represented by the NPX.

7.3.5.3 FINAL ADJUSTMENTS

It is possible that the power function (Get_Power_10) could produce a scaling value such that it forms a scaled result larger than the ASCII field could allow. For example, scaling

$9.999999999999999 \times 10^{4900}$ by $1.00000000000000010 \times 10^{-4883}$ produces $1.00000000000000009 \times 10^{18}$. The scale factor is within the accuracy of the NPX and the result is within the conversion accuracy, but it cannot be represented in BCD format. This is why there is a post-scaling test on the magnitude of the result. The result can be multiplied or divided by 10, depending on whether the result was too small or too large, respectively.

7.3.6 Output Format

For maximum flexibility in output formats, the position of the decimal point is indicated by a binary integer called the power value. If the power value is zero, then the decimal point is assumed to be at the right of the rightmost digit. Power values greater than zero indicate how many trailing zeros are not shown. For each unit below zero, move the decimal point to the left in the string.

The last step of the conversion is storing the result in BCD and indicating where the decimal point lies. The BCD string is then unpacked into ASCII decimal characters. The ASCII sign is set corresponding to the sign of the original value.

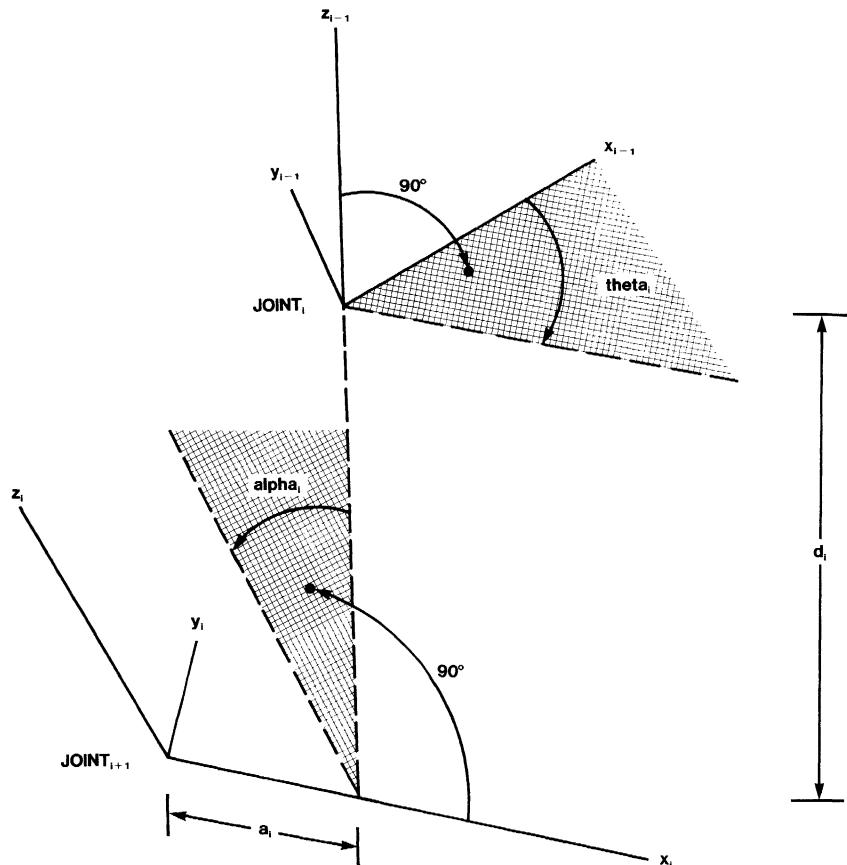
7.4 TRIGONOMETRIC CALCULATION EXAMPLES (NOT TESTED)

In this example, the kinematics of a robot arm is modeled with the 4×4 homogeneous transformation matrices proposed by Denavit and Hartenberg^{1,2}. The translational and rotational relationships between adjacent links are described with these matrices using the D-H matrix method. For each link, there is a 4×4 homogeneous transformation matrix that represents the link's coordinate system (\mathbf{L}_i) at the joint (\mathbf{J}_i) with respect to the previous link's coordinate system (\mathbf{J}_{i-1} , \mathbf{L}_{i-1}). The following four geometric quantities completely describe the motion of any rigid joint/link pair (\mathbf{J}_i , \mathbf{L}_i), as Figure 7-7 illustrates.

- θ_i = The angular displacement of the \mathbf{x}_i axis from the \mathbf{x}_{i-1} axis by rotating around the \mathbf{z}_{i-1} axis (anticlockwise).
- d_i = The distance from the origin of the $(i-1)^{th}$ coordinate system along the \mathbf{z}_{i-1} axis to the \mathbf{x}_i axis.
- a_i = The distance of the origin of the i^{th} coordinate system from the \mathbf{z}_{i-1} axis along the $-\mathbf{x}_i$ axis.
- α_i = The angular displacement of the \mathbf{z}_i axis from the \mathbf{z}_{i-1} about the \mathbf{x}_i axis (anticlockwise).

1. J. Denavit and R.S. Hartenberg, "A Kinematic Notation for Lower-Pair Mechanisms Based on Matrices," *J. Applied Mechanics*, June 1955, pp. 215-221.

2. C.S. George Lee, "Robot Arm Kinematics, Dynamics, and Control," *IEEE Computer*, Dec. 1982.



G40003

Figure 7-7. Relationships between Adjacent Joints

The D-H transformation matrix A_{i-1}^i for adjacent coordinate frames (from joint_{i-1} to joint_i) is calculated as follows:

$$A_{i-1}^i = T_{z,d} \times T_{z,\theta} \times T_{x,a} \times T_{x,\alpha}$$

...where...

$T_{z,d}$ represents a translation along the z_{i-1} axis

$T_{z,\theta}$ represents a rotation of angle θ about the z_{i-1} axis

$T_{x,a}$ represents a translation along the x_i axis

$T_{x,\alpha}$ represents a rotation of angle α about the x_i axis

$$A_{i-1}^i = \begin{vmatrix} \cos \theta_i & -\cos \alpha_i \sin \theta_i & \sin \alpha_i \sin \theta_i & \cos \theta_i \\ \sin \theta_i & \cos \alpha_i \cos \theta_i & -\sin \alpha_i \cos \theta_i & \sin \theta_i \\ 0 & \sin \alpha_i & \cos \alpha_i & d_i \\ 0 & 0 & 0 & 1 \end{vmatrix}$$

The composite homogeneous matrix T which represents the position and orientation of the joint/link pair with respect to the base system is obtained by successively multiplying the D-H transformation matrices for adjacent coordinate frames.

$$T_0^i = A_0^1 \times A_1^2 \times \dots \times A_{i-1}^i$$

This example in Figure 7-8 illustrates how the transformation process can be accomplished using the 80387. The program consists of two major procedures. The first procedure TRANS_PROC is used to calculate the elements in each D-H matrix, A_{i-1}^i . The second procedure MATRIXMUL_PROC finds the product of two successive D-H matrices.

```
XENIX286 80386 MACRO ASSEMBLER V1.0, ASSEMBLY OF MODULE TOS_STATUS
OBJECT MODULE PLACED IN tos.obj
ASSEMBLER INVOKED BY: asm386 tos.asm

LOC      OBJ          LINE      SOURCE
                                         1  +1 $title(Determine TOS register contents)
                                         2  ;
                                         3  ;      This subroutine will return a value
                                         4  ;      from 0-15 in eax corresponding
                                         5  ;      to the contents of NPX TOS. All
                                         6  ;      registers are transparent and no
                                         7  ;      errors are possible. The return
                                         8  ;      value corresponds to c3,c2,c1,c0
                                         9  ;      of FXAM instruction.
                                         10 ;
                                         11      name    tos_status
                                         12      public   tos_status
                                         13
-----                                         14      stack     stackseg    6
-----                                         15
-----                                         16      code      segment public er
                                         17
00000000                                         18      tos_status    proc
                                         19
                                         20      fxam           ; Get status of TOS register
00000002 9BDFE0                           21      fstsw  ax      ; Get current status
00000005 88E0                           22      mov    al,ah   ; Put bit 10-8 into bits 2-0
00000007 2507400000                      23      and   eax,4007h ; Mask out bits c3,c2,c1,c0
0000000C COEC03                           24      shr    ah,3    ; Put bit c3 into bit 11
0000000F 0BEO                           25      or    al,ah   ; Put c3 into bit 3
00000011 B400                           26      mov    ah,0    ; Clear return value
00000013 C3                           27      ret
                                         28
00000014                                         29      tos_status    endp
-----                                         30
-----                                         31      code      ends
                                         32      end

ASSEMBLY COMPLETE,  NO WARNINGS,  NO ERRORS.
```

Figure 7-8. Robot Arm Kinematics Example

LOC	OBJ	LINE	SOURCE
		37	
000000A9 C3		38	ret ; and fraction
		39	; OK to leave fextract running
		40	;
		41	Calculate the value using the
		42	exponentiate instruction. The following
		43	relations are used:
		44	$10^{**}x = 2^{**}(\log_2(10)*x)$
		45	$2^{**}(I+F) = 2^{**}I * 2^{**}F$
		46	if st(1) = I and st(0) = 2**F then
		47	fscale produces $2^{**}(I+F)$
000000AA		48	;
		49	out_of_range:
000000AA D9E9		50	fldl2t ; TOS = LOG2(10)
000000AC C8040000		51	enter 4,0
		52	;
000000B0 8945FC		53	; save power of 10 value, P
		54	mov [ebp-4],eax
		55	;
000000B3 DA4DFC		56	; TOS,X = LOG2(10)*P = LOG2($10^{**}P$)
000000B6 D9E8		57	fimul dword ptr [ebp-4]
000000B8 D9E0		58	fld1 ; Set TOS = -1.0
000000BA D9C1		59	fchs
		60	fld st(1) ; Copy power value
		61	; in base two
000000BC D9FC		62	frndint ; TOS = I: -inf < I <= X
		63	; where I is an integer
		64	; Rounding mode does
		65	; not matter
000000BE D9CA		66	fxch st(2) ; TOS = X, ST(1) = -1.0
		67	; ST(2) = I
000000C0 D8E2		68	fsub st,st(2) ; TOS,F = X-I:
		69	; -1.0 < TOS <= 1.0
		70	;
		71	; Restore original rounding control
000000C2 58		72	pop eax
000000C3 D9F0		73	f2xm1 ; TOS = $2^{**}(F) - 1.0$
000000C5 C9		74	leave ; Restore stack
000000C6 DEE1		75	fsubr ; Form $2^{**}(F)$
000000C8 C3		76	ret ; OK to leave fsnbr running
		77	;
000000C9		78	get_power_10 endp
-----		79	
		80	code ends
		81	end

ASSEMBLY COMPLETE, NO WARNINGS, NO ERRORS.

Figure 7-8. Robot Arm Kinematics Example (Cont'd.).

```

XENIX286 80386 MACRO ASSEMBLER V1.0, ASSEMBLY OF MODULE ROT_MATRIX_CAL
OBJECT MODULE PLACED IN transx.obj
ASSEMBLER INVOKED BY: asm386 transx.asm

LOC      OBJ          LINE      SOURCE
1        Name ROT_MATRIX_CAL
2
3
4
5      ; This example illustrates the use
6      ; of the 80387 floating point
7      ; instructions, in particular, the
8      ; FSINCOS function which gives both
9      ; the SIN and COS values.
10     ; The program calculates the
11     ; composite matrix for base to end-
12     ; effector transformation.
13     ;
14     ; Only the kinematics is considered in
15     ; this example.
16     ;
17     ; If the composite matrix mentioned above
18     ; is given by:
19     ;  $T_{in} = A_1 \times A_2 \times \dots \times A_n$ 
20     ;  $T_{in}$  is found by successively calling
21     ; trans_proc and matrixmul_proc until
22     ; all matrices have been exhausted.
23     ;
24     ; trans_proc calculates entries in each
25     ;  $A(A_1, \dots, A_n)$  while matrixmul_proc
26     ; performs the matrix multiplication for
27     ;  $A_i$  and  $A_{i+1}$ . matrixmul_proc in turn
28     ; calls matrix_row and matrix_elem to
29     ; do the multiplication.
30
31     ; Define stack space
32
33     .....
34     trans_stack stackseg 400
35
36     ; Define the matrix structure for
37     ; 4x4 transformational matrices
38
39     a_matrix struc
40         a11    dq    ?
41         a12    dq    ?
42         a13    dq    ?
43         a14    dq    ?
44         a21    dq    ?
45         a22    dq    ?
46         a23    dq    ?
47         a24    dq    ?
48         a31    dq    0h
49         a32    dq    ?
50         a33    dq    ?
51         a34    dq    ?
52         a41    dq    0h
53         a42    dq    0h
54         a43    dq    0h
55         a44    dq    1h
56     a_matrix ends
57

```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.)

```
58 ; Assume One joint in the storage
59 ; allocation and hence for
60 ; two sets of parameters; however,
61 ; more joints are possible
62 ;
63 alp_deg struc
64     alpha_deg1 dd ?
65     alpha_deg2 dd ?
66 alp_deg ends
67
68 tht_deg struc
69     theta_deg1 dd ?
70     theta_deg2 dd ?
71 tht_deg ends
72
73 A_array struc
74     A1      dq ?
75     A2      dq ?
76 A_array ends
77
78 D_array struc
79     D1      dq ?
80     D2      dq ?
81 D_array ends
82
83 ; trans_data is the data segment
84 ;
85
86 trans_data    segment rw public
87
88         Amx          a_matrix<>
89         Bmx          a_matrix<>
00000000 ??????????????
00000008 ??????????????
00000010 ??????????????
00000018 ??????????????
00000020 ??????????????
00000028 ??????????????
00000030 ??????????????
00000038 ??????????????
00000040 0000000000000000
00000048 ??????????????
00000050 ??????????????
00000058 ??????????????
00000060 0000000000000000
00000068 0000000000000000
00000070 0000000000000000
00000078 0100000000000000
00000080 ??????????????
00000088 ??????????????
00000090 ??????????????
00000098 ??????????????
000000A0 ??????????????
000000A8 ??????????????
000000B0 ??????????????
000000B8 ??????????????
000000C0 0000000000000000
000000C8 ??????????????
000000D0 ??????????????
000000D8 ??????????????
000000E0 0000000000000000
000000E8 0000000000000000
000000F0 0000000000000000
000000F8 0100000000000000
```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.).

Figure 7-8. Robot Arm Kinematics Example (Cont'd.)

```

135      ; theta(radians) in ST and
136      ; alpha(radians) in ST(1)
137
138      ; Calculate matrix elements
139      ; a11 = cos theta
140      ; a12 = -cos alpha * sin theta
141      ; a13 = sin alpha * sin theta
142      ; a14 = A * cos theta
143      ; a21 = sin theta
144      ; a22 = cos alpha * cos theta
145      ; a23 = -sin alpha * cos theta
146      ; a24 = A * sin theta
147      ; a32 = sin alpha
148      ; a33 = cos alpha
149      ; a34 = D
150      ; a31 = a41 = a42 = a43 = 0.0
151      ; a44 = 1
152
153      ; ebx contains the offset for the matrix
154
0000001A D9FB          155      fsincos           ;cos theta in ST
156                          ;sin theta in ST(1)
0000001C D9C0          157      fld    st             ;duplicate cos theta
0000001E DD13          158      fst    [ebx].a11 ;cos theta in a11
00000020 DC0CCD90010000   R     159      fmul   qword ptr A_VECTOR[ecx*8]
00000027 DD5B18          160      fstp   [ebx].a14 ;A * cos thetan a14
0000002A D9C9          161      fxch   st(1)        ;sin theta in ST
0000002C DD5320          162      fst    [ebx].a21 ;sin theta in a21
0000002F D9C0          163      fld    st             ;duplicate sin theta
00000031 DC0CCD90010000   R     164      fmul   qword ptr A_VECTOR[ecx*8]
00000038 DD5B38          165      fstp   [ebx].a24 ;A * sin theta in a24
0000003B D9C2          166      fld    st(2)        ;alpha in ST
0000003D D9FB          167      fsincos           ;cos alpha in ST
168                          ;sin alpha in ST(1)
169                          ;sin theta in ST(2)
170                          ;cos theta in ST(3)
0000003F DD5350          171      fst    [ebx].a33 ;cos alpha in a33
00000042 D9C9          172      fxch   st(1)        ;sin alpha in ST
00000044 DD5348          173      fst    [ebx].a32 ;sin alpha in a32
00000047 D9C2          174      fld    ST(2)        ;sin theta in ST
175                          ;sin alpha in ST(1)
00000049 D8C9          176      fmul   st,st(1) ;sin alpha * sin theta
0000004B DD5B10          177      fstp   [ebx].a13 ;stored in a13
0000004E D8CB          178      fmul   st,st(3) ;cos theta * sin alpha
00000050 D9E0          179      fchs   st,st(3) ;-cos theta * sin alpha
00000052 DD5B30          180      fstp   [ebx].a23 ;stored in a23
00000055 D9C2          181      fld    st(2)        ;cos theta in ST
182                          ;cos alpha in ST(1)
183                          ;sin theta in ST(2)
184                          ;cos theta in ST(3)
00000057 D8C9          185      fmul   st,st(1) ;cos theta * cos alpha
00000059 DD5B28          186      fstp   [ebx].a22 ;stored in a22
0000005C D8C9          187      fmul   st,st(1) ;cos alpha * sin theta
188
189      ; To take advantage of parallel operations
190      ; between the CPU and NPX
191
0000005E 50          192      push   eax ; save eax
193
194      ; also move D into a34 in a faster way
0000005F 8804CDA0010000   R     195      mov    eax, dword ptr D_VECTOR[ecx*8]
00000066 894358          196      mov    dword ptr [ebx + 88], eax

```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.)

```

00000069 8B04CDA4010000      R   197      mov    eax, dword ptr D_VECTOR[ecx*8 + 4]
00000070 89435C                198      mov    dword ptr [ebx + 92], eax
00000073 58                  199      pop    eax ; restore eax
00000074 D9E0                200      fchs   ;-cos alpha * sin theta
00000076 DD5B08                201      fstp   [ebx].a12 ;stored in a12
                                         202      ;and all nonzero elements
                                         203      ;have been calculated
00000079 CB                  204      ret
                                         205
0000007A                      206      trans_proc endp
                                         207
                                         208
0000007A                      209      matrix_elem proc far
                                         210
                                         ; This procedure calculate the dot product
                                         ; of the ith row of the first matrix and
                                         ; the jth column of the second matrix:
                                         211
                                         ;
                                         212      ; Tij Where Tij = sum of Aik x Bkj over k
                                         ;
                                         213      ; parameters passed from the calling routine,
                                         ; matrix_ROW;
                                         214      ; ESI = (i-1)*8
                                         ; EDI = (j-1)*8
                                         ; local register, EBP = (k-1)*8
                                         ;
                                         215      push   ebp     ; save ebp
                                         216      push   ecx     ; ecx to be used as a tmp reg
                                         217      mov    ecx, esi; save it for later indexing
                                         218
                                         ; locating the element in the first matrix, A
                                         219      imul  ecx, NUM_COL ; ecx contains offset due
                                         220                          ; to preceding rows; the
                                         221                          ; offset is from the
                                         222                          ; beginning of the matrix
                                         ;
                                         223      xor    ebp, ebp ; clear ebp, which will be
                                         224                          ; used a temp reg to index( k)
                                         225                          ; across the ith row of the first
                                         226                          ; matrix as well as down the jth
                                         227                          ; column of the second matrix
                                         ;
                                         228      xor    ebx, ebx ; clear ebx for accumulating Aik*Bkj
                                         229      mov    dword ptr [ecx+edi],ebp
                                         230      mov    dword ptr [ecx][edi+4], ebp
                                         ;
                                         231      push   ecx     ; save on stack: esi * num.col =
                                         232                          ; the offset of the beginning
                                         233                          ; of the ith row from the
                                         234                          ; beginning of the A matrix
                                         ;
                                         235      add    ecx, ebp ; get to the kth column entry
                                         236                          ; of the ith row of the A matrix
                                         ;
                                         237      add    ebx, ebp ; load Aik into 80387
                                         238      fld    qword ptr [eax][ecx]
                                         ;
                                         239      add    ebx, ebp ; locating Bkj
                                         240      imul  ecx, NUM_ROW ; ecx contains the offset
                                         241                          ; of the beginning of the
                                         242                          ; kth row from the
                                         ;
                                         243      add    ebx, ebp ; get to the kth column entry
                                         244                          ; of the ith row of the A matrix
                                         245
                                         246      add    ebx, ebp ; load Bkj into 80387
                                         247
                                         248      NXT_k:
                                         249      add    ecx, ebp ; get to the kth column entry
                                         250                          ; of the ith row of the A matrix
                                         ;
                                         251      add    ebx, ebp ; load Aik into 80387
                                         252      fld    qword ptr [eax][ecx]
                                         ;
                                         253      add    ebx, ebp ; locating Bkj
                                         254
                                         255      add    ebx, ebp ; get to the kth column entry
                                         256      imul  ecx, NUM_ROW ; ecx contains the offset
                                         257                          ; of the beginning of the
                                         258                          ; kth row from the
                                         259

```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.)

```

00000095 01F9
entry
00000097 DC0C0B
0000009A 59
0000009B 51
0000009C 01F9
0000009E DC040A
000000A1 DD1COA
000000A4 83C508
000000A7 83FD20
000000AA 7CDF
000000AC 59
000000AD 59
000000AE 5D
000000AF CB
000000B0
000000B0
000000B0 31FF
000000B2 9A7A000000---- R
000000B9 83C708
000000BC 83FF20
000000BF 7CF1
000000C1 CB
000000C2
000000C2
000000C2 proc far
matrix_elem endp
matrix_row proc far
    xor     edi, edi
    ; scan across a row
NXT_COL:
    call    matrix_elem
    add    edi, 8
    cmp    edi, NUM_COL*8
    jl     NXT_COL
    ret
matrix_row endp
matrixmul_proc proc far
    ; This procedure does the matrix
    ; multiplication by calling matrix_row
    ; to calculate entries in each row
    ;
    ; The matrix multiplication is
    ; performed in the following manner,
    ;  $T_{ij} = A_{ik} \times B_{kj}$ 
    ; where i and j denote the row and column
    ; respectively and k is the index for
    ; scanning across the ith row of the
    ; first matrix and the jth column of the
    ; second matrix.
320

```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.).

```

000000C2 5A          321      pop    edx ; offset Tmx in edx
000000C3 5B          322      pop    ebx ; offset Bmx in ebx
000000C4 5B          323      pop    eax ; offset Amx in eax
000000C5 31F6        324
000000C7             325      ; setup esi and edi
000000C7 9AB0000000---- R 326      ; edi points to the column
000000CE 83C60B       327      ; esi points to the row
000000D1 83FE20
000000D4 7CF1
000000D6 C8
000000D7             328
-----             329      xor    esi, esi ; clear esi
000000C7             330
000000C7 9AB0000000---- R 331      NXT_ROW:
000000CE 83C60B       332      call   matrix_row
000000D1 83FE20       333      add    esi, 8
000000D4 7CF1         334      cmp    esi, NUM_ROW*8
000000D6 C8           335      jl    NXT_ROW
000000D7             336      ret
000000D7             337
-----             338      matrixmul_proc endp
000000D7             339
000000D7             340
-----             341      trans_code ends
000000D7             342
000000D7             343      ;*****
000000D7             344      ;
000000D7             345      ;
000000D7             346      ;
000000D7             347      ; Main program
000000D7             348      ;
000000D7             349      ;
000000D7             350      ;
000000D7             351      ;*****
000000D7             352
-----             353      main_code segment er
000000D7             354
000000D7             355      START:
000000D7             356
000000D7  BC00000000 R 357      mov    esp, stackstart trans_stack
000000D7             358      ; save all registers
000000D7             359
000000D7  60           R 360      pushad
000000D7             361
000000D7             362      ; ECX denotes the number of joints
000000D7             363      ; where no of matrices = NUM_JOINT + 1
000000D7             364      ; Find the first matrix( from the base
000000D7             365      ; of the system to the first joint)
000000D7             366      ; and call it Bmx
000000D7             367      xor    ecx, ecx ; 1st matrix
000000D7  31C9          R 368      mov    ebx, offset Bmx ;
000000D7  BB80000000---- R 369      call   trans_proc ; is Bmx
000000D7  9A00000000---- R 370      inc    ecx
000000D7             371
000000D7             372      NXT_MATRIX:
000000D7             373      ; From the 2nd matrix and on, it
000000D7             374      ; will be stored in Amx.
000000D7             375      ; The result from the first matrix mult.
000000D7             376      ; is stored in Tmx but will be accessed
000000D7             377      ; as Bmx in the next multiplication.
000000D7             378      ; As a matter of fact, the roles of Bmx
000000D7             379      ; and Tmx alternate in successive
000000D7             380      ; multiplications. This is achieved by
000000D7             381      ; reversing the order of the Bmx and Tmx
000000D7             382      ; pointers being passed onto the program

```

Figure 7-8. Robot Arm Kinematics Example (Cont'd.).

```

00000015 BB00000000 R 383 ; stack. Thus, this is invisible to the
0000001A 9A00000000---- R 384 ; matrix multiplication procedure.
00000021 41 R 385 ; REVERSE serves as the indicator;
00000022 80358801000001 R 386 ; REVERSE = 0 means that the result
00000029 7511 R 387 ; is to placed in Tmx.
00000029 7511 R 388

00000015 BB00000000 R 389 mov ebx, offset Amx ; find Amx
0000001A 9A00000000---- R 390 call trans_proc
00000021 41 R 391 inc ecx
00000022 80358801000001 R 392 xor REVERSE, 1h
00000029 7511 R 393 jnz Bmx_as_Tmx
00000029 7511 R 394

0000002B 6800000000 R 395 ; no reversing. Bmx as the second input
00000030 6880000000 R 396 ; matrix while Tmx as the output matrix.
00000035 6800010000 R 397 push offset Amx
0000003A EBOF R 398 push offset Bmx
0000003A EBOF R 399 push offset Tmx
0000003A EBOF R 400 jmp CONTINUE
0000003C 6800000000 R 401
00000041 6800010000 R 402 ; reversing. Tmx as the second input
00000046 6880000000 R 403 ; matrix while Bmx as the output matrix.
00000046 6880000000 R 404 Bmx_as_Tmx:
00000046 6880000000 R 405 push offset Amx
00000041 6800010000 R 406 push offset Tmx ;reversing the
00000046 6880000000 R 407 push offset Bmx ;pointers passed
00000048 R 408

00000048 9AC2000000---- R 409 CONTINUE:
00000052 83F901 R 410 call matrixmul_proc
00000055 7E8E R 411 cmp ecx, NUM_JOINT
00000057 61 R 412 jle NXT_MATRIX
00000057 61 R 413

00000057 61 R 414 ; if REVERSE = 1 then the final answer
00000057 61 R 415 ; will be in Bmx otherwise, in Tmx.
00000057 61 R 416

00000057 61 R 417 popad
00000057 61 R 418
00000057 61 R 419 main_code ends
00000057 61 R 420
00000057 61 R 421 end START, ds:trans_data, ss:trans_stack

```

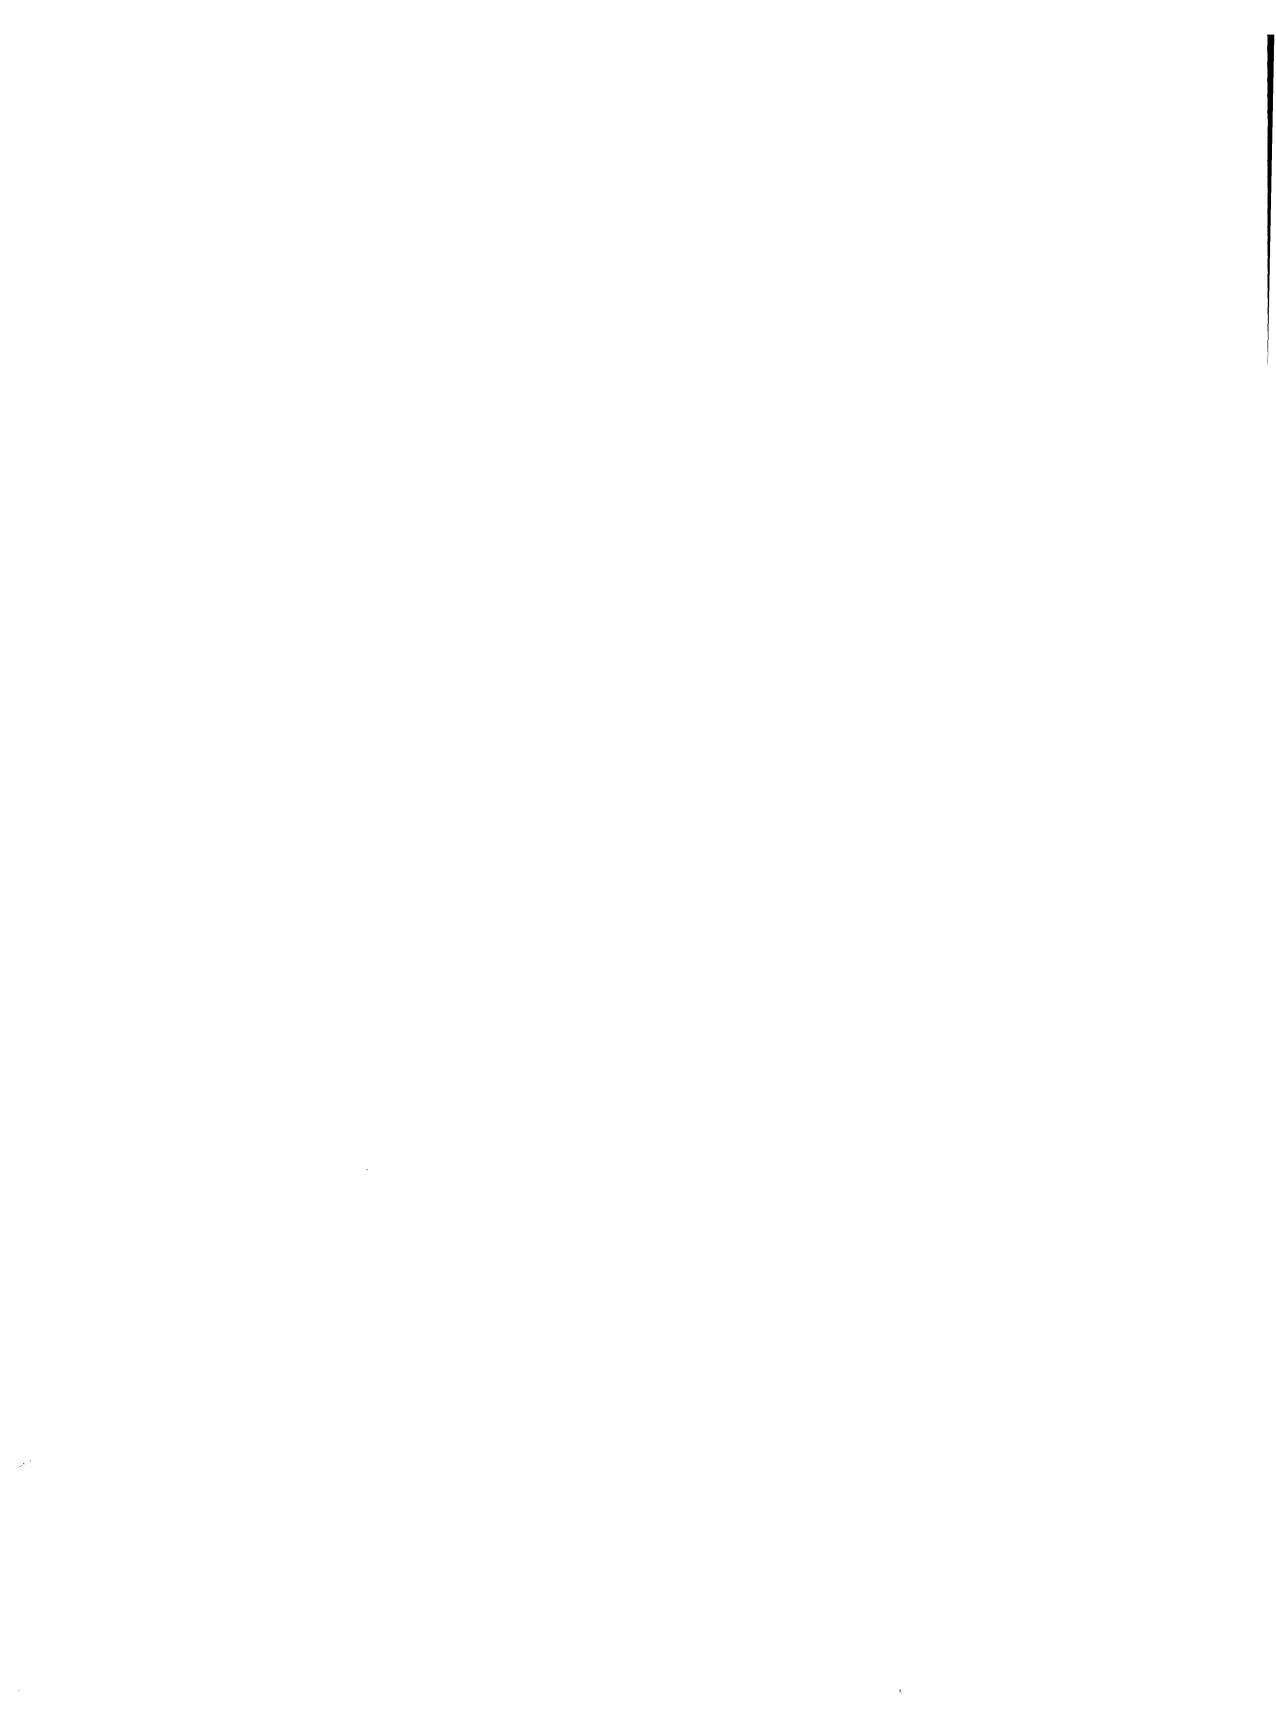
ASSEMBLY COMPLETE, NO WARNINGS, NO ERRORS.

Figure 7-8. Robot Arm Kinematics Example (Cont'd.)



Machine Instruction Encoding and Decoding

A



APPENDIX A

MACHINE INSTRUCTION ENCODING AND DECODING

1st Byte		2nd Byte	Bytes 3-7	ASM386 Instruction Format	
Hex	Binary				
D8	1101 1000	MOD 000 R/M	SIB, displ	FADD	single-real
D8	1101 1000	MOD 001 R/M	SIB, displ	FMUL	single-real
D8	1101 1000	MOD 010 R/M	SIB, displ	FCOM	single-real
D8	1101 1000	MOD 011 R/M	SIB, displ	FCOMP	single-real
D8	1101 1000	MOD 100 R/M	SIB, displ	FSUB	single-real
D8	1101 1000	MOD 101 R/M	SIB, displ	FSUBR	single-real
D8	1101 1000	MOD 110 R/M	SIB, displ	FDIV	single-real
D8	1101 1000	MOD 111 R/M	SIB, displ	FDIVR	single-real
D8	1101 1000	1100 0 REG		FADD	ST,ST(i)
D8	1101 1000	1100 1 REG		FMUL	ST,ST(i)
D8	1101 1000	1101 0 REG		FCOM	ST(i)
D8	1101 1000	1101 1 REG		FCOMP	ST(i)
D8	1101 1000	1110 0 REG		FSUB	ST,ST(i)
D8	1101 1000	1110 1 REG		FSUBR	ST,ST(i)
D8	1101 1000	1111 0 REG		FDIV	ST,ST(i)
D8	1101 1000	1111 1 REG		FDIVR	ST,ST(i)
D9	1101 1001	MOD 000 R/M	SIB, displ	FLD	single-real
D9	1101 1001	MOD 001 R/M		reserved	
D9	1101 1001	MOD 010 R/M	SIB, displ	FST	single-real
D9	1101 1001	MOD 011 R/M	SIB, displ	FSTP	single-real
D9	1101 1001	MOD 100 R/M	SIB, displ	FLDENV	14 or 28 bytes***
D9	1101 1001	MOD 101 R/M	SIB, displ	FLDCW	2 bytes
D9	1101 1001	MOD 110 R/M	SIB, displ	FSTENV	14 or 28 bytes***
D9	1101 1001	MOD 111 R/M	SIB, displ	FSTCW	2 bytes
D9	1101 1001	1100 0 REG		FLD	ST(i)
D9	1101 1001	1100 1 REG		FXCH	ST(i)
D9	1101 1001	1101 0000		FNOP	
D9	1101 1001	1101 0001		reserved	
D9	1101 1001	1101 001-		reserved	
D9	1101 1001	1101 01-		reserved	
D9	1101 1001	1101 1 REG		reserved	
D9	1101 1001	1110 0000		FCHS	
D9	1101 1001	1110 0001		FABS	
D9	1101 1001	1110 001-		reserved	
D9	1101 1001	1110 0100		FTST	
D9	1101 1001	1110 0101		FXAM	
D9	1101 1001	1110 011-		reserved	
D9	1101 1001	1110 1000		FLD1	
D9	1101 1001	1110 1001		FLDL2T	
D9	1101 1001	1110 1010		FLDL2E	
D9	1101 1001	1110 1011		FLDP1	
D9	1101 1001	1110 1100		FLDLG2	
D9	1101 1001	1110 1101		FLDLN2	
D9	1101 1001	1110 1110		FLDZ	
D9	1101 1001	1110 1111		reserved	
D9	1101 1001	1111 0000		F2XM1	
D9	1101 1001	1111 0001		FYL2X	
D9	1101 1001	1111 0010		FPTAN	

1st Byte		2nd Byte	Bytes 3-7	ASM386 Instruction Format	
Hex	Binary				
D9	1101 1001	1111 0011		FPATAN	
D9	1101 1001	1111 0100		FXTRACT	
D9	1101 1001	1111 0101		FPREM1	
D9	1101 1001	1111 0110		FDECSTP	
D9	1101 1001	1111 0111		FINCSTP	
D9	1101 1001	1111 1000		FPREM	
D9	1101 1001	1111 1001		FYL2XP1	
D9	1101 1001	1111 1010		FSQRT	
D9	1101 1001	1111 1011		FSIN COS	
D9	1101 1001	1111 1100		FRNDINT	
D9	1101 1001	1111 1101		FSCALE	
D9	1101 1001	1111 1110		FSIN	
D9	1101 1001	1111 1111		FCOS	
DA	1101 1010	MOD 000 R/M	SIB, displ	FIADD	short-integer
DA	1101 1010	MOD 001 R/M	SIB, displ	FIMUL	short-integer
DA	1101 1010	MOD 010 R/M	SIB, displ	FICOM	short-integer
DA	1101 1010	MOD 011 R/M	SIB, displ	FICOMP	short-integer
DA	1101 1010	MOD 100 R/M	SIB, displ	FISUB	short-integer
DA	1101 1010	MOD 101 R/M	SIB, displ	FISUBR	short-integer
DA	1101 1010	MOD 110 R/M	SIB, displ	FIDIV	short-integer
DA	1101 1010	MOD 111 R/M	SIB, displ	FIDIVR	short-integer
DA	1101 1010	110- ----		reserved	
DA	1101 1010	1110 0---		reserved	
DA	1101 1010	1110 1000		reserved	
DA	1010 1010	1110 1001		FUCOMPP	
DA	1101 1010	1110 101-		reserved	
DA	1101 1010	1110 11--		reserved	
DA	1101 1010	1111 ----		reserved	
DB	1101 1011	MOD 000 R/M	SIB, displ	FILD	short-integer
DB	1101 1011	MOD 001 R/M	SIB, displ	reserved	
DB	1101 1011	MOD 010 R/M	SIB, displ	FIST	short-integer
DB	1101 1011	MOD 011 R/M	SIB, displ	FISTP	short-integer
DB	1101 1011	MOD 100 R/M	SIB, displ	reserved	
DB	1101 1011	MOD 101 R/M	SIB, displ	FLD	extended-real
DB	1101 1011	MOD 110 R/M	SIB, displ	reserved	
DB	1101 1011	MOD 111 R/M	SIB, displ	FSTP	extended-real
DB	1101 1011	110- ----		reserved	
DB	1101 1011	1110 0000		**(1)	
DB	1101 1011	1110 0001		**(2)	
DB	1101 1011	1110 0010		FCLEX	
DB	1101 1011	1110 0011		FINIT	
DB	1101 1011	1110 0100		**(3)	
DB	1101 1011	1110 0101		reserved	
DB	1101 1011	1110 011-		reserved	
DB	1101 1011	1110 1---		reserved	
DB	1101 1011	1111 ----		reserved	
DC	1101 1100	MOD 000 R/M	SIB, displ	FADD	double-real
DC	1101 1100	MOD 001 R/M	SIB, displ	FMUL	double-real
DC	1101 1100	MOD 010 R/M	SIB, displ	FCOM	double-real
DC	1101 1100	MOD 011 R/M	SIB, displ	FCOMP	double-real
DC	1101 1100	MOD 100 R/M	SIB, displ	FSUB	double-real
DC	1101 1100	MOD 101 R/M	SIB, displ	FSUBR	double-real
DC	1101 1100	MOD 110 R/M	SIB, displ	FDIV	double-real
DC	1101 1100	MOD 111 R/M	SIB, displ	FDIVR	double-real
DC	1101 1100	1100 0 REG		FADD	ST(i),ST

1st Byte		2nd Byte	Bytes 3-7	ASM386 Instruction Format	
Hex	Binary				
DC	1101 1100	1100 1 REG		FMUL	ST(i),ST
DC	1101 1100	1101 0 REG		reserved	
DC	1101 100	1101 1 REG		reserved	
DC	1101 1100	1110 0 REG		FSUBR	ST(i),ST
DC	1101 1100	1110 1 REG		FSUB	ST(i),ST
DC	1101 1100	1111 0 REG		FDIVR	ST(i),ST
DC	1101 1100	1111 1 REG		FDIV	ST(i),ST
DD	1101 1101	MOD 000 R/M	SIB, displ	FLD	double-real
DD	1101 1101	MOD 001 R/M		reserved	
DD	1101 1101	MOD 010 R/M	SIB, displ	FST	double-real
DD	1101 1101	MOD 011 R/M	SIB, displ	FSTP	double-real
DD	1101 1101	MOD 100 R/M	SIB, displ	FRSTOR	94 or 108 bytes***
DD	1101 1101	MOD 101 R/M	SIB, displ	reserved	
DD	1101 1101	MOD 110 R/M	SIB, displ	FSAVE	94 or 108 bytes***
DD	1101 1101	MOD 111 R/M	SIB, displ	FSTSW	2 bytes
DD	1101 1101	1100 0 REG		FFREE	ST(i)
DD	1101 1101	1100 1 REG		reserved	
DD	1101 1101	1101 0 REG		FST	ST(i)
DD	1101 1101	1101 1 REG		FSTP	ST(i)
DD	1101 1101	1110 0 REG		FUCOM	ST(i)
DD	1101 1101	1110 1 REG		FUCOMP	ST(i)
DD	1101 1101	1111 ----		reserved	
DE	1101 1110	MOD 000 R/M	SIB, displ	FIADD	word-integer
DE	1101 1110	MOD 001 R/M	SIB, displ	FIMUL	word-integer
DE	1101 1110	MOD 010 R/M	SIB, displ	FICOM	word-integer
DE	1101 1110	MOD 011 R/M	SIB, displ	FICOMP	word-integer
DE	1101 1110	MOD 100 R/M	SIB, displ	FISUB	word-integer
DE	1101 1110	MOD 101 R/M	SIB, displ	FISUBR	word-integer
DE	1101 1110	MOD 110 R/M	SIB, displ	FIDIV	word-integer
DE	1101 1110	MOD 111 R/M	SIB, displ	FIDIVR	word-integer
DE	1101 1110	1100 0 REG		FADDP	ST(i),ST
DE	1101 1110	1100 1 REG		FMULP	ST(i),ST
DE	1101 1110	1101 0 ---		reserved	
DE	1101 1110	1101 1000		reserved	
DE	1101 1110	1101 1001		FCOMPP	
DE	1101 1110	1101 101-		reserved	
DE	1101 1110	1101 11--		reserved	
DE	1101 1110	1110 0 REG		FSUBRP	ST(i),ST
DE	1101 1110	1110 1 REG		FSUBP	ST(i),ST
DE	1101 1110	1111 0 REG		FDIVRP	ST(i),ST
DE	1101 1110	1111 1 REG		FDIVP	ST(i),ST
DF	1101 1111	MOD 000 R/M	SIB, displ	FILD	word-integer
DF	1101 1111	MOD 001 R/M	SIB, displ	reserved	
DF	1101 1111	MOD 010 R/M	SIB, displ	FIST	word-integer
DF	1101 1111	MOD 011 R/M	SIB, displ	FISTP	word-integer
DF	1101 1111	MOD 100 R/M	SIB, displ	FBLD	packed-decimal
DF	1101 1111	MOD 101 R/M	SIB, displ	FILD	long-integer
DF	1101 1111	MOD 110 R/M	SIB, displ	FBSTP	packed-decimal
DF	1101 1111	MOD 111 R/M	SIB, displ	FISTP	long-integer
DF	1101 1111	1100 0 REG		reserved	
DF	1101 1111	1100 1 REG		reserved	
DF	1101 1111	1101 0 REG		reserved	
DF	1101 1111	1101 1 REG		reserved	
DF	1101 1111	1110 0000		FSTSW AX	
DF	1101 1111	1110 0001		reserved	

1st Byte		2nd Byte	Bytes 3-7	ASM386 Instruction Format
Hex	Binary			
DF	1101 1111	1110 001-		reserved
DF	1101 1111	1110 01--		reserved
DF	1101 1111	1110 1---		reserved
DF	1101 1111	1111 ----		reserved

** The marked encodings can be generated by the language translators; however, the 80387 treats them as FNOP. They correspond to the following 8087 or 80287 instructions.

- (1) FENI
- (2) FDISI
- (3) FSETPM

*** The size of operand transferred depends on the 80386 operand-size attribute in effect for the instruction.

Exception Summary

B



APPENDIX B

EXCEPTION SUMMARY

The following table lists the instruction mnemonics in alphabetical order. For each mnemonic, it summarizes the exceptions that the instruction may cause. When writing 80387 programs that may be used in an environment that employs numerics exception handlers, assembly-language programmers should be aware of the possible exceptions for each instruction in order to determine the need for exception synchronization. Chapter 4 explains the need for exception synchronization.

Mnemonic	Instruction	IS	I	D	Z	O	U	P
F2XM1	$2^x - 1$	Y	Y	Y			Y	Y
FABS	Absolute value	Y						
FADD(P)	Add real	Y	Y	Y		Y	Y	Y
FBLD	BCD load	Y						
FBSTP	BCD store and pop	Y	Y					Y
FCHS	Change sign	Y						
FCLEX	Clear exceptions							
FCOM(P)(P)	Compare real	Y	Y	Y				
FCOS	Cosine	Y	Y	Y			Y	Y
FDECSTP	Decrement stack pointer							
FDIV(R)(P)	Divide real	Y	Y	Y	Y	Y	Y	Y
FFREE	Free register							
FIADD	Integer add	Y	Y	Y		Y	Y	Y
FICOM(P)	Integer compare	Y	Y	Y				
FIDIV	Integer divide	Y	Y	Y	Y		Y	Y
FIDIVR	Integer divide reversed	Y	Y	Y	Y	Y	Y	Y
FILD	Integer load	Y						
FIMUL	Integer multiply	Y	Y	Y		Y	Y	Y
FINCSTP	Increment stack pointer							
FINIT	Initialize processor							
FIST(P)	Integer store	Y	Y					Y
FISUB(R)	Integer subtract	Y	Y	Y		Y	Y	Y
FLD extended or stack	Load real	Y						
FLD single or double	Load real	Y	Y	Y				
FLD1	Load + 1.0	Y						
FLDCW	Load Control word	Y	Y	Y	Y	Y	Y	Y
FLDENV	Load environment	Y	Y	Y	Y	Y	Y	Y
FLDL2E	Load log _e	Y						
FLDL2T	Load log ₂ 10	Y						
FLDLG2	Load log ₁₀ 2	Y						
FLDLN2	Load log _e 2	Y						
FLDPI	Load π	Y						

IS—Invalid operand due to stack overflow/underflow

I—Invalid operand due to other cause

D—Denormal operand

Z—Zero-divide

O—Overflow

U—Underflow

P—Inexact result (precision)

Mnemonic	Instruction	IS	I	D	Z	O	U	P
FLDZ	Load + 0.0	Y						
FMUL(P)	Multiply real	Y	Y	Y		Y	Y	Y
FNOP	No operation							
FPATAN	Partial arctangent	Y	Y	Y			Y	Y
FPREM	Partial remainder	Y	Y	Y			Y	Y
FPREM1	IEEE partial remainder	Y	Y	Y			YY	
FPTAN	Partial tangent	Y	Y	Y			Y	Y
FRNDINT	Round to integer	Y	Y	Y			Y	Y
FRSTOR	Restore state	Y	Y	Y	Y	Y	Y	Y
FSAVE	Save state							
FSCALE	Scale	Y	Y	Y		Y	Y	Y
FSIN	Sine	Y	Y	Y			Y	Y
FSINCOS	Sine and cosine	Y	Y	Y			Y	Y
FSQRT	Square root	Y	Y	Y				Y
FST(P) stack or extended	Store real	Y						
FST(P) single or double	Store real	Y	Y	Y		Y	Y	Y
FSTCW	Store control word							
FSTENV	Store Environment							
FSTSW (AX)	Store status word							
FSUB(R)(P)	Subtract real	Y	Y	Y		Y	Y	Y
FTST	Test	Y	Y	Y				
FUCOM(P)(P)	Unordered compare real	Y	Y	Y				
FWAIT	CPU Wait							
FXAM	Examine							
FXCH	Exchange registers	Y						
FXTRACT	Extract	Y	Y	Y	Y	Y	Y	Y
FYL2X	$Y \cdot \log_2 X$	Y	Y	Y	Y			
FYL2XP1	$Y \cdot \log_2(X + 1)$	Y	Y	Y	Y			

IS—Invalid operand due to stack overflow/underflow

I—Invalid operand due to other cause

D—Denormal operand

Z—Zero-divide

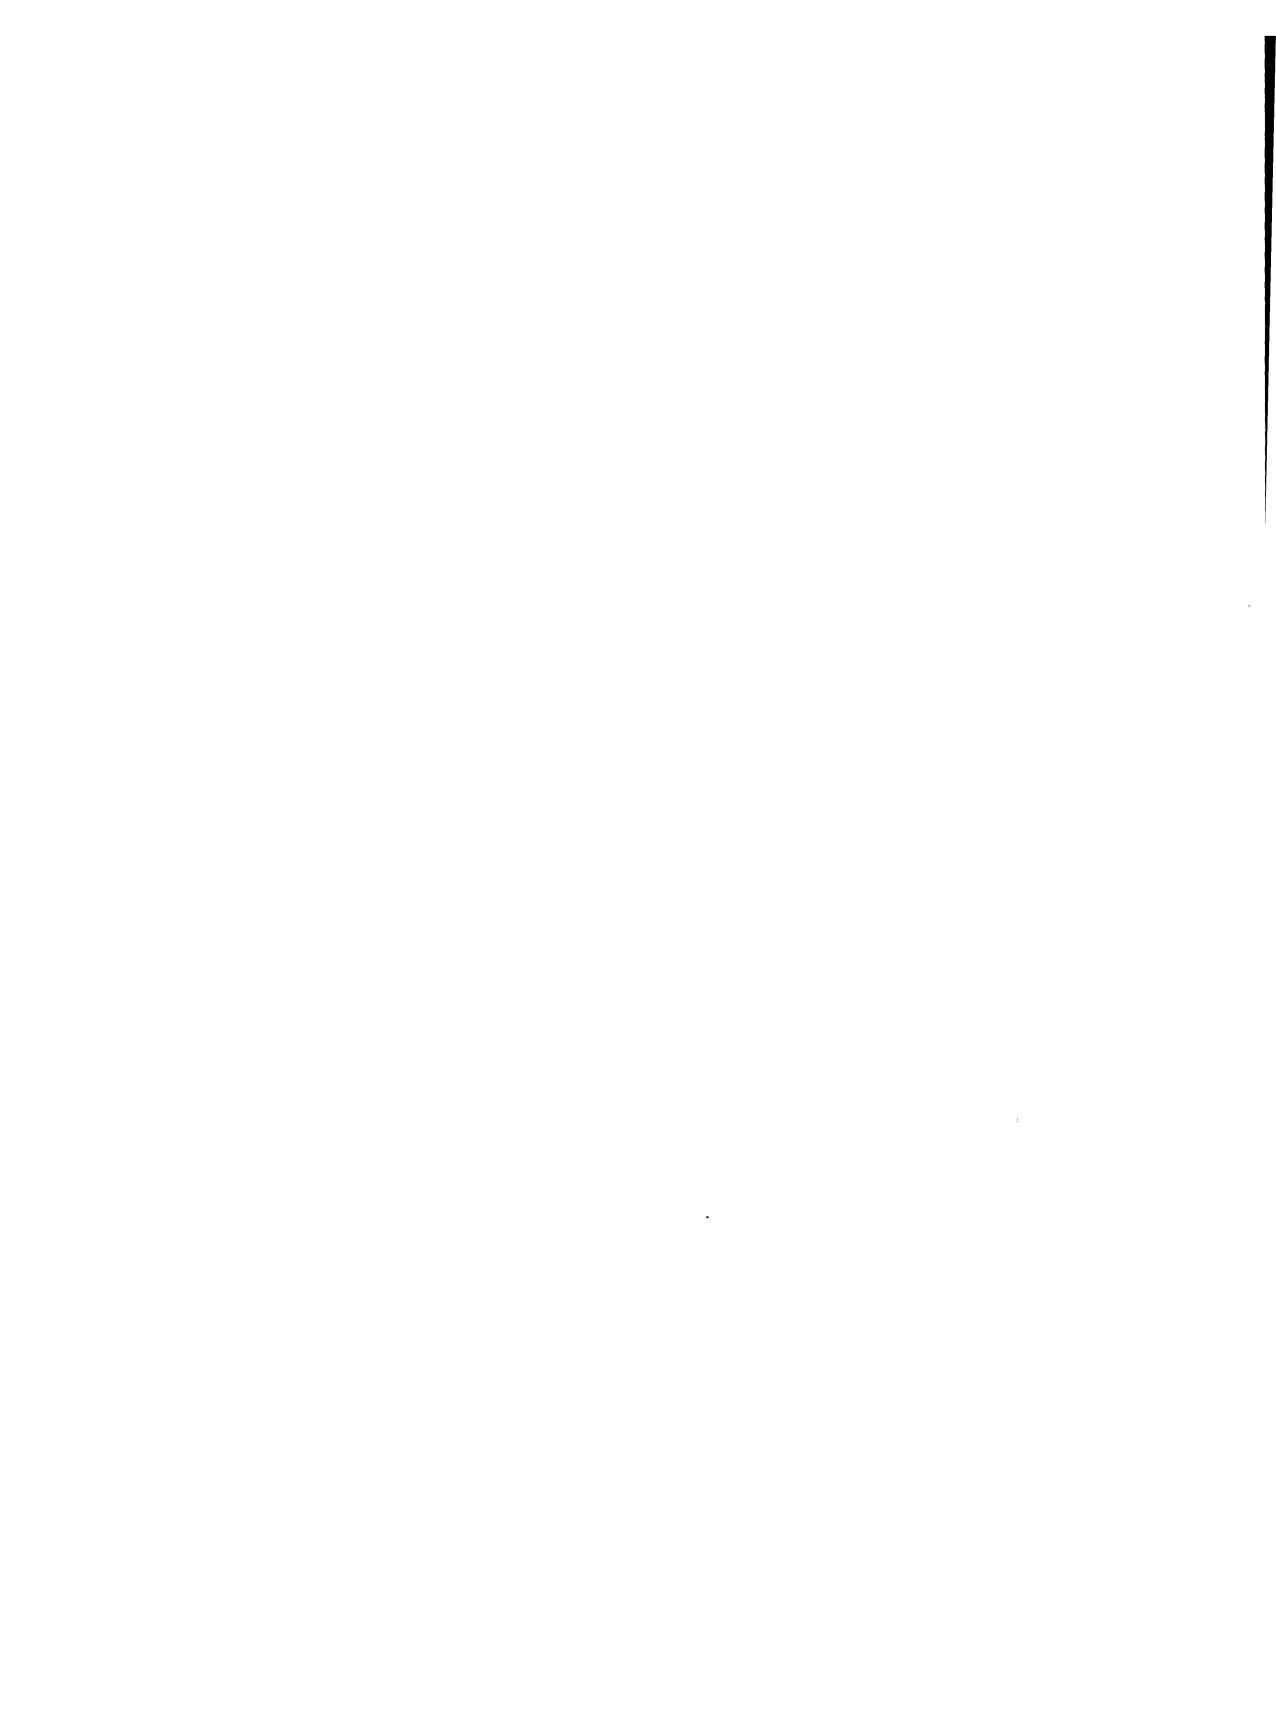
O—Overflow

U—Underflow

P—Inexact result (precision)

*Compatibility Between the
80387 and the 80287/8087*

C



APPENDIX C

COMPATIBILITY BETWEEN THE 80387 AND THE 80287/8087

This appendix summarizes the differences between the 80387 and its predecessors the 80287 and the 8087, and analyzes the impact of these differences on software that must be transported from the 80287 or 8087 to the 80387. Any migration from the 8087 directly to the 80387 must also take into account the additional differences between the 8087 and the 80387 as listed in Appendix D of this manual.

C.1 INITIALIZATION SEQUENCE

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087 / 80287 Behavior		
RESET, FINIT, and ERROR# PIN	After a hardware RESET, the ERROR# output is asserted to indicate that an 80387 is present. To accomplish this, the IE and ES bits of the status word are set, and the IM bit in the control word is reset. After FINIT, the status word and the control word have the same values as in an 80287/8087 after RESET.	No difference between RESET and FINIT.	80387 initialization software must execute an FNINIT instruction to clear ERROR#. The FNINIT is not required for 80287/8087 software, though Intel documentation recommends its use (refer to the Numerics Supplement to the <i>iAPX 286 Programmer's Reference Manual</i>).	Permits the 80386 to differentiate between the 80287 and the 80387.

C.2 DATA TYPES AND EXCEPTION HANDLING

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087 / 80287 Behavior		
NaN	The 80387 distinguishes between signaling NaNs and quiet NaNs. The 80387 only generates quiet NaNs. An invalid-operation exception is raised only upon encountering a signaling NaN (except for FCOM, FIST, and FBSTP which also raise IE for quiet NaNs).	The 80287/8087 only generates one kind of NaN (the equivalent of a quiet NaN) but raises an invalid-operation exception upon encountering any kind of NaN.	Uninitialized memory locations that contain QNaNs should be changed to SNaNs to cause the 80387 to fault when uninitialized memory locations are referenced.	IEEE Standard 754 compatibility.
Pseudozero, Pseudo-NaN, Pseudoinfinity, and Unnormal Formats	The 80387 neither generates nor supports these formats; it raises an invalid-operation exception whenever it encounters them in an arithmetic operation.	The 80287/8087 defines and supports special handling for these formats.	None. The 80387 does not generate these formats, and therefore will not encounter them unless a programmer deliberately enters them.	IEEE Standard 754 compatibility.

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
Tag Word Bits for Unsupported Data Formats	The encoding in the tag word for the unsupported data formats mentioned in Section C.2.2 is "special data" (type 10).	The encoding for pseudo-zero and unnormal is "valid" (type 00); the others are "special data" (type 10).	The exception handler may need to be changed if programmers use such data types.	IEEE Standard 754 compatibility.
Invalid-Operation Exception	No invalid-operation exception is raised upon encountering a denormal in FSQRT, FDIV, or FPREM or upon conversion to BCD or to integer. The operation proceeds by first normalizing the value.	Upon encountering a denormal in FSQRT, FDIV, or FPREM or upon conversion to BCD or to integer, the invalid-operation exception is raised.	None. Software on the 80387 will continue to execute in cases where the 80287/8087 would trap.	Upgrade, to eliminate exception.
Denormal Exception	The denormal exception is raised in transcendental instructions and FXTRACT.	The denormal exception is not raised in transcendental instructions and FXTRACT.	The exception handler needs to be changed only if it gives special treatment to different opcodes.	Performance enhancement for normal case.
Overflow Exception	<p>Overflow exception masked.</p> <p>If the rounding mode is set to chop (toward zero), the result is the most positive or most negative number.</p> <p>Overflow exception not masked.</p> <p>The precision exception is flagged. When the result is stored in the stack, the significand is rounded according to the precision control (PC) bit of the control word or according to the opcode.</p>	<p>Overflow exception masked.</p> <p>The 80287/8087 does not signal the overflow exception when the masked response is not infinity; i.e., it signals overflow only when the rounding control is not set to round to zero. If rounding is set to chop (toward zero), the result is positive or negative infinity.</p> <p>Overflow exception not masked.</p> <p>The precision exception is not flagged and the significand is not rounded.</p>	<p>Overflow exception masked.</p> <p>Under the most common rounding modes, no impact. If rounding is toward zero (chop), a program on the 80387 produces under overflow conditions a result that is different in the least significant bit of the significand, compared to the result on the 80287.</p> <p>Overflow exception not masked.</p> <p>If the result is stored on the stack, a program on the 80387 produces a different result under overflow conditions than on the 80287/8087. The difference is apparent only to the exception handler.</p>	IEEE Standard 754 compatibility.

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
Underflow Exception Two related events contribute to underflow: 1. The creation of a tiny result. A tiny number, because it is so small, may cause some other exception later (such as overflow upon division). 2. Loss of accuracy during the denormalization of a tiny number. Which of these events triggers the underflow exception depends on whether the underflow exception is masked.	Conditions for underflow. When the underflow exception is masked, the underflow exception is signaled when both the result is tiny and denormalization results in a loss of accuracy. Response to underflow. When the underflow exception is unmasked and the instruction is supposed to store the result on the stack, the significand is rounded to the appropriate precision (according to the precision control (PC) bit of the control word, for those instructions controlled by PC, otherwise to extended precision).	Conditions for underflow. When the underflow exception is masked and rounding is toward zero, the underflow exception flag is raised on tiness, regardless of loss of accuracy. Response to underflow. When the underflow exception is not masked and the destination is the stack, the significand is not rounded but rather is left as is.	Underflow exception masked. No impact. The underflow exception occurs less often when rounding is toward zero. Underflow exception not masked. A program on the 80387 produces a different result during underflow conditions than on the 80287/8087 if the result is stored on the stack. The difference is only in the least significant bit of the significand and is apparent only to the exception handler.	IEEE Standard 754 compatibility.
Exception Precedence	There is no difference in the precedence of the denormal exception, whether it be masked or not.	When the denormal exception is not masked, it takes precedence over all other exceptions.	None, but some unneeded normalization of denormal operands is prevented on the 80387.	Operational improvement.

C.3 TAG, STATUS, AND CONTROL WORDS

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
Bits C3-C0 of Status Word	After FINIT, incomplete FPREM, and hardware reset, the 80387 sets these bits to zero.	After FINIT, incomplete FPREM, and hardware reset, the 80287/8087 leaves these bits intact (they contain the prior value).	None.	Upgrade, to provide consistent state after reset.
Bit C2 of Status Word	Bit 10 (C2) serves as an incomplete bit for FPTAN.	This bit is undefined for FPTAN.	None. Programs don't check C2 after FPTAN.	Upgrade to allow fast checking of operand range.
Infinity Control	Only affine closure is supported. Bit 12 remains programmable but has no effect on 80387 operation.	Both affine and projective closures are supported. After RESET, the default value in the control word is projective.	Software that requires projective infinity arithmetic may give different results.	IEEE Standard 754 compatibility.

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
Status Word Bit 6 for Stack Fault	When an invalid-operation exception occurs due to stack overflow or underflow, not only is bit 0 (IE) of the status word set, but also bit 6 is set to indicate a stack fault and bit 9 (C1) specifies overflow or underflow. Bit 6 is called SF and serves to distinguish invalid exceptions caused by stack overflow/underflow from those caused by numeric operations.	When an invalid-operation exception occurs due to stack overflow or underflow, only bit 0 (IE) of the status word is set. Bit 6 is RESERVED.	None. Existing exception handlers need not change, but may be upgraded to take advantage of the additional information. Newly written handlers will be more effective.	Upgrade and performance improvement.
Tag Word	When loading the tag word with an FLDENV or FRSTOR instruction, the only interpretations of tag values used by the 80387 are <i>empty</i> (value 11) and <i>nonempty</i> (values 00, 01, and 10). Subsequent operations on a nonempty register always examine the value in the register, not the value in its tag. The FSTENV and FSAVE instructions examine the nonempty registers and put the correct values in the tags before storing the tag word.	The corresponding tag is checked before each register access to determine the class of operand in the register; the tag is updated after every change to a register so that the tag always reflects the most recent status of the register. Programmers can load a tag with a value that disagrees with the contents of a register (for example, the register contains valid contents, but the tag says <i>special</i> ; the 80287/8087, in this case, honors the tag and does not examine the register).	Software may not operate correctly if it uses FLDENV or FRSTOR to change tags to values (other than empty) that are different from actual register contents.	Performance improvement.

C.4 INSTRUCTION SET

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
FBSTP, FDIV, FIST(P), FPREM, FSQRT	Operation on denormal operand is supported. An underflow exception can occur.	Operation on denormal operand raises invalid-operation exception. Underflow is not possible.	The exception handler for underflow may require change only if it gives different treatment to different opcodes. Possibly fewer invalid-operation exceptions will occur.	IEEE Standard 754 compatibility.
FSCALE	The range of the scaling operand is not restricted. If $0 < ST(1) < 1$, the scaling factor is zero; therefore, ST(0) remains unchanged. If the rounded result is not exact or if there was a loss of accuracy (masked underflow), the precision exception is signaled.	The range of the scaling operand is restricted. If $0 < ST(1) < 1$, the result is undefined and no exception is signaled.	Different result when $0 < ST(1) < 1$.	Upgrade.



COMPATIBILITY BETWEEN THE 80387 AND THE 80287/8087

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
FPREM1	Performs partial remainder according to IEEE Standard 754 standard.	Does not exist.	None.	IEEE Standard 754 compatibility and upgrade.
FPREM	Bits C0, C3, C1 of the status word, correctly reflect the three low-order bits of the quotient.	The quotient bits are incorrect when performing a reduction of $64^N + M$ when $N \geq 1$ and $M=1$ or $M=2$.	None. Software that works around the bug should not be affected.	Upgrade.
FUCOM, FUCOMP, FUCOMPP	Perform unordered compare according to IEEE Standard 754 standard.	Do not exist.	None.	IEEE Standard 754 compatibility.
FPTAN	Range of operand is much less restricted ($ ST(0) < 2^{63}$), reduces operand internally using an internal $\pi/4$ constant that is more accurate. After a stack overflow when the invalid-operation exception is masked, both ST and ST(1) contain quiet NaNs.	Range of operand is restricted ($ ST(0) < \pi/4$); operand must be reduced to range using FPREM. After a stack overflow when the invalid-operation exception is masked, the original operand remains unchanged, but is pushed to ST(1).	None.	Upgrade. IEEE Standard 754 compatibility.
FSIN, FCOS, FSINCOS	Perform three common trigonometric functions.	Do not exist.	None.	Upgrade.
FPATAN	Range of operands is unrestricted.	$ ST(0) $ must be smaller than $ ST(1) $.	None.	Upgrade.
F2XM1	Wider range of operand ($-1 \leq ST(0) \leq +1$).	The supported operand range is $0 \leq ST(0) \leq 0.5$.	None.	Upgrade.
FLD extended-real	Does not report denormal exception because the instruction is not arithmetic.	Reports denormal exception.	None.	Upgrade.
FXTRACT	If the operand is zero, the zero-divide exception is reported and ST(1) is $-\infty$. If the operand is $+\infty$, no exception is reported.	If the operand is zero, ST(1) is zero and no exception is reported. If the operand is $+\infty$, the invalid-operation exception is reported.	None. Software usually bypasses zero and ∞ .	IEEE 754 recommendation to fully support the logb function.
FLD constant	Rounding control is in effect.	Rounding control is not in effect.	Results are the same as for the 8087/80287 when rounding control is set to round to zero, round to $-\infty$, and (in the case of FLDL2T) round to nearest. Results are different by one in the least significant bit of the significand in round to $+\infty$ and round to nearest (excluding FLDL2T). FLD1 and FLDZ are always the same.	IEEE 754 recommendation.

Issue	Difference Description		Impact on Software	Reason for the Difference
	80387 Behavior	8087/80287 Behavior		
FLD single/double precision	Loading a denormal causes the number to be converted to extended precision (because it is put on the stack).	Loading a denormal causes the number to be converted to an unnormal.	If the next instruction is FXTRACT or FXAM, the 80387 will give a different result than the 80287/8087.	IEEE Standard 754 compatibility.
FLD single/double precision	When loading a signaling NaN, raises invalid exception.	Does not raise an exception when loading a signaling NaN.	The exception handler need to be updated to handle this condition.	IEEE Standard 754 compatibility.
FSETPM	Treated as FNOP (no operation).	Informs the 80287 that the system is in protected mode.	None.	The 80386 handles all addressing and exception-pointer information, whether in protected mode or not.
FXAM	When encountering an empty register, the 80387 will not generate combinations of C3-C0 equal to 1101 or 1111.	May generate these combinations, among others.	None.	Upgrade, to provide repeatable results.
All Transcendental Instructions	May generate different results in round-up bit of status word is undefined for these instructions.	Round-up bit of status word is undefined for these instructions.	None.	Upgrade, to signal rounding status.

Compatibility Between the 80387 and the 8087

D



APPENDIX D

COMPATIBILITY BETWEEN THE 80387 AND THE 8087

The 80386/80387 operating in real-address mode will execute 8087 programs without major modification. However, because of differences in the handling of numeric exceptions between the 80387 NPX and the 8087 NPX, exception-handling routines *may* need to be changed.

This appendix summarizes the additional differences between the 80387 NPX and the 8087 NPX (other than those already included in Appendix B), and provides details showing how 8087 programs can be ported to the 80387.

1. The 80387 signals exceptions through a dedicated ERROR# line to the 80386; no interrupt controller is needed for this purpose. The 8087 requires an interrupt controller (8259A) to interrupt the CPU when an unmasked exception occurs. Therefore, any interrupt-controller-oriented instructions in numeric exception handlers for the 8087 should be deleted.
2. The 8087 instructions FENI/FNENI and FDISI/FNDISI perform no useful function in the 80387. If the 80387 encounters one of these opcodes in its instruction stream, the instruction will effectively be ignored—none of the 80387 internal states will be updated. While 8087 code containing these instructions may be executed on the 80387, it is unlikely that the exception-handling routines containing these instructions will be completely portable to the 80387.
3. In real mode and protected mode (not including virtual 8086 mode), interrupt vector 16 must point to the numeric exception handling routine. In virtual 8086 mode, the V86 monitor can be programmed to accommodate a different location of the interrupt vector for numeric exceptions.
4. The ESC instruction address saved in the 80386/80387 or 80386/80287 includes any leading prefixes before the ESC opcode. The corresponding address saved in the 8086/8087 does not include leading prefixes.
5. In protected mode (not including virtual 8086 mode), the format of the 80387's saved instruction and address pointers is different than for the 8087. The instruction opcode is not saved in protected mode—exception handlers will have to retrieve the opcode from memory if needed.
6. Interrupt 7 will occur in the 80386 when executing ESC instructions with either TS (task switched) or EM (emulation) of the 80386 MSW set (TS=1 or EM=1). If TS is set, then a WAIT instruction will also cause interrupt 7. An exception handler should be included in 80387 code to handle these situations.
7. Interrupt 9 will occur if the second or subsequent words of a floating-point operand fall outside a segment's size. Interrupt 13 will occur if the starting address of a numeric operand falls outside a segment's size. An exception handler should be included to report these programming errors.

8. Except for the processor control instructions, all of the 80387 numeric instructions are automatically synchronized by the 80386 CPU—the 80386 automatically waits until all operands have been transferred between the 80386 and the 80387 before executing the next ESC instruction. No explicit WAIT instructions are required to assure this synchronization. For the 8087 used with 8086 and 8088 processors, explicit WAITs are required before each numeric instruction to ensure synchronization. Although 8087 programs having explicit WAIT instructions will execute perfectly on the 80387 without reassembly, these WAIT instructions are unnecessary.
9. Since the 80387 does not require WAIT instructions before each numeric instruction, the ASM386 assembler does not automatically generate these WAIT instructions. The ASM86 assembler, however, automatically precedes every ESC instruction with a WAIT instruction. Although numeric routines generated using the ASM86 assembler will generally execute correctly on the 80386/20, reassembly using ASM386 may result in a more compact code image and faster execution.

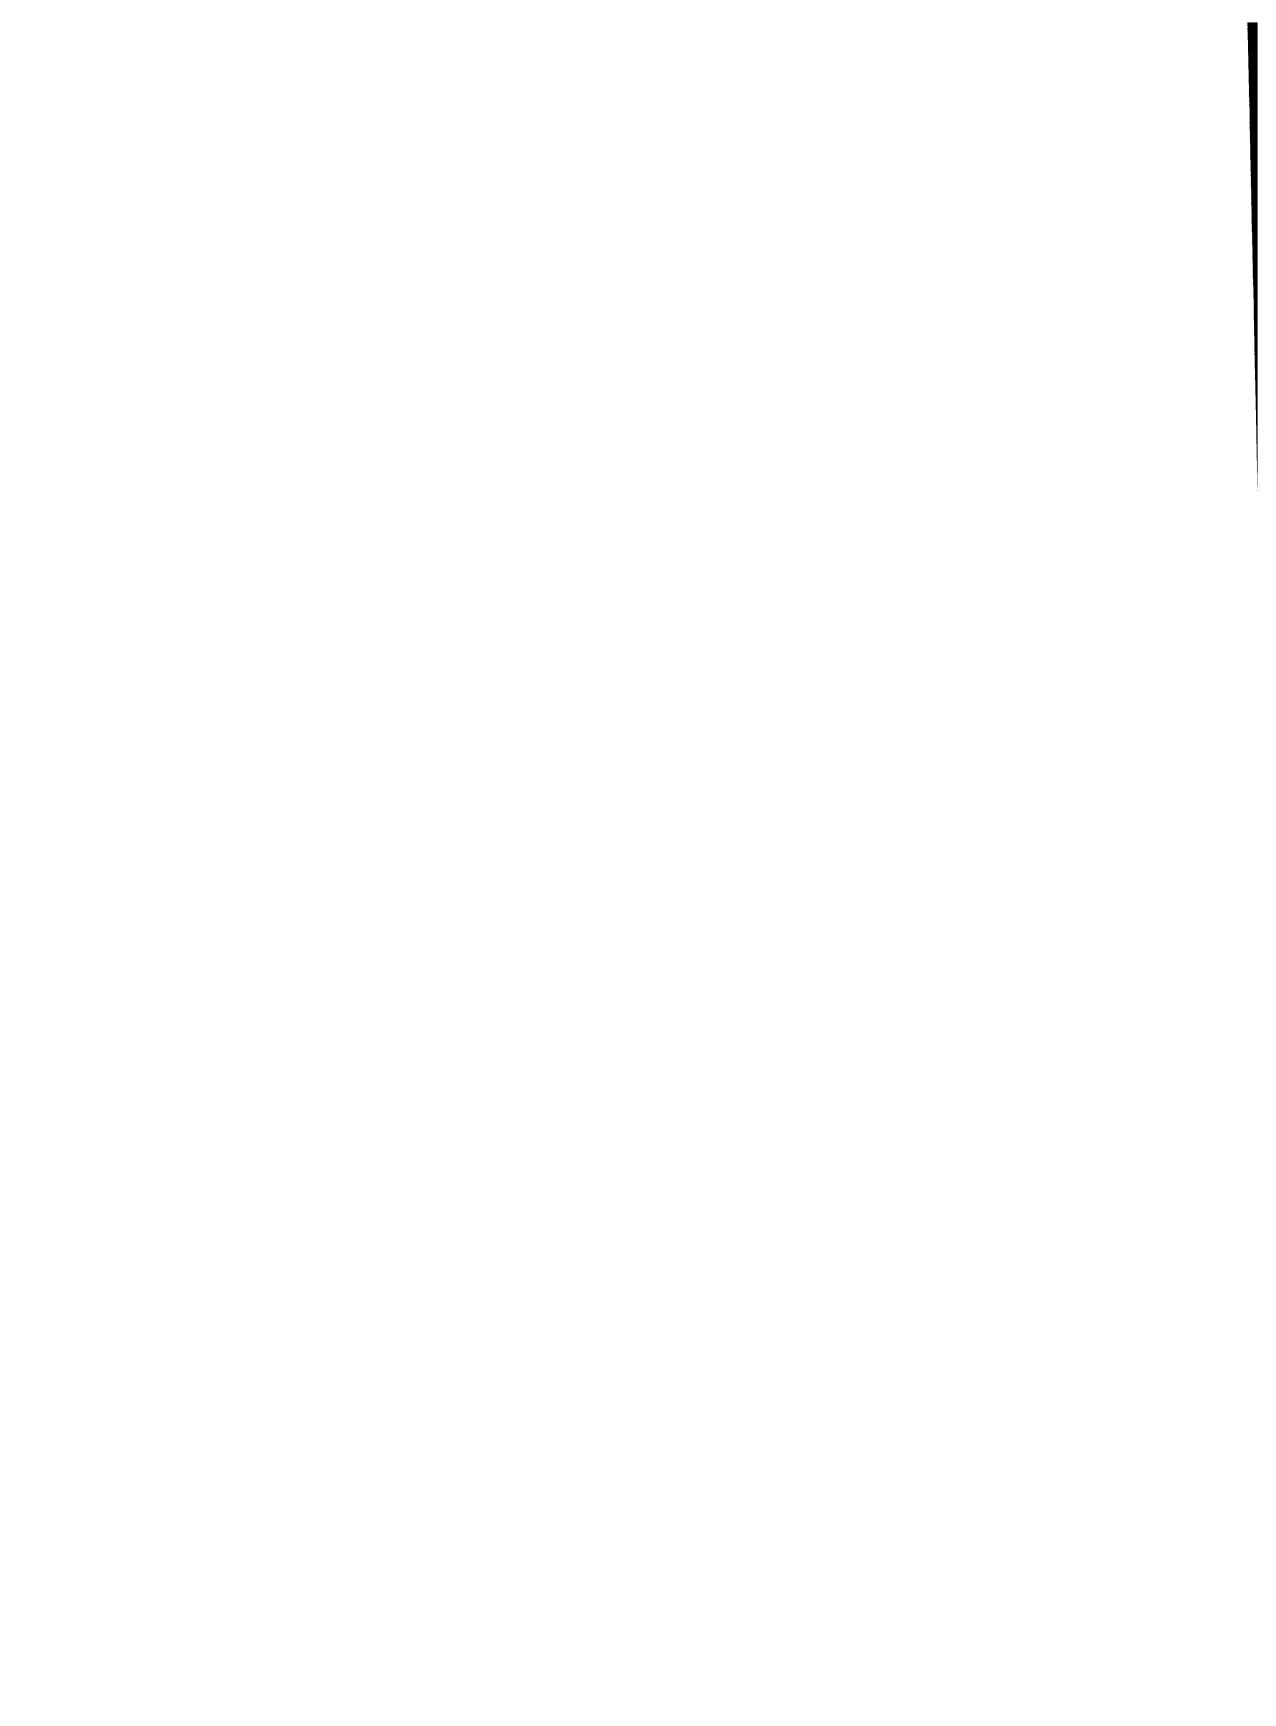
The processor control instructions for the 80387 may be coded using either a WAIT or No-WAIT form of mnemonic. The WAIT forms of these instructions cause ASM386 to precede the ESC instruction with a CPU WAIT instruction, in the identical manner as does ASM86.

10. The address of a memory operand stored by FSAVE or FSTENV is undefined if the previous ESC instruction did not refer to memory.
11. Because the 80387 automatically normalizes denormal numbers when possible, an 8087 program that uses the denormal exception solely to normalize denormal operands can run on an 80387 by masking the denormal exception. The 8087 denormal exception handler would not be used by the 80387 in this case. A numerics program runs faster when the 80387 performs normalization of denormal operands. A program can detect at run-time whether it is running on an 80387 or 8087/80287 and disable the denormal exception when an 80387 is used.

80387 80-Bit CHMOS III Numeric Processor Extension

E

This appendix is a copy of the 80387 Data Sheet, which is also available separately. (The AC specifications have been deliberately left out.) The specifications in data sheets are subject to change; consult the most recent data sheet for design-in information.



80387

80-BIT CHMOS III

NUMERIC PROCESSOR EXTENSION

- High Performance 80-Bit Internal Architecture
- Implements ANSI/IEEE Standard 754-1985 for Binary Floating-Point Arithmetic
- Five to Six Times 8087/80287 Performance
- Upward Object-Code Compatible from 8087 and 80287
- Expands 80386 Data Types to Include 32-, 64-, 80-Bit Floating Point, 32-, 64-Bit Integers and 18-Digit BCD Operands
- Directly Extends 80386 Instruction Set to Include Trigonometric, Logarithmic, Exponential and Arithmetic Instructions for All Data Types
- Full-Range Transcendental Operations for SINE, COSINE, TANGENT, ARCTANGENT and LOGARITHM
- Built-In Exception Handling
- Operates Independently of Real, Protected and Virtual-8086 Modes of the 80386
- Eight 80-Bit Numeric Registers, Usable as Individually Addressable General Registers or as a Register Stack
- Available in 68-Pin PGA Package
(See Packaging Spec: Order #231369)

The Intel 80387 is a high-performance numerics processor extension that extends the 80386 architecture with floating point, extended integer and BCD data types. The 80386/80387 computing system fully conforms to the ANSI/IEEE floating-point standard. Using a numerics oriented architecture, the 80387 adds over seventy mnemonics to the 80386/80387 instruction set, making the 80386/80387 a complete solution for high-performance numerics processing. The 80387 is implemented with 1.5 micron, high-speed CHMOS III technology and packaged in a 68-pin ceramic pin grid array (PGA) package. The 80386/80387 is upward object-code compatible from the 80386/80287, 80286/80287 and 8086/8087 computing systems.

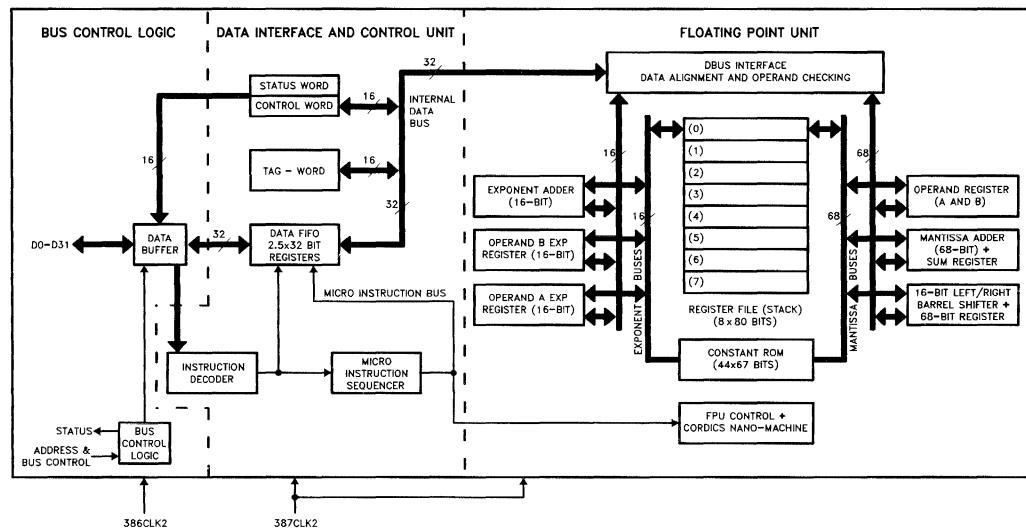


Figure 0.1. 80387 Block Diagram

231920-1

CONTENTS

1.0 Functional Description	4
2.0 Programming Interface	5
2.1 Data Types	5
2.2 Numeric Operands	5
2.3 Register Set	7
2.3.1 Data Registers	7
2.3.2 Tag Word	7
2.3.3 Status Word	8
2.3.4 Instruction and Data Pointers	11
2.3.5 Control Word	13
2.4 Interrupt Description	13
2.5 Exception Handling	14
2.6 Initialization	14
2.7 8087 and 80287 Compatibility	15
2.7.1 General Differences	15
2.7.2 Exceptions	16
3.0 Hardware Interface	16
3.1 Signal Description	16
3.1.1 80386 Clock 2 (386CLK2)	16
3.1.2 80387 Clock 2 (387CLK2)	16
3.1.3 80387 Clocking Mode (CKM)	18
3.1.4 System Reset (RESETIN)	18
3.1.5 Processor Extension Request (PEREQ)	18
3.1.6 Busy Status (BUSY #)	18
3.1.7 Error Status (ERROR #)	18
3.1.8 Data Pins (D31-D0)	18
3.1.9 Write/Read Bus Cycle (W/R #)	18
3.1.10 Address Strobe (ADS #)	19
3.1.11 Bus Ready Input (READY #)	19
3.1.12 Ready Output (READYO #)	19
3.1.13 Status Enable (STEN)	19
3.1.14 NPX Select #1 (NPS1#)	19
3.1.15 NPX Select #2 (NPS2)	19
3.1.16 Command (CMD0 #)	19
3.2 Processor Architecture	19
3.2.1 Bus Control Logic	19
3.2.2 Data Interface and Control Unit	20
3.2.3 Floating Point Unit	20
3.3 System Configuration	20
3.3.1 Bus Cycle Tracking	21
3.3.2 80387 Addressing	21
3.3.3 Function Select	21
3.3.4 CPU/NPX Synchronization	21
3.3.5 Synchronous or Asynchronous Modes	21
3.3.6 Automatic Bus Cycle Termination	22
3.4 Bus Operation	22
3.4.1 Nonpipelined Bus Cycles	23
3.4.1.1 Write Cycle	23
3.4.1.2 Read Cycle	23
3.4.2 Pipelined Bus Cycles	24
3.4.3 Bus Cycles of Mixed Type	25
3.4.4 BUSY # and PEREQ Timing Relationship	25
4.0 Mechanical Data	27

5.0 Electrical Data	28
5.1 Absolute Maximum Ratings	28
5.2 DC Characteristics	28
5.3 AC Characteristics	29
6.0 80387 Extensions to the 80386 Instruction Set	33
Appendix A—Compatibility Between the 80287 NPX and the 8087	37

FIGURES

Figure 0.1 80387 Block Diagram	1
Figure 1.1 80386/80387 Register Set	4
Figure 2.1 80387 Tag Word	7
Figure 2.2 80387 Status Word	8
Figure 2.3 Protected Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format	11
Figure 2.4 Real Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format	12
Figure 2.5 Protected Mode 80387 Instruction and Data Pointer Image in Memory, 16-Bit Format	12
Figure 2.6 Real Mode 80387 Instruction and Data Pointer Image in Memory, 16-Bit Format	12
Figure 2.7 80387 Control Word	13
Figure 3.1 80387 Pin Configuration	18
Figure 3.2 80386/80387 System Configuration	20
Figure 3.3 Bus State Diagram	22
Figure 3.4 Nonpipelined Read and Write Cycles	24
Figure 3.5 Fastest Transitions to and from Pipelined Cycles	25
Figure 3.6 Pipelined Cycles with Wait States	26
Figure 3.7 STEN, BUSY # and PEREQ Timing Relationship	26
Figure 4.1 Package Description	27
Figure 5.1 386CLK2/387CLK2 Waveform	30
Figure 5.2 Output Signals	30
Figure 5.3 Input and I/O Signals	31
Figure 5.4 RESET Signal	31
Figure 5.5 Float from STEN	31
Figure 5.6 Other Parameters	32

TABLES

Table 2.1 80387 Data Type Representation in Memory	6
Table 2.2 Condition Code Interpretation	9
Table 2.3 Condition Code Interpretation after FPREM and FPREM1 Instructions	10
Table 2.4 Condition Code Resulting from Comparison	10
Table 2.5 Condition Code Defining Operand Class	10
Table 2.6 80386 Interrupt Vectors Reserved for NPX	14
Table 2.7 Exceptions	15
Table 3.1 80387 Pin Summary	17
Table 3.2 80387 Pin Cross-Reference	17
Table 3.3 Output Pin Status after Reset	18
Table 3.4 Bus Cycles Definition	21
Table 5.1 DC Specifications	28
Table 5.2 Timing Requirements	29
Table 5.3 Other Parameters	32

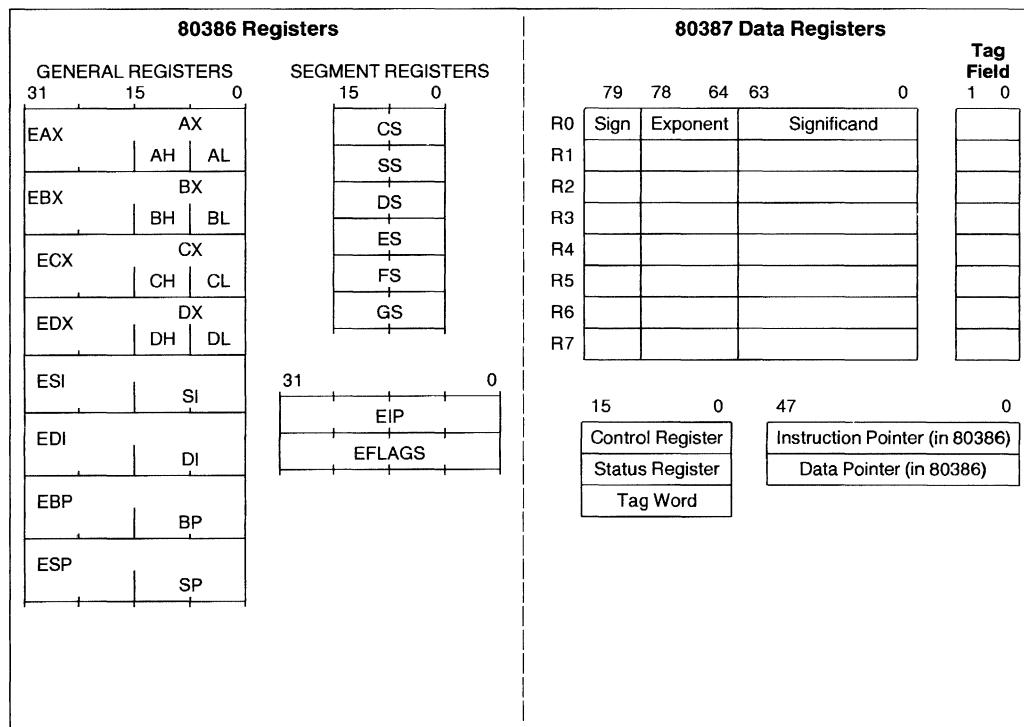


Figure 1.1. 80386/80387 Register Set

1.0 FUNCTIONAL DESCRIPTION

The 80387 Numeric Processor Extension (NPX) provides arithmetic instructions for a variety of numeric data types in 80386/80387 systems. It also executes numerous built-in transcendental functions (e.g. tangent, sine, cosine, and log functions). The 80387 effectively extends the register and instruction set of an 80386 system for existing data types and adds several new data types as well. Figure 1.1 shows the model of registers visible to 80386/80387 programs. Essentially, the 80387 can be treated as an additional resource or an extension to the 80386. The 80386 together with an 80387 can be used as a single unified system, the 80386/80387.

The 80387 works the same whether the 80386 is executing in real-address mode, protected mode, or virtual-8086 mode. All memory access is handled by the 80386; the 80387 merely operates on instructions and values passed to it by the 80386. Therefore, the 80387 is not sensitive to the processing mode of the 80386.

In real-address mode and virtual-8086 mode, the 80386/80387 is completely upward compatible with software for 8086/8087, 80286/80287 real-address mode, and 80386/80287 real-address mode systems.

In protected mode, the 80386/80387 is completely upward compatible with software for 80286/80287 protected mode, and 80386/80287 protected mode systems.

The only differences of operation that may appear when 8086/8087 programs are ported to a protected-mode 80386/80387 system (*not* using virtual-8086 mode), is in the format of operands for the administrative instructions FLDENV, FSTENV, FRSTOR and FSAVE. These instructions are normally used only by exception handlers and operating systems, not by applications programs.

The 80387 contains three functional units that can operate in parallel to increase system performance. The 80386 can be transferring commands and data to the 80387 *bus control logic* for the next instruction while the 80387 *floating-point unit* is performing the current numeric instruction.

2.0 PROGRAMMING INTERFACE

The 80387 adds to an 80386 system additional data types, registers, instructions, and interrupts specifically designed to facilitate high-speed numerics processing. To use the 80387 requires no special programming tools, because all new instructions and data types are directly supported by the 80386 assembler and compilers for high-level languages. All 8086/8088 development tools that support the 8087 can also be used to develop software for the 80386/80387 in real-address mode or virtual-8086 mode. All 80286 development tools that support the 80287 can also be used to develop software for the 80386/80387.

All communication between the 80386 and the 80387 is transparent to applications software. The CPU automatically controls the 80387 whenever a numerics instruction is executed. All physical memory and virtual memory of the CPU are available for storage of the instructions and operands of programs that use the 80387. All memory addressing modes, including use of displacement, base register, index register, and scaling, are available for addressing numerics operands.

Section 6 at the end of this data sheet lists by class the instructions that the 80387 adds to the instruction set of an 80386 system.

2.1 Data Types

Table 2.1 lists the seven data types that the 80387 supports and presents the format for each type. Operands are stored in memory with the least significant digit at the lowest memory address. Programs retrieve these values by generating the lowest address. For maximum system performance, all operands should start at physical-memory addresses evenly divisible by four (doubleword boundaries); operands may begin at any other addresses, but will require extra memory cycles to access the entire operand.

Internally, the 80387 holds all numbers in the extended-precision real format. Instructions that load operands from memory automatically convert operands represented in memory as 16-, 32-, or 64-bit integers, 32- or 64-bit floating-point numbers, or 18-digit packed BCD numbers into extended-precision real format. Instructions that store operands in memory perform the inverse type conversion.

2.2 Numeric Operands

A typical NPX instruction accepts one or two operands and produces a single result. In two-operand instructions, one operand is the contents of an NPX register, while the other may be a memory location. The operands of some instructions are predefined; for example FSQRT always takes the square root of the number in the top stack element.

Table 2.1. 80387 Data Type Representation in Memory

Data Formats	Range	Precision	Most Significant Byte										HIGHEST ADDRESSED BYTE									
			7	0	7	0	7	0	7	0	7	0	7	0	7	0	7	0	7	0	7	0
Word Integer	10^4	16 Bits																				
Short Integer	10^9	32 Bits																				
Long Integer	10^{19}	64 Bits																				
Packed BCD	10^{18}	18 Digits																				
Single Precision	$10^{\pm 38}$	24 Bits																				
Double Precision	$10^{\pm 308}$	53 Bits																				
Extended Precision	$10^{\pm 4932}$	64 Bits																				

231920-2

NOTES:

- (1) S = Sign bit (0 = positive, 1 = negative)
- (2) d_n = Decimal digit (two per byte)
- (3) X = Bits have no significance; 80387 ignores when loading, zeros when storing
- (4) ▲ = Position of implicit binary point
- (5) I = Integer bit of significand; stored in temporary real, implicit in single and double precision
- (6) Exponent Bias (normalized values):
 - Single: 127 (7FH)
 - Double: 1023 (3FFH)
 - Extended Real: 16383 (3FFFH)
- (7) Packed BCD: $(-1)^S (D_{17} \dots D_0)$
- (8) Real: $(-1)^S (2^E\text{-BIAS}) (F_0 F_1 \dots)$

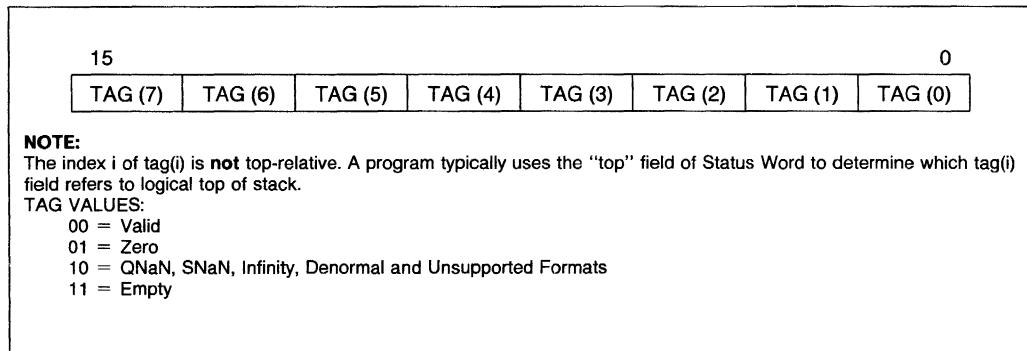


Figure 2.1. 80387 Tag Word

2.3 Register Set

Figure 1.1 shows the 80387 register set. When an 80387 is present in a system, programmers may use these registers in addition to the registers normally available on the 80386.

2.3.1 DATA REGISTERS

80387 computations use the 80387's data registers. These eight 80-bit registers provide the equivalent capacity of twenty 32-bit registers. Each of the eight data registers in the 80387 is 80 bits wide and is divided into "fields" corresponding to the NPXs extended-precision real data type.

The 80387 register set can be accessed either as a stack, with instructions operating on the top one or two stack elements, or as a fixed register set, with instructions operating on explicitly designated registers. The TOP field in the status word identifies the current top-of-stack register. A "push" operation decrements TOP by one and loads a value into the new top register. A "pop" operation stores the value from the current top register and then increments

TOP by one. Like 80386 stacks in memory, the 80387 register stack grows "down" toward lower-addressed registers.

Instructions may address the data registers either implicitly or explicitly. Many instructions operate on the register at the TOP of the stack. These instructions implicitly address the register at which TOP points. Other instructions allow the programmer to explicitly specify which register to use. This explicit register addressing is also relative to TOP.

2.3.2 TAG WORD

The tag word marks the content of each numeric data register, as Figure 2.1 shows. Each two-bit tag represents one of the eight numerics registers. The principal function of the tag word is to optimize the NPXs performance and stack handling by making it possible to distinguish between empty and nonempty register locations. It also enables exception handlers to check the contents of a stack location without the need to perform complex decoding of the actual data.

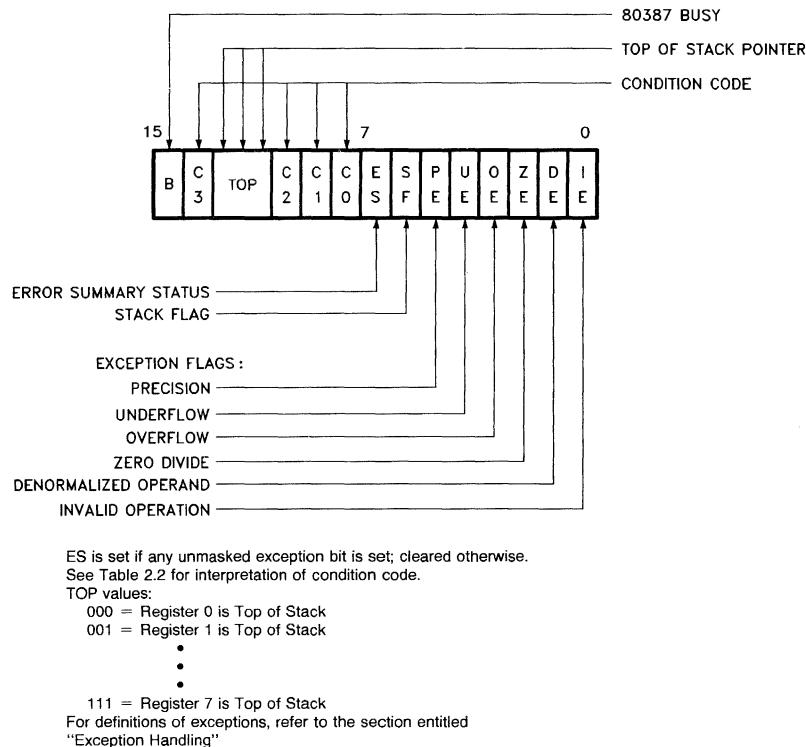


Figure 2.2. 80387 Status Word

2.3.3 STATUS WORD

The 16-bit status word (in the status register) shown in Figure 2.2 reflects the overall state of the 80387. It may be read and inspected by CPU code.

Bit 15, the B-bit (busy bit) is included for 8087 compatibility only. It reflects the contents of the ES bit (bit 7 of the status word), not the status of the BUSY# output of 80387/80287.

Bits 13–11 (TOP) point to the 80387 register that is the current top-of-stack.

The four numeric condition code bits (C₃–C₀) are similar to the flags in a CPU; instructions that perform arithmetic operations update these bits to reflect the outcome. The effects of these instructions on the condition code are summarized in Tables 2.2 through 2.5.

Bit 7 is the error summary (ES) status bit. This bit is set if any unmasked exception bit is set; it is clear otherwise. If this bit is set, the ERROR# signal is asserted.

Bit 6 is the stack flag (SF). This bit is used to distinguish invalid operations due to stack overflow or underflow from other kinds of invalid operations. When SF is set, bit 9 (C₁) distinguishes between stack overflow (C₁ = 1) and underflow (C₁ = 0).

Figure 2.2 shows the six exception flags in bits 5–0 of the status word. Bits 5–0 are set to indicate that the 80387 has detected an exception while executing an instruction. A later section entitled "Exception Handling" explains how they are set and used.

Note that when a new value is loaded into the status word by the FLDENV or FRSTOR instruction, the value of ES (bit 7) and its reflection in the B-bit (bit 15) are not derived from the values loaded from memory but rather are dependent upon the values of the exception flags (bits 5–0) in the status word and their corresponding masks in the control word. If ES is set in such a case, the ERROR# output of the 80387 is activated immediately.

Table 2.2. Condition Code Interpretation

Instruction	C0 (S)	C3 (Z)	C1 (A)	C2 (C)
FPREM, FPREM1 (see Table 2.3)	Three least significant bits of quotient Q2	Q0	Q1 or O/U#	Reduction 0 = complete 1 = incomplete
FCOM, FCOMP, FCOMPP, FTST, FUCOM, FUCOMP, FUCOMPP, FICOM, FICOMP	Result of comparison (see Table 2.4)		Zero or O/U#	Operand is not comparable (Table 2.4)
FXAM	Operand class (see Table 2.5)		Sign or O/U#	Operand class (Table 2.5)
FCHS, FABS, FXCH, FINCTOP, FDECSTOP, Constant loads, FXTRACT, FLD, FILD, FBLD, FSTP (ext real)	UNDEFINED		Zero or O/U#	UNDEFINED
FIST, FBSTP, FRNDINT, FST, FSTP, FADD, FMUL, FDIV, FDIVR, FSUB, FSUBR, FSCALE, FSQRT, FPATAN, F2XM1, FYLN2X, FYL2XP1	UNDEFINED		Roundup or O/U#	UNDEFINED
FPTAN, FSIN FCOS, FSINCOS	UNDEFINED		Roundup or O/U#, undefined if C2 = 1	Reduction 0 = complete 1 = incomplete
FLDENV, FRSTOR	Each bit loaded from memory			
FLDCW, FSTENV, FSTCW, FSTSW, FCLEX, FINIT, FSAVE	UNDEFINED			
O/U#	When both IE and SF bits of status word are set, indicating a stack exception, this bit distinguishes between stack overflow (C1 = 1) and underflow (C1 = 0).			
Reduction	If FPREM or FPREM1 produces a remainder that is less than the modulus, reduction is complete. When reduction is incomplete the value at the top of the stack is a partial remainder, which can be used as input to further reduction. For FPTAN, FSIN, FCOS, and FSINCOS, the reduction bit is set if the operand at the top of the stack is too large. In this case the original operand remains at the top of the stack.			
Roundup	When the PE bit of the status word is set, this bit indicates whether the last rounding in the instruction was upward.			
UNDEFINED	Do not rely on finding any specific value in these bits.			

Table 2.3. Condition Code Interpretation after FPREM and FPREM1 Instructions

Condition Code				Interpretation after FPREM and FPREM1			
C2	C3	C1	C0				
1	X	X	X	Incomplete Reduction: further iteration required for complete reduction			
0	Q1	Q0	Q2	Q MOD8	Complete Reduction: C0, C3, C1 contain three least significant bits of quotient		
	0	0	0	0			
	0	1	0	1			
	1	0	0	2			
	1	1	0	3			
	0	0	1	4			
	0	1	1	5			
	1	0	1	6			
	1	1	1	7			

Table 2.4. Condition Code Resulting from Comparison

Order	C3	C2	C0
TOP > Operand	0	0	0
TOP < Operand	0	0	1
TOP = Operand	1	0	0
Unordered	1	1	1

Table 2.5. Condition Code Defining Operand Class

C3	C2	C1	C0	Value at TOP
0	0	0	0	+ Unsupported
0	0	0	1	+ NaN
0	0	1	0	- Unsupported
0	0	1	1	- NaN
0	1	0	0	+ Normal
0	1	0	1	+ Infinity
0	1	1	0	- Normal
0	1	1	1	- Infinity
1	0	0	0	+ 0
1	0	0	1	+ Empty
1	0	1	0	- 0
1	0	1	1	- Empty
1	1	0	0	+ Denormal
1	1	1	0	- Denormal

2.3.4 INSTRUCTION AND DATA POINTERS

Because the NPX operates in parallel with the CPU, any errors detected by the NPX may be reported after the CPU has executed the ESC instruction which caused it. To allow identification of the failing numeric instruction, the 80386/80387 contains two pointer registers that supply the address of the failing numeric instruction and the address of its numeric memory operand (if appropriate).

The instruction and data pointers are provided for user-written error handlers. These registers are actually located in the 80386, but appear to be located in the 80387 because they are accessed by the ESC instructions FLDENV, FSTENV, FSAVE, and FRSTOR. (In the 8086/8087 and 80286/80287, these registers are located in the NPX.) Whenever the 80386 decodes a new ESC instruction, it saves

the address of the instruction (including any prefixes that may be present), the address of the operand (if present), and the opcode.

The instruction and data pointers appear in one of four formats depending on the operating mode of the 80386 (protected mode or real-address mode) and depending on the operand-size attribute in effect (32-bit operand or 16-bit operand). When the 80386 is in virtual-8086 mode, the real-address mode formats are used. (See Figures 2.3 through 2.6.) The ESC instructions FLDENV, FSTENV, FSAVE, and FRSTOR are used to transfer these values between the 80386 registers and memory. Note that the value of the data pointer is *undefined* if the prior ESC instruction did not have a memory operand.

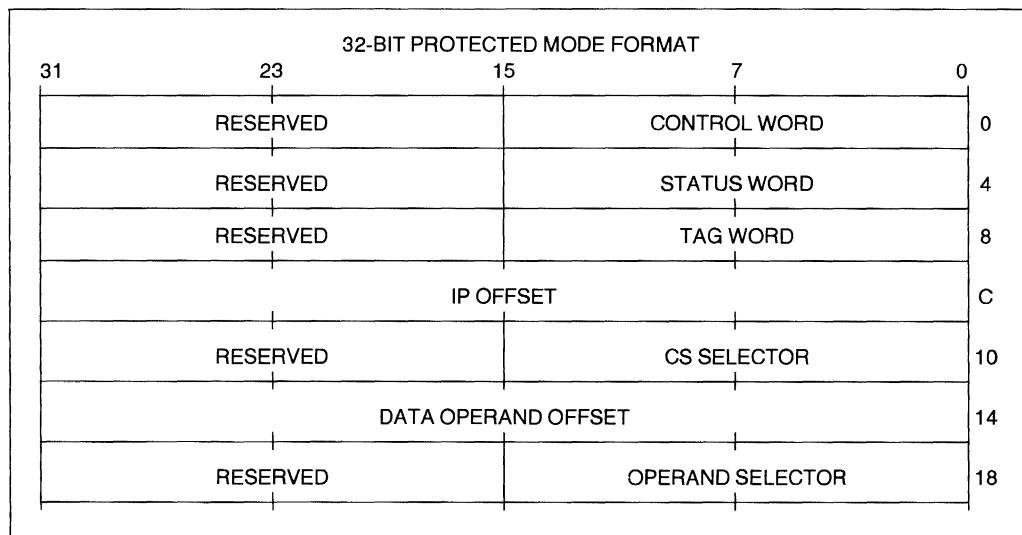


Figure 2.3. Protected Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format

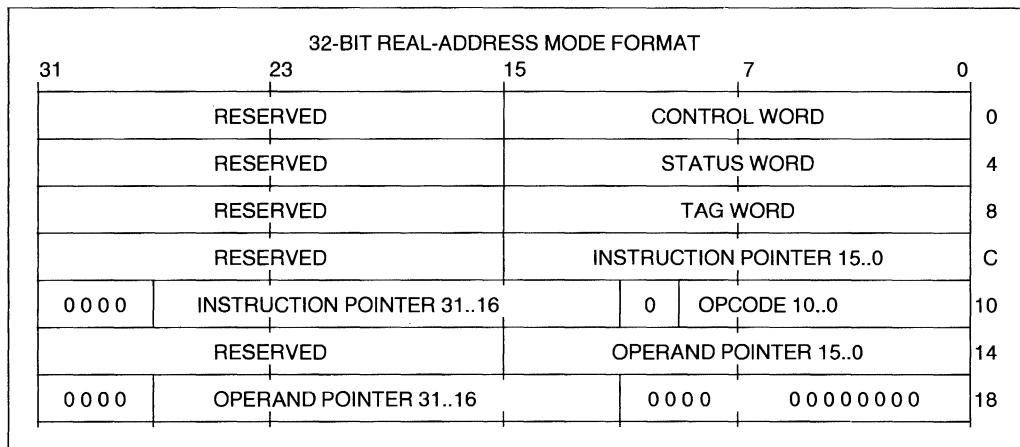


Figure 2.4. Real Mode 80387 Instruction and Data Pointer Image in Memory, 32-Bit Format

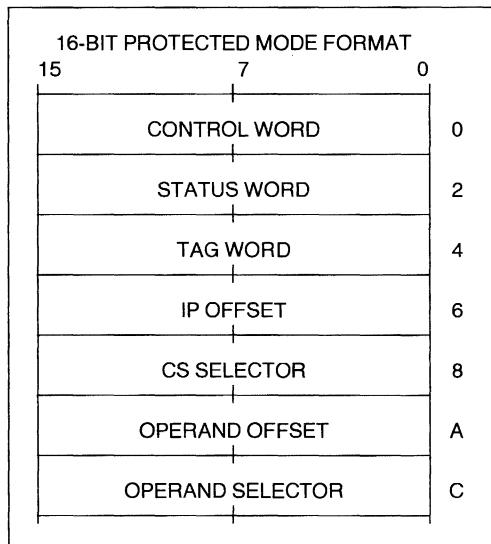


Figure 2.5. Protected Mode 80387
Instruction and Data Pointer
Image in Memory, 16-Bit Format

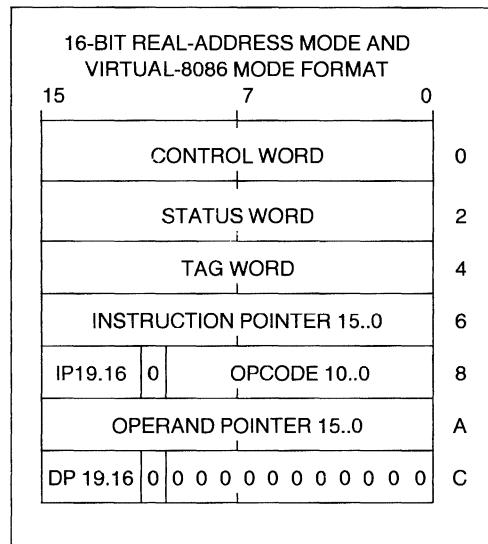


Figure 2.6. Real Mode 80387
Instruction and Data Pointer
Image in Memory, 16-Bit Format

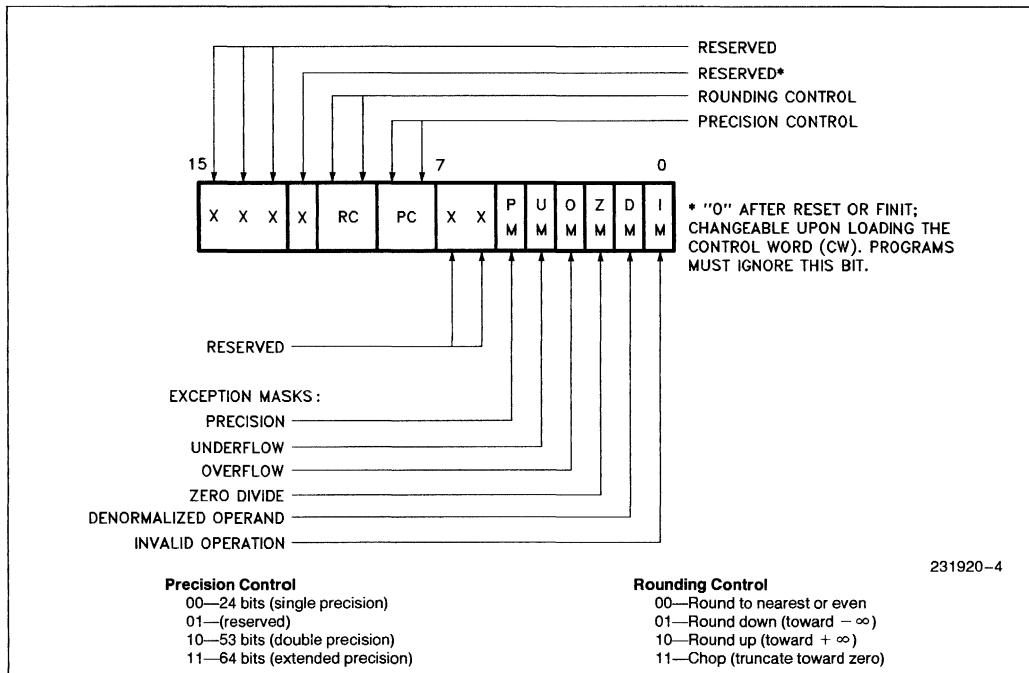


Figure 2.7. 80387 Control Word

2.3.5 CONTROL WORD

The NPX provides several processing options that are selected by loading a control word from memory into the control register. Figure 2.7 shows the format and encoding of fields in the control word.

The low-order byte of this control word configures the 80387 error and exception masking. Bits 5–0 of the control word contain individual masks for each of the six exceptions that the 80387 recognizes.

The high-order byte of the control word configures the 80387 operating mode, including precision and rounding.

- Bit 12 no longer defines infinity control and is a reserved bit. Only affine closure is supported for infinity arithmetic. The bit is initialized to zero after RESET or FINIT and is changeable upon loading the CW. Programs must ignore this bit.
- The rounding control (RC) bits (bits 11–10) provide for directed rounding and true chop, as well as the unbiased round to nearest even mode specified in the IEEE standard. Rounding control

affects only those instructions that perform rounding at the end of the operation (and thus can generate a precision exception); namely, FST, FSTP, FIST, all arithmetic instructions (except FPREM, FPREM1, FXTRACT, FABS, and FCHS), and all transcendental instructions.

- The precision control (PC) bits (bits 9–8) can be used to set the 80387 internal operating precision of the significand at less than the default of 64 bits (extended precision). This can be useful in providing compatibility with early generation arithmetic processors of smaller precision. PC affects only the instructions ADD, SUB, DIV, MUL, and SQRT. For all other instructions, either the precision is determined by the opcode or extended precision is used.

2.4 Interrupt Description

Several interrupts of the 80386 are used to report exceptional conditions while executing numeric programs in either real or protected mode. Table 2.6 shows these interrupts and their causes.

Table 2.6. 80386 Interrupt Vectors Reserved for NPX

Interrupt Number	Cause of Interrupt
7	An ESC instruction was encountered when EM or TS of 80386 control register zero (CR0) was set. EM = 1 indicates that software emulation of the instruction is required. When TS is set, either an ESC or WAIT instruction causes interrupt 7. This indicates that the current NPX context may not belong to the current task.
9	An operand of a coprocessor instruction wrapped around an addressing limit (0FFFFH for small segments, 0FFFFFFFH for big segments, zero for expand-down segments) and spanned inaccessible addresses ^a . The failing numerics instruction is not restartable. The address of the failing numerics instruction and data operand may be lost; an FSTENV does not return reliable addresses. As with the 80286/80287, the segment overrun exception should be handled by executing an FNINIT instruction (i.e. an FINIT without a preceding WAIT). The return address on the stack does not necessarily point to the failing instruction nor to the following instruction. The interrupt can be avoided by never allowing numeric data to start within 108 bytes of the end of a segment.
13	The first word or doubleword of a numeric operand is not entirely within the limit of its segment. The return address pushed onto the stack of the exception handler points at the ESC instruction that caused the exception, including any prefixes. The 80387 has not executed this instruction; the instruction pointer and data pointer register refer to a previous, correctly executed instruction.
16	The previous numerics instruction caused an unmasked exception. The address of the faulty instruction and the address of its operand are stored in the instruction pointer and data pointer registers. Only ESC and WAIT instructions can cause this interrupt. The 80386 return address pushed onto the stack of the exception handler points to a WAIT or ESC instruction (including prefixes). This instruction can be restarted after clearing the exception condition in the NPX. FNINIT, FNCLEX, FNSTSW, FNSTENV, and FNSAVE cannot cause this interrupt.

a. An operand may wrap around an addressing limit when the segment limit is near an addressing limit and the operand is near the largest valid address in the segment. Because of the wrap-around, the beginning and ending addresses of such an operand will be at opposite ends of the segment. There are two ways that such an operand may also span inaccessible addresses: 1) if the segment limit is not equal to the addressing limit (e.g. addressing limit is FFFFH and segment limit is FFDH) the operand will span addresses that are not within the segment (e.g. an 8-byte operand that starts at valid offset FFFC will span addresses FFFC–FFF7 and 0000–0003; however addresses FFFE and FFFF are not valid, because they exceed the limit); 2) if the operand begins and ends in present and accessible pages but intermediate bytes of the operand fall in a not-present page or a page to which the procedure does not have access rights.

2.5 Exception Handling

The 80387 detects six different exception conditions that can occur during instruction execution. Table 2.7 lists the exception conditions in order of precedence, showing for each the cause and the default action taken by the 80387 if the exception is masked by its corresponding mask bit in the control word.

Any exception that is not masked by the control word sets the corresponding exception flag of the status word, sets the ES bit of the status word, and asserts the ERROR# signal. When the CPU attempts to execute another ESC instruction or WAIT, exception 16 occurs. The exception condition must be resolved via an interrupt service routine. The 80386/80387 saves the address of the floating-point instruction that caused the exception and the address of any memory operand required by that instruction.

2.6 Initialization

80387 initialization software must execute an FNINIT instruction (i.e. an FINIT without a preceding WAIT) to clear ERROR#. The FNINIT is not required for the 80287, though Intel documentation recommends its use (refer to the Numerics Supplement to the *iAPX 286 Programmer's Reference Manual*). After a hardware RESET, the ERROR# output is asserted to indicate that an 80387 is present. To accomplish this, the IE and ES bits of the status word are set, and the IM bit in the control word is reset. After FNINIT, the status word and the control word have the same values as in an 80287 after RESET.

2.7 8087 and 80287 Compatibility

This section summarizes the differences between the 80387 and the 80287. Any migration from the 8087 directly to the 80387 must also take into account the differences between the 8087 and the 80287 as listed in Appendix A.

Many changes have been designed into the 80387 to directly support the IEEE standard in hardware. These changes result in increased performance by eliminating the need for software that supports the standard.

2.7.1 GENERAL DIFFERENCES

The 80387 supports only affine closure for infinity arithmetic, not projective closure. Bit 12 of the Control Word (CW) no longer defines infinity control. It is a reserved bit; but it is initialized to zero after RESET or FINIT and is changeable upon loading the CW. Programs must ignore this bit.

Operands for FSCALE and FPATAN are no longer restricted in range (except for $\pm\infty$); F2XM1 and FPTAN accept a wider range of operands.

The results of transcendental operations may be slightly different from those computed by 80287.

In the case of FPTAN, the 80387 supplies a true tangent result in ST(1), and (always) a floating point 1 in ST.

Rounding control is in effect for FLD *constant*.

Software cannot change entries of the tag word to values (other than empty) that do not reflect the actual register contents.

After reset, FINIT, and incomplete FPREM, the 80387 resets to zero the condition code bits C₃-C₀ of the status word.

In conformance with the IEEE standard, the 80387 does not support the special data formats: pseudozero, pseudo-NaN, pseudoinfinity, and unnormal.

Table 2.7. Exceptions

Exception	Cause	Default Action (if exception is masked)
Invalid Operation	Operation on a signaling NaN, unsupported format, indeterminate form ($0^*\infty$, $0/0$, $(+\infty) + (-\infty)$, etc.), or stack overflow/underflow (SF is also set).	Result is a quiet NaN, integer indefinite, or BCD indefinite
Denormalized Operand	At least one of the operands is denormalized, i.e. it has the smallest exponent but a nonzero significand.	Normal processing continues
Zero Divisor	The divisor is zero while the dividend is a noninfinite, nonzero number.	Result is ∞
Overflow	The result is too large in magnitude to fit in the specified format.	Result is largest finite value or ∞
Underflow	The true result is nonzero but too small to be represented in the specified format, and, if underflow exception is masked, denormalization causes loss of accuracy.	Result is denormalized or zero
Inexact Result (Precision)	The true result is not exactly representable in the specified format (e.g. 1/3); the result is rounded according to the rounding mode.	Normal processing continues

2.7.2 EXCEPTIONS

When the overflow or underflow exception is masked, one difference from the 80287 is in rounding when overflow or underflow occurs. The 80387 produces results that are consistent with the rounding mode. The other difference is that the 80387 sets its underflow flag only if there is also a loss of accuracy during denormalization.

A number of differences exist due to changes in the IEEE standard and to functional improvements to the architecture of the 80387:

1. Fewer invalid-operation exceptions due to denormal operands, because the instructions FSQRT, FDIV, FPREM and conversions to BCD or to integer normalize denormal operands before proceeding.
2. The FSQRT, FBSTP, and FPREM instructions may cause underflow, because they support denormal operands.
3. The denormal exception can occur during the transcendental instructions and the FXTRACT instruction.
4. The denormal exception no longer takes precedence over all other exceptions.
5. When the operand is zero, the FXTRACT instruction reports a zero-divide exception and leaves $-\infty$ in ST(1).
6. The status word has a new bit (SF) that signals when invalid-operation exceptions are due to stack underflow or overflow.
7. FLD *extended precision* no longer reports denormal exceptions, because the instruction is not numeric.
8. FLD *single/double precision* when the operand is denormal converts the number to extended precision and signals the denormalized operand exception. When loading a signaling NaN, FLD *single/double precision* signals an invalid-operation exception.
9. The 80387 only generates quiet NaNs (as on the 80287); however, the 80387 distinguishes between quiet NaNs and signaling NaNs. Signaling NaNs trigger exceptions when they are used as operands; quiet NaNs do not (except for FCOM, FIST, and FBSTP which also raise IE for quiet NaNs).

3.0 HARDWARE INTERFACE

In the following description of hardware interface, the # symbol at the end of a signal name indicates that the active or asserted state occurs when the

signal is at a low voltage. When no # is present after the signal name, the signal is asserted when at the high voltage level.

3.1 Signal Description

In the following signal descriptions, the 80387 pins are grouped by function as follows:

1. Execution control—386CLK2, 387CLK2, CKM, RESETIN
2. NPX handshake—PEREQ, BUSY#, ERROR#
3. Bus interface pins—D31-D0, W/R#, ADS#, READY#, READYO#
4. Chip/Port Select—STEN, NPS1#, NPS2, CMD0#
5. Power supplies—V_{CC}, V_{SS}

Table 3.1 lists every pin by its identifier, gives a brief description of its function, and lists some of its characteristics. All output signals are tristate; they leave floating state only when STEN is active. The output buffers of the bidirectional data pins D31-D0 are also tristate; they leave floating state only in read cycles when the 80387 is selected (i.e. when STEN, NPS1#, and NPS2 are all active).

Figure 3.1 and Table 3.2 together show the location of every pin in the pin grid array.

3.1.1 80386 CLOCK 2 (386CLK2)

This input uses the 80386 CLK2 signal to time the bus control logic. Several other 80387 signals are referenced to the rising edge of this signal. When CKM = 1 (synchronous mode) this pin also clocks the data interface and control unit and the floating-point unit of the 80387. This pin requires MOS-level input. The signal on this pin is divided by two to produce the internal clock signal CLK.

3.1.2 80387 CLOCK 2 (387CLK2)

When CKM = 0 (asynchronous mode) this pin provides the clock for the data interface and control unit and the floating-point unit of the 80387. In this case, the ratio of the frequency of 387CLK2 to the frequency of 386CLK2 must lie within the range 10:16 to 16:10. When CKM = 1 (synchronous mode) this pin is ignored; 386CLK2 is used instead for the data interface and control unit and the floating-point unit. This pin requires TTL-level input.

Table 3.1. 80387 Pin Summary

Pin Name	Function	Active State	Input/Output	Referenced To
386CLK2	80386 CLock 2		—	
387CLK2	80387 CLock 2		—	
CKM	80387 CLocking Mode		—	
RESETIN	System reset	High	—	386CLK2
PEREQ	Processor Extension REQuest	High	O	386CLK2/STEN
BUSY #	Busy status	Low	O	386CLK2/STEN
ERROR #	Error status	Low	O	387CLK2/STEN
D31-D0	Data pins	High	I/O	386CLK2
W/R #	Write/Read bus cycle	Hi/Lo	—	386CLK2
ADS #	ADdress Strobe	Low	—	386CLK2
READY #	Bus ready input	Low	—	386CLK2
READYO #	Ready output	Low	O	386CLK2/STEN
STEN	STatus ENable	High	—	386CLK2
NPS1 #	NPX select # 1	Low	—	386CLK2
NPS2	NPX select # 2	High	—	386CLK2
CMD0 #	CoMmand	Low	—	386CLK2
V _{CC}			—	
V _{SS}			—	

NOTE:

STEN is referenced to only when getting the output pins into or out of tristate mode.

Table 3.2. 80387 Pin Cross-Reference

A2	—	D9	C11	—	V _{SS}	J10	—	V _{SS}
A3	—	D11	D1	—	D5	J11	—	CKM
A4	—	D12	D2	—	D4	K1	—	PEREQ
A5	—	D14	D10	—	D24	K2	—	BUSY #
A6	—	V _{CC}	D11	—	D25	K3	—	Tie High
A7	—	D16	E1	—	V _{CC}	K5	—	W/R #
A8	—	D18	E2	—	V _{SS}	K5	—	V _{CC}
A9	—	V _{CC}	E10	—	D26	K6	—	NPS2
A10	—	D21	E11	—	D27	K7	—	ADS #
B1	—	D8	F1	—	V _{CC}	K8	—	READY #
B2	—	V _{SS}	F2	—	V _{SS}	K9	—	No Connect
B3	—	D10	F10	—	V _{CC}	K10	—	386CLK2
B4	—	V _{CC}	F11	—	V _{SS}	K11	—	387CLK2
B5	—	D13	G1	—	D3	L2	—	ERROR #
B6	—	D15	G2	—	D2	L3	—	READYO #
B7	—	V _{SS}	G10	—	D28	L4	—	STEN
B8	—	D17	G11	—	D29	L5	—	V _{SS}
B9	—	D19	H1	—	D1	L6	—	NPS1 #
B10	—	D20	H2	—	D0	L7	—	V _{CC}
B11	—	D22	H10	—	D30	L8	—	CMD0 #
C1	—	D7	H11	—	D31	L9	—	Tie High
C2	—	D6	J1	—	V _{SS}	L10	—	RESETIN
C10	—	D23	J2	—	V _{CC}			

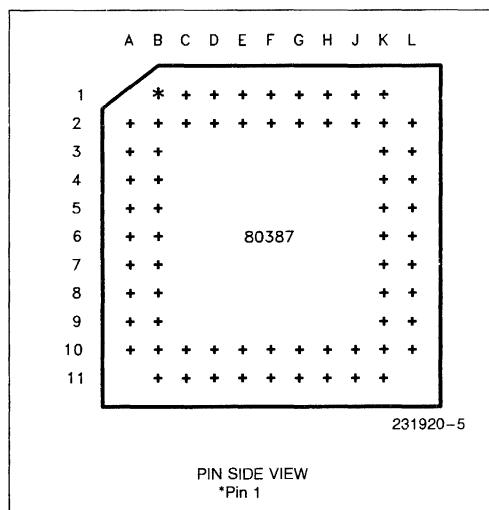


Figure 3.1. 80387 Pin Configuration

3.1.3 80387 CLOCKING MODE (CKM)

This pin is a strapping option. When it is strapped to V_{CC}, the 80387 operates in synchronous mode; when strapped to V_{SS}, the 80387 operates in asynchronous mode. These modes relate to clocking of the data interface and control unit and the floating-point unit only; the bus control logic always operates synchronously with respect to the 80386.

3.1.4 SYSTEM RESET (RESETIN)

A LOW to HIGH transition on this pin causes the 80387 to terminate its present activity and to enter a dormant state. RESETIN must remain HIGH for at least 40 387CLK2 periods. The HIGH to LOW transitions of RESETIN must be synchronous with 386CLK2, so that the phase of the internal clock of the bus control logic (which is the 386CLK2 divided by 2) is the same as the phase of the internal clock of the 80386. After RESETIN goes LOW, at least 50 387CLK2 periods must pass before the first NPX instruction is written into the 80387. This pin should be connected to the 80386 RESET pin. Table 3.3 shows the status of other pins after a reset.

Table 3.3. Output Pin Status during Reset

Pin Value	Pin Name
HIGH	READYO#, BUSY#
LOW	PREQ, ERROR#
Tri-State OFF	D31-D0

3.1.5 PROCESSOR EXTENSION REQUEST (PREQ)

When active, this pin signals to the 80386 CPU that the 80387 is ready for data transfer to/from its data FIFO. When all data is written to or read from the data FIFO, PREQ is deactivated. This signal always goes inactive before BUSY# goes inactive. This signal is referenced to 386CLK2. It should be connected to the 80386 PREQ input. Refer to Figure 3.7 for the timing relationships between this and the BUSY# and ERROR# pins.

3.1.6 BUSY STATUS (BUSY#)

When active, this pin signals to the 80386 CPU that the 80387 is currently executing an instruction. This signal is referenced to 386CLK2. It should be connected to the 80386 BUSY# pin. Refer to Figure 3.7 for the timing relationships between this and the PREQ and ERROR# pins.

3.1.7 ERROR STATUS (ERROR#)

This pin reflects the ES bits of the status register. When active, it indicates that an unmasked exception has occurred (except that, immediately after a reset, it indicates to the 80386 that an 80387 is present in the system). This signal can be changed to inactive state only by the following instructions (without a preceding WAIT): FNINIT, FNCLEX, FNSTENV, and FNSAVE. This signal is referenced to 386CLK2. It should be connected to the 80386 ERROR# pin. Refer to Figure 3.7 for the timing relationships between this and the PREQ and BUSY# pins.

3.1.8 DATA PINS (D31-D0)

These bidirectional pins are used to transfer data and opcodes between the 80386 and 80387. They are normally connected directly to the corresponding 80386 data pins. HIGH state indicates a value of one. D0 is the least significant data bit. Timings are referenced to 386CLK2.

3.1.9 WRITE/READ BUS CYCLE (W/R#)

This signal indicates to the 80387 whether the 80386 bus cycle in progress is a read or a write cycle. This pin should be connected directly to the 80386 W/R# pin. HIGH indicates a write cycle; LOW, a read cycle. This input is ignored if any of the signals STEN, NPS1#, or NPS2 is inactive. Setup and hold times are referenced to 386CLK2.

3.1.10 ADDRESS STROBE (ADS#)

This input, in conjunction with the READY# input indicates when the 80387 bus-control logic may sample W/R# and the chip-select signals. Setup and hold times are referenced to 386CLK2. This pin should be connected to the 80386 ADS# pin.

3.1.11 BUS READY INPUT (READY#)

This input indicates to the 80387 when an 80386 bus cycle is to be terminated. It is used by the bus-control logic to trace bus activities. Bus cycles can be extended indefinitely until terminated by READY#. This input should be connected to the same signal that drives the 80386 READ# input. Setup and hold times are referenced to 386CLK2.

3.1.12 READY OUTPUT (READYO#)

This pin is activated at such a time that write cycles are terminated after two clocks and read cycles after three clocks. In configurations where no extra wait states are required, it can be used to directly drive the 80386 READY# input. Refer to section 3.4 "Bus Operation" for details. This pin is activated only during bus cycles that select the 80387. This signal is referenced to 386CLK2.

3.1.13 STATUS ENABLE (STEN)

This pin serves as a chip select for the 80387. When inactive, this pin forces BUSY#, PEREQ, ERROR#, and READYO# outputs into floating state. D31-D0 are normally floating and leave floating state only if STEN is active and additional conditions are met. STEN also causes the chip to recognize its other chip-select inputs. STEN makes it easier to do on-board testing (using the overdrive method) of other chips in systems containing the 80387. STEN should be pulled up with a resistor so that it can be pulled down when testing. In boards that do not use on-board testing, STEN should be connected to V_{CC}. Setup and hold times are relative to 386CLK2. Note that STEN must maintain the same setup and hold times as NPS1#, NPS2, and CMD0# (i.e. if STEN changes state during an 80387 bus cycle, it should change state during the same CLK period as the NPS1#, NPS2, and CMD0# signals).

3.1.14 NPX Select #1 (NPS1#)

When active (along with STEN and NPS2) in the first period of an 80386 bus cycle, this signal indicates that the purpose of the bus cycle is to communicate with the 80387. This pin should be connected directly to the 80386 M/IO# pin, so that the 80387 is selected only when the 80386 performs I/O cycles. Setup and hold times are referenced to 386CLK2.

3.1.15 NPX SELECT #2 (NPS2)

When active (along with STEN and NPS1#) in the first period of an 80386 bus cycle, this signal indicates that the purpose of the bus cycle is to communicate with the 80387. This pin should be connected directly to the 80386 A31 pin, so that the 80387 is selected only when the 80386 uses one of the I/O addresses reserved for the 80387 (800000F8 or 800000FC). Setup and hold times are referenced to 386CLK2.

3.1.16 COMMAND (CMD0#)

During a write cycle, this signal indicates whether an opcode (CMD0# active) or data (CMD0# inactive) is being sent to the 80387. During a read cycle, it indicates whether the control or status register (CMD0# active) or a data register (CMD0# inactive) is being read. CMD0# should be connected directly to the A2 output of the 80386. Setup and hold times are referenced to 386CLK2.

3.2 Processor Architecture

As shown by the block diagram on the front page, the NPX is internally divided into three sections: the bus control logic (BCL), the data interface and control unit, and the floating point unit (FPU). The FPU (with the support of the control unit which contains the sequencer and other support units) executes all numerics instructions. The data interface and control unit is responsible for the data flow to and from the FPU and the control registers, for receiving the instructions, decoding them, and sequencing the microinstructions, and for handling some of the administrative instructions. The BCL is responsible for 80386 bus tracking and interface. The BCL is the only unit in the 80387 that must run synchronously with the 80386; the rest of the 80387 can run asynchronously with respect to the 80386.

3.2.1 BUS CONTROL LOGIC

The BCL communicates solely with the CPU using I/O bus cycles. The BCL appears to the CPU as a special peripheral device. It is special in two respects: the CPU initiates I/O automatically when it encounters ESC instructions, and the CPU uses reserved I/O addresses to communicate with the BCL. The BCL does not communicate directly with memory. The CPU performs all memory access, transferring input operands from memory to the 80387 and transferring outputs from the 80387 to memory.

3.2.2 DATA INTERFACE AND CONTROL UNIT

The data interface and control unit latches the data and, subject to BCL control, directs the data to the FIFO or the instruction decoder. The instruction decoder decodes the ESC instructions sent to it by the CPU and generates controls that direct the data flow in the FIFO. It also triggers the microinstruction sequencer that controls execution of each instruction. If the ESC instruction is FINIT, FCLEX, FSTSW, FSTSW AX, or FSTCW, the control executes it independently of the FPU and the sequencer. The data interface and control unit is the one that generates the BUSY#, PEREQ and ERROR# signals that synchronize 80387 activities with the 80386. It also supports the FPU in all operations that it cannot perform alone (e.g. exceptions handling, transcendental operations, etc.).

3.2.3 FLOATING POINT UNIT

The FPU executes all instructions that involve the register stack, including arithmetic, logical, transcen-

dental, constant, and data transfer instructions. The data path in the FPU is 84 bits wide (68 significant bits, 15 exponent bits, and a sign bit) which allows internal operand transfers to be performed at very high speeds.

3.3 System Configuration

As an extension to the 80386, the 80387 can be connected to the CPU as shown by Figure 3.2. A dedicated communication protocol makes possible high-speed transfer of opcodes and operands between the 80386 and 80387. The 80387 is designed so that no additional components are required for interface with the 80386. The 80387 shares the 32-bit wide local bus of the 80386 and most control pins of the 80387 are connected directly to pins of the 80386.

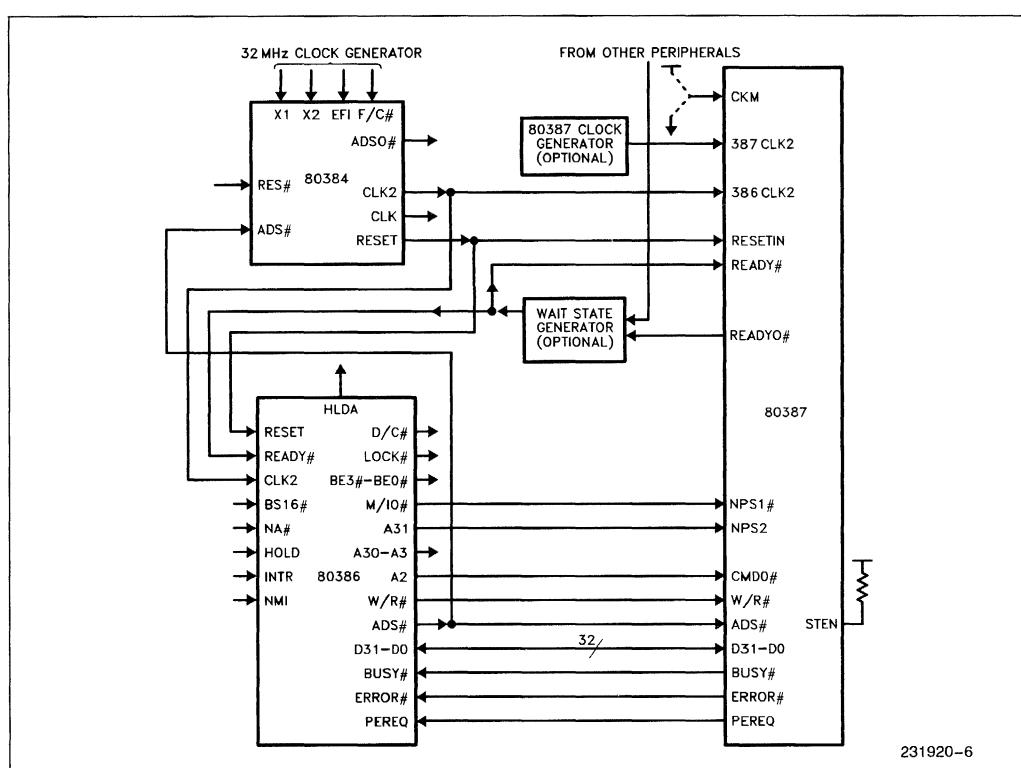


Figure 3.2. 80386/80387 System Configuration

231920-6

Table 3.4. Bus Cycles Definition

STEN	NPS1 #	NPS2	CMD0 #	W/R #	Bus Cycle Type
0	x	x	x	x	80387 not selected and all outputs in floating state
1	1	x	x	x	80387 not selected
1	x	0	x	x	80387 not selected
1	0	1	0	0	CW or SW read from 80387
1	0	1	0	1	Opcode write to 80387
1	0	1	1	0	Data read from 80387
1	0	1	1	1	Data write to 80387

3.3.1 BUS CYCLE TRACKING

The ADS# and READY# signals allow the 80387 to track the beginning and end of 80386 bus cycles, respectively. When ADS# is asserted at the same time as the 80387 chip-select inputs, the bus cycle is intended for the 80387. To signal the end of a bus cycle for the 80387, READY# may be asserted directly or indirectly by the 80387 or by other bus-control logic. Refer to Table 3.4 for definition of the types of 80387 bus cycles.

3.3.2 80387 ADDRESSING

The NPS1#, NPS2 and STEN signals allow the NPX to identify which bus cycles are intended for the NPX. The NPX responds only to I/O cycles when bit 31 of the I/O address is set. In other words, the NPX acts as an I/O device in a reserved I/O address space.

Because A₃₁ is used to select the 80387 for data transfers, it is not possible for a program running on the 80386 to address the 80387 with an I/O instruction. Only ESC instructions cause the 80386 to communicate with the 80387. The 80386 BS16# input must be inactive during I/O cycles when A₃₁ is active.

3.3.3 FUNCTION SELECT

The CMD0# and W/R# signals identify the four kinds of bus cycle: control or status register read, data read, opcode write, data write.

3.3.4 CPU/NPX Synchronization

The pin pairs BUSY#, PEREQ, and ERROR# are used for various aspects of synchronization between the CPU and the NPX.

BUSY# is used to synchronize instruction transfer from the 80386 to the 80387. When the 80387 recognizes an ESC instruction, it asserts BUSY#. For most ESC instructions, the 80386 waits for the 80387 to deassert BUSY# before sending the new opcode.

The NPX uses the PEREQ pin of the 80386 CPU to signal that the NPX is ready for data transfer to or from its data FIFO. The NPX does not directly access memory; rather, the 80386 provides memory access services for the NPX. Thus, memory access on behalf of the NPX always obeys the rules applicable to the mode of the 80386, whether the 80386 be in real-address mode or protected mode.

Once the 80386 initiates an 80387 instruction that has operands, the 80386 waits for PEREQ signals that indicate when the 80387 is ready for operand transfer. Once all operands have been transferred (or if the instruction has no operands) the 80386 continues program execution while the 80387 executes the ESC instruction.

In 8086/8087 systems, WAIT instructions may be required to achieve synchronization of both commands and operands. In 80286/80287 and 80386/80387 systems, WAIT instructions are required only for operand synchronization; namely, after NPX stores to memory (except FSTSW and FSTCW) or loads from memory. Used this way, WAIT ensures that the value has already been written or read by the NPX before the CPU reads or changes the value.

Once it has started to execute a numerics instruction and has transferred the operands from the 80386, the 80387 can process the instruction in parallel with and independent of the host CPU. When the NPX detects an exception, it asserts the ERROR# signal, which causes an 80386 interrupt.

3.3.5 SYNCHRONOUS OR ASYNCHRONOUS MODES

The internal logic of the 80387 (the FPU) can either operate directly from the CPU clock (synchronous mode) or from a separate clock (asynchronous mode). The two configurations are distinguished by the CKM pin. In either case, the bus control logic (BCL) of the 80387 is synchronized with the CPU clock. Use of asynchronous mode allows the 80386 and the FPU section of the 80387 to run at different speeds. In this case, the ratio of the frequency of

387CLK2 to the frequency of 386CLK2 must lie within the range 10:16 to 16:10. Use of synchronous mode eliminates one clock generator from the board design.

3.3.6 AUTOMATIC BUS CYCLE TERMINATION

In configurations where no extra wait states are required, READYO# can be used to drive the 80386 READY# input. If this pin is used, it should be connected to the logic that ORs all READY outputs from peripherals on the 80386 bus. READYO# is asserted by the 80387 only during I/O cycles that select the 80387. Refer to section 3.4 "Bus Operation" for details.

3.4 Bus Operation

With respect to the bus interface, the 80387 is fully synchronous with the 80386. Both operate at the same rate, because each generates its internal CLK signal by dividing 386CLK2 by two.

The 80386 initiates a new bus cycle by activating ADS#. The 80387 recognizes a bus cycle, if, during the cycle in which ADS# is activated, STEN, NPS1#, and NPS2 are all activated. Proper operation is achieved if NPS1# is connected to the M/IO# output of the 80386, and NPS2 to the A31 output. The 80386's A31 output is guaranteed to be inactive in all bus cycles that do not address the 80387 (i.e. I/O cycles to other devices, interrupt acknowledge, and reserved types of bus cycles). System logic must not signal a 16-bit bus cycle via the 80386 BS16# input during I/O cycles when A31 is active.

During the CLK period in which ADS# is activated, the 80387 also examines the W/R# input signal to determine whether the cycle is a read or a write cycle and examines the CMD0# input to determine whether an opcode, operand, or control/status register transfer is to occur.

The 80387 supports both pipelined and nonpipelined bus cycles. A nonpipelined cycle is one for which the 80386 asserts ADS# when no other 80387 bus cycle is in progress. A pipelined bus cycle is one for which the 80386 asserts ADS# and provides valid next-address and control signals as soon as in the second CLK period after the ADS# assertion for the previous 80386 bus cycle. Pipelining increases the availability of the bus by at least one CLK period. The 80387 supports pipelined bus cycles in order to optimize address pipelining by the 80386 for memory cycles.

Bus operation is described in terms of an abstract *state machine*. Figure 3.3 illustrates the states and state transitions for 80387 bus cycles:

- T_I is the idle state. This is the state of the bus logic after RESET, the state to which bus logic returns after every nonpipelined bus cycle, and the state to which bus logic returns after a series of pipelined cycles.
- T_{RS} is the READY# sensitive state. Different types of bus cycle may require a minimum of one or two successive T_{RS} states. The bus logic remains in T_{RS} state until READY# is sensed, at which point the bus cycle terminates. Any number of wait states may be implemented by delaying READY#, thereby causing additional successive T_{RS} states.
- T_P is the first state for every pipelined bus cycle.

The READYO# output of the 80387 indicates when a bus cycle for the 80387 may be terminated if no extra wait states are required. For all write cycles (except those for the instructions FLDENV and FRSTOR), READYO# is always asserted in the first T_{RS} state, regardless of the number of wait states. For all read cycles and write cycles for FLDENV and FRSTOR, READYO# is always asserted in the second T_{RS} state, regardless of the number of wait states. These rules apply to both pipelined and non-pipelined cycles. Systems designers may use READYO# in one of three ways:

1. Leave it disconnected and use external logic to generate READY# signals. When choosing this option, 80387 requirements for wait states in read cycles and write cycles of FLDENV and FRSTOR must be obeyed.
2. Connect it (directly or through logic that ORs READY signals from other devices) to the READY# inputs of the 80386 and 80387.
3. Use it as one input to a wait-state generator.

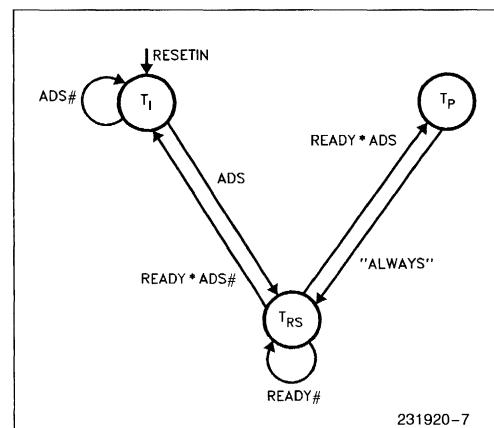


Figure 3.3. Bus State Diagram

The following sections illustrate different types of 80387 bus cycles.

Because different instructions have different amounts of overhead before, between, and after operand transfer cycles, it is not possible to represent in a few diagrams all of the combinations of successive operand transfer cycles. The following bus-cycle diagrams show memory cycles between 80387 operand-transfer cycles. Note however that, during the instructions FLDENV, FSTENV, FSAVE, and FRSTOR, some consecutive accesses to the NPX do not have intervening memory accesses. For the timing relationship between operand transfer cycles and opcode write or other overhead activities, see Figure 3.7.

3.4.1 NONPIPELINED BUS CYCLES

Figure 3.4 illustrates bus activity for consecutive nonpipelined bus cycles.

3.4.1.1 Write Cycle

At the second clock of the bus cycle, the 80387 enters the T_{RS} (READY#-sensitive) state. During this state, the 80387 samples the READY# input and stays in this state as long as READY# is inactive.

In write cycles, the 80387 drives the READYO# signal for one CLK period beginning with the second CLK of the bus cycle; therefore, the fastest write cycle takes two CLK cycles (see cycle 2 of Figure 3.4). For the instructions FLDENV and FRSTOR, however, the 80387 forces a wait state by delaying the activation of READYO# to the second T_{RS} cycle (not shown in Figure 3.4).

When READY# is asserted the 80387 returns to the idle state, in which ADS# could be asserted again by the 80386 for the next cycle.

3.4.1.2 Read Cycle

At the second clock of the bus cycle, the 80387 enters the T_{RS} state. See Figure 3.4. In this state, the 80387 samples the READY# input and stays in this state as long as READY# is inactive.

At the rising edge of CLK in the second clock period of the cycle, the 80387 starts to drive the D31-D0 outputs and continues to drive them as long as it stays in T_{RS} state.

In read cycles that address the 80387, at least one wait state must be inserted to insure that the 80386 latches the correct data. Since the 80387 starts driving the system data bus only at the rising edge of CLK in the second clock period of the bus cycle, not enough time is left for the data signals to propagate and be latched by the 80386 at the falling edge of the same clock period. The 80387 drives the READYO# signal for one CLK period in the third CLK of the bus cycle. Therefore, if the READYO# output is used to drive the 80386 READY# input, one wait state is inserted automatically.

Because one wait state is required for 80387 reads, the minimum is three CLK cycles per read, as cycle 3 of Figure 3.4 shows.

When READY# is asserted the 80387 returns to the idle state, in which ADS# could be asserted again by the 80386 for the next cycle. The transition from T_{RS} state to idle state causes the 80387 to put the tristate D31-D0 outputs into the floating state, allowing another device to drive the system data bus.

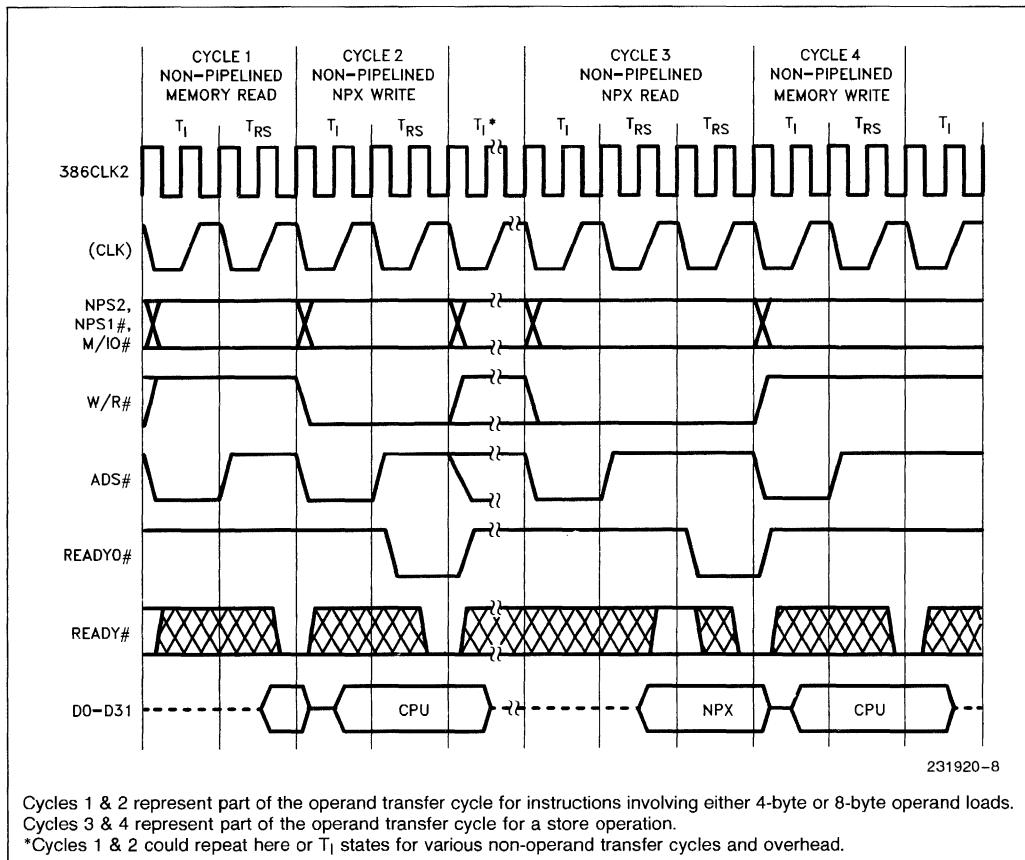


Figure 3.4. Nonpipelined Read and Write Cycles

3.4.2 PIPELINED BUS CYCLES

Because all the activities of the 80387 bus interface occur either during the T_{RS} state or during the transitions to or from that state, the only difference between a pipelined and a nonpipelined cycle is the manner of changing from one state to another. The exact activities in each state are detailed in the previous section "Nonpipelined Bus Cycles".

When the 80386 asserts ADS# before the end of a bus cycle, both ADS# and READY# are active during a T_{RS} state. This condition causes the 80387 to change to a different state named T_P. The 80387 activities in the transition from a T_{RS} state to a T_P state are exactly the same as those in the transition from a T_{RS} state to a T_I state in nonpipelined cycles.

T_P state is metastable; therefore, one clock period later the 80387 returns to T_{RS} state. In consecutive pipelined cycles, the 80387 bus logic uses only T_{RS} and T_P states.

Figure 3.5 shows the fastest transition into and out of the pipelined bus cycles. Cycle 1 in this figure represents a nonpipelined cycle. (Nonpipelined write cycles with only one T_{RS} state (i.e. no wait states) are always followed by another nonpipelined cycle, because READY# is asserted before the earliest possible assertion of ADS# for the next cycle.)

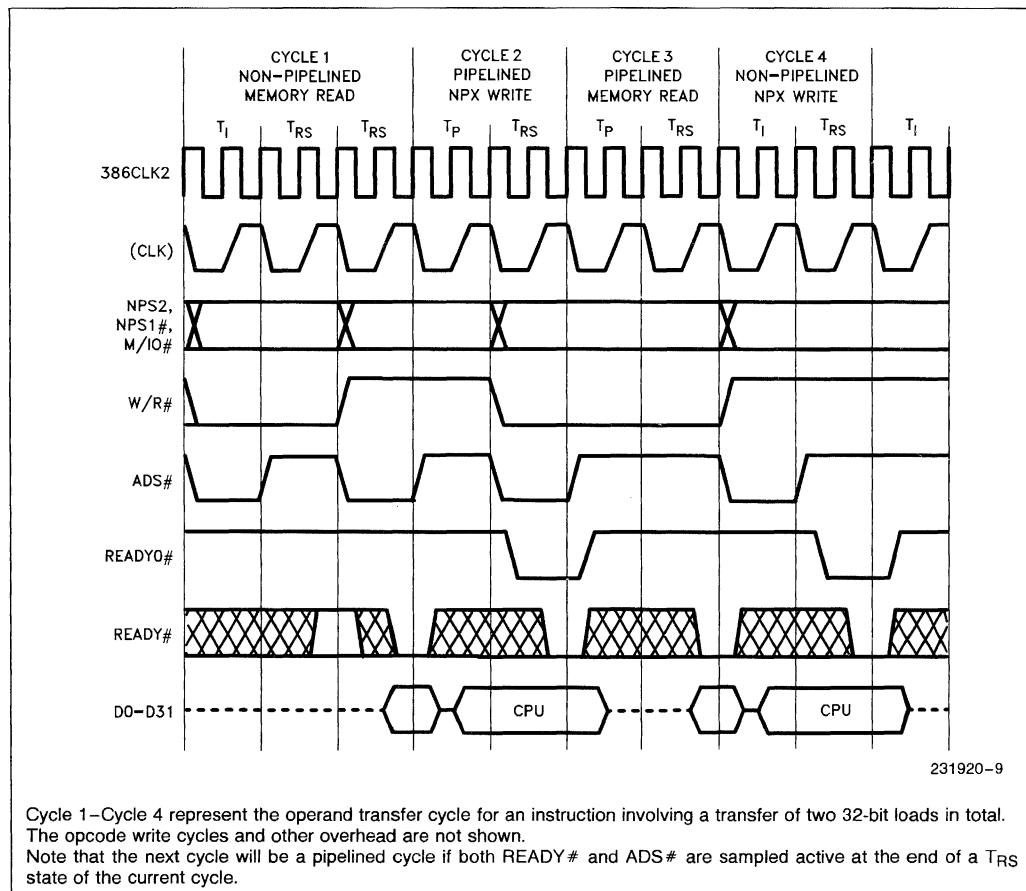
Figure 3.6 shows the pipelined write and read cycles with one additional T_{RS} states beyond the minimum required. To delay the assertion of READY# requires external logic.

3.4.3 BUS CYCLES OF MIXED TYPE

When the 80387 bus logic is in the T_{RS} state, it distinguishes between nonpipelined and pipelined cycles according to the behavior of ADS# and READY#. In a nonpipelined cycle, only READY# is activated, and the transition is from T_{RS} to idle state. In a pipelined cycle, both READY# and ADS# are active and the transition is first from T_{RS} state to T_P state then, after one clock period, back to T_{RS} state.

3.4.4 BUSY# AND PEREQ TIMING RELATIONSHIP

Figure 3.7 shows the activation of BUSY# at the beginning of instruction execution and its deactivation after execution of the instruction is complete. PEREQ is activated in this interval. If ERROR# (not shown in the diagram) is ever asserted, it would occur at least six 386CLK2 periods after the deactivation of PEREQ and at least six 386CLK2 periods before the deactivation of BUSY#. Figure 3.7 shows also that STEN is activated at the beginning of a bus cycle.



Cycle 1–Cycle 4 represent the operand transfer cycle for an instruction involving a transfer of two 32-bit loads in total. The opcode write cycles and other overhead are not shown.

Note that the next cycle will be a pipelined cycle if both READY# and ADS# are sampled active at the end of a T_{RS} state of the current cycle.

Figure 3.5. Fastest Transitions to and from Pipelined Cycles

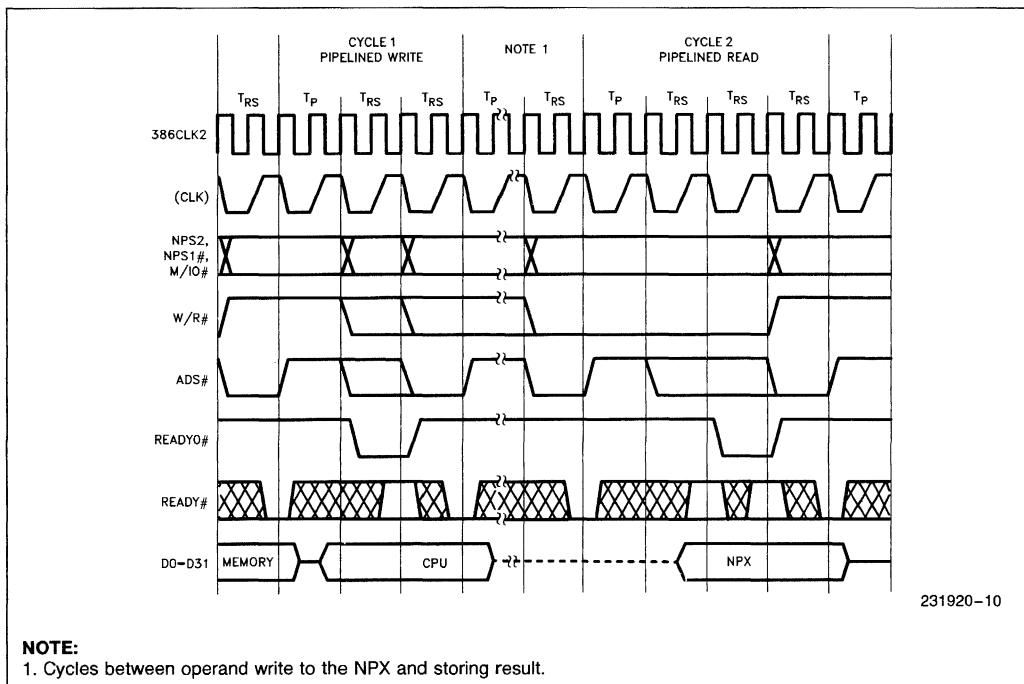


Figure 3.6. Pipelined Cycles with Wait States

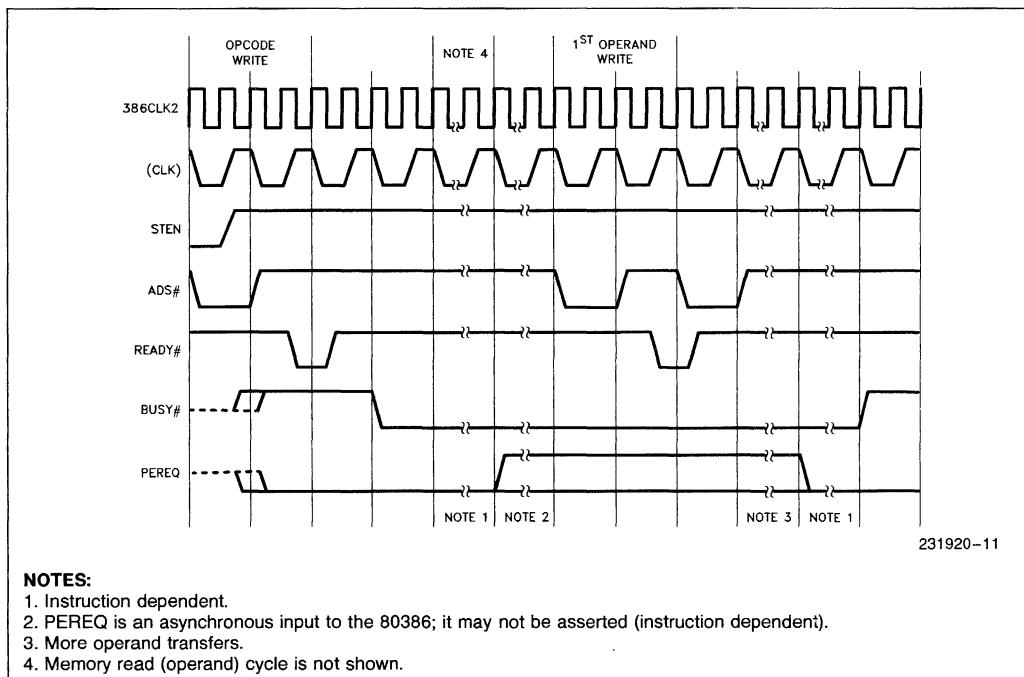
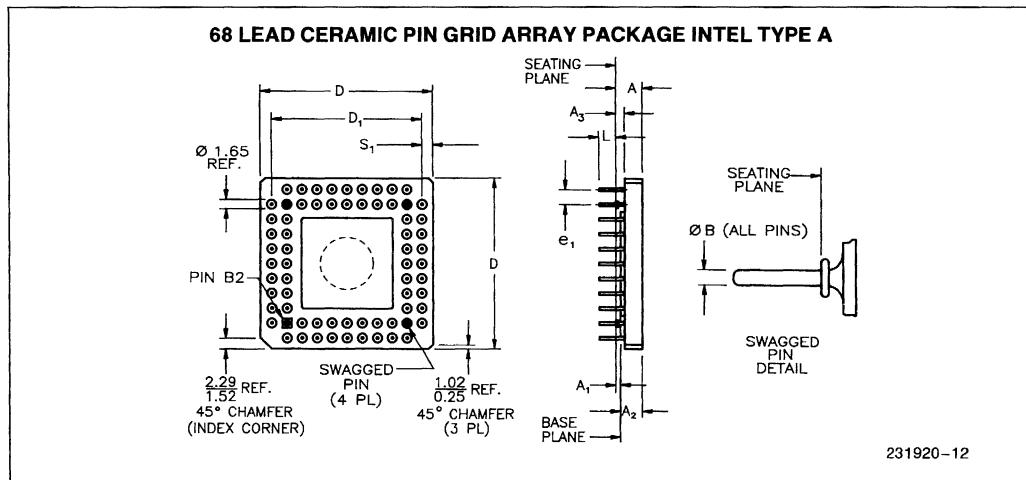


Figure 3.7. STEN, BUSY# and PEREQ Timing Relationship

4.0 MECHANICAL DATA



Family: Ceramic Pin Grid Array Package						
Symbol	Millimeters			Inches		
	Min	Max	Notes	Min	Max	Notes
A	3.56	4.57		0.140	0.180	
A ₁	0.76	1.27	Solid Lid	0.030	0.050	Solid Lid
A ₁		0.41	EPROM Lid		0.016	EPROM Lid
A ₂	2.72	3.43	Solid Lid	0.107	0.135	Solid Lid
A ₂	3.43	4.32	EPROM Lid	0.135	0.170	EPROM Lid
A ₃	1.14	1.40		0.045	0.055	
B	0.43	0.51		0.017	0.020	
D	28.83	29.59		1.135	1.165	
D ₁	25.27	25.53		0.995	1.005	
e ₁	2.29	2.79		0.090	0.110	
L	2.29	3.30		0.090	0.130	
N	68			68		
S ₁	1.27	2.54		0.050	0.100	
ISSUE	IWS REV 7 3/26/86					

Figure 4.1. Package Description

Consult the most recent 80387 data sheet for AC specifications.

Consult the most recent 80387 data sheet for AC specifications.

Consult the most recent 80387 data sheet for AC specifications.

Consult the most recent 80387 data sheet for AC specifications.

Consult the most recent 80387 data sheet for AC specifications.

Instruction										Optional Fields	
First Byte				Second Byte							
1	11011	OPA		1	MOD		1	OPB	R/M	SIB	DISP
2	11011	MF		OPA	MOD		OPB		R/M	SIB	DISP
3	11011	d	P	OPA	1	1	OPB		ST(i)		
4	11011	0	0	1	1	1	1	OP			
5	11011	0	1	1	1	1	1	OP			
	15-11	10	9	8	7	6	5	4	3	2	1 0

6.0 80387 EXTENSIONS TO THE 80386 INSTRUCTION SET

Instructions for the 80387 assume one of the five forms shown in the following table. In all cases, instructions are at least two bytes long and begin with the bit pattern 11011B, which identifies the ESCAPE class of instruction. Instructions that refer to memory operands specify addresses using the 80386 addressing modes.

OP = Instruction opcode, possible split into two fields OPA and OPB

MF = Memory Format
 00—32-bit real
 01—32-bit integer
 10—64-bit real
 11—16-bit integer

P = Pop
 0—Do not pop stack
 1—Pop stack after operation

ESC = 11011

d = Destination
 0—Destination is ST(0)
 1—Destination is ST(i)

R XOR d = 0—Destination (op) Source
 R XOR d = 1—Source (op) Destination

ST(i) = Register stack element i

000 = Stack top

001 = Second stack element

-

-

-

111 = Eighth stack element

MOD (Mode field) and R/M (Register/Memory specifier) have the same interpretation as the corresponding fields of 80386 instructions (refer to *80386 Programmer's Reference Manual*)

SIB (Scale Index Base) byte and DISP (displacement) are optionally present in instructions that have MOD and R/M fields. Their presence depends on the values of MOD and R/M, as for 80386 instructions.

The instruction summaries that follow assume that the instruction has been prefetched, decoded, and is ready for execution; that bus cycles do not require wait states; that there are no local bus HOLD request delaying processor access to the bus; and that no exceptions are detected during instruction execution. If the instruction has MOD and R/M fields that call for both base and index registers, add one clock.

80387 Extensions to the 80386 Instruction Set

Instruction	Encoding			Clock Count Range			
	Byte 0	Byte 1	Optional Bytes 2-6	32-Bit Real	32-Bit Integer	64-Bit Real	16-Bit Integer
DATA TRANSFER							
FLD = Load a							
Integer/real memory to ST(0)	ESC MF 1	MOD 000 R/M	SIB/DISP	20	45-52	25	61-65
Long integer memory to ST(0)	ESC 111	MOD 101 R/M	SIB/DISP		56-67		
Extended real memory to ST(0)	ESC 011	MOD 101 R/M	SIB/DISP		44		
BCD memory to ST(0)	ESC 111	MOD 100 R/M	SIB/DISP		266-275		
ST(i) to ST(0)	ESC 001	11000 ST(i)			14		
FST = Store							
ST(0) to integer/real memory	ESC MF 1	MOD 010 R/M	SIB/DISP	44	79-93	45	82-95
ST(0) to ST(i)	ESC 101	11010 ST(i)			11		
FSTP = Store and Pop							
ST(0) to integer/real memory	ESC MF 1	MOD 011 R/M	SIB/DISP	44	79-93	45	82-95
ST(0) to long integer memory	ESC 111	MOD 111 R/M	SIB/DISP		80-97		
ST(0) to extended real	ESC 011	MOD 111 R/M	SIB/DISP		53		
ST(0) to BCD memory	ESC 111	MOD 110 R/M	SIB/DISP		512-534		
ST(0) to ST(i)	ESC 101	11001 ST(i)			12		
FXCH = Exchange							
ST(i) and ST(0)	ESC 001	11001 ST(i)			18		
COMPARISON							
FCOM = Compare							
Integer/real memory to ST(0)	ESC MF 0	MOD 010 R/M	SIB/DISP	26	56-63	31	71-75
ST(i) to ST(0)	ESC 000	11010 ST(i)			24		
FCOMP = Compare and pop							
Integer/real memory to ST	ESC MF 0	MOD 011 R/M	SIB/DISP	26	56-63	31	71-75
ST(i) to ST(0)	ESC 000	11011 ST(i)			26		
FCOMPP = Compare and pop twice							
ST(1) to ST(0)	ESC 110	1101 1001			26		
FTST = Test ST(0)							
ESC 001	1110 0100				28		
FUCOM = Unordered compare							
ESC 101	11100 ST(i)				24		
FUCOMP = Unordered compare and pop							
ESC 101	11101 ST(i)				26		
FUCOMPP = Unordered compare and pop twice							
ESC 010	1110 1001				26		
FXAM = Examine ST(0)							
ESC 001	11100101				30-38		
CONSTANTS							
FLDZ = Load +0.0 into ST(0)	ESC 001	1110 1110			20		
FLD1 = Load +1.0 into ST(0)	ESC 001	1110 1000			24		
FLDPi = Load pi into ST(0)	ESC 001	1110 1011			40		
FLDL2T = Load log ₂ (10) into ST(0)	ESC 001	1110 1001			40		

Shaded areas indicate instructions not available in 8087/80287.

NOTE:

- a. When loading single- or double-precision zero from memory, add 5 clocks.

80387 Extensions to the 80386 Instruction Set (Continued)

Instruction	Encoding			Clock Count Range			
	Byte 0	Byte 1	Optional Bytes 2-6	32-Bit Real	32-Bit Integer	64-Bit Real	16-Bit Integer
CONSTANTS (Continued)							
FLDL2E = Load log ₂ (e) into ST(0)	ESC 001	1110 1010				40	
FLDLG2 = Load log ₁₀ (2) into ST(0)	ESC 001	1110 1100				41	
FLDLN2 = Load log _e (2) into ST(0)	ESC 001	1110 1101				41	
ARITHMETIC							
FADD = Add							
Integer/real memory with ST(0)	ESC MF 0	MOD 000 R/M	SIB/DISP	24-32	57-72	29-37	71-85
ST(i) and ST(0)	ESC d P 0	11000 ST(i)			23-31 ^b		
FSUB = Subtract							
Integer/real memory with ST(0)	ESC MF 0	MOD 10 R R/M	SIB/DISP	24-32	57-82	28-36	71-83 ^c
ST(i) and ST(0)	ESC d P 0	1110 R R/M			26-34 ^d		
FMUL = Multiply							
Integer/real memory with ST(0)	ESC MF 0	MOD 001 R/M	SIB/DISP	27-35	61-82	32-57	76-87
ST(i) and ST(0)	ESC d P 0	1100 1 R/M			29-57 ^e		
FDIV = Divide							
Integer/real memory with ST(0)	ESC MF 0	MOD 11 R R/M	SIB/DISP	89	120-127 ^f	94	136-140 ^g
ST(i) and ST(0)	ESC d P 0	1111 R R/M			88 ^h		
FSQRTⁱ = Square root	ESC 001	1111 1010				122-129	
FSCALE = Scale ST(0) by ST(1)	ESC 001	1111 1101				67-86	
FPREM = Partial remainder	ESC 001	1111 1000				74-155	
FPREM1 = Partial remainder (IEEE)	ESC 001	1111 0101				95-185	
FRNDINT = Round ST(0) to integer	ESC 001	1111 1100				66-80	
FXTRACT = Extract components of ST(0)	ESC 001	1111 0100				70-76	
FABS = Absolute value of ST(0)	ESC 001	1110 0001				22	
FCHS = Change sign of ST(0)	ESC 001	1110 0000				24-25	

Shaded areas indicate instructions not available in 8087/80287.

NOTES:

- b. Add 3 clocks to the range when d = 1.
- c. Add 1 clock to **each** range when R = 1.
- d. Add 3 clocks to the range when d = 0.
- e. typical = 52 (When d = 0, 46-54, typical = 49).
- f. Add 1 clock to the range when R = 1.
- g. 135-141 when R = 1.
- h. Add 3 clocks to the range when d = 1.
- i. $-0 \leq \text{ST}(0) \leq +\infty$.

80387 Extensions to the 80386 Instruction Set (Continued)

Instruction	Encoding			Clock Count Range
	Byte 0	Byte 1	Optional Bytes 2-6	
TRANSCENDENTAL				
FCOS^k = Cosine of ST(0)	ESC 001	1111 1111		123-772i
FPTAN^k = Partial tangent of ST(0)	ESC 001	1111 0010		191-497i
FPATAN = Partial arctangent	ESC 001	1111 0011		314-487
FSIN^k = Sine of ST(0)	ESC 001	1111 1110		122-771i
FSINCOS^k = Sine and cosine of ST(0)	ESC 001	1111 1011		194-809i
F2XM1^l = 2 ^{ST(0) - 1}	ESC 001	1111 0000		211-476
FYL2X^m = ST(1) * log ₂ (ST(0))	ESC 001	1111 0001		120-538
FYL2XP1ⁿ = ST(1) * log ₂ (ST(0) + 1.0)	ESC 001	1111 1001		257-547
PROCESSOR CONTROL				
FINIT = Initialize NPX	ESC 011	1110 0011		33
FSTSW AX = Store status word	ESC 111	1110 0000		13
FLDCW = Load control word	ESC 001	MOD 101 R/M	SIB/DISP	19
FSTCW = Store control word	ESC 101	MOD 111 R/M	SIB/DISP	15
FSTSW = Store status word	ESC 101	MOD 111 R/M	SIB/DISP	15
FCLEX = Clear exceptions	ESC 011	1110 0010		11
FSTENV = Store environment	ESC 001	MOD 110 R/M	SIB/DISP	103-104
FLDENV = Load environment	ESC 001	MOD 100 R/M	SIB/DISP	71
FSAVE = Save state	ESC 101	MOD 110 R/M	SIB/DISP	375-376
FRSTOR = Restore state	ESC 101	MOD 100 R/M	SIB/DISP	308
FINCSTP = Increment stack pointer	ESC 001	1111 0111		21
FDECSTP = Decrement stack pointer	ESC 001	1111 0110		22
FFREE = Free ST(i)	ESC 101	1100 0 ST(i)		18
FNOP = No operations	ESC 001	1101 0000		12

Shaded areas indicate instructions not available in 8087/80287.

NOTES:

- j. These timings hold for operands in the range $|x| < \pi/4$. For operands not in this range, up to 76 additional clocks may be needed to reduce the operand.
- k. $0 \leq |ST(0)| < 2^{63}$.
- l. $-1.0 \leq ST(0) \leq 1.0$.
- m. $0 \leq ST(0) < \infty, -\infty < ST(1) < +\infty$.
- n. $0 \leq |ST(0)| < (2 - SQRT(2))/2, -\infty < ST(1) < +\infty$.

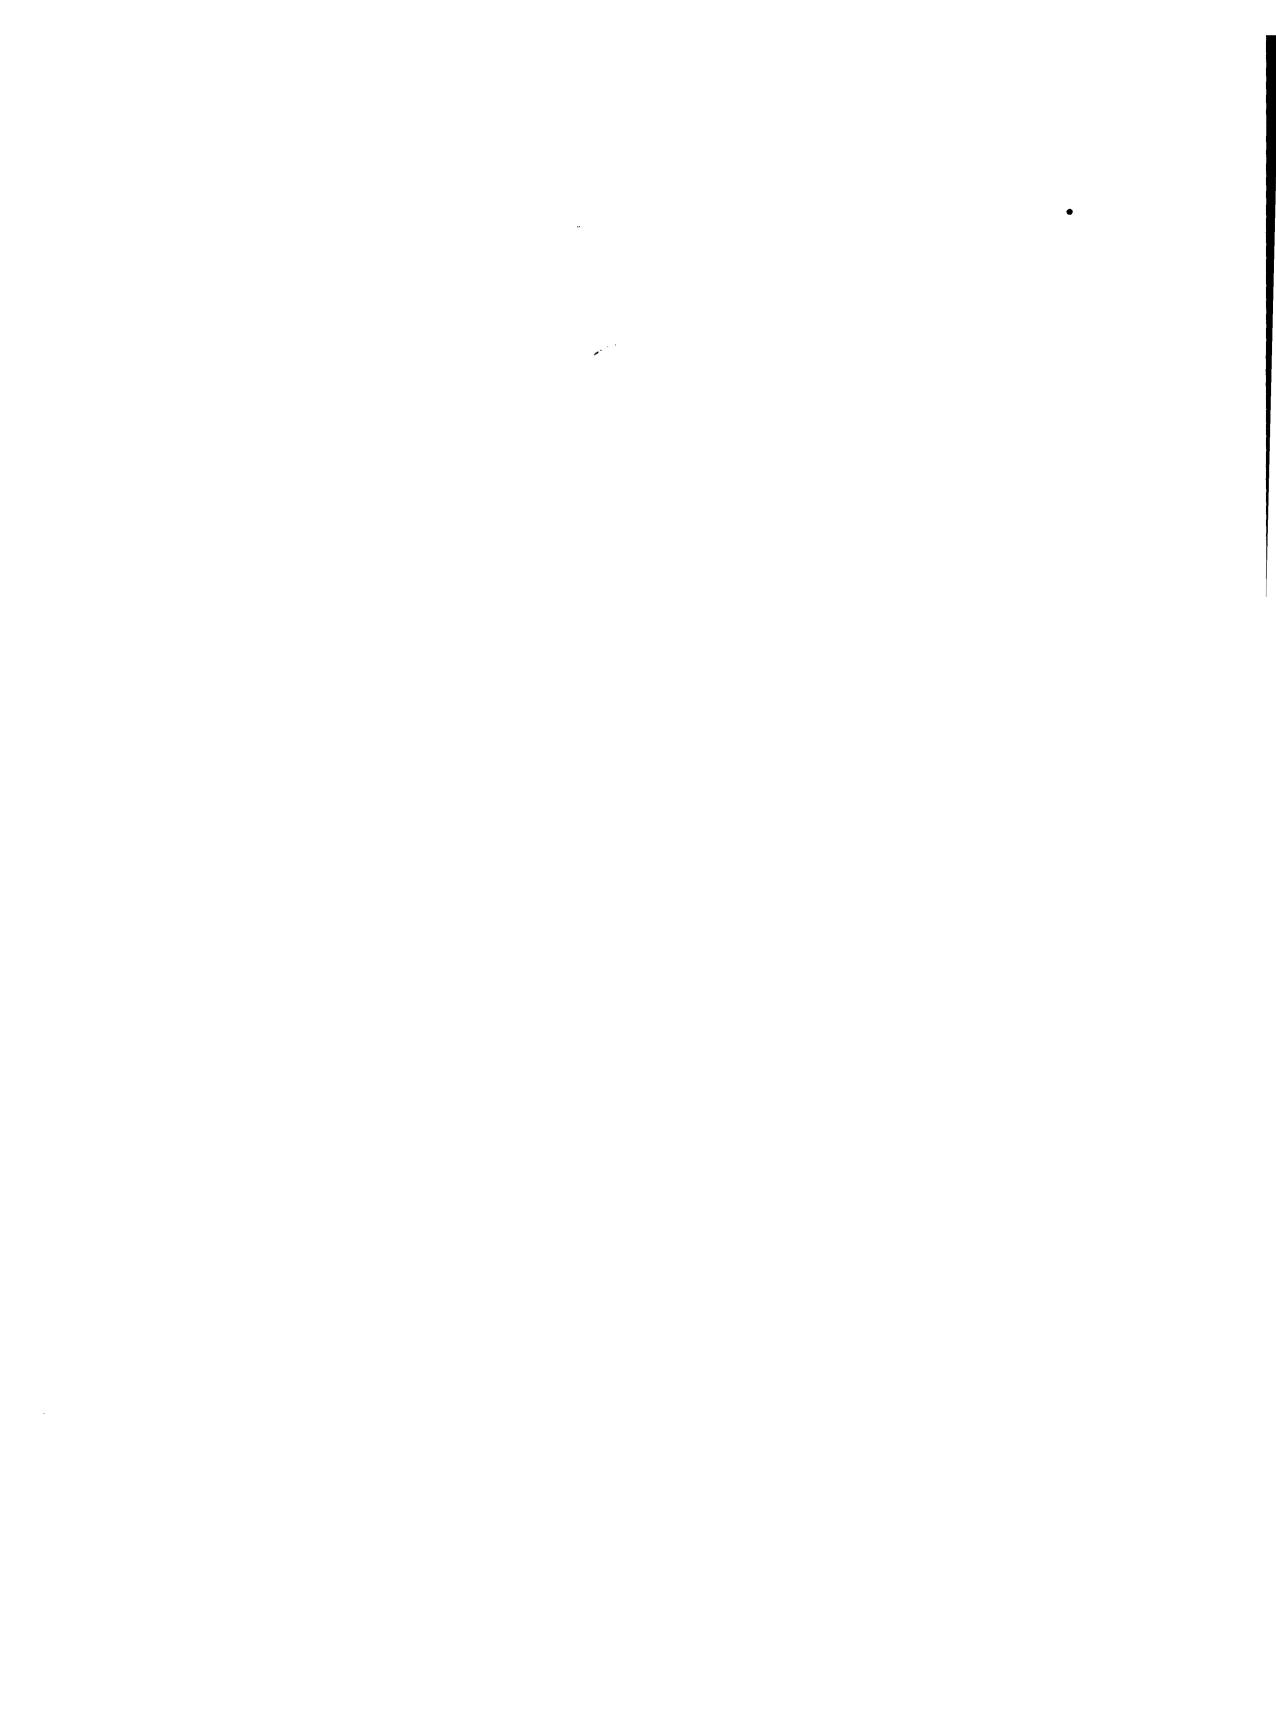
APPENDIX A COMPATIBILITY BETWEEN THE 80287 AND THE 8087

The 80286/80287 operating in Real-Address mode will execute 8086/8087 programs without major modification. However, because of differences in the handling of numeric exceptions by the 80287 NPX and the 8087 NPX, exception-handling routines *may* need to be changed.

This appendix summarizes the differences between the 80287 NPX and the 8087 NPX, and provides details showing how 8086/8087 programs can be ported to the 80286/80287.

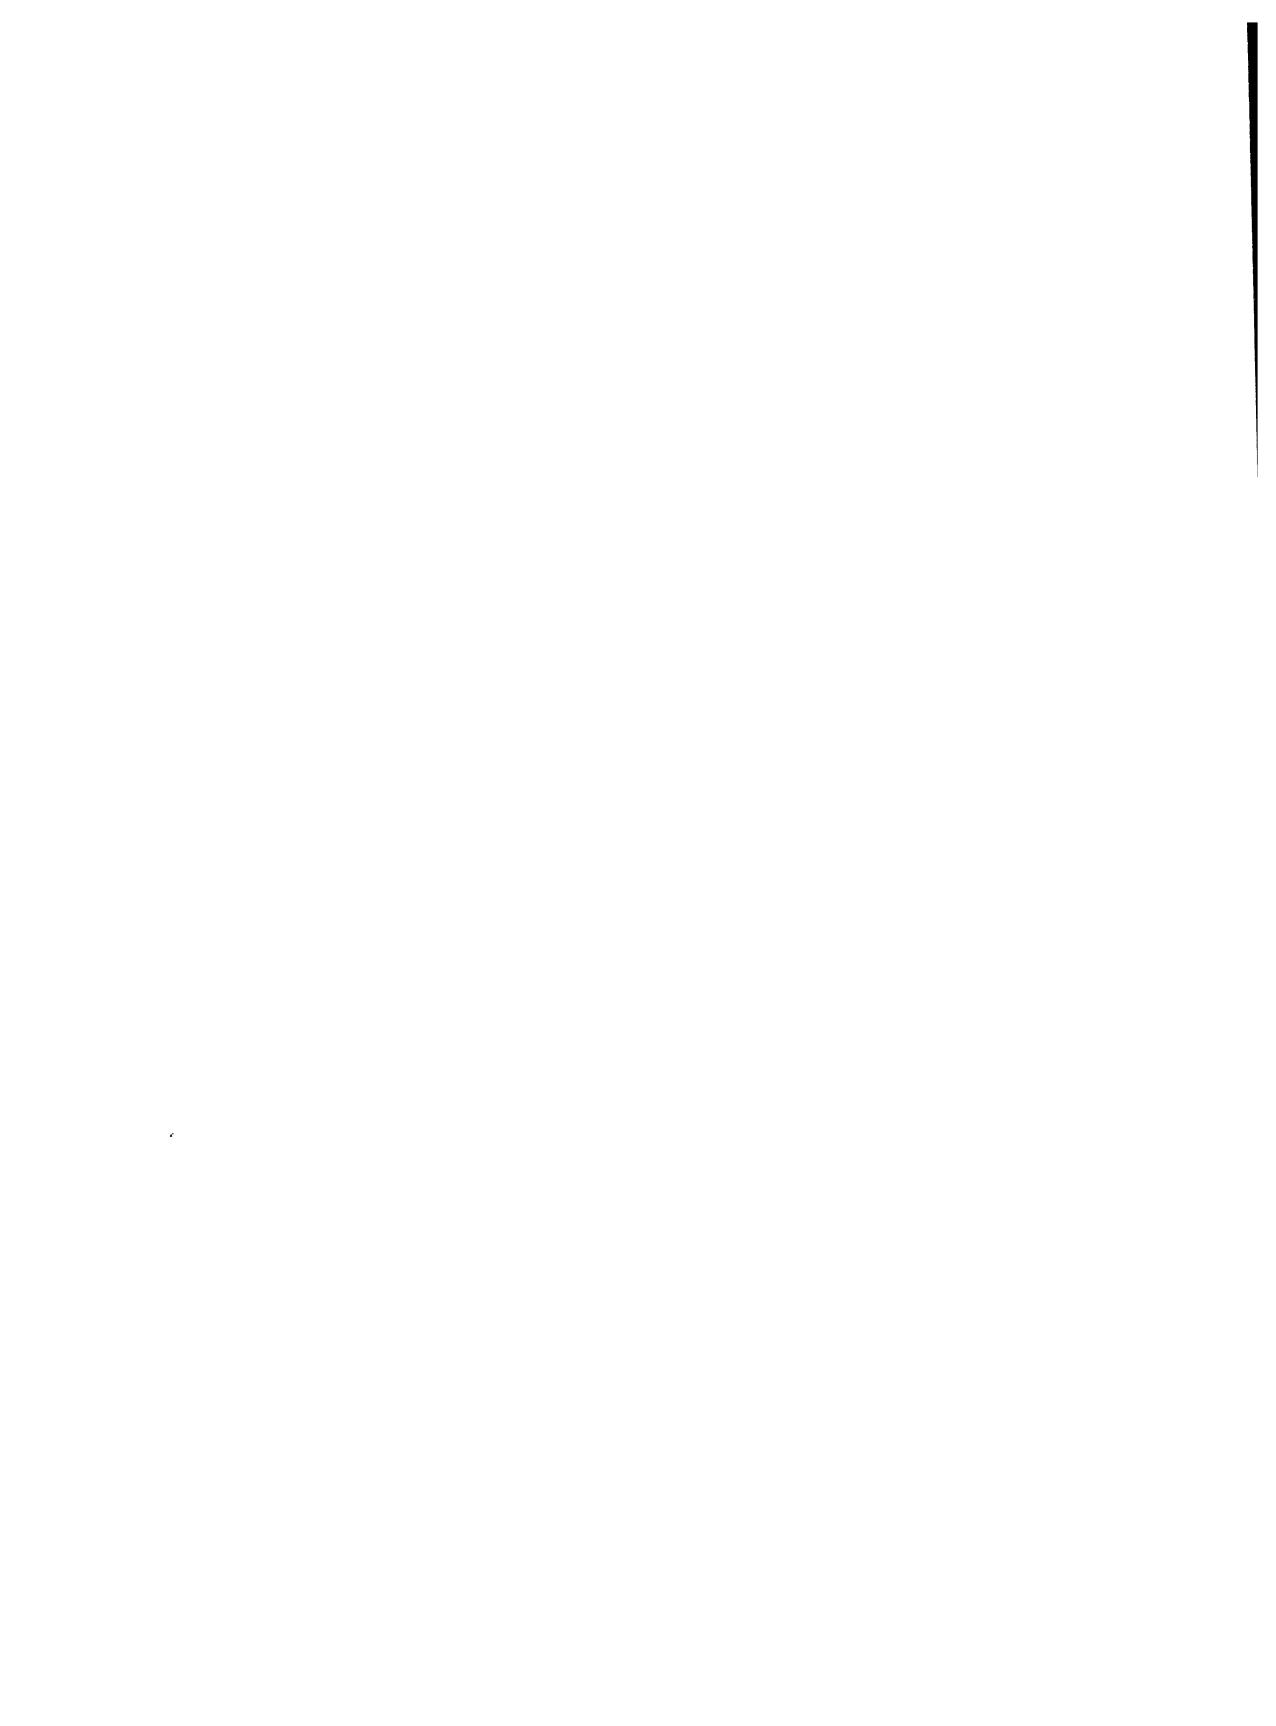
1. The NPX signals exceptions through a dedicated ERROR line to the 80286. The NPX error signal does not pass through an interrupt controller (the 8087 INT signal does). Therefore, any interrupt-controller-oriented instructions in numeric exception handlers for the 8086/8087 should be deleted.
2. The 8087 instructions FENI/FNENI and FDISI/FNDISI perform no useful function in the 80287. If the 80287 encounters one of these opcodes in its instruction stream, the instruction will effectively be ignored—none of the 80287 internal states will be updated. While 8086/8087 containing these instructions may be executed on the 80286/80287, it is unlikely that the exception-handling routines containing these instructions will be completely portable to the 80287.
3. Interrupt vector 16 must point to the numeric exception handling routine.
4. The ESC instruction address saved in the 80287 includes any leading prefixes before the ESC opcode. The corresponding address saved in the 8087 does not include leading prefixes.
5. In Protected-Address mode, the format of the 80287's saved instruction and address pointers is different than for the 8087. The instruction opcode is not saved in Protected mode—exception handlers will have to retrieve the opcode from memory if needed.
6. Interrupt 7 will occur in the 80286 when executing ESC instructions with either TS (task switched) or EM (emulation) of the 80286 MSW set (TS = 1 or EM = 1). If TS is set, then a WAIT instruction will also cause interrupt 7. An exception handler should be included in 80286/80287 code to handle these situations.
7. Interrupt 9 will occur if the second or subsequent words of a floating-point operand fall outside a segment's size. Interrupt 13 will occur if the starting address of a numeric operand falls outside a segment's size. An exception handler should be included in 80286/80287 code to report these programming errors.
8. Except for the processor control instructions, all of the 80287 numeric instructions are automatically synchronized by the 80286 CPU—the 80286 automatically tests the BUSY line from the 80287 to ensure that the 80287 has completed its previous instruction before executing the next ESC instruction. No explicit WAIT instructions are required to assure this synchronization. For the 8087 used with 8086 and 8088 processors, explicit WAITs are required before each numeric instruction to ensure synchronization. Although 8086/8087 programs having explicit WAIT instructions will execute perfectly on the 80286/80287 without reassembly, these WAIT instructions are unnecessary.
9. Since the 80287 does not require WAIT instructions before each numeric instruction, the ASM286 assembler does not automatically generate these WAIT instructions. The ASM86 assembler, however, automatically precedes every ESC instruction with a WAIT instruction. Although numeric routines generated using the ASM86 assembler will generally execute correctly on the 80286/80287, reassembly using ASM286 may result in a more compact code image.

The processor control instructions for the 80287 may be coded using either a WAIT or No-WAIT form of mnemonic. The WAIT forms of these instructions cause ASM286 to precede the ESC instruction with a CPU WAIT instruction, in the identical manner as does ASM86.



*PC/AT-Compatible
80387 Connection*

F



APPENDIX F

PC/AT*-COMPATIBLE 80387 CONNECTION

The PC/AT uses a nonstandard scheme to report 80287 exceptions to the 80286. When replicating the PC/AT coprocessor interface in 80386-based systems, the PC/AT interface cannot be used in exactly the same way; however, this appendix outlines a similar interface that works on 80386/80387 systems and maintains compatibility with the nonstandard PC/AT scheme.

Note that the interface outlined here does not represent a new interface standard; it needs to be incorporated in AT-compatible designs only because the 80286 and 80287 in the PC/AT are not connected according to the standards defined by Intel. The standard 80386/80387 connection recommended by Intel in the 80387 Data Sheet functions properly; the 80386 implementation has not been and will not be altered.

F.1 THE PC/AT INTERFACE

In the PC/AT, the ERROR# input to the 80286 is tied inactive (high) permanently. The ERROR# output of the 80287 is tied to an interrupt port (IRQ13). This interrupt replaces exception signaling via the 80286's ERROR# input. To guarantee (in the case of an 80287 exception) that INTR 13 will be serviced prior to the execution of any further 80287 instructions, an edge-triggered flip-flop latches BUSY# using ERROR# as a clock. The output of this latch is ORed with the BUSY# output of the 80287 and drives the BUSY# input of the 80286. This PC/AT scheme effectively delays deactivation of BUSY# at the 80286 whenever an 80287 ERROR# is signaled.

Since the 80286 BUSY# input remains active after an exception, the 80286 interrupt 13 handler is guaranteed to execute before any other 80287 instructions may begin. The interrupt 13 handler clears the BUSY# latch (via a write to a special I/O port), thus allowing execution of 80287 instructions to proceed. The interrupt 13 handler then branches to the NMI handler, where the user-defined numerics exception handler resides in PC-compatible systems.

The use of an interrupt guarantees that an exception from a coprocessor instruction will be detected. Latching BUSY# guarantees that any coprocessor instruction (except FINIT, FSETPM, and FCLEX) following the instruction that raised the exception will not be executed before the NMI handler is executed.

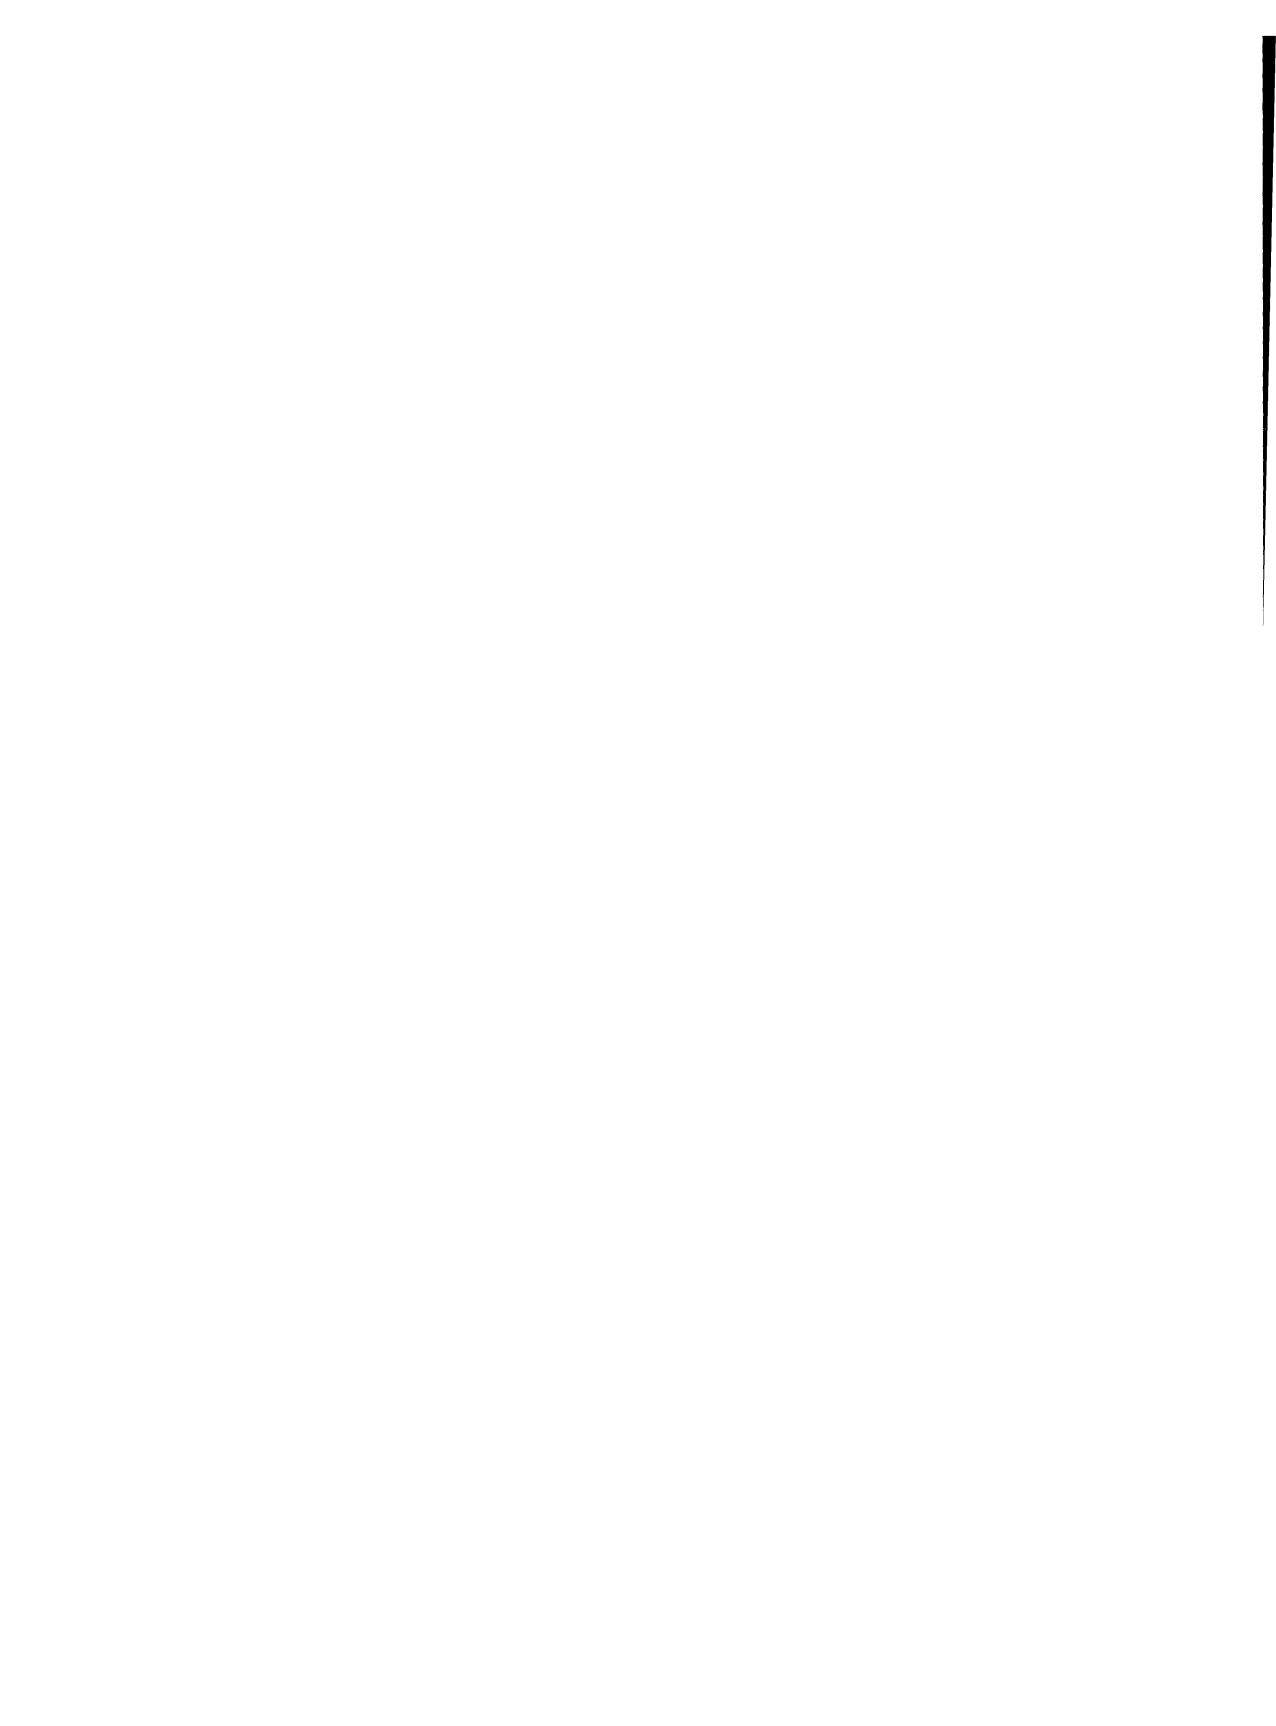
This PC/AT scheme approximates the exception reporting scheme between the 8087 and 8088 in the original PC.

F.2 HOW TO ACHIEVE THE SAME EFFECT IN AN 80386 SYSTEM

The 80386 can use a PC/AT-compatible interface to communicate with an 80387 provided that, when an NPX exception occurs, BUSY# active time is extended and PEREQ is reactivated only after 80387 BUSY# has gone inactive. The 80387 is left active (tying STEN high) at all times. Also, the 80386 and 80387 must be reset by the same RESET signal.

The reactivation of PEREQ for the 80386 is needed for store instructions (for example, FST *mem*) because the 80387 drops PEREQ once it signals an exception. While the 80386 has not yet recognized the occurrence of the exception, it still expects the data transfers to complete via PEREQ reactivation. It is permissible for the 80386 to receive undefined data during such I/O read cycles. Disabling the 80387 is not necessary, because the dummy data-transfer cycles directed to the 80387 when PEREQ is externally reactivated for the 80386 will not disturb the operation of the 80387. The interrupt 13 handler should remove the extension of BUSY# and reactivation of PEREQ via a write to PC/AT-compatible hardware at I/O port F0H.

Glossary of 80387 and Floating-Point Terminology



GLOSSARY OF 80387 AND FLOATING-POINT TERMINOLOGY

This glossary defines many terms that have precise technical meanings as specified in the IEEE 754 Standard or as specified in this manual. Where these terms are used, they have been italicized to emphasize the precision of their meanings. In reading these definitions, you may therefore interpret any italicized terms or phrases as cross-references.

Base: (1) a term used in logarithms and exponentials. In both contexts, it is a number that is being raised to a power. The two equations ($y = \log$ base b of x) and ($b^y = x$) are the same.

Base: (2) a number that defines the representation being used for a string of digits. *Base 2* is the binary representation; *base 10* is the decimal representation; *base 16* is the hexadecimal representation. In each case, the *base* is the factor of increased significance for each succeeding digit (working up from the bottom).

Bias: a constant that is added to the true exponent of a real number to obtain the *exponent* field of that number's *floating-point* representation in the 80387. To obtain the true *exponent*, you must subtract the *bias* from the given *exponent*. For example, the *single real* format has a *bias* of 127 whenever the given *exponent* is nonzero. If the 8-bit *exponent* field contains 10000011, which is 131, the true *exponent* is $131 - 127$, or +4.

Biased Exponent: the *exponent* as it appears in a *floating-point* representation of a number. The *biased exponent* is interpreted as an unsigned, positive number. In the above example, 131 is the *biased exponent*.

Binary Coded Decimal: a method of storing numbers that retains a *base 10* representation. Each decimal digit occupies 4 full bits (one hexadecimal digit). The hexadecimal values A through F (1010 through 1111) are not used. The 80387 supports a *packed decimal* format that consists of 9 bytes of *binary coded decimal* (18 decimal digits) and one sign byte.

Binary Point: an entity just like a decimal point, except that it exists in binary numbers. Each binary digit to the right of the *binary point* is multiplied by an increasing negative power of two.

C3—C0: the four "condition code" bits of the 80387 *status word*. These bits are set to certain values by the compare, test, examine, and remainder functions of the 80387.

Characteristic: a term used for some non-Intel computers, meaning the *exponent* field of a *floating-point* number.

Chop: to set one or more low-order bits of a real number to zero, yielding the nearest representable number in the direction of zero.

Condition Code: the four bits of the 80387 *status word* that indicate the results of the compare, test, examine, and remainder functions of the 80387.

Control Word: a 16-bit 80387 register that the user can set, to determine the modes of computation the 80387 will use and the exception interrupts that will be enabled.

Denormal: a special form of *floating-point* number. On the 80387, a *denormal* is defined as a number that has a *biased exponent* of zero. By providing a *significand* with leading zeros, the range of possible negative *exponents* can be extended by the number of bits in the *significand*. Each leading zero is a bit of lost accuracy, so the extended *exponent* range is obtained by reducing significance.

Double Extended: the Standard's term for the 80387's *extended format*, with more *exponent* and *significand* bits than the *double format* and an explicit *integer bit* in the *significand*.

Double Format: a *floating-point* format supported by the 80387 that consists of a sign, an 11-bit *biased exponent*, an *implicit integer bit*, and a 52-bit *significand*—a total of 64 explicit bits.

Environment: the 14 or 28 (depending on addressing mode) bytes of 80387 registers affected by the FSTENV and FLDENV instructions. It encompasses the entire state of the 80387, except for the 8 registers of the 80387 stack. Included are the *control word*, *status word*, *tag word*, and the instruction, opcode, and operand information provided by interrupts.

Exception: any of the six conditions (invalid operand, denormal, numeric overflow, numeric underflow, zero-divide, and precision) detected by the 80387 that may be signaled by status flags or by traps.

Exception Pointers: The data maintained by the 80386 to help exception handlers identify the cause of an exception. This data consists of a pointer to the most recently executed ESC instruction and a pointer to the memory operand of this instruction, if it had a memory operand. An exception handler can use the FSTENV and FSAVE instructions to access these pointers.

Exponent: (1) any number that indicates the power to which another number is raised.

Exponent: (2) the field of a *floating-point* number that indicates the magnitude of the number. This would fall under the above more general definition (1), except that a *bias* sometimes needs to be subtracted to obtain the correct power.

Extended Format: the 80387's implementation of the Standard's *double extended* format. *Extended format* is the main *floating-point* format used by the 80387. It consists of a sign, a 15-bit *biased exponent*, and a *significand* with an explicit *integer bit* and 63 fractional-part bits.

Floating-Point: of or pertaining to a number that is expressed as base, a sign, a significand, and a signed exponent. The value of the number is the signed product of its significand and the base raised to the power of the exponent. *Floating-point* representations are more versatile than *integer* representations in two ways. First, they include fractions. Second, their *exponent* parts allow a much wider range of magnitude than possible with fixed-length *integer* representations.

Gradual Underflow: a method of handling the *underflow* error condition that minimizes the loss of accuracy in the result. If there is a *denormal* number that represents the correct result, that *denormal* is returned. Thus, digits are lost only to the extent of denormalization. Most computers return zero when *underflow* occurs, losing all significant digits.

Implicit Integer Bit: a part of the *significand* in the *single real* and *double real* formats that is not explicitly given. In these formats, the entire given *significand* is considered to be to the right of the *binary point*. A single *implicit integer bit* to the left of the *binary point* is always one, except in one case. When the *exponent* is the minimum (*biased exponent* is zero), the *implicit integer bit* is zero.

Indefinite: a special value that is returned by functions when the inputs are such that no other sensible answer is possible. For each *floating-point* format there exists one *quiet NaN* that is designated as the *indefinite* value. For binary *integer* formats, the negative number furthest from zero is often considered the *indefinite* value. For the 80387 *packed decimal* format, the *indefinite* value contains all 1's in the sign byte and the uppermost digits byte.

Inexact: The Standard's term for the 80387's *precision exception*.

Infinity: a value that has greater magnitude than any *integer* or any *real* number. It is often useful to consider *infinity* as another number, subject to special rules of arithmetic. All three Intel *floating-point* formats provide representations for $+\infty$ and $-\infty$.

Integer: a number (positive, negative, or zero) that is finite and has no fractional part. *Integer* can also mean the computer representation for such a number: a sequence of data bytes, interpreted in a standard way. It is perfectly reasonable for *integers* to be represented in a *floating-point* format; this is what the 80387 does whenever an *integer* is pushed onto the 80387 stack.

Integer Bit: a part of the *significand* in *floating-point* formats. In these formats, the *integer bit* is the only part of the *significand* considered to be to the left of the *binary point*. The *integer bit* is always one, except in one case: when the *exponent* is the minimum (*biased exponent* is zero), the *integer bit* is zero. In the *extended format* the *integer bit* is explicit; in the *single format* and *double format* the *integer bit* is implicit; i.e., it is not actually stored in memory.

Invalid Operation: the exception condition for the 80387 that covers all cases not covered by other exceptions. Included are 80387 stack overflow and underflow, NaN inputs, illegal infinite inputs, out-of-range inputs, and inputs in unsupported formats.

Long Integer: an *integer* format supported by the 80387 that consists of a 64-bit *two's complement* quantity.

Long Real: an older term for the 80387's 64-bit *double format*.

Mantissa: a term used with some non-Intel computers for the *significand* of a *floating-point* number.

Masked: a term that applies to each of the six 80387 *exceptions* I,D,Z,O,U,P. An exception is *masked* if a corresponding bit in the 80387 *control word* is set to one. If an exception is *masked*, the 80387 will not generate an interrupt when the exception condition occurs; it will instead provide its own exception recovery.

Mode: One of the *status word* fields “rounding control” and “precision control” which programs can set, sense, save, and restore to control the execution of subsequent arithmetic operations.

NaN: an abbreviation for “Not a Number”; a *floating-point* quantity that does not represent any numeric or infinite quantity. *NaNs* should be returned by functions that encounter serious errors. If created during a sequence of calculations, they are transmitted to the final answer and can contain information about where the error occurred.

Normal: the representation of a number in a *floating-point* format in which the *significand* has an *integer bit* one (either *explicit* or *implicit*).

Normalize: convert a denormal representation of a number to a normal representation.

NPX: Numeric Processor Extension. This is the 80387, 80287, or 8087.

Overflow: an exception condition in which the correct answer is finite, but has magnitude too great to be represented in the destination format. This kind of overflow (also called numeric overflow) is not to be confused with stack overflow.

Packed Decimal: an *integer* format supported by the 80387. A *packed decimal* number is a 10-byte quantity, with nine bytes of 18 *binary coded decimal* digits and one byte for the sign.

Pop: to remove from a stack the last item that was placed on the stack.

Precision: The effective number of bits in the significand of the *floating-point* representation of a number.

Precision Control: an option, programmed through the 80387 *control word*, that allows all 80387 arithmetic to be performed with reduced precision. Because no speed advantage results from this option, its only use is for strict compatibility with the *standard* and with other computer systems.

Precision Exception: an 80387 *exception* condition that results when a calculation does not return an exact answer. This exception is usually *masked* and ignored; it is used only in extremely critical applications, when the user must know if the results are exact. The *precision exception* is called *inexact* in the *standard*.

Pseudozero: one of a set of special values of the *extended real* format. The set consists of numbers with a zero *significand* and an *exponent* that is neither all zeros nor all ones. *Pseudozeros* are not created by the 80387 but are handled correctly when encountered as operands.

Quiet NaN: a *NaN* in which the most significant bit of the fractional part of the *significand* is one. By convention, these *NaNs* can undergo certain operations without causing an exception.

Real: any finite value (negative, positive, or zero) that can be represented by a (possibly infinite) decimal expansion. *Reals* can be represented as the points of a line marked off like a ruler. The term *real* can also refer to a *floating-point* number that represents a *real* value.

Short Integer: an *integer* format supported by the 80387 that consists of a 32-bit *two's complement* quantity. *short integer* is not the shortest 80387 *integer* format—the 16-bit *word integer* is.

Short Real: an older term for the 80387's 32-bit *single format*.

Signaling NaN: a *NaN* that causes an *invalid-operation exception* whenever it enters into a calculation or comparison, even a nonordered comparison.

Significand: the part of a *floating-point* number that consists of the most significant nonzero bits of the number, if the number were written out in an unlimited binary format. The *significand* is composed of an *integer bit* and a *fraction*. The *integer bit* is implicit in the *single format* and *double format*. The *significand* is considered to have a *binary point* after the *integer bit*; the *binary point* is then moved according to the value of the *exponent*.

Single Extended: a *floating-point* format, required by the *standard*, that provides greater precision than *single*; it also provides an explicit *integer bit* in the *significand*. The 80387's *extended format* meets the *single extended* requirement as well as the *double extended* requirement.

Single Format: a *floating-point* format supported by the 80387, which consists of a sign, an 8-bit *biased exponent*, an *implicit integer bit*, and a 23-bit *significand*—a total of 32 explicit bits.

Stack Fault: a special case of the *invalid-operation exception* which is indicated by a one in the SF bit of the *status word*. This condition usually results from stack underflow or overflow.

Standard: “IEEE Standard for Binary Floating-Point Arithmetic,” ANSI/IEEE Std 754-1985.

Status Word: A 16-bit 80387 register that can be manually set, but which is usually controlled by side effects to 80387 instructions. It contains condition codes, the 80387 stack pointer, busy and interrupt bits, and exception flags.

Tag Word: a 16-bit 80387 register that is automatically maintained by the 80387. For each space in the 80387 stack, it tells if the space is occupied by a number; if so, it gives information about what kind of number.

Temporary Real: an older term for the 80387's 80-bit *extended format*.

Tiny: of or pertaining to a floating-point number that is so close to zero that its exponent is smaller than smallest exponent that can be represented in the destination format.

TOP: The three-bit field of the status word that indicates which 80387 register is the current top of stack.

Transcendental: one of a class of functions for which polynomial formulas are always approximate, never exact for more than isolated values. The 80387 supports trigonometric, exponential, and logarithmic functions; all are *transcendental*.

Two's Complement: a method of representing *integers*. If the uppermost bit is zero, the number is considered positive, with the value given by the rest of the bits. If the uppermost bit is one, the number is negative, with the value obtained by subtracting ($2^{\text{bit count}}$) from all the given bits. For example, the 8-bit number 11111100 is -4 , obtained by subtracting 2^8 from 252.

Unbiased Exponent: the true value that tells how far and in which direction to move the *binary point* of the *significand* of a *floating-point* number. For example, if a *single-format exponent* is 131, we subtract the Bias 127 to obtain the *unbiased exponent* +4. Thus, the *real* number being represented is the *significand* with the *binary point* shifted 4 bits to the right.

Underflow: an exception condition in which the correct answer is nonzero, but has a magnitude too small to be represented as a *normal* number in the destination *floating-point* format. The Standard specifies that an attempt be made to represent the number as a *denormal*. This denormalization may result in a loss of significant bits from the significand. This kind of underflow (also called numeric overflow) is not to be confused with stack underflow.

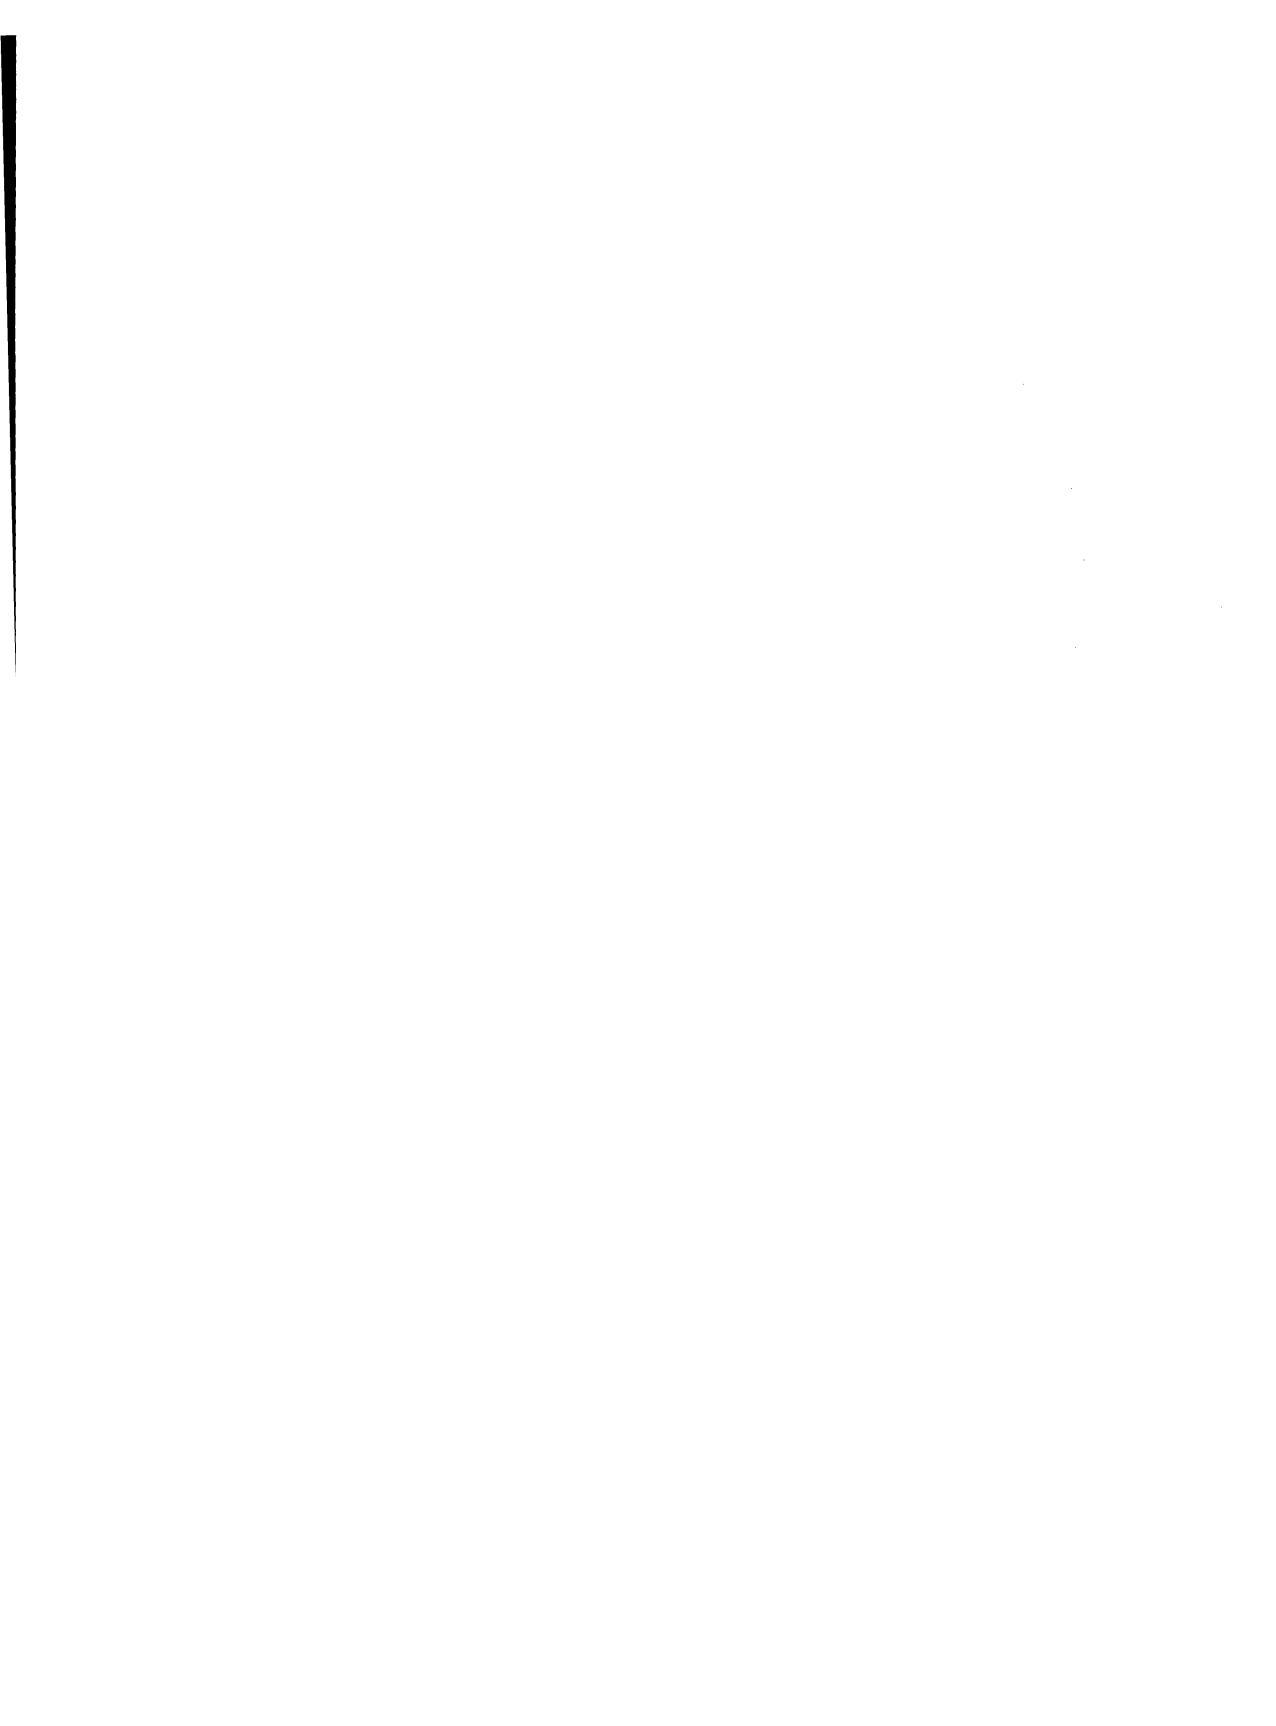
Unmasked: a term that applies to each of the six 80387 *exceptions*: I,D,Z,O,U,P. An exception is *unmasked* if a corresponding bit in the 80387 *control word* is set to zero. If an exception is *unmasked*, the 80387 will generate an interrupt when the exception condition occurs. You can provide an interrupt routine that customizes your exception recovery.

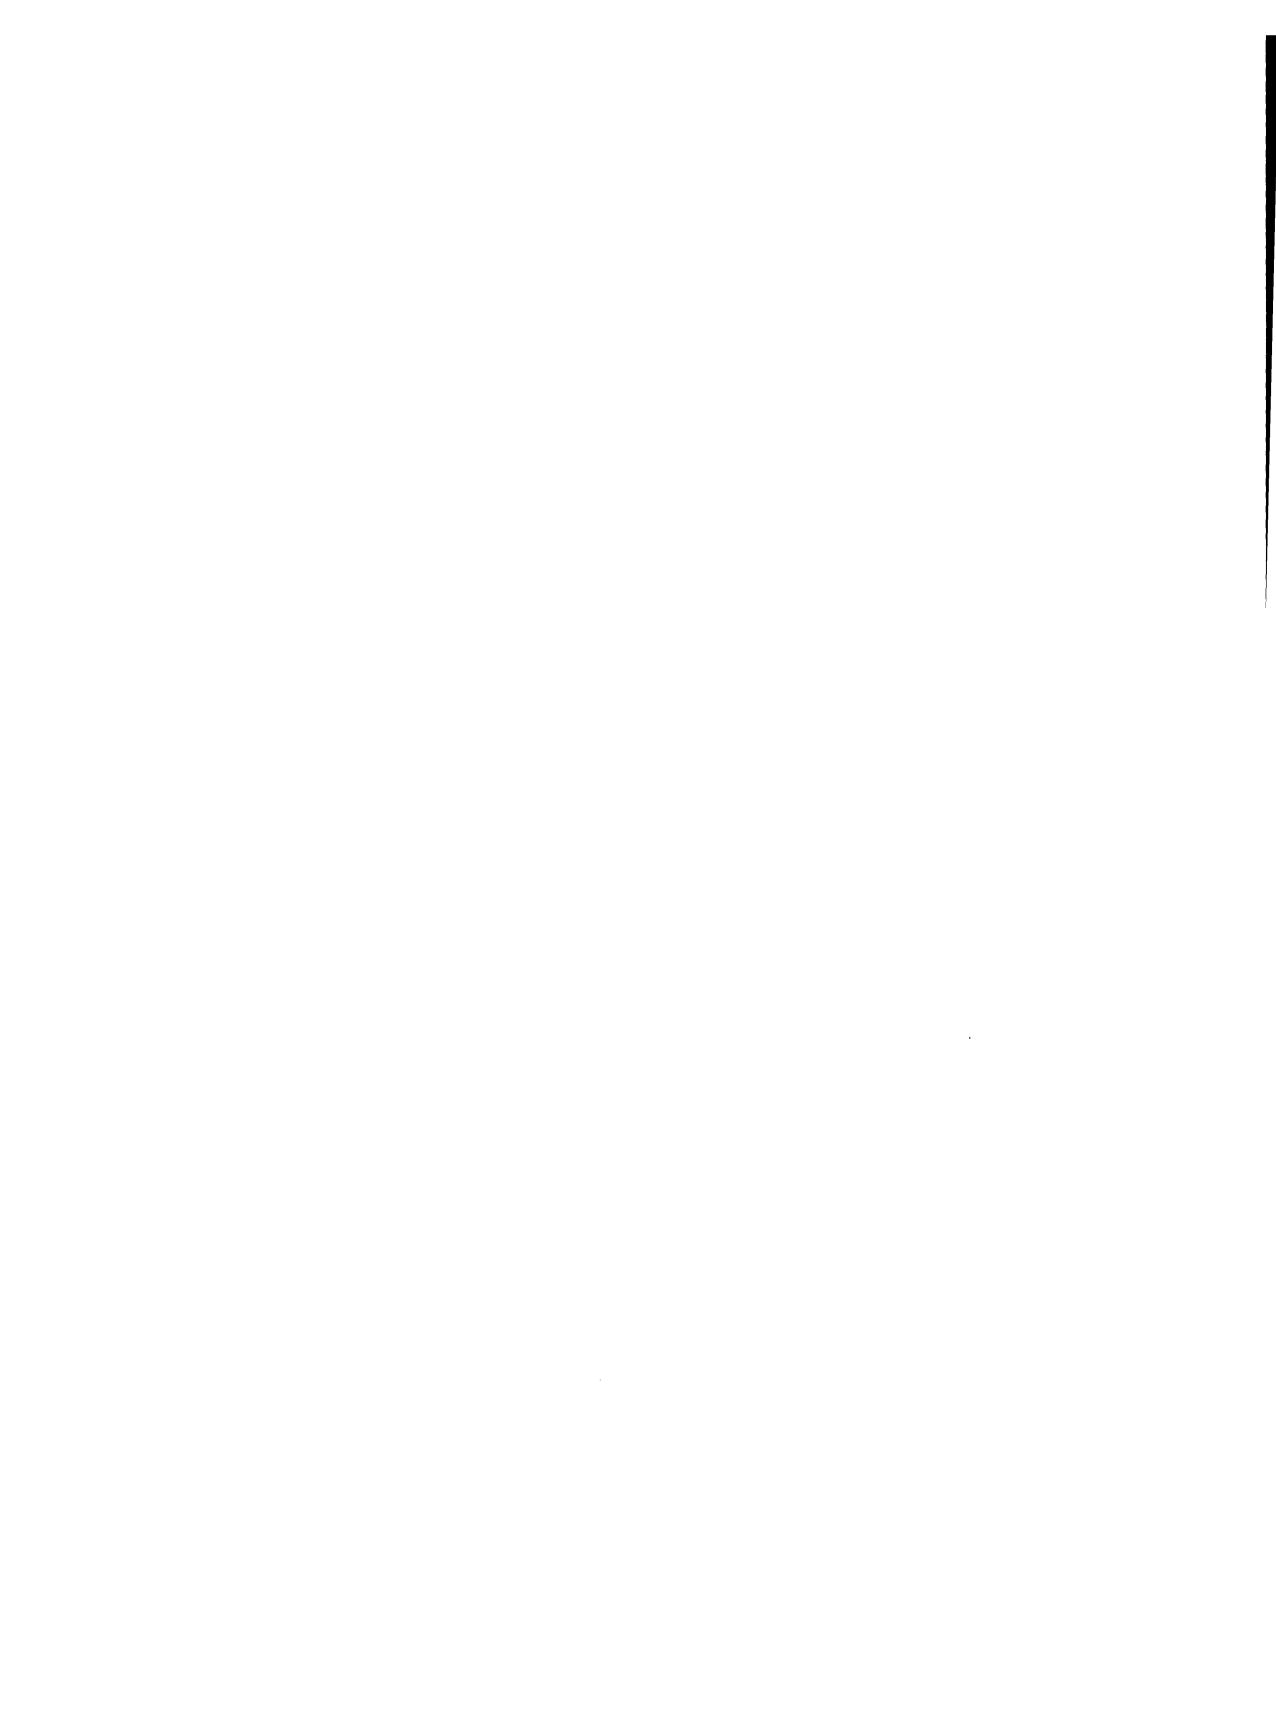
Unnormal: a *extended real* representation in which the explicit *integer* bit of the *significand* is zero and the exponent is nonzero. Unnormal values are not supported by the 80387; they cause the invalid-operation exception when encountered as operands.

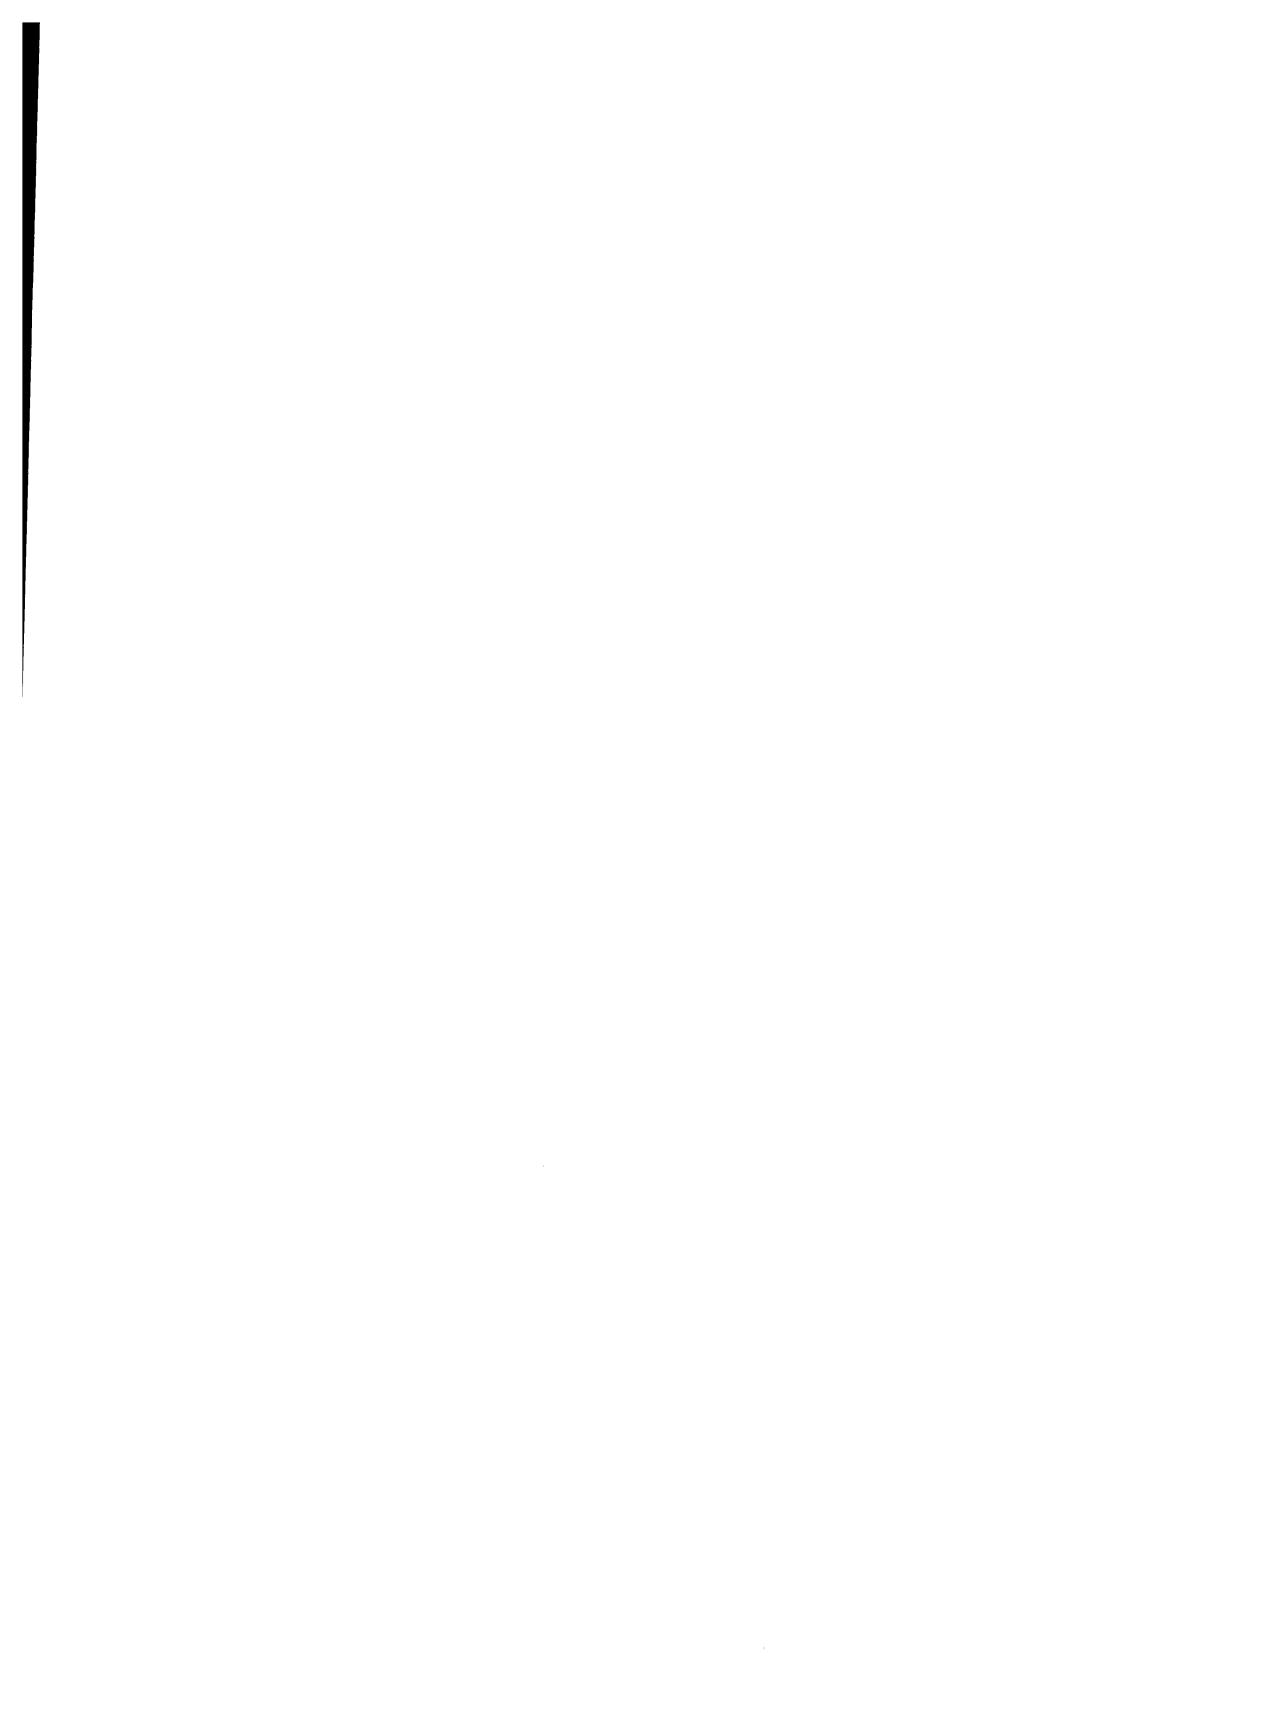
Unsupported Format: Any number representation that is not recognized by the 80387. This includes several formats that are recognized by the 8087 and 80287; namely: pseudo-NaN, pseudoinfinity, and unnormal.

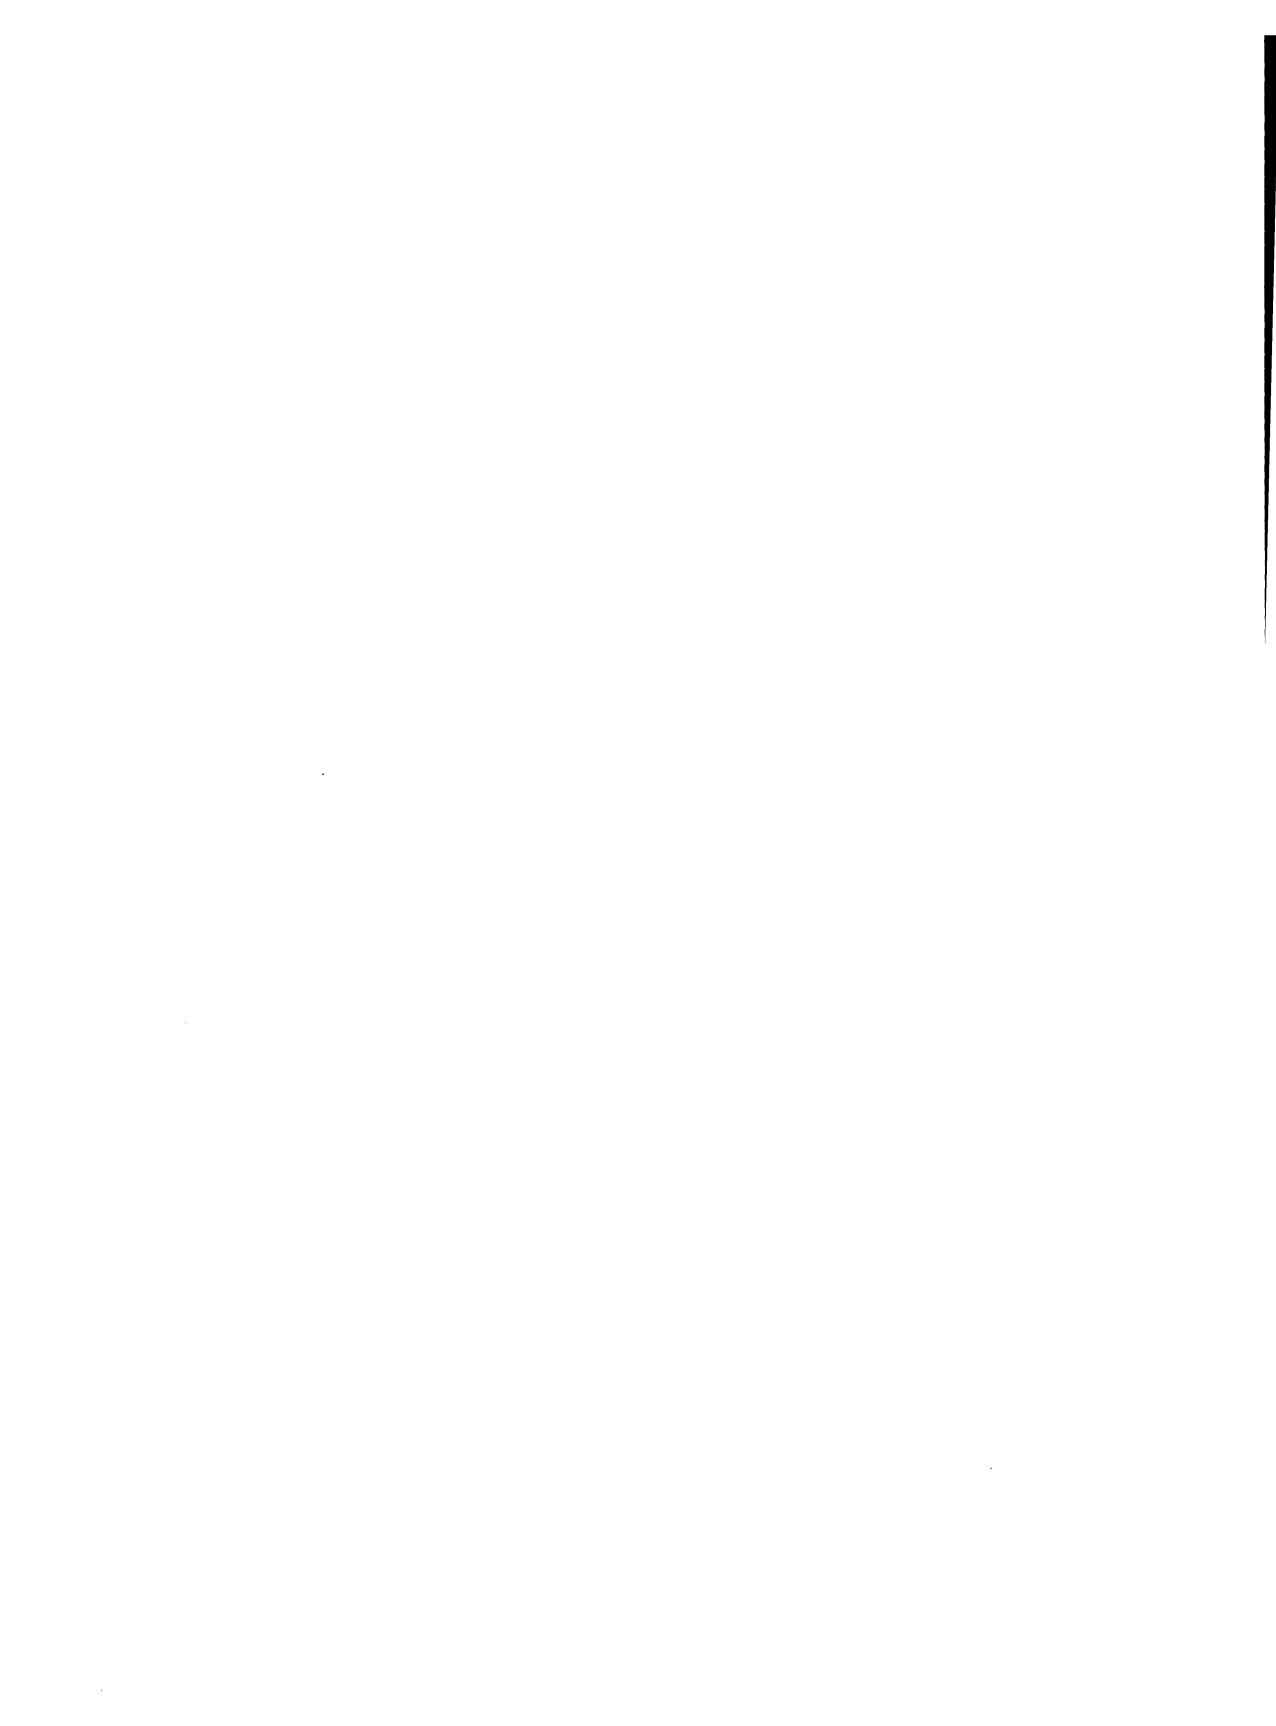
Word Integer: an *integer* format supported by both the 80386 and the 80387 that consists of a 16-bit *two's complement* quantity.

Zero divide: an *exception* condition in which the inputs are finite, but the correct answer, even with an unlimited *exponent*, has infinite magnitude.











DOMESTIC SALES OFFICES

ALABAMA

Intel Corp.
5015 Bradford Drive
Suite 200
Huntsville 35805
Tel: (205) 830-4010

ARIZONA

Intel Corp.
11225 N. 28th Drive
Suite 214D
Phoenix 85029
Tel: (602) 869-4980

Intel Corp.
1161 N. El Dorado Place
Suite 301
Tucson 85715
Tel: (602) 299-6815

CALIFORNIA

Intel Corp.
21515 Vanowen Street
Suite 116
Woodland Park 91303
Tel: (818) 704-8500

Intel Corp.
2250 Imperial Highway
Suite 218
El Segundo 90245
Tel: (213) 640-6040

Intel Corp.
1510 Arden Way, Suite 101
Sacramento 95815
Tel: (916) 920-8096

Intel Corp.
4350 Executive Drive
Suite 108
San Diego 92121
Tel: (619) 452-5980

Intel Corp.
400 N. Martin Avenue
Suite 450
Santa Ana 92705
Tel: (714) 835-9642
TWX: 910-595-1114

Intel Corp.
San Tomas 4
2700 San Tomas Expressway
Santa Clara, CA 95051
Tel: (408) 986-8086
TWX: 910-338-0255

COLORADO

Intel Corp.
4445 Northpark Drive
Suite 100
Colorado Springs 80907
Tel: (303) 594-6622

Intel Corp.
650 S. Cherry St., Suite 915
Denver 80222
Tel: (303) 321-8086
TWX: 910-931-2289

CONNECTICUT

Intel Corp.
26 Mill Plain Road
Danbury 06811
Tel: (203) 748-3130
TWX: 710-456-1199

FLORIDA

Intel Corp.
242 N. Westmonte Dr., Suite
100
Altamonte Springs 32714
Tel: (305) 869-5588
TWX: 710-456-1199

Intel Corp.
8385 N.W. 6th Way, Suite 100
P.O. Lauderdale 33309
Tel: (305) 771-0600
TWX: 510-956-9407

Intel Corp.
11300 4th Street North
Suite 170
St. Petersburg 33702
Tel: (813) 577-2413

GEORGIA

Intel Corp.
3250 Peachtree Parkway
Suite 200
Norcross 30092
Tel: (404) 449-0541

ILLINOIS

Intel Corp.*
300 N. Martingale Road, Suite 400
Schaumburg 60173
Tel: (312) 310-0311

INDIANA

Intel Corp.
8777 Purdue Road
Suite 125
Indianapolis 46228
Tel: (317) 875-0623

IOWA

Intel Corp.
St. Andrews Building
1930 St. Andrews Drive N.E.
Cedar Rapids 52402
Tel: (319) 393-5510

KANSAS

Intel Corp.
8400 W. 110th Street
Suite 170
Overland Park 66210
Tel: (913) 345-2727

MARYLAND

Intel Corp.*
7321 Parkway Drive South
Suite C
Hanover 21076
Tel: (301) 795-7500
TWX: 710-862-1944

MASSACHUSETTS

Intel Corp.*
Westford Corp. Center
3 Carlisle Road
Westford 01886
Tel: (317) 692-3222
TWX: 710-343-6333

MICHIGAN

Intel Corp.
7071 Orchard Lake Road
Suite 100
West Bloomfield 48033
Tel: (313) 851-8096

MINNESOTA

Intel Corp.
3500 W. 80th St., Suite 360
Bloomington 55431
Tel: (612) 833-6722
TWX: 910-576-2867

MISSOURI

Intel Corp.
4203 Earth City Expressway
Suite 131
Earth City 63045
Tel: (314) 291-1990

NEW JERSEY

Intel Corp.*
Parkway 109 Office Center
328 Newman Springs Road
Rte. Bank 07701
Tel: (201) 747-2233

Intel Corp.
280 Corporate Center
75 Livingston Avenue
First Floor
Roseland 07068
Tel: (201) 740-0111

NEW MEXICO

Intel Corp.
8500 Mutual Boulevard N.E.
Suite B 295
Albuquerque 87112
Tel: (505) 292-8086

NEW YORK

Intel Corp.*
850 Cross Keys Office Park
Fairport 14450
Tel: (716) 425-2750
TWX: 510-253-7391

INDIANA

Intel Corp.*
300 Motor Parkway
Hauppauge 11787
Tel: (516) 231-3300
TWX: 510-227-6236

KANSAS

Intel Corp.
Suite 2B Holbrook Park
15 Myers Corners Road
Wappingers Falls 12590
Tel: (914) 297-6161
TWX: 510-248-0060

NORTH CAROLINA

Intel Corp.
5700 Executive Center Drive
Suite 213
Charlotte 28212
Tel: (704) 568-8966

OHIO

Intel Corp.*
3401 Park Center Drive
Suite 220
Dayton 45414
Tel: (937) 890-5350
TWX: 810-450-2528

PENNSYLVANIA

Intel Corp.*
25700 Science Park Dr., Suite 100
Bensalem 19022
Tel: (215) 425-2736
TWX: 810-427-9298

OKLAHOMA

Intel Corp.
6801 N. Broadway
Suite 115
Oklahoma City 73116
Tel: (405) 848-8086

OREGON

Intel Corp.
15254 N.W. Greenbrier Parkway, Bldg. B
Beaverton 97006
Tel: (503) 645-8051
TWX: 910-467-8741

PENNSYLVANIA

Intel Corp.
1513 Cedar Cliff Drive
Camp Hill 17011
Tel: (717) 737-3035

Intel Corp.*
455 Pennsylvania Avenue
Fort Washington 19034
Tel: (215) 641-1000
TWX: 510-661-2077

Intel Corp.*
400 Penn Center Blvd., Suite 610
Pittsburgh 15235
Tel: (412) 823-4970

PUEBLO RICO

Intel Microprocessor Corp.
South Industrial Park
P.O. Box 910
Las Piedras 00671
Tel: (609) 733-6616

TEXAS

Intel Corp.
313 E. Anderson Lane
Suite 390
Austin 78752
Tel: (512) 454-3628

Intel Corp.*
1250 Concord Road
Suite 390
Dallas 75234
Tel: (214) 241-8087
TWX: 910-860-5617

Intel Corp.*
7322 S.W. Freeway
Suite 1490
Houston 77074
Tel: (713) 988-8086
TWX: 910-881-2490

UTAH

Intel Corp.
5201 Green Street
Suite 290
Murray 84123
Tel: (801) 263-8051

VIRGINIA

Intel Corp.
1603 Santa Rosa Road
Suite 109
Richmond 23288
Tel: (804) 282-5668

WASHINGTON

Intel Corp.
151-108 Avenue N.E.
Suite 300
Bellevue 98004
Tel: (206) 453-8086
TWX: 910-443-3002

WISCONSIN

Intel Corp.
409 N. Mullion Road
Suite 102
Spokane 99206
Tel: (509) 928-8086

ONTARIO

Intel Semiconductor of Canada, Ltd.
301-2245 W. Broadway
Vancouver V6K 2E4
Tel: (604) 738-6522

CANADA

BRITISH COLUMBIA
Intel Semiconductor of Canada, Ltd.
2850 Seminoleview Drive
Suite 250
Ottawa K2B 8H6
Tel: (613) 829-9714
TLX: 053-4115

ONTARIO
Intel Semiconductor of Canada, Ltd.
190 Atwell Drive
Suite 500
Reddale M9W 6H8
Tel: (416) 675-2105
TLX: 05983574

QUEBEC

Intel Semiconductor of Canada, Ltd.
620 St. Jean Boulevard
Pointe Claire H9R 3K3
Tel: (514) 694-9130
TWX: 514-694-9134

*Field Application Location



DOMESTIC DISTRIBUTORS

ALABAMA

Arrow Electronics, Inc.
1015 Henderson Road
Huntsville 35805
Tel: (205) 837-6955
†Hamilton/Avnet Electronics
4940 Research Drive
Huntsville 35805
Tel: (205) 837-7210
TWX: 810-26-2162

Pioneer Technologies Group Inc.
4825 University Square
Huntsville 35805
Tel: (205) 837-9300
TWX: 810-26-2197

ARIZONA

†Hamilton/Avnet Electronics
505 S. Madison Drive
Tempe 85281
Tel: (602) 231-5100
TWX: 910-950-0077

Kierulff Electronics, Inc.
4134 E. Wood Street
Phoenix 85040
Tel: (602) 437-0750
TWX: 910-951-1550

Wyle Distribution Group
17855 N. Black Canyon Highway
Phoenix 85024
Tel: (602) 866-2888

CALIFORNIA

Arrow Electronics, Inc.
19748 Dearborn Street
Chatsworth 91311
Tel: (818) 761-7500
TWX: 910-933-2086

Arrow Electronics, Inc.
1500 Crocker Avenue
Hayward 94544
Tel: (408) 487-4600

Arrow Electronics, Inc.
9511 Ridgehaven Court
San Diego 92123
Tel: (619) 565-4090
TWX: 880085

†Arrow Electronics, Inc.
521 Wedell Drive
Sunrise Park
Tel: (408) 745-5600
TWX: 910-339-9371

Arrow Electronics, Inc.
2961 Dow Avenue
Torrance 90503
Tel: (714) 638-5422
TWX: 910-595-2860

†Arrow Electronics,
550 McCormick Avenue
Costa Mesa 92626
Tel: (714) 754-6051
TWX: 910-905-1928

Hamilton/Avnet Electronics
175 Bonanza Drive
Sunvalley 94086
Tel: (408) 743-3300
TWX: 910-339-9332

†Hamilton/Avnet Electronics
4545 Viewridge Avenue
San Diego 92123
Tel: (619) 571-7500
TWX: 910-595-2638

†Hamilton/Avnet Electronics
20501 Plummer Street
Chatsworth 91311
Tel: (818) 700-6271
TWX: 910-494-2207

†Hamilton/Avnet Electronics
4103 Northgate Boulevard
Sacramento 95834
Tel: (916) 920-3150

†Hamilton/Avnet Electronics
3002 G Street
Ontario 91311
Tel: (714) 989-9411

Hamilton/Avnet Electronics
1951 So. Vermont Avenue
Torrance 90502
Tel: (213) 615-3909
TWX: 910-349-6263

Hamilton Electro Sales
9650 De Soto Avenue
Chatsworth 91311
Tel: (818) 700-6550

CALIFORNIA (Cont'd)

†Hamilton Electro Sales
10950 W. Washington Blvd.
Culver City 90230
Tel: (213) 558-6368
TWX: 910-340-6364

Hamilton Electro Sales
1361 B West 190th Street
Gardena 90248
Tel: (213) 558-2131

†Hamilton Electro Sales
3170 Pullman Street
Costa Mesa 92626
Tel: (714) 641-4470
TWX: 910-595-2638

Kierulff Electronics, Inc.
10824 Hope Street
Cypress 90630
Tel: (714) 220-6300

†Kierulff Electronics, Inc.
1180 Murphy Avenue
San Jose 95131
Tel: (408) 971-2600

†Kierulff Electronics, Inc.
1180 Franklin Avenue
Tustin 92680
Tel: (714) 731-5711

TWX: 910-595-2599

†Kierulff Electronics, Inc.
5650 Wilshire Street
Commerce 90040
Tel: (213) 725-0325

Wyle Distribution Group
26560 Apgus Street
Calabasas 91302
Tel: (818) 880-9000

TWX: 818-372-0232

†Wyle Distribution Group
124 Maryland Street
El Segundo 90245
Tel: (213) 322-8100

TWX: 910-348-7140 or 7111

†Wyle Distribution Group
17872 Cowan Avenue
Irvine 92714
Tel: (714) 863-9953

TWX: 910-595-1572

Wyle Distribution Group
11151 Sun Center Drive
Rancho Cordova 95670
Tel: (916) 638-5282

†Wyle Distribution Group
3000 Bowers Avenue
Santa Clara 95051
Tel: (408) 727-2500

Wyle Military
18910 Teller Avenue
Irvine 92750
Tel: (714) 851-9958

TWX: 910-335-1590

†Wyle Distribution Group
9525 Chesapeake Drive
San Diego 92123
Tel: (619) 585-1717

TWX: 910-335-1927

Wyle Systems
7382 Lampson Avenue
Garden Grove 92641
Tel: (714) 691-1717

TWX: 910-595-2642

Arrow Electronics, Inc.
1380 S. Potomac Street
Suite 136
Aurora 80012
Tel: (303) 696-1111

†Hamilton/Avnet Electronics
8765 E. Orchard Road
Suite 708
Englewood 80111
Tel: (303) 740-1017

TWX: 910-935-0787

†Wyle Distribution Group
451 E. 124th Avenue
Thornton 80241
Tel: (303) 457-9953

TWX: 910-936-0770

CONNECTICUT

†Arrow Electronics, Inc.
12 Beaumont Road
Wallingford 06492
Tel: (203) 661-0161
TWX: 710-476-0162

Hamilton/Avnet Electronics
Commerce Industrial Park
Commerce Drive
Danbury 06810
Tel: (203) 971-2800
TWX: 710-456-9974

†Pioneer Northeast Electronics
112 Main Street
Norwalk 06851
Tel: (203) 853-1515
TWX: 710-468-3373

FLORIDA

†Arrow Electronics, Inc.
350 Farway Drive
Deerfield Beach 33441
Tel: (305) 429-8200
TWX: 510-955-9456

Arrow Electronics, Inc.
1001 N.W. 62nd St., Ste. 108
Ft. Lauderdale 33309
Tel: (305) 776-7790
TWX: 510-955-9456

†Arrow Electronics, Inc.
50 Woodlake Drive W., Bldg. B
Palm Bay 32901
Tel: (305) 725-1480
TWX: 510-953-6337

†Hamilton/Avnet Electronics
6801 N.W. 5th Way
Ft. Lauderdale 33309
Tel: (305) 971-2900
TWX: 510-956-3097

Hamilton/Avnet Electronics
3197 Tech Drive North
St. Petersburg 33707
Tel: (813) 576-3930
TWX: 810-863-0374

Hamilton/Avnet Electronics
6947 University Boulevard
Winterpark 32792
Tel: (305) 628-3888
TWX: 810-853-0322

†Pioneer Electronics
337 N. Lake Blvd., Ste. 1000
Alta Monte Springs 32701
Tel: (305) 834-9050
TWX: 810-853-0284

Pioneer Electronics
674 S. Military Trail
Deerfield Beach 33442
Tel: (305) 428-8877
TWX: 510-955-9653

GEORGIA

†Arrow Electronics, Inc.
1555 Northwoods Parkway
Suite 100
Norcross 30071
Tel: (404) 449-8252

Hamilton/Avnet Electronics
5026 D. Peasey Corners
Norcross 30092
Tel: (404) 447-7500
TWX: 710-828-9702

†Arrow Electronics, Inc.
3100 Northwoods Place
Norcross 30071
Tel: (404) 448-7111
TWX: 810-766-0432

Pioneer Electronics
3100 F. Northwoods Place
Norcross 30071
Tel: (404) 448-7111
TWX: 810-766-4515

ILLINOIS

†Arrow Electronics, Inc.
2000 E. Alonquist Street
Schaumburg 60195
Tel: (312) 397-3440
TWX: 910-291-3544

†Hamilton/Avnet Electronics
1130 Thordale Avenue
Bensenville 60106
Tel: (312) 869-7780
TWX: 910-227-0060

Kierulff Electronics, Inc.
1140 W. Thordale
Itasca 60143
Tel: (312) 250-0500

ILLINOIS (Cont'd)

MTI Systems Sales
1100 West Thordale
Itasca 60143
Tel: (312) 773-2300

†Pioneer Electronics
1551 Carmen Drive
Elk Grove Village 60007
Tel: (312) 437-5680
TWX: 810-222-1834

INDIANA

†Arrow Electronics, Inc.
2405 Directors Row, Suite H
Indianapolis 46241
Tel: (317) 243-9333
TWX: 810-341-3119

Hamilton/Avnet Electronics
485 Grade Drive
Carmel 46032
Tel: (317) 844-9333
TWX: 810-260-3966

†Pioneer Electronics
6408 Castleplace Drive
Indianapolis 46250
Tel: (317) 849-7300
TWX: 810-260-1794

KANSAS

†Hamilton/Avnet Electronics
9129 Quivira Road
Overland Park 66215
Tel: (913) 869-5900
TWX: 910-743-0005

Pioneer Electronics
10551 Lackman Rd.
Lenexa 66215
Tel: (913) 452-0500

KENTUCKY

Hamilton/Avnet Electronics
10203 Bren Road East
Minnetonka 55343
Tel: (612) 850-1000
TWX: 910-576-3125

MARYLAND

Arrow Electronics, Inc.
8300 Gulford Road #H
Rivers Center
Columbia 21046
Tel: (301) 995-0003

†Hamilton/Avnet Electronics
6822 Oak Hall Lane
Columbia 21045
Tel: (301) 995-3000
TWX: 710-862-1861

†Mesa Technology Corp.
9720 Patuxentwood Dr.
Columbia 21046
Tel: (301) 720-5020
TWX: 710-828-9702

Hamilton/Avnet Electronics
911 Gulf Rd.
Gainesburg 20877
Tel: (301) 921-0660
TWX: 710-828-0545

MASSACHUSETTS

†Arrow Electronics, Inc.
1 Arrow Drive
Woburn 01801
Tel: (617) 933-6130
TWX: 710-393-6770

†Hamilton/Avnet Electronics
100 Centennial Drive
Peabody 01960
Tel: (617) 532-3701
TWX: 710-393-0382

Kierulff Electronics, Inc.
13 Fortune Dr.
Billerica 01821
Tel: (617) 667-8331

MTI Systems Sales
13 Fortune Drive
Billerica 01821
Tel: (617) 667-8331

Pioneer Northeast Electronics
44 Hartwell Avenue
Lexington 02173
Tel: (617) 861-9200
TWX: 710-326-6617

†Hamilton/Avnet Electronics
1 Keystone Avenue
Bldg. 36
Cherry Hill 08003
Tel: (609) 424-0110
TWX: 710-940-0262

MICHIGAN

Arrow Electronics, Inc.
755 Phoenix Drive
Ann Arbor 48104
Tel: (313) 997-0220
TWX: 810-223-6200

†Hamilton/Avnet Electronics
32487 Schoolcraft Road
Livonia 48150
Tel: (313) 522-4700
TWX: 810-242-8775

Hamilton/Avnet Electronics
215 29th Street S.E.
Spokane 59245
Grand Rapids 49508
Tel: (616) 555-1800

Pioneer Electronics
4505 Broadmoor Ave. S.E.
Grand Rapids 49508
Tel: (616) 555-1800

†Pioneer Electronics
13485 Stamford
Livonia 48150
Tel: (313) 526-3738

Hamilton/Avnet Electronics
12400 White Water Drive
Minnetonka 55343
Tel: (612) 852-0600
TWX: 910-576-2720

†Pioneer Electronics
12023 Shoreline Court
Earth City 63045
Tel: (314) 344-1200
TWX: 910-762-0684

Kierulff Electronics, Inc.
11804 Borman Dr.
St. Louis 63146
Tel: (314) 997-4956

NEW HAMPSHIRE

†Arrow Electronics, Inc.
3 Perimeter Road
Manchester 33104
Tel: (603) 624-9400

NEW JERSEY

†Arrow Electronics, Inc.
901 G. Gilbert
Marlton 08053
Tel: (609) 689-6968
TWX: 710-220-1694

Hamilton/Avnet Electronics
444 E. Industrial Drive
Manchester 33104
Tel: (603) 624-9400

NEW YORK

†Arrow Electronics, Inc.
2 Industrial Road
Fairfield 07006
Tel: (201) 575-5300
TWX: 710-998-2206

†Hamilton/Avnet Electronics
1 Bldg. 36
Cherry Hill 08003
Tel: (609) 424-0110
TWX: 710-940-0262



DOMESTIC DISTRIBUTORS

NEW JERSEY (Cont'd)

†Hamilton/Avnet Electronics
10 Industrial
Fairfield 07006
Tel: (201) 575-3390
TWX: 710-734-4388

†Pioneer Northeast Electronics
45 Route 46
Pinelawn 06508
Tel: (201) 575-3510
TWX: 710-734-4382

†MTI Systems Sales
383 Route 46 W
Fairfield 07006
Tel: (201) 227-5552

NEW MEXICO

Alliance Electronics Inc.
11030 Cochiti S.E.
Albuquerque 87123
Tel: (505) 292-3360
TWX: 910-989-1151

Hamilton/Avnet Electronics
2524 Baylor Drive S.E.
Albuquerque 87106
Tel: (505) 765-1500
TWX: 910-989-0614

NEW YORK

Arrow Electronics, Inc.
25 Hub Drive
Melville 11747
Tel: (516) 694-6800
TWX: 510-224-6126

†Arrow Electronics, Inc.
3375 Brighton-Henrietta Townline Rd.
Rochester 14623
Tel: (716) 427-0300
TWX: 510-253-4766

Arrow Electronics, Inc.
7705 Mallgate Drive
Livermore 94550
Tel: (415) 652-1000
TWX: 710-545-0230

Arrow Electronics, Inc.
20 Averitt Avenue
West Seneca 14288
Tel: (716) 231-1000
TWX: 510-227-6623

Hamilton/Avnet Electronics
333 Mario Park
Rochester 14623
Tel: (716) 475-9130
TWX: 510-253-5470

†Hamilton/Avnet Electronics
100 Twin Oaks Drive
Syracuse 13206
Tel: (315) 437-2641
TWX: 710-541-1560

†Hamilton/Avnet Electronics
933 Motor Parkway
Hauppauge 11788
Tel: (516) 231-9800
TWX: 510-224-6166

†MTI Systems Sales
38 Harbor Park Drive
P.O. Box 271
Port Washington 11050
Tel: (516) 621-6200
TWX: 510-223-0846

†Pioneer Northeast Electronics
1806 Vestal Parkway East
Vestal 13850
Tel: (607) 748-8211
TWX: 510-252-0893

†Pioneer Northeast Electronics
60 Crossway Park West
Woodbury, Long Island 11797
Tel: (516) 991-2100
TWX: 510-221-2184

NEW YORK (Cont'd)

†Pioneer Northeast Electronics
840 Fairport Park
Fairport 14450
Tel: (518) 876-7070
TWX: 510-253-7001

NORTH CAROLINA

†Arrow Electronics, Inc.
5240 Greenberry Road
Raleigh 27604
Tel: (919) 876-3132
TWX: 710-878-1856

†Hamilton/Avnet Electronics
3510 Spring Forest Drive
Raleigh 27604
Tel: (919) 878-0819
TWX: 510-928-1836

Pioneer Electronics
9801 A-Southern Pine Blvd.
Charlotte 28210
Tel: (704) 527-8188
TWX: 810-621-0366

OHIO

Arrow Electronics, Inc.
7620 McEwen Road
Centerville 45459
Tel: (513) 435-5563
TWX: 810-427-9409

†Pioneer Electronics
261 Giraltar Road
Horsham 19044
Tel: (215) 674-4000
TWX: 510-665-6778

†Arrow Electronics, Inc.
3220 Commander Drive
Carrollton 75006
Tel: (214) 360-5464
TWX: 910-860-5377

†Arrow Electronics, Inc.
10899 Kinghurst
Suite 100
Houston 77099
Tel: (713) 530-4700
TWX: 910-880-4439

†Hamilton/Avnet Electronics
954 Senate Drive
Westerville 43081
Tel: (614) 882-7004
TWX: 810-450-2531

†Hamilton/Avnet Electronics
4588 Energy Industrial Parkway
Worrellswood Heights 44128
Tel: (216) 831-3500
TWX: 810-427-9452

†Pioneer Electronics
4433 Interpoint Blvd.
Dayton 45424
Tel: (513) 236-9900
TWX: 910-459-1622

†Pioneer Electronics
4900 E. 31st Street
Cleveland 44105
Tel: (216) 587-3600
TWX: 810-422-2211

OKLAHOMA

Arrow Electronics, Inc.
4719 S. Memorial Drive
Dallas 75243
Tel: (918) 665-7700

OREGON

†Almac Electronics Corpora-
tion
1895 N.W. 169th Place
Beaverton 97006

Tel: (503) 629-8090
TWX: 910-467-8743

†Hamilton/Avnet Electronics
6024 S.W. Jean Road
Bldg. C, Suite 10
Lake Oswego 97034
Tel: (503) 635-7848
TWX: 910-455-8179

OREGON (Cont'd)

Wyle Distribution Group
5250 N.E. Elam Young Parkway
Suite 600
Portland 97124
Tel: (503) 640-5000
TWX: 910-460-2203

PENNSYLVANIA

Arrow Electronics, Inc.
650 Seco Road
Monroeville 15146
Tel: (724) 281-7000

Hamilton/Avnet Electronics
2800 Liberty Ave., Bldg. E
Pittsburgh 15238
Tel: (412) 281-4150

Pioneer Electronics
259 Kapoor Drive
Pittsburgh 15238
Tel: (412) 782-2300
TWX: 710-795-3122

†Pioneer Electronics
261 Giraltar Road
Horsham 19044
Tel: (215) 674-4000
TWX: 510-665-6778

TEXAS

†Arrow Electronics, Inc.
3220 Commander Drive
Carrollton 75006
Tel: (214) 360-5464
TWX: 910-860-5377

†Arrow Electronics, Inc.
10899 Kinghurst
Suite 100
Houston 77099
Tel: (713) 530-4700
TWX: 910-880-4439

†Arrow Electronics, Inc.
1025 Metropolitan
Austin 78758
Tel: (512) 835-4180
TWX: 910-874-1348

Hamilton/Avnet Electronics
2401 Rutledge
Austin 78758
Tel: (512) 837-8911
TWX: 910-874-1319

†Hamilton/Avnet Electronics
2111 W. Walnut Hill Lane
Irving 75062
Tel: (918) 659-4100
TWX: 910-860-5929

†Hamilton/Avnet Electronics
4850 Wright Road #190
Stafford 77477
Tel: (713) 780-1771
TWX: 910-881-5523

Kierulff Electronics, Inc.
9610 Skilled Lane
Dallas 75243
Tel: (214) 343-2401

†Pioneer Electronics
1826 D. Kramer Lane
Austin 78758
Tel: (512) 835-4000
TWX: 910-874-1323

†Pioneer Electronics
13710 Omega Road
Dallas 75234
Tel: (214) 386-7300
TWX: 910-850-5563

†Pioneer Electronics
5853 Point West Drive
Houston 77036
Tel: (713) 988-5555
TWX: 910-881-1606

UTAH

†Hamilton/Avnet Electronics
1058 West 2100 South
Salt Lake City 84119
Tel: (801) 466-2200
TWX: 910-925-4018

Kierulff Electronics, Inc.
1945 W. Parkway Blvd.
Salt Lake City 84119
Tel: (801) 973-6913

Wyle Distribution Group
1325 West 2200 South
Suite E
Salt Lake City 84119
Tel: (801) 974-9953

WASHINGTON

†Almac Electronics Corp.
14360 S.E. Eastgate Way
Bellevue 98007
Tel: (206) 643-9992
TWX: 910-444-2067

Arrow Electronics, Inc.
14320 N.E. 21st Street
Bellevue 98005
Tel: (206) 453-5874
TWX: 910-444-2469

Wyle Distribution Group
1750 132nd Ave., N.E.
Bellevue 98005
Tel: (206) 453-8300

†Hamilton/Avnet Electronics
14212 N.E. 21st Street
Bellevue 98005
Tel: (206) 453-8300

WISCONSIN

†Arrow Electronics, Inc.
430 W. Rauson Avenue
Oak Creek 53154
Tel: (414) 764-6600
TWX: 910-262-1193

Hamilton/Avnet Electronics
2975 Woodland Road
Madison 53713
Tel: (414) 232-5110
TWX: 910-262-1192

Kierulff Electronics, Inc.
2238-B W. Bluemound Rd.
Waukesha 53185
Tel: (414) 784-8160

CANADA

ALBERTA

Hamilton/Avnet Electronics
2816 21st Street N.E.
Calgary T2A 5Z2
Tel: (403) 230-3586
TWX: 03-827-642

Hamilton/Avnet Electronics
6045 Rexwood Road Unit 6
Mississauga, Ontario L4V1R2
Tel: (416) 677-0484

ZENTRONICS

BURMALAY

BRITISH COLUMBIA

Hamilton/Avnet Electronics
105-2550 Boundary Road
Burnaby V5M 3Z3
Tel: (604) 272-4242

BRITISH COLUMBIA (Cont'd)

Zentronics
108-1140 Bridgeport Road
Richmond V6V 1T2
Tel: (604) 272-5575
TWX: 04-507-89

MANITOBA

Zentronics
550 Berry Street
Winnipeg R3H 0S1
Tel: (204) 775-8861

ONTARIO

Arrow Electronics Inc.
24 Martin Ross Avenue
Downsvie MJ3 2K9
Tel: (416) 661-0220
TLX: 06-218213

Arrow Electronics Inc.
148 Colonnade Road
Nepean K2E 7J5
Tel: (613) 226-6903

†Hamilton/Avnet Electronics
210 Colonnade Road South
Nepean K2E 7L5
Tel: (613) 226-1700
TWX: 05-349-71

†Zentronics
8 Tilbury Court
Brantford L8S 1T4
Tel: (419) 451-9800
TWX: 05-976-78

Zentronics
564/10 Weber Street North
Waterloo N2L 5C6
Tel: (519) 845-5700

†Zentronics
155 Colonnade Road
Unit 11
Nepean K2E 7K1
Tel: (613) 225-8840
TWX: 05-976-78

QUEBEC

†Arrow Electronics Inc.
4050 Jean Talon Quay
Montreal H4P 1W1
Tel: (514) 735-5511
TLX: 05-2559

Arrow Electronics Inc.
909 Charest Blvd.
Quebec G1N 2B9
Tel: (418) 687-4231
TLX: 05-13388

Hamilton/Avnet Electronics
2795 Rue Halpern
St. Laurent H4P 1W8
Tel: (514) 335-1000
TWX: 610-421-3731

Zentronics
505 Locke Street
St. Laurent H4L 1X7
Tel: (514) 335-5361
TWX: 05-827-535



EUROPEAN SALES OFFICES

BELGIUM

Intel Corporation S.A.
Rue de Cottages 65
B-1180 Brussels
Tel: (02) 347-0666

DENMARK

Intel Denmark A/S*
Glentevej 61 - 3rd Floor
DK-2400 Copenhagen
Tel: (01) 19-80-33
TLX: 19567

FINLAND

Intel Finland OY
Pohjois-Piikkiö
00390 Helsinki
Tel: (0) 0544-644
TLX: 12332

FRANCE

Intel Paris
1 Rue Edison, BP 303
78054 Saint-Quentin-en-Yvelines Cedex
Tel: (33) 7-342-57-7000
TLX: 6990167

Intel Corporation, S.A.R.L.
Immeuble BBC
4 Quai des Etroits
69005 Lyon
Tel: (7) 742-4089
TLX: 30513

WEST GERMANY

Intel Semiconductor GmbH*
Seidestrasse 27
D-8000 Muenchen 2
Tel: (09) 53891
TLX: 05-23177 INTL D

Intel Semiconductor GmbH
Verkaufsbüro Wiesbaden
Abraham-Lincoln Str. 16-18
6200 Wiesbaden
Tel: (6121) 76050
TLX: 04186183 INTW D

Intel Semiconductor GmbH
Verkaufsbüro Stuttgart
Bräuerstrasse 91
7012 Fellbach
Tel: (711) 58-00-82
TLX: 7254826 INTS D

Intel Semiconductor GmbH
Verkaufsbüro Hannover
Hohenholznerstrasse 5
3000 Hannover 1
Tel: (611) 34-40-81
TLX: 923625 INT H D

ISRAEL

Intel Semiconductor Ltd*
Attidim Industrial Park
Neve Sharet
Dvora Hanavia
Bldg. No. 13, 4th Floor
P.O. Box 100-002
Tel: (03) 611430

Tel: (03) 491-099, 491-098
TLX: 371215

ITALY

Intel Corporation S.P.A.*
Milanofiori, Palazzo E/4
20090 Assago (Milano)
Tel: (02) 824-4071
TLX: 341286 INT MIL

NETHERLANDS

Intel Semiconductor (Nederland) B.V.*
Alexander Airport Building
Meerweg 10, 1990 ZH
3068 Rotterdam
Tel: (10) 21-23-77
TLX: 22283

NORWAY

Intel Norway A/S
P.O. Box 92
Hvammeien 4
N-2013, Skjetten
Tel: (2) 742-420
TLX: 78018

SPAIN

Intel Iberia
Calle Zurbaran 28-IZQDA
28010 Madrid
Tel: (1) 410-4004
TLX: 46880

SWEDEN

Intel Sweden A.B.*
Dalvagen 24
S-171 36 Solna
Tel: (6) 734-0100
TLX: 12261

SWITZERLAND

Intel Semiconductor A.G.*
Tiefenbachstrasse 17
8152 Giattello
CH-8065 Zurich
Tel: (01) 829-2977
TLX: 57989 ICH CH

UNITED KINGDOM

Intel Corporation (U.K.) Ltd.*
Pipers Way
Swindon, Wiltshire SN1 1RJ
Tel: (0793) 696-0000
TLX: 444447 INT SWN

*Field Application Location

EUROPEAN DISTRIBUTORS/REPRESENTATIVES

AUSTRIA

Bacher Elektronics Ges.m.b.H.
Rotenmühlgasse 26
A-1120 Vienna
Tel: (222) 835-6460
TLX: 131532

BELGIUM

Ineoce Belgium S.A.
Ave. des Croix de Guerre, 94
Bruxelles 1120
Tel: (02) 216-01-60
TLX: 64475

BENELUX

Koning en Hartman Electrotechniek B.V.
Postbus 125
2600 AH Delft
Tel: (15) 69-90-906
TLX: 38250

DENMARK

ITT MultiKomponent
Naverland 29
DK-2600 Glostrup
Tel: (02) 456-66-45
TLX: 33355 ITTCG DK

FINLAND

Oy Fintronik AB
Mekonkatu 24A
SF-00210 Helsinki 21
Tel: (0) 692-50-22
TLX: 124224 FINRON SF

FRANCE

Generim
Zone d'Activité de Courtabeauf
Avenue de la Batiqne
91943 Les Ulis Cedex
Tel: (1) 69-07-78-78
TLX: 691700

Jermyn
73-75 Rue des Solets
Silic 585
94663 Rungis Cedex
Tel: (1) 45-50-04-00
TLX: 230967

Metrologie
Tour d'Aspières
4 Avenue Laurent Cely
92606 Asnières
Tel: (1) 47-90-62-40
TLX: 611448

FRANCE (Cont'd)

Tekelec Artronic
Cite des Bruyeres
Rue Paul Vermet BP 2
92310 Suresnes
Tel: (1) 45-34-75-35
TLX: 20452

WEST GERMANY

Electronic 2000 Vertriebs AG
Stahlgruberring 12
8000 Muenchen 82
Tel: (089) 42-00-10
TLX: 522561 ELEC D

Metrologie GmbH
Meglingerstr. 49
8000 Muenchen 71
Tel: (089) 570-940
TLX: 523185

Metrologie GmbH
Rheinstr. 94-96
6100 Frankfurt am Main
Tel: (0611) 23661
TLX: 176151820

Proelectron Vertriebs AG
Max-Planck-Straße 1-3
6000 Frankfurt 1
Tel: (06103) 3040
TLX: 417972

ITT-MultiKomponent
Bebelstrasse 44
7114 Metzingen
Tel: (07141) 4879
TLX: 7264399 MUKO D

ISRAEL

Eastronics Ltd.
11 Ronasan Street
P.O. Box 39300
Tel Aviv 61392
Tel: (03) 541-51
TLX: 342610 DATIX IL or
33638 RONIX IL

ITALY

Eledra Componenti S.P.A.
V - Giacomo Watt, 37
20143 Milano
Tel: (02) 82821
TLX: 332332

ITALY (Cont'd)

Intesi
Milanofiori E5
20090 Assago
Tel: (02) 824701
TLX: 311351

NORWAY

Nordisk Elektronik A/S
Postboks 130
N-1364 Hvalstad
6277 Bad Camberg
Tel: (2) 846-210
TLX: 77546 NENAS N

PORUGAL

Ditram
Avenida Marques de Tomar, 46A
Lisboa P-1000
Tel: (351-1) 734-834
TWX: (0404) 14182

SPAIN

ITT
21-3 Miquel Angel
Madrid 28010
Tel: (1) 419-54-00
TWX: 27461

SWEDEN

A.T.D. Electronica S.A.
Pl. Ciudad de Viena 6
28040 Madrid
Tel: (1) 234-4000
TWX: 42477

SWEDEN

Nordisk Elektronik AB
Box 4057
S-171 27 Solna
Tel: (8) 734-97-70
TLX: 10547

SWITZERLAND

Industrade AG
Herristrasse 31
CH-8304 Wallisellen
Tel: (01) 830-5040
TLX: 56788

UNITED KINGDOM

Accent Electronic Components Ltd.
Jubilee House, Jubilee Way
Leatherhead, Surrey KT6 1QH
England
Tel: (0462) 686666
TLX: 626923

BYTECH LTD.

Unit 2 Western Centre
Western Industrial Estate
Bracknell, Berkshire RG12 1RW
England
Tel: (0344) 482211
TLX: 846215

CONWAY MICROSYSTEMS LTD.

John Scott House, Market St.
Bracknell, Berkshire RG12 1QP
England
Tel: (0344) 55333
TLX: 847201

IBRI MICROCOMPUTERS LTD.

Unit 2 Western Centre
Western Industrial Estate
Bracknell, Berkshire RG12 1RW
England
Tel: (0344) 486-555
TLX: 849381

JERMYN INDUSTRIES

Vestry Estate, Oxford Road
Sevenoaks, Kent TN14 5EU
England
Tel: (0732) 450144
TLX: 95142

RAPID SILICON

Rapid House, Denmark St.
High Wycombe, Bucks HP11 2ER
England

Tel: (094) 442266
TLX: 837931

RAPID SYSTEMS

Rapid House, Denmark St.
High Wycombe, Bucks HP11 2ER
England
Tel: (094) 450244
TLX: 837931

MICRO MARKETING

Glenageary Office Park
Glenageary, Co. Dublin
Ireland
Tel: (001) 856288
TLX: 31584

YUGOSLAVIA

H.R. Microelectronics Corp.
2005 La La Cruz Blvd., Ste. 223
Santa Clara, CA 95050 U.S.A.

Tel: (408) 988-0286
TLX: 367452



INTERNATIONAL SALES OFFICES

AUSTRALIA

Intel Australia Pty. Ltd.
Spectrum Building
200 Pacific Hwy., Level 6
Crows Nest, NSW, 2065
Tel: (02) 9571-2744
TLX: 20097
FAX: (02) 923-2632

CHINA

Intel PRC Corporation
15/F, Office 1, Citic Bldg.
Jing An, Chongming Wai Street
Beijing, PRC
Tel: (1) 500-4850
TLX: 22947 INTEL CN
FAX: (1) 500-2953

HONG KONG

Intel Semiconductor Ltd.
1701-3 Connaught Centre
1 Connaught Road
Tel: (5) 801-3555
TLX: 50982 INTEL HK
FAX: (5) 294-589

JAPAN

Intel Japan K.K.
5-6 Tohoku Toyosato-machi
Tsukuba-ku, Ibaraki-ken 300-26
Tel: (02) 97-47-8511
TLX: 03656-160

Intel Japan K.K.
Daishi Mitsugi Bldg.
1-8889 Fuchu-cho
Fuchu-shi, Tokyo 183
Tel: (04) 23-60-7871

Intel Japan K.K.
Flower Hill Shin-machi Bldg.
1-23-9 Shinmachi
Setagaya-ku, Tokyo 154
Tel: (03) 426-2231

Intel Japan K.K.
Kumagaya Bldg.
2-69 Honcho
Kumagaya, Saitama 360
Tel: (04) 85-24-6871

Intel Japan K.K.
Mishima Tokyo-Kaijo Bldg.
1-1 Shibahon-cho, Mishima-shi
Shizuoka-ken 411
Tel: (05) 59-72-2141

JAPAN (Cont'd)

Intel Japan K.K.
Mitaka Seisaku Mutsashi-Kosugi Bldg.
915-20 Shimmuku, Nakahara-ku
Kawasaki-shi, Kanagawa 211
Tel: (04) 47-33-7011

Intel Japan K.K.
Nihon Seimei Bldg.
1-12 Asahi-cho
Atsugi, Kanagawa 243
Tel: (04) 62-23-3511

Intel Japan K.K.
Pyokyo-Station Bldg.
2-4-1 Terauchi
Toyonaka, Osaka 560
Tel: (06) 863-1091

Intel Japan K.K.
Shinmaru Bldg.
1-5-1 Marunouchi
Chiyoda-ku, Tokyo 100
Tel: (03) 201-3621

Okaya Koki
2-14-6 Kita-ku
Naka-ku, Nagoya-shi 460
Tel: 052-204-2911
FAX: 052-204-2901

Ryovo Electro Corp.
Konwa Bldg.
1-12-22 Tsukiji
Chuo-ku, Tokyo 104
Tel: 03-546-5011
FAX: 03-546-5044

Samsung Semiconductor &
Telecommunications Co., Ltd.
150, 2-Ka, Taipyung-ro, Chung-ku
Seoul 100
Tel: 82-10-51-3987
TLX: 27970 KORST
FAX: 82-2-784-8391

J-Tek Corporation
8th Floor, Government Pension Bldg.
2-3 Yodo-Dong
Youngdeungpo-ku
Seoul 150
Tel: 82-2-782-8039
TLX: 25299 KODIGIT
FAX: 82-2-784-8391

Dicopel S.A.
Tochili 388 Fracc. Ind. San Antonio
Azcapotzalco
C.P. 02760-Mexico, D.F.
Tel: 52-5-561-3211
TLX: 1773790 DICOME

KOREA

Intel Technology Asia Ltd.
Room 1008, Singsong Bldg.
25-1, Yoido-Dong, Youngdeungpo-ku
Seoul 150
Tel: (2) 784-8186
TLX: 29312 INTELKO
FAX: (2) 784-8096

SINGAPORE

Intel Singapore Technology, Ltd.
1-1 Thomson Road
#21-06 Golden Square
Singapore 03130
Tel: 250-7811
TLX: 39921 INTEL
FAX: 250-9256

TAIWAN

Intel Technology (Far East) Ltd.
Taiwan Branch
10/F., No. 205, Tun Hua N. Road
Taipei, R.O.C.
Tel: 886-2-29660
TLX: 13159 INTELTWN
FAX: (02) 717-2455

*Field Application Location

INTERNATIONAL DISTRIBUTORS/REPRESENTATIVES

ARGENTINA

VLC S.R.L. Bartolome Mitre 1711
3 Piso
1037 Buenos Aires
Tel: 54-1-49-2092
TLX: 17575 EDARG-AR

AUSTRALIA

Total Electronics
Private Bag 250
9 Harker Street
Burwood, Victoria 3125
Tel: 61-3-288-4044
TLX: 23297

Total Electronics
P.O. Box 139
Artamon, N.S.W. 2064
Tel: 61-02-438-1855
TLX: 26297

BRAZIL

Eletra Microelectronic S/A
Geraldo Flausino Gomes, 78
9 Ander
04575 - Sao Paulo - S.P.
Tel: 55-1-534-9600
TLX: 3911125131 ELBR BR
FAX: 55-1-534-9424

CHILE

DIN Instruments
Suecia 2222
Casilla 6055, Correo 22
Santiago
Tel: 56-2-225-8139
TLX: 440422 RUDY CZ

CHINA

Novel Precision Machinery Co., Ltd.
Flat D, 20 Kingsford Ind. Bldg.
Phase 1, 26 Kwai Hei Street
N.T., Nowloon
Hong Kong
Tel: 852-0-223-222
TWX: 39114 JINMI HX
FAX: 852-0-261-602

CHINA (Cont'd)

Schmidt & Co. Ltd.
19/F Great Eagle Centre
23 Harbour Road
Wanchai, Hong Kong
Tel: 852-5-833-0222
TWX: 74766 SCHMCN HX
FAX: 852-5-891-8754

INDIA

Micronic Devices
Arun Chambers
No. 10, D.V.S. Road
Basavangudi
Bangalore 560 004
Tel: 91-81-600-631
TLX: 0845-8332 MD BG IN

Micronic Devices
403, Gagan Deep
12, Rajendra Place
New Delhi 110 008
Tel: 91-11-97-71
TLX: 03163235 MDND IN

Micronic Devices
No. 516 5th Floor
Swastik Chambers
Sion, Chambar Road
Bombay 400 071
Tel: 91-91-52-03-63
TLX: 9531 171447 MDEV IN

JAPAN

Asahi Electronics Co. Ltd.
KMM Bldg. 2-14-1 Asano
Kokurakita-ku
Kitakyushu-shi 802
Tel: 093-511-6471
FAX: 093-551-7861

C. Itoh Techno-Science Co., Ltd.
C. Itoh Bldg., 2-5-1 Kita-Aoyama
Minato-ku, Tokyo 107
Tel: 03-3477-4959
FAX: 03-347-4959

JAPAN (Cont'd)

Da Semicon Systems, Inc.
Wasacore 64, 1-37-8 Sangenjaya
Setagaya-ku, Tokyo 154
Tel: 03-487-0389
FAX: 03-487-8088

Nakao Koki
2-14-6 Kita-ku
Naka-ku, Nagoya-shi 460
Tel: 052-204-2911
FAX: 052-204-2901

Ryovo Electro Corp.
Konwa Bldg.
1-12-22 Tsukiji
Chuo-ku, Tokyo 104
Tel: 03-546-5011
FAX: 03-546-5044

J-Tek Corporation
8th Floor, Government Pension Bldg.
2-3 Yodo-Dong
Youngdeungpo-ku
Seoul 150
Tel: 82-2-782-8039
TLX: 25299 KODIGIT
FAX: 82-2-784-8391

Samsung Semiconductor &
Telecommunications Co., Ltd.
150, 2-Ka, Taipyung-ro, Chung-ku
Seoul 100
Tel: 82-10-51-3987
TLX: 27970 KORST
FAX: 82-2-784-8391

Dicopel S.A.
Tochili 388 Fracc. Ind. San Antonio
Azcapotzalco
C.P. 02760-Mexico, D.F.
Tel: 52-5-561-3211
TLX: 1773790 DICOME

NEW ZEALAND

Northrup Instruments & Systems Ltd.
459 Kabel Pass Road
P.O. Box 9464, Newmarket
Auckland 1
Tel: 64-9-501-219, 501-801
TLX: 21570 THERMAL

Northrup Instruments & Systems Ltd.
P.O. Box 2406
Wellington 85668
Tel: 64-4-856-658
TLX: NZ3380
FAX: 64-4-857-276

SINGAPORE

Francotron Electronics Pte. Ltd.
17 Harvey Road #04-01
Singapore 069600
Tel: 65-283-0888, 289-1618
TWX: 65541 FRELS
FAX: 2895327

SOUTH AFRICA

Electronic Building Elements, Pty. Ltd.
P.O. Box 4609
Pine Square, 18th Street
Hazelwood, Pretoria 0001
Tel: 27-12-465921
TLX: 3-227786 SA

TAIWAN

Mitac Corporation
No. 588, Ming Shen East Rd.
Taipei, R.O.C.
Tel: 886-2-501-8231
FAX: 886-2-501-4265

VENEZUELA

P. Benavides S/A
Avilanes a Rio
Residencias Kamarata
Locales 4 A17
La Guaira, Caracas
Tel: 58-2-571-0396
TLX: 28450 PBVEN VC
FAX: 58-2-572-3321

*Field Application Location



DOMESTIC SERVICE OFFICES

ALABAMA

Intel Corp.
5015 Bradford Drive, #2
Huntsville 35805
Tel: (205) 830-4010

ARIZONA

Intel Corp.
11225 N. 28th Dr. #D214
Phoenix 85029
Tel: (602) 869-9800

Intel Corp.
500 E. Fry Blvd., Suite M-15
Sierra Vista 85635
Tel: (602) 459-5010

ARKANSAS

Intel Corp.
P.O. Box 206
Ulm 72170
Tel: (501) 241-3264

CALIFORNIA

Intel Corp.
21515 Vanowen
Suite 100
Canoga Park 91303
Tel: (818) 704-8500

Intel Corp.
2250 E. Imperial Highway
Suite 118
El Segundo 90245
Tel: 1-800-468-3548

Intel Corp.
1900 Prairie City Rd
Folsom 95630-5597
Tel: (916) 351-6143

Intel Corp.
2000 E. 10th Street
Suite 110
Santa Ana 92705
Tel: (714) 835-5789
TWX: 910-595-2475

Intel Corp.
2700 San Tomas Expressway
Santa Clara 95051
Tel: (408) 970-1740

Intel Corp.
4250 Executive Drive
Suite 150
San Diego 92121
Tel: (619) 452-5880

COLORADO

Intel Corp.
650 South Cherry
Suite 915
Denver 80222
Tel: (303) 321-8086
TWX: 910-931-2289

CONNECTICUT

Intel Corp.
26 Mill Plain Road
Danbury 06811
Tel: (203) 748-3130

FLORIDA

Intel Corp.
1500 N.W. 62, Suite 104
 Ft. Lauderdale 33309
Tel: (305) 771-0600
TWX: 510-956-9407

Intel Corp.
242 N. Westmonte Drive
Suite 105
Altamonte Springs 32714
Tel: (305) 869-5588

GEORGIA

Intel Corp.
3200 Peachtree Parkway
Suite 200
Norcross 30092
Tel: (404) 441-1171

ILLINOIS

Intel Corp.
300 N. Martingale Rd.
Suite 300
Schaumburg 60194
Tel: (312) 310-5733

INDIANA

Intel Corp.
8777 Purdue Rd., #125
Indianapolis 46288
Tel: (317) 875-0623

KANSAS

Intel Corp.
8400 W. 110th Street
Suite 170
Overland Park 66210
Tel: (913) 345-2727

KENTUCKY

Intel Corp.
3525 Tatescreek Road,
#51
Lexington 40502
Tel: (502) 272-6745

MARYLAND

Intel Corp.
5th Floor
7833 Walker Drive
Greenbelt 20770
Tel: (301) 441-1020

MASSACHUSETTS

Intel Corp.
3 Carlisle Road
Westford 01886
Tel: (617) 692-1060

MICHIGAN

Intel Corp.
7071 Orchard Lake Road
Suite 100
West Bloomfield 48033
Tel: (313) 851-8905

MISSOURI

Intel Corp.
4203 Earth City Expressway
Suite 143
Earth City 63045
Tel: (314) 291-2015

NEW JERSEY

Intel Corp.
305 Syntex Avenue
Englewood Cliffs 07632
Tel: (201) 567-0821
TWX: 710-991-8593

NORTH CAROLINA

Intel Corp.
2306 W. Meadowview Road
Suite 206
Greensboro 27407
Tel: (919) 294-1541

OHIO

Intel Corp.
One Chagrin Boulevard
Suite 305
28001 Chagrin Boulevard
Cleveland 44122
Tel: (216) 464-6915
TWX: 810-427-9298

OREGON

Intel Corp.
1525 N.W. Greenbrier
Beaverton 97005
Tel: (503) 645-8051
TWX: 910-467-8741

PENNSYLVANIA

Intel Corp.
5200 N.E. Elam Young Parkway
Hillsboro 97123
Tel: (503) 681-8080

PENNSYLVANIA

Intel Corp.
201 Penn Center Boulevard
Suite 301 W
Pittsburgh 15235
Tel: (313) 354-1540

TEXAS

Intel Corp.
313 E. Anderson Lane
Suite 314
Austin 78752
Tel: (512) 454-3628
TWX: 910-874-1347

Intel Corp.
12300 Ford Road
Suite 100
Dallas 75234
Tel: (214) 241-2820
TWX: 910-860-5617

Intel Corp.
8815 Dyer St., Suite 225
El Paso 79904
Tel: (915) 751-0186

Intel Corp.
1603 Santa Rosa Rd., #109
Richmond 23288
Tel: (804) 282-5668

WASHINGTON

Intel Corp.
110 110th Avenue N.E.
Suite 5110
Bellevue 98004
Tel: 1-800-468-3548
TWX: 910-443-3002

WISCONSIN

Intel Corp.
450 N. Sunnyslope Road
Suite 130
Brookfield 53005
Tel: (414) 784-8087

CANADA

Intel Corp.
190 Atwell Drive, Suite 103
Bendale, Ontario
Canada K2H 8R2
Tel: (416) 675-2105

Intel Corp.
620 St. Jean Blvd.
Pointe Claire, Quebec
Canada H9R 3K2
Tel: (514) 694-9130

Intel Corp.
2650 Queensview Drive, #250
Ottawa, Ontario,
Canada K2B 8H6
Tel: (613) 829-9714

CUSTOMER TRAINING CENTERS

CALIFORNIA

2700 San Tomas Expressway
Santa Clara 95051
Tel: (408) 970-1700

ILLINOIS

300 N. Martingale, #300
Schaumburg 60173
Tel: (312) 310-5700

MASSACHUSETTS

3 Carlisle Road
Westford 01886
Tel: (617) 692-1000

MARYLAND

7833 Walker Dr., 4th Floor
Greenbelt 20770
Tel: (301) 220-3380

CALIFORNIA

2700 San Tomas Expressway
Santa Clara 95051
Tel: (408) 986-8086

ILLINOIS

300 N. Martingale, #300
Schaumburg 60173
Tel: (312) 310-8031

MASSACHUSETTS

3 Carlisle Road
Westford 01886
Tel: (617) 692-3222

NEW YORK

300 Motor Parkway
Hauppauge 11788
Tel: (516) 231-3300

SYSTEMS ENGINEERING OFFICES



UNITED STATES

Intel Corporation
3065 Bowers Avenue
Santa Clara, CA 95051

JAPAN

Intel Japan K.K.
5-6 Tokodai Tovosato-machi
Tsukuba-gun, Ibaraki-ken 300-26
Japan

FRANCE

Intel Paris
1 Rue Edison, BP 303
78054 Saint-Quentin en Yvelines
France

UNITED KINGDOM

Intel Corporation (U.K.) Ltd.
Piper's Way
Swindon
Wiltshire, England SN3 1RJ

WEST GERMANY

Intel Semiconductor GmbH
Seidlstrasse 27
D-8000 Munchen 2
West Germany

HONG KONG

Intel Semiconductor Ltd.
1701 Connaught Centre
1 Connaught Road
Hong Kong