Infinite Rules for Rounding
Binary Floating Point Arithmetic
Ignoring Over/underflow.

The following rules would have to be amended only slightly to accommodate overflow, which is a nearly independent and much more complicated topic. For simplicity here we consider the "representable numbers" to be an infinite discrete subset of the continuum of real numbers.

(1) The representable numbers must include 0, 1 and, if x then -x too.

(2) Each representable number must be represented uniquely by a finite symbol string that represents nothing else.

(3) Each "arithmetic operation" which, when executed without roundoff error would produce a representable number, must actually be executed without error.

(4) Never discard information unnecessarily.

(5) Each arithmetic operation which must be executed with roundoff error must result in a representable number that is as close as possible to what would have been produced in the absence of roundoff error.

(6) The preceding rule is ambiguous when two representable numbers are nearest the unrounded result. This ambiguity is resolved in a systematic way which preserves sign symmetry (e.g. $x-y = -(y-x)$) and is "unbiased" in the sense that "drift" cannot occur; e.g. the sequence $x_0, x_1, \ldots$ defined for arbitrary x_0 and y by $x_{n+1} := (x_n + y) - y$ $= x_2 = x_3 = \cdots$

(7) The basic operations include $+$, $-$, $\times$, $/$, $|\cdot|$, and conversion; they are to be extended to include $\sqrt{\ }$ and other FORTRAN functions.

(8) The rules above were slightly relaxed.

W. Kahan
Univ of Calif @ Berkeley
Sept. 20, 1977